

## CS411 Selected Lecture Notes

### This is one big WEB page, used for printing

These are not intended to be complete lecture notes.  
Complicated figures or tables or formulas are included here  
in case they were not clear or not copied correctly in class.  
Computer commands, directory names and file names are included.  
Specific help may be included here yet not presented in class.  
Source code may be included in line or by a link.

Lecture numbers correspond to the syllabus numbering.

### Contents

- [Lecture 1, Introduction, terminology](#)
- [Lecture 2, Benchmarks](#)
- [Lecture 3, Performance](#)
- [Lecture 4, CPU Operation](#)
- [Lecture 5, Instructions and Registers](#)
- [Lecture 6, VHDL introduction](#)
- [Lecture 7, Arithmetic](#)
- [Lecture 8, ALU](#)
- [Lecture 9, Multiply](#)
- [Lecture 10, Divide](#)
- [Lecture 11, Floating Point](#)
- [Lecture 12, VHDL - circuits and debugging](#)
- [Lecture 13, Microporgramming - Review](#)
- [Lecture 14, mid-term exam](#)
- [Lecture 15, Control Unit](#)
- [Lecture 16, Pipelining 1](#)
- [Lecture 17, Pipelining 2](#)
- [Lecture 18, Project outline and VHDL](#)
- [Lecture 19, Pipelining Data Forwarding](#)
- [Lecture 20, Hazards and Stalls](#)
- [Lecture 21, Cache](#)
- [Lecture 22, Cache Performance](#)
- [Lecture 23, Virtual Memory 1](#)
- [Lecture 24, Virtual Memory 2](#)
- [Lecture 25, I/O types and performance](#)
- [Lecture 26, DVR, DVD-RW, CDR, CD-RW](#)
- [Lecture 27, Busses, I/O-processor connection](#)
- [Lecture 28, Multiprocessors](#)
- [Lecture 29, Review](#)
- [Lecture 30, Final Exam](#)
- [Other Links](#)

### Lecture 1, Introduction, terminology

#### Introduction:

Hello, my name is Jon Squire and I have been programming computers since 1959. I have served my time in corporate management for 25 years. This course covers a little history of computer architecture through some of the latest advances and practical information you may use in buying, upgrading or building your own computer.

After this course, you can say that you have performed "modeling and simulation" possibly a valuable asset in finding a job. You will be skilled in converting graphical and schematic

information to textual information and the reverse.

#### Some Brief History:

The ISA card slots were replaced by PCI card slots that are replaced by external USB devices. The serial port for RS232 devices is replaced by the USB port. Floppy disk are disappearing along with that connector on the motherboard. RAM still uses DIMM's and the slots have grown to handle 4, 8 and 16 gigabytes of memory. ATA hard drives are replaced by SATA hard drives, 5TB and more available. Some rotating hard drives are being replaced by SSD, solid state drives. The printer port will be going as will the AGP graphics connector. HDMI and now DP. That expensive graphics card you bought will probably not work in your new computer.

[I have been saving architecture news.](#)

#### Overview:

This course will present detailed information on the internal working of the CPU, cache, memory, busses and peripheral devices such as disk drives and DVD's. The course five part project will have each student simulate a small computer using the VHDL digital simulation language. Either Cadence VHDL or free GHDL.

[Read the syllabus.](#)

Much of the lectures are covered in these WEB pages.

Lecture notes are often updated just before class.

And, sometimes corrected after class. :)

Some information is still presented on the blackboard/whiteboard.

Check UMBC "Blackboard" for announcements and grades.

The Top 500 Multiprocessor systems are evaluated about every six months. These are not your typical home computers.

The Top 10 are shown [www.top500.org/list/2018/06](http://www.top500.org/list/2018/06)  
Over 2 million cores!

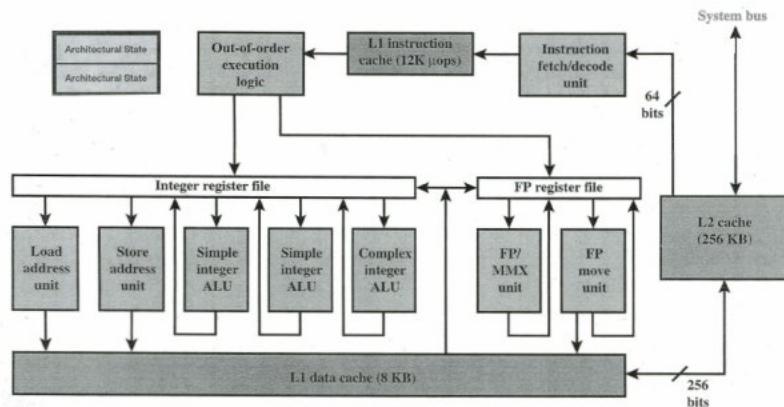
[More Lecture 1, pdf format](#)

The free market system and resulting competition, provide better and more economical products to consumers. Expect flip-flop between vendors for best or most economical products.

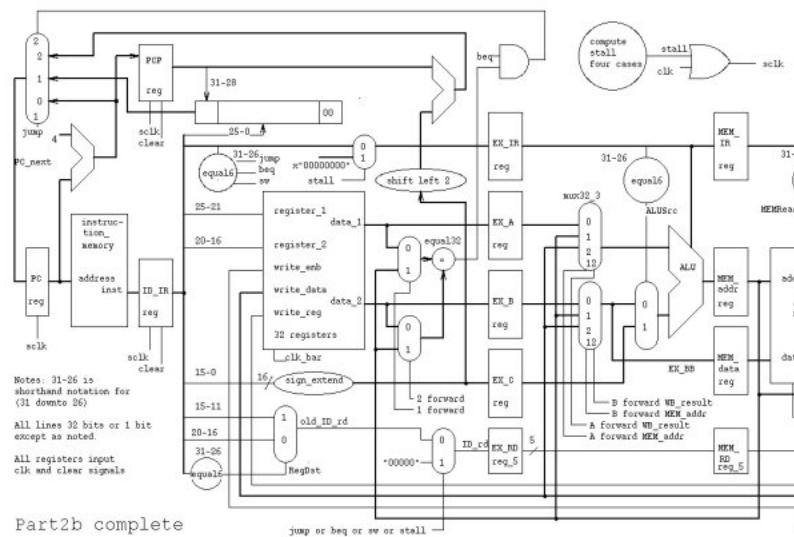
A standard engineering statement is:  
Fast, Cheap, Reliable - pick any two.

Monopolies: Ford Motor Company, Standard Oil of New Jersey, IBM, AT&T, ...  
Microsoft.

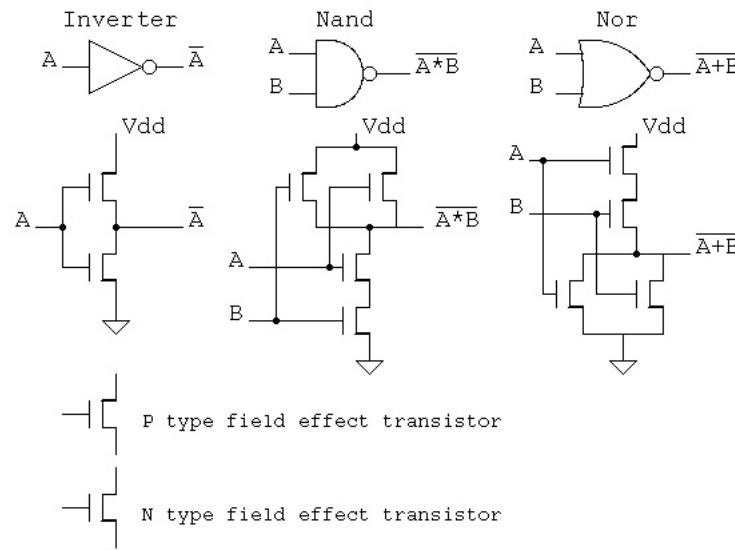
Computer Architecture Development:  
System Architecture



### Logic Design



### Circuit Design



Basic CMOS circuits

### The 6 transistor Memory Cell

The basic cell for static memory design is based on 6 transistors, with two pass gates instead of one. The corresponding schematic diagram is given in Figure 10-16. The circuit consists again of the 2 cross-coupled inverters, but uses two pass transistors instead of one. The cell has been designed to be duplicated in X and Y in order to create a large array of cells. Usual sizes for Megabit SRAM memories are 256 column x 256 rows or higher. A modest arrangement of 4x4 RAM cells is proposed in figure 10-16. The selection lines  $WL$  concerns all the cells of one row. The bit lines  $BL$  and  $\sim BL$  concern all the cells of one column.

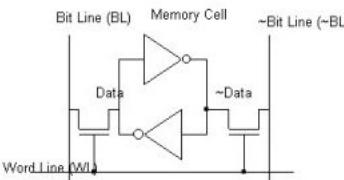
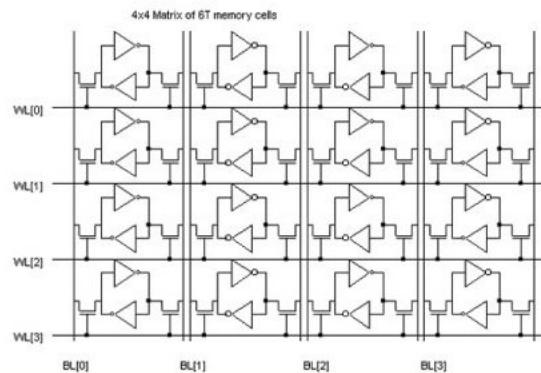


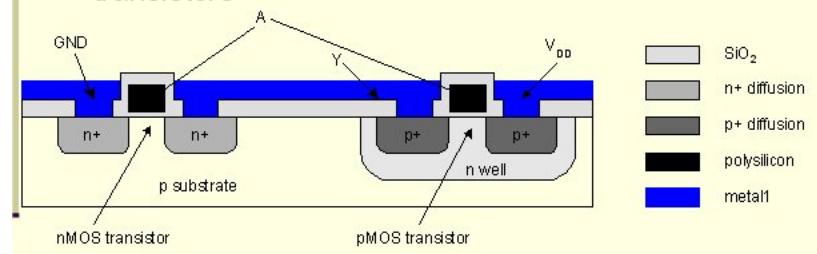
Figure 10-16: The layout of the 6 transistor static memory cell



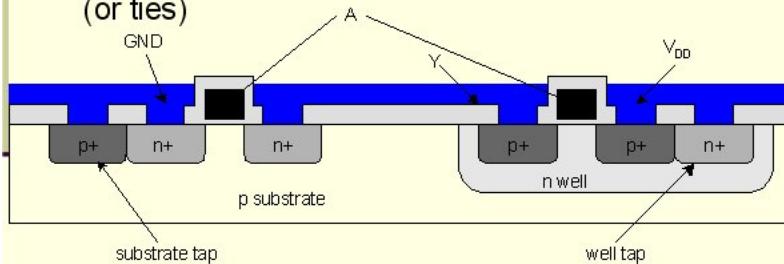
### Device Physics

For the inverter above, a chip cross section is:

- Typically use p-type substrate for nMOS transistors
- Requires to make an n-well for body of pMOS transistors



- Substrate must be tied to GND and n-well to  $V_{DD}$
- Metal to lightly-doped semiconductor forms poor connection called Schottky Diode
- Use heavily doped well and substrate contacts/taps (or ties)

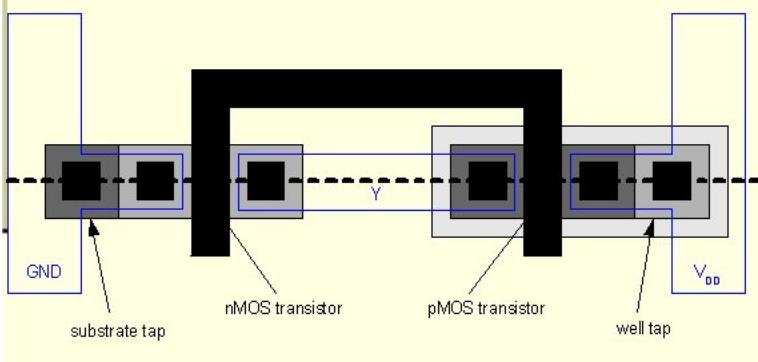


N type and P type impurities are diffused into the silicon substrate through a mask, typically in a high temperature vacuum process.

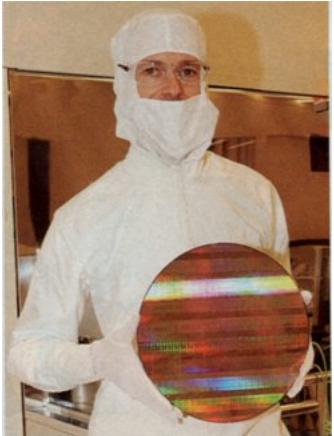
Oh! Oh! It is now predicted that Moore's Law:  
The gate width of transistors will halve every 18 months,  
will end in 2021. Prior estimates ended in 2028.  
Never fear, monolithic 3D is here.

#### Mask Making and Processing

- Transistors and wires are defined by masks
- Cross-section taken along dashed line



The black would be a metalization mask, here showing the transistor input connection. Other masks are for P+, N+, N well and via (the etch through the SiO<sub>2</sub> to allow electrical connection to metal.)



The large round wafer, after processing with all the masks, is broke up into many rectangular dies. Each die is placed in a package and the input and output pads on the die are connected to the pins on the package. The die in the package is called a chip or IC chip or Integrated Circuit Chip.

"Feature size" is the smallest dimension of metal width, gate width, metal spacing, etc. coming 12 nanometers is 0.000 000 012 meter or less than 1 millionth of an inch.

# AMD Ships Power-Saving 65-Nanometer Chips

ADVANCED Micro Devices is shipping chips produced using the 65-nanometer (nm) process, a move that's rekindling its manufacturing rivalry with Intel. The first chips produced are for desktop PCs; notebook and server chips will follow in the near future.

AMD says chips made with the 65nm process will consume about 30 percent less energy than those produced with the 90nm process at the same speed. With the first 65nm chips shipped, virtually all of the benefit comes from reduced power consumption. Later, the company will balance energy conservation with performance gains, depending on what the designers want to achieve with various desktop-, notebook-, and server-chip designs.

AMD made improvements,

in part, by straining the silicon in the transistors inside the chips, using a silicon germanium film to improve performance. While that's a first for AMD, Intel already uses sil-

icon germanium extensively. Straining silicon improves the performance of transistors because the larger germanium atoms slightly rearrange the silicon atoms, allowing electrical carriers to move more rapidly.



AMD's new chips once

again put it in the middle of a manufacturing battle with Intel, which first began shipping 65nm chips in October 2005. Besides providing more performance and/or consum-

ing less power than 90nm chips, the more-advanced 65nm chips also cost less to produce.

Intel's 14-month lead in manufacturing has been one of the primary reasons it has been able to undercut AMD in some segments. AMD reiterated the company's goal of shipping chips produced using the

45nm process in 18 months, a shorter-than-normal time period between manufacturing nodes. If successful, this will trim Intel's manufacturing advantage to about six or seven months. —Michael Kanellos

cnet News.com

AMD Opteron™ 6300 Series Processors						
Model Number	Cores	Core Speed	AMD Turbo CORE Max Frequency	L3 Cache	TDP	Socket Type
6386SE	16	2.8GHz	3.5GHz	16MB	140W	G34
6380	16	2.5GHz	3.4GHz	16MB	115W	G34
6378	16	2.4GHz	3.3GHz	16MB	115W	G34
6376	16	2.3GHz	3.2GHz	16MB	115W	G34
6370P	16	2.0GHz	2.5GHz	16MB	99W	G34
6348	12	2.8GHz	3.4GHz	16MB	115W	G34
6344	12	2.6GHz	3.2GHz	16MB	115W	G34
6338P	12	2.3GHz	2.8GHz	16MB	99W	G34
6328	8	3.2GHz	3.8GHz	16MB	115W	G34
6320	8	2.8GHz	3.3GHz	16MB	115W	G34
6308	4	3.5GHz	N/A	16MB	115W	G34
6366 HE	16	1.8GHz	3.1GHz	16MB	85W	G34

This gets smaller every year or so.

## BYOC

Build your own computer

Everything you need to build your own PC!

Check our website for latest Prices

**Gigabyte/Intel® Ivy Bridge Unlocked Barebones Kit**

- Gigabyte Intel Z77 CrossfireX ATX Motherboard
- Intel® 3rd Gen Core® i5-3570K 3.40 GHz CPU
- Thermaltake Commander MS-I Snow Edition Mid-Tower Case with USB 3.0
- Seagate 1TB ATA Hard Drive
- ADATA Premier Series 8GB DDR3 Memory (2x4GB)
- Sony Optiarc 24X DVDRW SATA OEM
- Ultra 650-Watt Power Supply

**Complete Bundle**  
**\$619<sup>99</sup>\***  
B69-1412  
*\*After Mail-In Rebate*

## Cooler Master HAF XM Mid Tower Case

Mid Tower Performance Case - ATX, 8x 3.5" Bays, 3x 5.25" Bays, 9x Expansion Slots, USB 3.0, SATA Dock

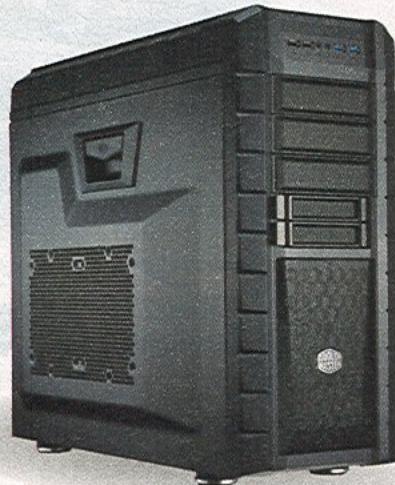
**ATX**  
**MID-TOWER**

**3** EXTERNAL  
5.25" BAYS

**6** INTERNAL  
3.5" BAYS

**2** EXTERNAL  
3.5" BAYS

**\$129<sup>99</sup>**



C283-2080

## Corsair Carbide Series 500R

Mid Tower Gaming Case - ATX, mATX, 4x Ext 5.25" Bays, 6x Int 3.5" Bays, 1x 200mm and 2x 120mm White LED Fans, 1x 120mm Fan, 2x USB 3.0 and 1x FireWire Ports

**ATX**  
**MID-TOWER**

**4** EXTERNAL  
5.25" BAYS

**6** INTERNAL  
3.5" BAYS

**\$134<sup>99</sup>**

C13-6500



## ThermalTake Overseer RX-1

Full Tower Gaming Case - ATX, Micro-ATX, Extended ATX, 3 Ext 5.25", 1 Ext 3.5", 5 Int 3.5", 2x 200mm LED Fans, 1x 120mm Fan, 2 USB 3.0, 2 USB 2.0, 1 eSATA, Window

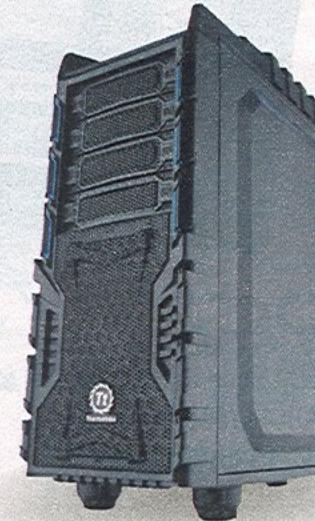
**ATX**  
**FULL-TOWER**

**3** 5.25" BAYS

**2** USB 3.0

**\$119<sup>99</sup>**

T925-7006



## Cooler Master ATX Mid Tower Case

ATX, Micro ATX, 4x 3.5" Bays, 4x 5.25" Bays, 4x USB, 200mm Red LED Fan

**\$79<sup>99</sup>**

C283-2071



## Cooler Master Mid-Tower ATX Case

ATX, MicroATX, 2, USB, Audio, 5x 3.5" Bays, 3x 5.25" Bays, 120mm Blue LED Fan

**\$49<sup>99</sup>**

C283-3122

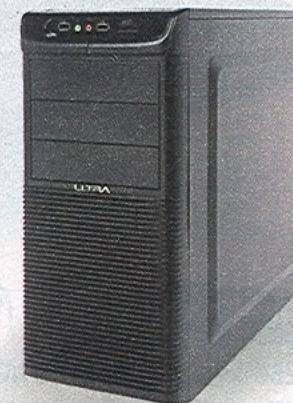


## Ultra X-Blaster Mid-Tower V2 Case

ATX, MicroATX, USB 2.0, Audio, 6x 3.5" Bays, 4x 5.25" Bays, 7x PCI Slots, Lifetime Warranty with Registration

**\$29<sup>99</sup>**

U12-41570

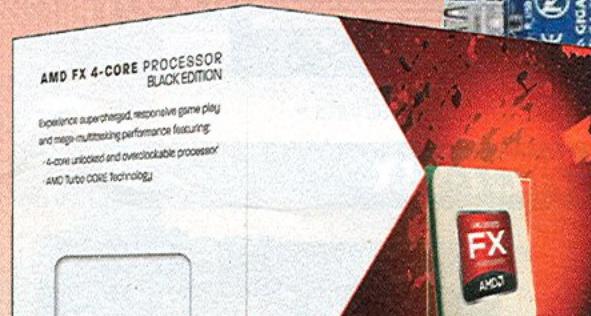


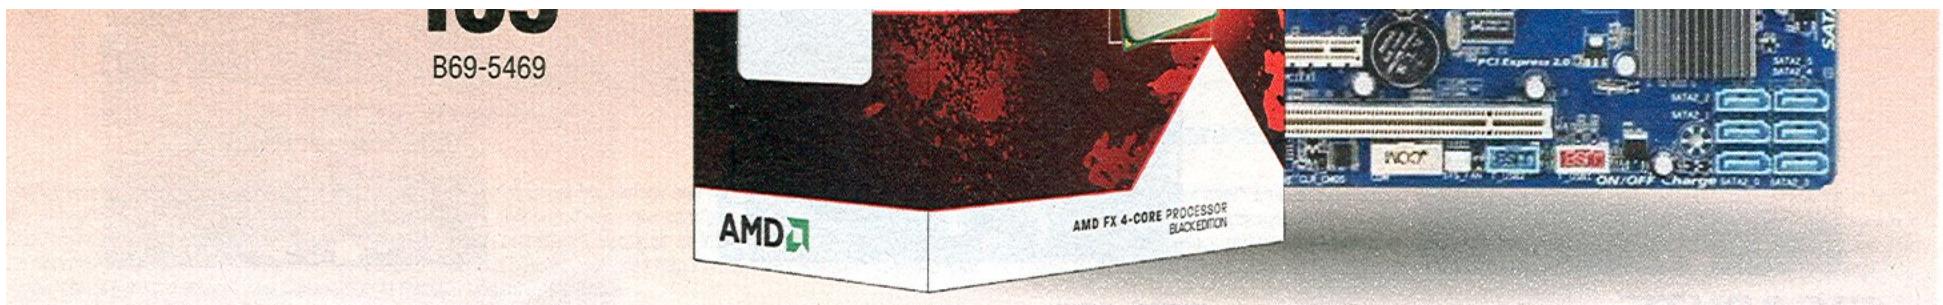
# Performance Motherboards

### GIGABYTE GA-78LMT-S2P AMD 760 AM3+ MOTHERBOARD AND AMD FX-4100 3.60 GHZ QUAD CORE AM3+ UNLOCKED PROCESSOR BUNDLE

- AMD FD4100WMGUSBX FX-4100
  - 8MB L3 Cache - 2MB L2 Cache - 3.60GHz (3.80GHz Max Turbo) - Quad Core
  - Socket Am3+ - Fan Included - Unlocked
- GIGABYTE GA-78LMT-S2P AMD 760 Motherboard
  - Micro ATX - Socket AM3+ - 1333MHz DDR3
  - SATA II (3Gb/s) - Gigabit LAN - USB 2.0

**\$169<sup>99</sup>**

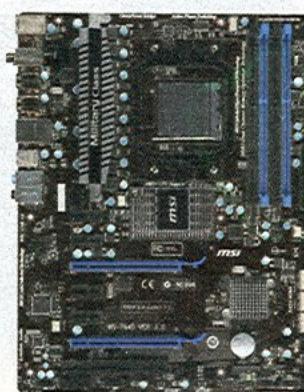




**ASUS**  
**Sabertooth  
990FX AMD  
AM3+ TUF  
Motherboard**

**\$189<sup>99</sup>**  
A455-3019

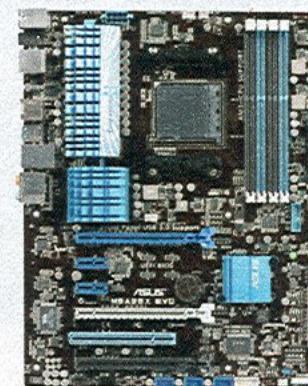
- 1866MHz-DDR3 • SATA 6.0 Gb/s • Gigabit LAN
- SuperSpeed USB 3.0 • SLI/CrossFireX Ready



**msi**  
**990FXA-GD65  
V2 AMD 990  
Motherboard**

**\$129<sup>99</sup>**  
M452-8440

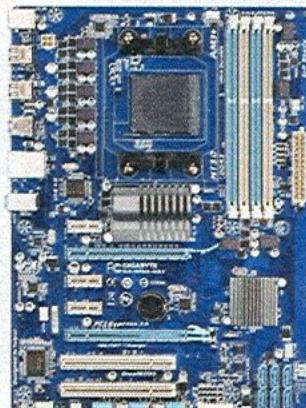
- 2133MHz-DDR3 • 6x SATA III Ports • Socket AM3+ • USB 3.0 • ATX • LAN



**ASUS**  
**M5A99X EVO  
AMD 990X  
AM3+  
Motherboard**

**\$149<sup>99</sup>**  
A455-3058

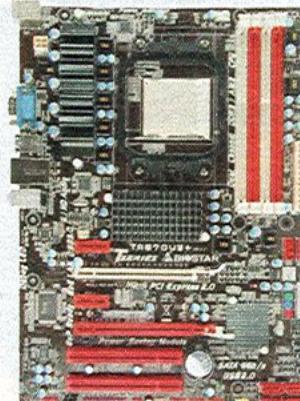
- 2133MHz-DDR3 • SATA 6.0 Gb/s • Gigabit LAN
- USB 3.0 • SLI/CrossFireX Ready



**GIGABYTE**  
**GA-970A-DS3  
AMD 970/SB950  
AM3+ FX ATX  
Motherboard**

**\$79<sup>99</sup>**  
G452-8442

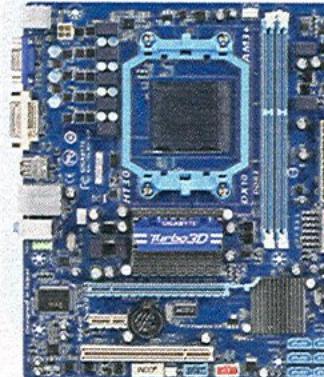
- Socket AM3 • 2000MHz DDR3 • SATA 6.0 Gb/s



**BIOSTAR**  
**TA870U3+  
AMD 870  
Motherboard**

**\$89<sup>99</sup>**  
B450-2425

- ATX • Socket Am3 • 1600MHz-DDR3



**GIGABYTE**  
**GA-78LMT-S2P  
AMD 760  
Motherboard**

**\$64<sup>99</sup>**  
G452-2609

- Micro ATX • Socket AM3+ • 1333MHz DDR3

- Gigabit LAN • USB 3.0 • CrossFire X Ready

- SATA 6.0 Gb/s • Gigabit LAN • USB 3.0

- SATA II (3Gb/s) • Gigabit LAN • USB 2.0

# Performance Processors



- 2x Corsair 4GB DDR3 RAM
- Seagate Barracude 500GB Hard Drive
- Sony Optiarc 24x DVDRW Drive



## MSI 970A-G46 AMD 9 SERIES AM3+ MOTHERBOARD BUNDLE

- AMD FX-8120 Processor
- Eight Core - 8MB L3 Cache
- 8MB L2 Cache - 3.10GHz (4.00 Max Turbo)
- MSI 970-A-G46 AMD 9 Series Motherboard
- ATX - AMD 970 Chipset - 2133MHz DDR3
- SATA III (6Gb/s) - Socket AM3+ - USB 3.0 - SLI/CrossFireX Ready - Gigabit LAN

**\$649<sup>99</sup>**

B69-0558

**FX-8150**



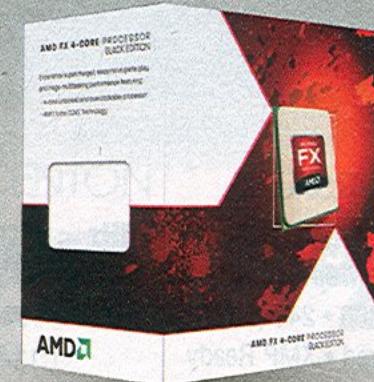
- Eight Core • 8MB L3 Cache • Fan Included
- 3.60GHz (4.20GHz Max Turbo) • Socket AM3+

**FX-8120**



- Eight Core • 8MB L3 Cache • 8MB L2 Cache
- Fan Included • 3.60GHz (4.00GHz Max Turbo)
- Socket AM3+

**FX-4100**



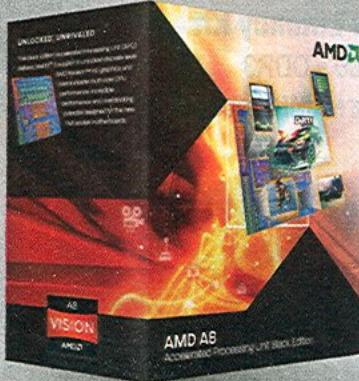
- Quad Core • 8MB L3 Cache • 2MB L2 Cache
- Fan Included • 3.60GHz (3.80GHz Max Turbo)
- Socket AM3+

**\$199<sup>99</sup>**  
A79-8150

**\$169<sup>99</sup>**  
A79-8120

**\$109<sup>99</sup>**  
A79-4100

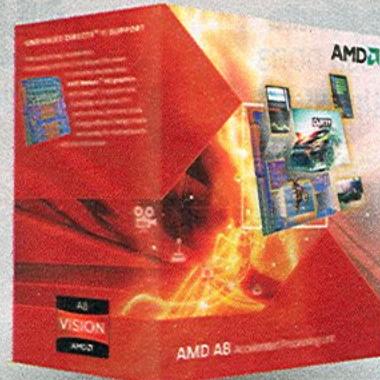
**A8-3870K BLACK EDITION**



- Quad Core • 4MB L2 Cache • Fan Included
- 3.0GHz • Socket FM1 • Radeon HD 6550D
- Dual Graphics Ready

**\$119<sup>99</sup>**  
A79-3870

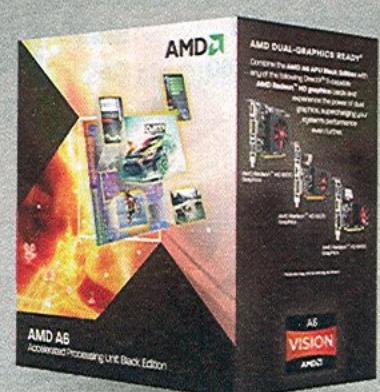
**A8-3850 APU**



- Quad Core • 4MB L2 Cache • Fan Included
- 2.9GHz • Socket FM1 • Radeon HD 6550D
- Dual Graphics Ready

**\$99<sup>99</sup>**  
A79-3850

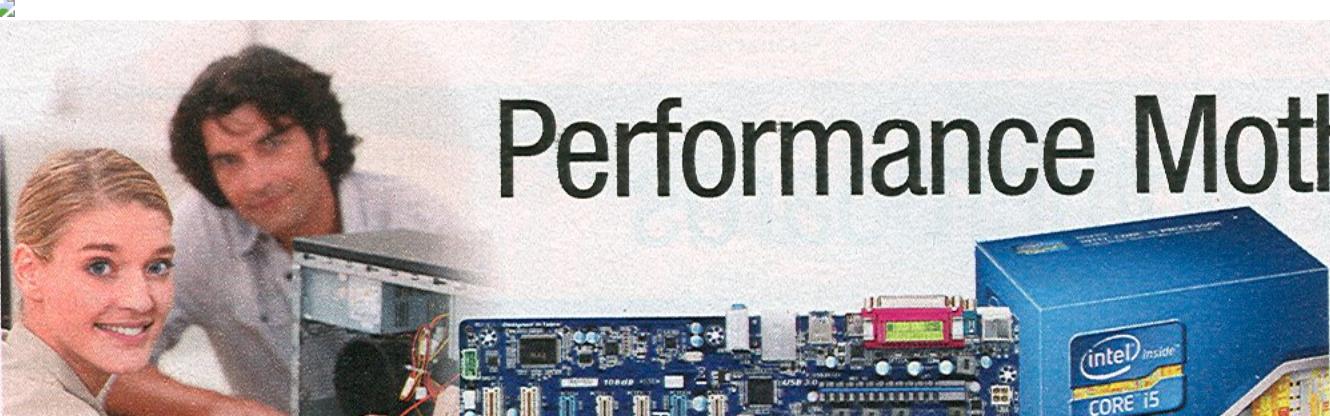
**A6-3670K BLACK EDITION**



- Quad Core • 4MB L2 Cache • Fan Included
- 2.7GHz • Socket FM1 • Radeon HD 6530D
- Dual Graphics Ready

**\$89<sup>99</sup>**  
A79-3670

# Performance Motherboards

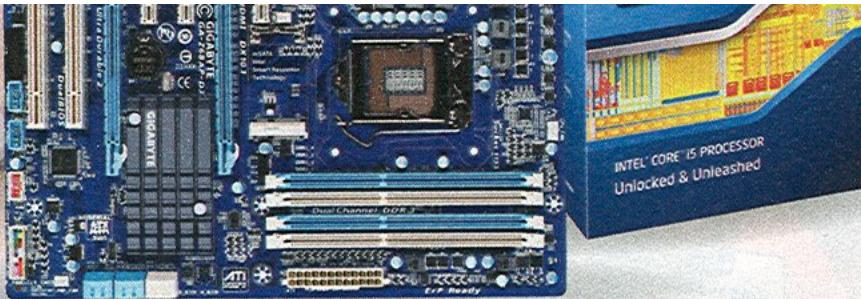


**GIGABYTE GA-Z68AP-D3  
MOTHERBOARD AND INTEL CORE  
I5-2500K PROCESSOR BUNDLE**

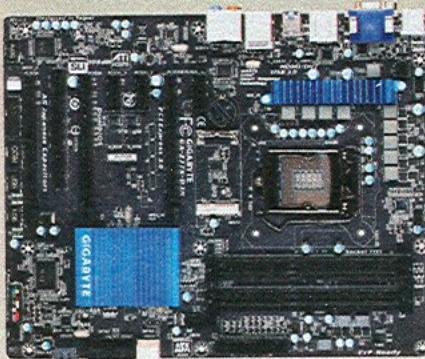
- 2nd Gen Intel® Core™ i5-2500K 3.30GHz



Scan to  
learn more.



### GA-Z77X-D3H



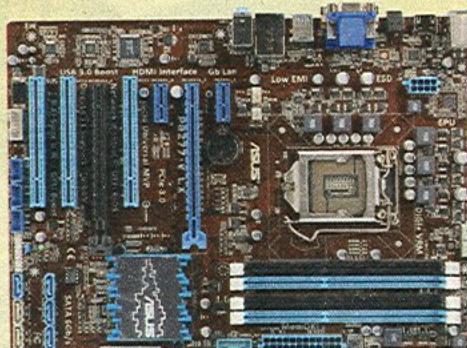
**GIGABYTE™**

- LGA115 socket H2; supports 2nd generation Intel® Core™ processors
- Intel Z77 Express • 1600MHz DDR3 • USB 3.0
- SATA III (6Gb/s) • 7.1-CH Audio • Gigabit LAN

**\$149<sup>99</sup>**

G452-2364

### P8Z77-V LX



**ASUS®**

- LGA115 socket H2; supports Intel® processors
- Intel Z77 Express • 2400MHz DDR3 • USB 3.0
- SATA III (6Gb/s) • 8-CH Audio • Gigabit LAN

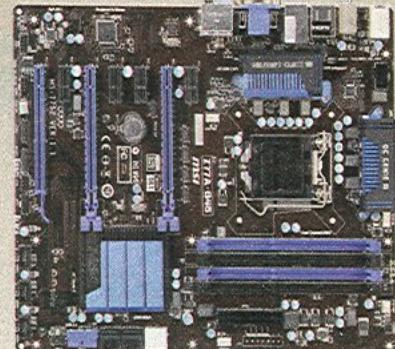
**\$139<sup>99</sup>**

A455-2042

### GA-Z77-DS3H



### Z77A-G45



**msi®**

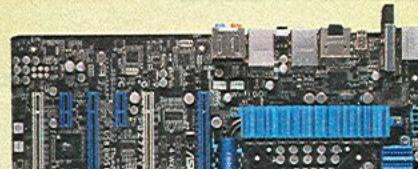
- LGA115 socket H2; supports Intel® processors
- Intel Z77 Express • 2667MHz DDR3 • USB 3.0
- SATA III (6Gb/s) • 8-CH Audio • Gigabit LAN

**\$139<sup>99</sup>**

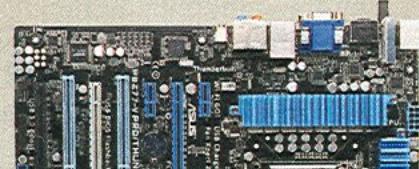
**SLI Ready**

M452-8420

### P8Z77-V PREMIUM



### P8Z77-V PRO/THUNDERBOLT



**GIGABYTE**

- LGA115 socket H2; supports 2nd generation Intel® Core™ processors
- Intel Z77 Express • 1600MHz DDR3 • USB 3.0
- SATA III (6Gb/s) • 7.1-CH Audio • Gigabit LAN

**\$119<sup>99</sup>**

G452-2366

**ASUS**

- LGA115 socket H2; supports 2nd generation Intel® Core™ processors
- Intel Z77 Express • 2600MHz DDR3 (O.C.) • RAID • USB 3.0
- SATA III (6Gb/s) • 8-CH Audio • Gigabit LAN

**\$449<sup>99</sup>**

A455-8425

SLI Ready

**ASUS**

- LGA115 socket H2; supports 2nd generation Intel® Core™ processors
- Intel Z77 Express • 2600MHz DDR3 (O.C.) • RAID • USB 3.0
- SATA III (6Gb/s) • 8-CH Audio • Gigabit LAN

**\$259<sup>99</sup>**

A455-8427

**NEW**  
3rd Generation  
Intel® Core™ i5 Processors  
with Intel® TURBO BOOST  
Technology 2.0

**intel inside™**  
**CORE™ i5**

**Introducing 3rd generation Intel® Core™ i5 processors.**

- Get an automatic burst of speed whenever you need it with Intel® Turbo Boost Technology 2.0<sup>1</sup>
- Experience smooth, beautiful movies, games, and more with Built-in Visuals<sup>2</sup>
- Effortlessly move between applications with Intel® Hyper-Threading Technology<sup>3</sup>

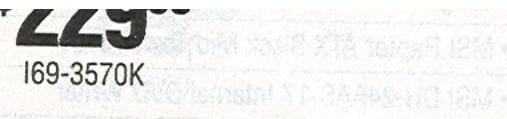
**Intel® Core™ i5-3570K Processor Unlocked**

- 3.40GHz (3.80Ghz Max Turbo)  
Unlocked Processor
- Quad Core

**\$220<sup>99</sup>**



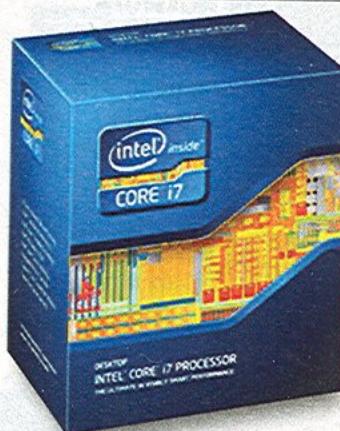
- Socket H2
- Fan Included



**Look for PCs with Intel® Core™ i5 processors**

**Sponsors of Tomorrow™ Intel®**

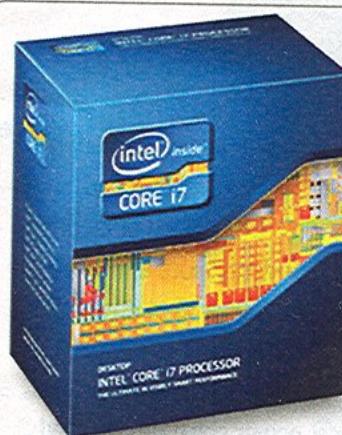
'Requires a system with Intel® Turbo Boost Technology. Intel® Turbo Boost Technology and Intel® Turbo Boost Technology 2.0 are only available on select Intel® processors. Consult your PC manufacturer. Performance varies depending on hardware, software and system configuration. For more information, visit <http://www.intel.com/go/turbo>. \*Built-in Visual features are not enabled on all PCs and optimized software may be required. Check with your system manufacturer. Learn more information including details on which processors support HT Technology, visit <http://www.intel.com/info/hyperthreading>. Copyright © 2012, Intel Corporation. All rights reserved. Intel Inside, Intel Sponsors of Tomorrow, Intel Core, Intel, and the Intel logo are trademarks of Intel Corporation in the U.S. and other countries.



**Intel® Core™  
i7-3770K  
Unlocked  
Quad-Core  
Processor**

**\$349<sup>99</sup>**  
I69-3770K

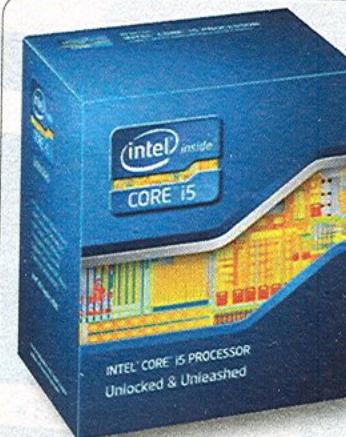
- 8MB L3 Cache • Quad Core CPU • Fan Included
- 3.50GHz (3.90GHz Max) • Socket R (LGA1155)



**Intel® Core™  
i7-3820  
Quad-Core  
Processor**

**\$299<sup>99</sup>**  
I69-3820

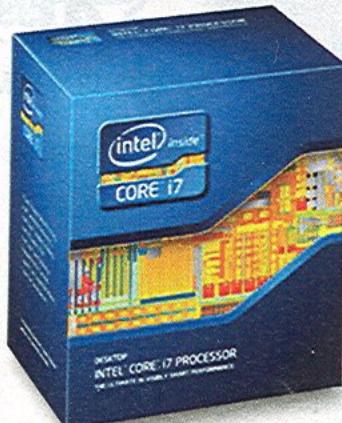
- 10MB L3 Cache • Quad Core CPU
- 3.60GHz (3.90GHz Max) • Socket R (LGA2011)



**Intel® Core™  
i5-3470  
Quad Core  
Processor**

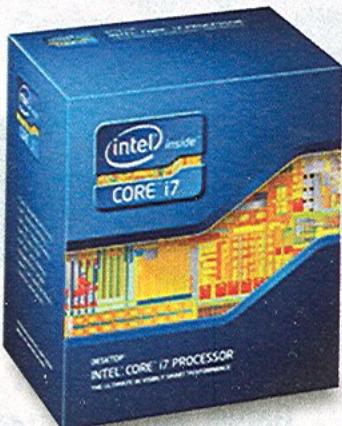
**\$199<sup>99</sup>**  
I69-3470

- 6MB L3 Cache • Quad Core
- 3.20GHz (3.60GHz Max) • Socket H2 (LGA1155)



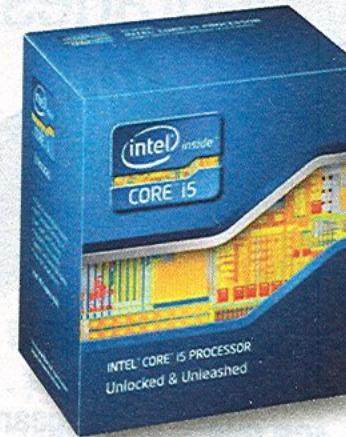
**Intel® Core™  
i7-2700K  
Unlocked  
Quad-Core  
Processor**

**\$309<sup>99</sup>**  
I69-2700K



**Intel® Core™  
i7-2600K  
Unlocked  
Quad-Core  
Processor**

**\$289<sup>99</sup>**  
I69-2600K



**Intel® Core™  
i5-2500K  
Unlocked  
Quad Core  
Processor**

**\$219<sup>99</sup>**  
I69-2500K

- 8MB L3 Cache • Quad Core CPU • Fan Included
- 3.50GHz (3.90GHz Max) • Socket R (LGA1155)

- 8MB L3 Cache • Quad Core CPU • Fan Included
- 3.40GHz (3.90GHz Max) • Socket R (LGA1155)

- 6MB L3 Cache • Quad Core CPU • Fan Included
- 3.30GHz (3.90GHz Max) • Socket R (LGA1155)

# High-Performance Memory Modules

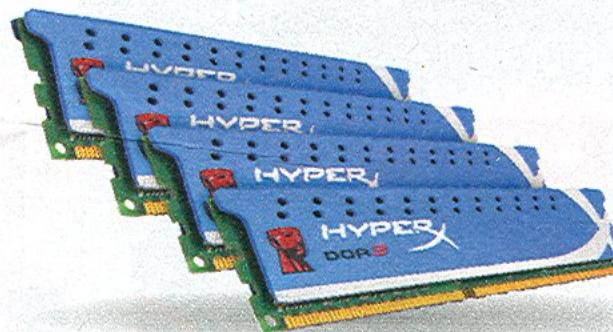
The best value for PC gamers seeking maximum system performance



**CORSAIR™**

**32GB (4x 8GB) DDR3-1600MHz Desktop Memory Kit** **\$259<sup>99</sup>**  
 • PC3-12800 • 240-pin DIMM  
 • Unbuffered • XMP Ready

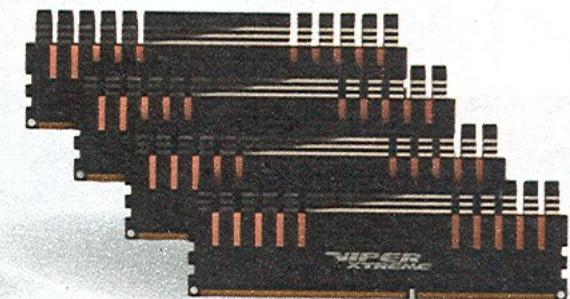
C13-8118



**Kingston® TECHNOLOGY**

**16GB (4x 4GB) DDR3-1866MHz Desktop Memory Module** **\$129<sup>99</sup>**  
 • 240-pin DIMM • PC3-14900  
 • XMP Ready

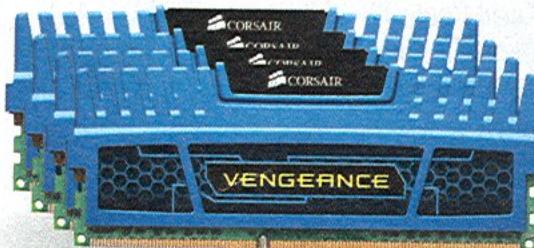
K24-8843



**PATRIOT® MEMORY**

**16GB Low Latency Quad Kit** **\$189<sup>99</sup>**  
 • PC3-17000 • DDR3  
 • 2133MHz

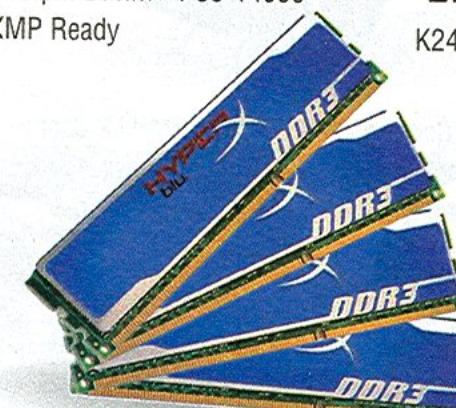
P33-5305



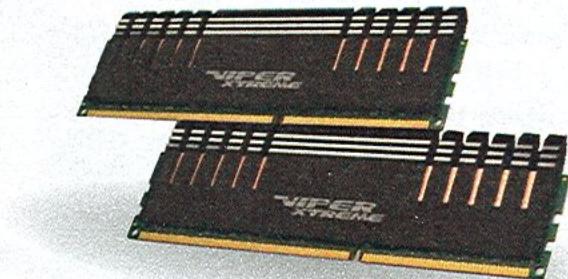
**CORSAIR™**

**16GB (4x 4GB) DDR3-1666MHz**

[https://userpages.umbc.edu/~squire/cs411\\_lect.html](https://userpages.umbc.edu/~squire/cs411_lect.html)



**Kingston® TECHNOLOGY**  
**16GB DDR3-1666MHz**



**PATRIOT® MEMORY**

**Desktop Memory Kit**

- PC3-12800 • Unbuffered
- XMP Ready

**\$109<sup>99</sup>**

C13-5724

**CORSAIR™****8GB (2x 4GB) DDR3-1866MHz Desktop Memory Kit**

- PC3-15000 • 240-pin DIMM
- Unbuffered • XMP Ready

**\$64<sup>99</sup>**

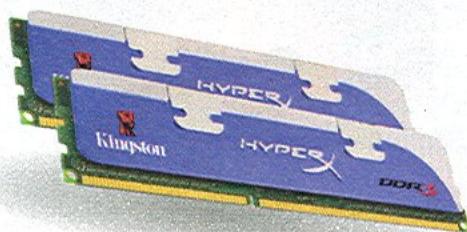
C13-5803

**1000 DDR3-1333MHz Memory Bundle**

- PC3-10600 DDR3
- DIMM Unbuffered

**\$94<sup>99</sup>**

K24-9978

**Kingston TECHNOLOGY****8GB (2x 4GB) DDR3-1600MHz Desktop Memory Kit**

- PC3-12800 • XMP Ready
- DIMM

**\$54<sup>99</sup>**

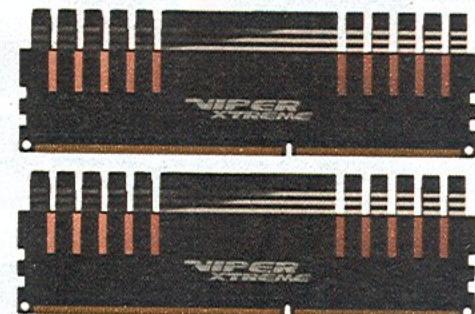
K24-9917

**16GB Low Latency Kit**

- PC3-15000 • DDR3
- 1866MHz • XMP Ready

**\$129<sup>99</sup>**

P33-1200

**PATRIOT MEMORY****8GB Enhanced Latency Kit**

- 2400MHz RoHS
- DDR3

**\$99<sup>99</sup>**

P33-2102

# Solid State Drives

These solid state drives provide the best selection for performance and price efficiency. A solid state drive offers you a faster response time and with no moving parts, are quieter and require less power.



See the difference and SSD can make on boot time!





**OCZ**  
Technology

- 128GB Solid State Drive
- SATA III (6Gb/s)
- 2.5" • 535MB/s Read

**\$114<sup>99</sup>**

0261-6408



**OCZ**  
Technology

- 256GB Solid State Drive
- SATA III (6Gb/s)
- 2.5" • 535MB/s Read

**\$209<sup>99</sup>**

0261-6409



**OCZ**  
Technology

- 512GB Solid State Drive
- SATA III (6Gb/s)
- 2.5" • 420MB/s Read

**\$379<sup>99</sup>**

0261-9803



**OCZ**  
Technology

- 60GB Solid State Drive
- SATA III (6Gb/s)
- 2.5" • 535MB/s Read

**\$69<sup>99</sup>**

0261-6379



**OCZ**  
Technology

- 120GB Solid State Drive
- SATA III
- 2.5" • 550MB/s Read

**\$109<sup>99</sup>**

0261-6380



**OCZ**  
Technology

- 240GB Solid State Drive
- SATA III
- 2.5" • 550MB/s Read

**\$199<sup>99</sup>**

0261-6381



 <ul style="list-style-type: none"> <li>• 128GB Solid State Drive</li> <li>• SATA • SATA III-6Gb/s</li> <li>• 2.5" • 320MB Write</li> </ul> <p><b>\$149<sup>99</sup></b> S203-1041</p>	 <ul style="list-style-type: none"> <li>• 256GB Solid State Drive</li> <li>• SATA • SATSA III</li> <li>• 2.5" • 520MB Read</li> </ul> <p><b>\$299<sup>99</sup></b> S203-1042</p>	 <ul style="list-style-type: none"> <li>• 512GB Solid State Drive</li> <li>• SATA • SATA III</li> <li>• 2.5" • 520MB Read</li> </ul> <p><b>\$699<sup>99</sup></b> S203-1043</p>
--	---	---

[A+ Products](#) ► [Motherboards](#) ► [ [H8GL-F](#) ]

#### Key Features



1. Single AMD Opteron™ 6100 Series processors (G34) Twelve/Eight-Core ready; HT3.0 Link support
2. AMD SR5650 + SP5100 Chipset
3. Up to 128GB DDR3 Registered ECC or 32GB DDR3 Unl. ECC/non-ECC 1333/1066/800 SDRAM in 8 DIMMs
4. Two Intel® 82574L controllers, Dual-Port Gigabit Ethernet
5. 6x SATA2 3.0 Gbps Ports via AMD SP5100 controller, RAID 0, 1, 10
6. 1/1/1/3 PCI-E 2.0 x8 (in x16) / x8 / x4 (in x8) / PCI
7. Integrated Matrox G200 Graphics
8. IPMI 2.0 with Dedicated LAN support

I have built several computers buying a case, motherboard, cpu, ram, drives, video, audio.

My older desktop, AMD FX 8-core is Cybertron G1244A  
16 GB ram, 1/2 TB SSD.  
(Replacing my old 12 core AMD that is acting up.)  
Now new Dell Precision 7920 Tower with 16 cores.

You want DDR3, SATA3, SSD we will cover these in future lectures.

[Look at Homework 1, it is assigned today.](#)

## Lecture 2, Benchmarks

The best method of measuring a computers performance is to use benchmarks. Some suggestions from my personal experience preparing a benchmark suite and several updates and personal benchmark experience are presented in pdf format.

#### [Lecture 2](#)

Smaller time is better, higher clock frequency is better.  
 $\text{time} = 1 / \text{frequency}$     $T = 1/F$    and    $F = 1/T$   
 1 nanosecond = 1 / 1 GHz  
 1 microsecond = 1 / 1 MHz

**Definitions:**

CPI    Clocks Per Instruction  
MHz    Megahertz, millions of cycles per second  
MIPS   Millions of Instructions Per Second = MHz / CPI  
MOPS   Millions of Operations Per Second  
MFLOPS Millions of Floating point Operations Per Second  
MIOPS Millions of Integer Operations Per Second

Do not trust your computers clock or the software  
that reads and processes the time.

First: Test the wall clock time against your watch.

[time\\_test.c](#)  
[time\\_test.java](#)  
[time\\_test.py](#)  
[time\\_test.f90](#)

The program displays 0, 5, 10, 15 ... at 0 seconds,  
5 seconds, 10 seconds etc.

demonstrate time\_test if possible

Note the use of <time.h> and 'time()'

Beware, midnight is zero seconds.  
Then  $60 \text{ sec/min} * 60 \text{ min/hr} * 24 \text{ hr/day} = 86,400 \text{ sec/day}$   
Just before midnight is 86,399 seconds.  
Running a benchmark across midnight may give a negative time.

Then: Test CPU time, this should be just the time  
used by the program that is running. With only  
this program running, checking against your watch  
should work.

[time\\_cpu.c](#)  
[time\\_cpu.java](#)  
[time\\_cpu.py](#)

The program displays 0, 5, 10, 15 ... at 0 seconds,  
5 seconds, 10 seconds etc.

Note the use of <time.h> and  
 $'(double)clock()/(double)CLOCKS\_PER\_SEC'$

I have found one machine with the constant  
CLOCKS\_PER\_SECOND completely wrong and  
another machine with a value 64 that should  
have been 100. A computer used for real time  
applications could have a value of 1,000,000  
or more.

[More graphs of FFT benchmarks](#)

The source code, C language, for the FFT benchmarks:

Note the check run to be sure the code works.

Note the non uniform data to avoid special cases.

[fft\\_time.c](#) main program  
[fftc.h](#) header file

FFT and inverse FFT for various numbers of complex data points  
The same source code was used for all benchmark measurements.  
These were optimized for embedded computer use where all  
constants were burned into rom.

[fft16.c](#)    [ifft16.c](#)  
[fft32.c](#)    [ifft32.c](#)  
[fft64.c](#)    [ifft64.c](#)  
[fft128.c](#)    [ifft128.c](#)  
[fft256.c](#)    [ifft256.c](#)  
[fft512.c](#)    [ifft512.c](#)  
[fft1024.c](#) [ifft1024.c](#)  
[fft2048.c](#) [ifft2048.c](#)  
[fft4096.c](#) [ifft4096.c](#)

Some of the result files:

[P1-166MHz](#)  
[P1-166MHz -O2](#)  
[P2-266MHz](#)  
[P2-266MHz -O2](#)  
[Celeron-500MHz](#)  
[P3-450MHz MS](#)  
[P3-450MHz Linux](#)  
[PPC-2.2GHz](#)  
[PPC-2.5GHz](#)  
[P4-2.53GHz XP](#)  
[Alpha-533MHz XP](#)  
[Xeon-2.8GHz](#)  
[Athlon-1.4GHz MS](#)  
[Athlon-1.4GHz XP](#)  
[Athlon-1.4GHz SuSe](#)  
[Laptop Win7](#)  
[Laptop Ubuntu](#)

What if you are benchmarking a multiprocessor?

For example, a two core or quad core, then use both CPU time  
and wall time to get average processor loading:

[time\\_mp2.c for two cores](#)  
[time\\_mp4.c for quad cores](#)  
[time\\_mp8.c for two quad cores](#)  
[time\\_mp12.c for two six cores](#)

The output from a two cores is:

[time\\_mp2.c.out for two core Xeon](#)

The output from four cores is:

[time\\_mp4.c.out for Mac quad G5](#)

The output from eight cores is:

[time\\_mp8.c.out for AMD 12-core](#)

The output from twelve cores is:

[time\\_mp12.c.out for AMD 12-core](#)

```
end of time_mp12_c.out file:
total CPU time is 342.970000 seconds
wall time is 29.000000 seconds
average number of processors used = 11.826552
time_mp12.c exiting
```

Similar tests in Java

[time\\_test.java](#)  
[time\\_cpu.java](#)  
[time\\_mp4.java for quad cores](#)  
[time\\_mp8.java for eight cores](#)  
[time\\_mp8.java for eight and twelve cores](#)  
[time\\_mp8.java.out for quad Xeon G5](#)  
[time\\_mp8.java.out for 8 thread Xeon G5](#)  
[time\\_mp8.java\\_fx.out for 8 core AMD FX](#)  
[time\\_mp12.java.out for 8 thread Xeon G5](#)  
[time\\_mp12\\_12.java.out for 12 core AMD](#)  
[matmul\\_thread4.java](#)  
[matmul\\_thread4.java.out](#)

Time\_test and threads in Python

[time\\_test.py](#)  
[time\\_cpu.py](#)  
[parallel\\_matmul.py](#)  
[parallel\\_matmul\\_py.out](#)

OK, since these were old and I did not want to change them,  
they give some indications of performance on various machines  
with various operating systems and compiler options.

To measure very short times, a higher quality, double-difference  
method is needed. The following program measures the time  
to do a double precision floating point add. This may be  
a time smaller than 1ns,  $10^{-9}$  seconds.

A test harness is needed to calibrate the loops and make sure  
dead code elimination can not be used by the compiler.

The the item to be tested is placed in a copy of the test harness  
to make the measurement.

The time of the test harness is the stop minus start time in seconds.

The time for the measurement is the stop minus start time in seconds.

The difference, thus double difference, between the harness and  
measurement is the time for the item being measured.  
Here A = A + B with B not known to be a constant by the compiler,  
is reasonably expected to be a single instruction to add B to  
a register. If not, we have timed the full statement.

The double difference time must be divided by the total  
number of iterations from the nested loops to get the  
time for the computer to execute the item once.

An attempt is made to get a very stable time measurement.  
Doubling the number of iterations should double the time.

Summary of double difference

```
t1 saved
run test harness
t2 saved

t3 saved
run measurement, test harness with item to be timed
t4 saved
tdiff = (t4-t3) - (t2-t1)
t_item = tdiff / number of iterations

check against previous time, if not close, double iterations
```

The source code is:

[time\\_fadd.c](#)  
[fadd on P4 2.53GHz](#)  
[fadd on Xeon 2.66GHz](#)  
[fadd on Mac 2.5GHz](#)

```
end of Mac output:
time_fadd.c
...
rep=16384, t measured=0.814363
rep=32768, t measured=1.62344
rep=65536, t measured=3.28666
tmeas=3.28666, t_prev=0, rep=65536
rep=65536, t measured=3.28829
tmeas=3.28829, t_prev=3.28666, rep=65536
time measured=3.28829, under minimum
raw time=3.28829, fadd time=5.01629e-10, rep=65536, stable=0.000497342
```

Some extra information for students wanting to explore their computer:

Windows OS

Linux OS

What is in my computer?

start	cd /proc
control panel	cat cpuinfo
system	
device manager	
processor	
etc.	

What processes are running in my computer?

ctrl-alt-del	ps -el
process	top

How do I easily time a program?	time prog < input > output
command prompt	
time	
prog < input > output	
time	

The time available through normal software calls may be updated less than 30 times per second to more than a million times per second. A general rule of thumb is to have the time being measured be 10 seconds or more. This will give a reasonable accurate time measurement on all computers. Just repeat what is being measured if it does not run 10 seconds.

Some history about computer time reporting.

There were time sharing system where you bought time on the computer by the cpu second. There is the cpu time your program requires that is usually called your process time. There is also operating system cpu time. When there are multiple processes running, the operating system time slices, running each job for a short time, called a quanta. The operating system must manage memory, devices, scheduling and related tasks. In the past we had to keep a very close eye on how cpu time was charged to the users process verses the systems processes and was "dead time" the idle process, charged to either. From a users point of view, the user did not request to be swapped out, thus the user does not want any of the operating system time for stopping and restarting the users process to be charged to the user.

Another historic tidbit, some Unix systems would add one microsecond to the time reported on each system request for the time. Never allowing the same time to be reported twice even if the clock had not updated. This was to ensure that all disk file times were unique and thus programs such as 'make' would be reliable.

For more recent SPEC benchmarks, [2006 is suit date, run 2015, 2016, 2017, 2018, 2019](#)

[see CPU integer benchmarks, SPECint, floating point benchmarks, SPECfp](#)  
[www.spec.org/cpu2006/Docs/](#)

Some times you just have to buy the top of the line and forget benchmarks.

**REVIEWS** COMPONENTS**AMD ATI Radeon X1950 Pro**

Impressive but pricey graphics

AMD'S ATI RADEON X1950 Pro graphics card is designed to replace the former midrange Radeon-line stalwart, the X1800 GTO. The card features some impressive numbers: a 575MHz core clock speed, and 256MB of DDR3 RAM running at 1,380MHz, faster than either the X1800 GTO or the competing nVidia GeForce 7900 GS. But you'll have to pony up for that performance, considering the X1950's \$299 price.

Like most current 3D cards, the Radeon X1950 Pro requires a connection to your PC's power supply, and AMD recommends at least a 550-watt power supply if it's used in dual-card CrossFire mode. In fact, the revamped CrossFire design on this card is probably more exciting than the Radeon



2,560x2,048, higher than the top-end resolution of 2,560x1,600 that nVidia's competing Scalable Link Interface (SLI) dual-card technology currently offers.

In our tests, the Radeon X1950 Pro stacked up well against the GeForce 7900 GS, in both single- and dual-card mode. On our high-resolution Oblivion test, a pair of Radeon X1950 Pros looked especially impressive, clocking nearly 45 frames per second on this notoriously difficult benchmark test. The GeForce took the lead on our Quake 4 test, however, which surprised us, since Radeon cards have done well on that test of late.

Overall, the Radeon X1950 Pro edges out the GeForce 7900 GS slightly on performance. But, given the steep cost of the Radeon, we can't recommend it over the less-expensive nVidia cards. —Rich Brown

CNET.com

**EDITORS' RATING** 7.0

BETTER...

**PROS** New CrossFire design is affordable, easy to set up; beats nVidia's competing card in certain tests

**CONS** nVidia's card wins out on some games; street price is high

Advanced Micro Devices  
888-974-6728  
at.amd.com

Direct Price \$299

Now find a display with 2,560 by 2,048 resolution!  
(other than the NASA display)

**NEW! 24" High-Definition Widescreen**

Gateway FPD2485W 24" Widescreen High-Definition LCD Flat-Panel Display

- 1920x1200 resolution
- 1000:1 contrast ratio
- 6ms response time
- Picture-in-picture

\$679.99

Newegg has an Acer 22 inch HDMI 1920 by 1080 for under \$100 in 2013  
HDMI replaces VGA connection from computer to display.

**Lecture 3, Performance**

Repeating some definitions:

CPI Clocks Per Instruction

MHz megahertz, millions of cycles per second

MIPS Millions of Instructions Per Second = MHz / CPI

MOPS Millions of Operations Per Second

MFLOPS Millions of Floating point Operations Per Second

MIOPS Millions of Integer Operations Per Second

(Classical, old, terms. Today would be billions.)

Amdahl's Law (many forms, understand the concept)

the part of time improved

`new time = ----- + the part of time not improved  
                  factor improved by`

`old time = the part of time improved + the part of time not improved`

`speedup = ----- (always bigger over smaller when faster)  
              old time  
              new time`

Given: on some program, the CPU takes 9 sec and the disk I/O takes 1 sec  
What is the speedup using a CPU 9 times faster?

9 sec

Answer: new time = ----- + 1 sec = 2 sec  
              9

old time = 9 + 1 = 10 sec

speedup = 10 / 2 = 5 a pure number

---

Amdahl's Law (many forms, understand the concept)

`new performance  
speedup = -----  
              old performance`

Given: Performance of M1 is 100 MFLOPS and 200 MIOPS  
Performance of M2 is 50 MFLOPS and 250 MIOPS  
On a program using 10% floating point and 90% integer  
Which is faster? What is the speedup?

Answer; .1 \* 100 + .9 \* 200 = 190 MIPS  
.1 \* 50 + .9 \* 250 = 230 MIPS (M2 is faster)

speedup = 230/190 = 1.21

---

`old performance  
new performance = -----  
                  fraction of old improved  
                  ----- + fraction of old unimproved  
                  improvement factor`

Given: half of a 100 MIPS machine is speeded up by a factor of 3  
what is the speedup relative to the original machine?

Answer: new performance = ----- \* 100 MIPS = ----- \* 100 MIPS = 150 MIPS  
              0.5                          .666  
              ----- + 0.5  
              3

speedup = 150 / 100 = 1.5 (same as -----)  
                  fraction improved  
                  ----- + fraction  
                  improvement factor    unimproved

speedup is a pure number, no units. The units must cancel.

---

SPEC Benchmarks

The benchmarks change infrequently, for example 2006 - 2016 same  
 The speed seems to increase every year.

SPEC Int2006, 9 in C, 3 in C++SPEC Flt2006, 17 in assorted Fortran, C, C++SPEC many rules to followrecent int resultsrecent flt results

Note number of core available, results seem to be using just one core.

---

CPI is average Clocks Per Instruction. units: clock/inst  
 MHz is frequency, we use millions of clocks per second. units: clock/sec  
 MIPS is millions of instruction per second. units: inst/sec  
 Note: MIPS=MHz/CPI because  $(\text{clock/sec}) / (\text{clock/inst}) = 10^6 \text{ inst/sec}$   
 ( 5/4 of people do not understand fractions. )

## Computing average CPI, Clocks Per Instruction

-----given----- -----compute-----

type	clocks	%use	product
RR	3	25%	3 * 25 = 75
RM	4	50%	4 * 50 = 200
MM	5	25%	5 * 25 = 125
<hr/>			
100%		400	sum

$$400/100 = 4 \text{ average CPI}$$

-----given----- -----compute-----

type	clocks	instructions	product
RR	3	25,000	3 * 25,000 = 75,000
RM	4	50,000	4 * 50,000 = 200,000
MM	5	25,000	5 * 25,000 = 125,000
<hr/>		100,000	400,000 sum

$$400,000/100,000 = 4 \text{ average CPI}$$

Find the faster sequence of instructions Prog1 vs Prog2

-----given-----

type	clocks
A	1
B	2
C	3
instruction counts for A B C	
Prog1	2 1 2

```

-----compute-----
Prog1
A   1   2      1 * 2 = 2
B   2   1      2 * 1 = 2
C   3   2      3 * 2 = 6
                  sum
                  10 clocks

Prog2
A   1   4      1 * 4 = 4
B   2   1      2 * 1 = 2
C   3   1      3 * 1 = 3
                  sum
                  9 clocks  more instructions yet faster

speedup = 10 clocks / 9 clocks = 1.111  a number (no units)

```

cs411\_OPCODES.txt different from Computer Organization and Design 1/8/2020

**rd** is register destination, the result, general register 1 through 31  
**rs** is the first register, A, source, general register 0 through 31  
**rt** is the second register, B, source, general register 0 through 31

--val--- generally a 16 bit number that gets sign extended  
--addr--- a 16 bit address, gets sign extended and added to (rx)  
"i" is generally immediate, operand value is in the instruction

Opcode	Operands	Machine code format									
		6	5	5	5	5	6	number of bits in field			
nop		RR	00	0	0	0	0	00			
add	rd,rs,rt	RR	00	rs	rt	rd	0	32			
sub	rd,rs,rt	RR	00	rs	rt	rd	0	34			
mul	rd,rs,rt	RR	00	rs	rt	rd	0	27			
div	rd,rs,rt	RR	00	rs	rt	rd	0	24			
and	rd,rs,rt	RR	00	rs	rt	rd	0	13			
or	rd,rs,rt	RR	00	rs	rt	rd	0	15			
srl	rd,rt,shf	RR	00	0	rt	rd	shf	03			
sll	rd,rt,shf	RR	00	0	rt	rd	shf	02			
cmpl	rd,rt	RR	00	0	rt	rd	0	11			
j	jadr	J	02	-----jadr-----							
lwim	rd,rs,val	M	15	rs	rd	---val---					
addi	rd,rs,val	M	12	rs	rd	---val---					
beq	rs,rt,adr	M	29	rs	rt	---adr---					
lw	rd,adr(rx)	M	35	rx	rd	---adr---					
sw	rt,adr(rx)	M	43	rx	rt	---adr---					

```

0 0 0 0 0 0 a a a a a b b b b b r r r r r -ignored- 0 1 1 0 1 1 mul r,a,b
0 0 0 0 0 0 a a a a a b b b b b r r r r r -ignored- 0 1 1 0 0 0 div r,a,b
0 0 0 0 0 0 a a a a a b b b b b r r r r r -ignored- 0 0 1 1 0 1 and r,a,b
0 0 0 0 0 0 a a a a a b b b b b r r r r r -ignored- 0 0 1 1 1 1 or r,a,b
0 0 0 0 0 0 0 0 0 0 0 b b b b b r r r r r s s s s s 0 0 0 0 1 1 srl r,b,s
0 0 0 0 0 0 0 0 0 0 0 b b b b b r r r r r s s s s s 0 0 0 0 1 0 sll r,b,s
0 0 0 0 0 0 0 0 0 0 0 b b b b b r r r r r -ignored- 0 0 1 0 1 1 cmpl r,b
0 0 0 1 0 -----address to bits (27:2) of PC----- j adr
0 0 1 1 1 1 x x x x x r r r r r ---2's complement value----- lwim r,val(x)
0 0 1 1 0 0 x x x x x r r r r r ---2's complement value----- addi r,val(x)
0 1 1 1 0 1 a a a a a b b b b b ---2's complement address----- beq a,b,adr
1 0 0 0 1 1 x x x x x r r r r r ---2's complement address----- lw r,adr(x)
1 0 1 0 1 1 x x x x x b b b b b ---2's complement address----- sw b,adr(x)

```

**Definitions:**

**nop** no operation, no programmer visible registers or memory  
 are changed, except PC gets PC+4

**j adr** bits 0 through 25 of the instruction are inserted into PC(27:2)  
 probably should zero bits PC(1:0) but should be zero already

**lw r,adr(x)** load word into register r from memory location (register x plus  
 sign extended adr field)

**sw b,adr(x)** store word from register b into memory location (register x plus  
 sign extended adr field)

**beq a,b,adr** branch on equal, if the contents of register a are equal  
 to the contents of register b, add the, shifted by two,  
 sign extended adr to the PC (The PC will have 4 added by then)

**lwim r,val(x)** load immediate, the contents of register x is added to the  
 sign extended value and the result put into register r

**addi r,val(x)** add immediate, the contents of register x is added to the  
 sign extended value and the result added to register r

**add r,a,b** add register a to register b and put result into register r

**sub r,a,b** subtract register b from register a and put result into register r

**mul r,a,b** multiply register a by register b and put result into register r

**div r,a,b** divide register a by register b and put result into register r

**and r,a,b** and register a to register b and put result into register r

```

or r,a,b    or register a to register b and put result into register r
srl r,b,s    shift the contents of register b by s places right and put
              result in register r
sll r,b,s    shift the contents of register b by s places left and put
              result in register r
cmpl r,b    one's complement of register b goes into register r

```

Also: no instructions are to have side effects or additional "features"

General register list (applies to MIPS ISA and project)  
 (note: project op codes may differ from MIPS/SGI)

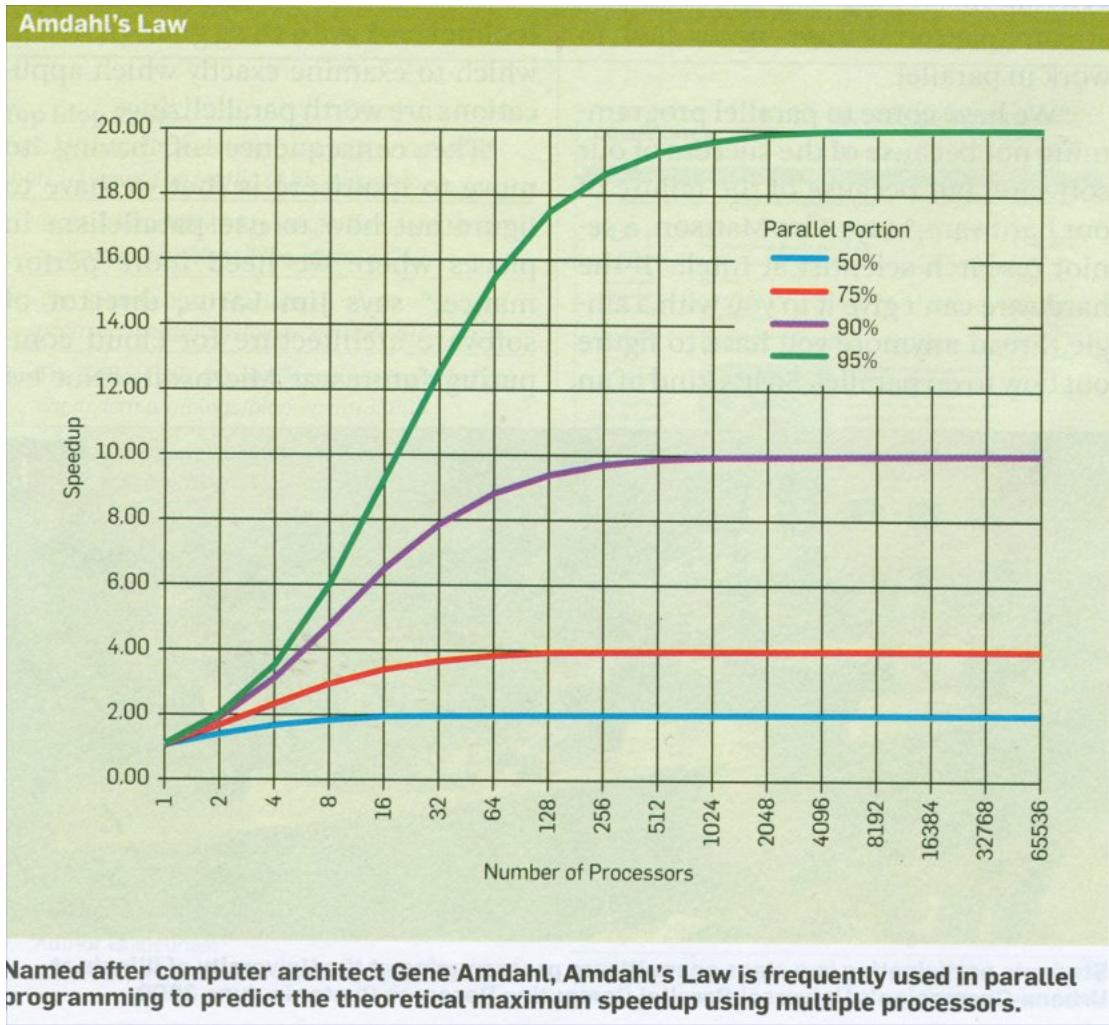
Register	notes
0 \$0	zero value, not writable
1 \$1	
2 \$2 \$v0	return values (convention, not constrained by hardware)
3 \$3 \$v1	
4 \$4 \$a0	arguments (convention, not constrained by hardware)
5 \$5 \$a1	
6 \$6 \$a2	
7 \$7 \$a3	
8 \$8 \$t0	temporaries(not saved by software convention over calls)
9 \$9 \$t1	
10 \$10 \$t2	
11 \$11 \$t3	
12 \$12 \$t4	
13 \$13 \$t5	
14 \$14 \$t6	
15 \$15 \$t7	
16 \$16 \$s0	saved by software convention over calls
17 \$17 \$s1	
18 \$18 \$s2	
19 \$19 \$s3	
20 \$20 \$s4	
21 \$21 \$s5	
22 \$22 \$s6	
23 \$23 \$s7	
24 \$24 \$t8	more temporaries
25 \$25 \$t9	
26 \$26	
27 \$27	
28 \$28 \$gp	global pointer ( more designations by software convention)
29 \$29 \$sp	stack pointer
30 \$30 \$fp	frame pointer
31 \$31 \$ra	return address

Remember: From a hardware view registers 1 through 31 are general purpose  
 and identical. The above table is just software conventions.  
 Register zero is always zero!

[Basic digital logic](#)

[IA-64 Itanium](#)

We will cover multicore and parallel processors later.  
 Amdahls law applies to them also.



[HW2 assignment](#)

## Lecture 4, CPU Operation

We now look at instructions in memory, how they got there and how they execute:

1. Start by using an editor to enter compiler language statements.  
The editor writes your source code to a disk file.
2. A compiler reads the source code disk file and produces

- assembly language instructions for a specific ISA that will perform your compiler language statements. The assembly language is written to a disk file.
3. An assembler reads the assembly language disk file and produces a relocatable binary version of your program and writes it to a disk file. This may be a main program or just a function or subroutine. Typical file name extension is .o or .obj
  4. A linkage editor or binder or loader combines the relocatable binary files into an executable file. Addresses are relocated and typically all instructions are put sequentially in a code segment, all constant data in another segment, variables and arrays in another segment and possibly making other segments. The addresses in all executable files for a specific computer start at the same address. These are virtual addresses and the operating system will place the segments into RAM at other real memory addresses. Windows file extension .exe
  5. A program is executed by having the operating system load the executable file into RAM and set the program counter to the address of the first instruction that is to be executed in the program. All programs might have the same starting address, yet the operating system has set up the TLB to translate the virtual instruction and data addresses to physical memory addresses. The physical addresses are not available to the program or to a debugger. This is part of the security an operating system provides to prevent one persons program from affecting another persons program.

A simple example:

```
Compiler input      int a, b=4, c=7;
                   a = b + c;

Assembly language fragment (not unique)
    lw    $2,12($fp)    b at 12 offset from frame pointer
    lw    $3,16($fp)    c at 16 offset from frame pointer
    add   $2,$2,$3      R format instruction
    sw    $2,8($fp)     a at 8 offset from frame pointer
```

Memory addresses in bytes, integer typically 4 bytes, 32 bits.

Loaded in machine	
virtual address	content 32-bits 8-hexadecimal digits
00000000	8FC2000C lw \$2,12(\$fp)
00000004	8FC30010 lw \$3,16(\$fp)
00000008	00000000 nop inserted for pipeline
0000000C	00431020 add \$2,\$2,\$3
00000010	AFC20008 sw \$2,8,\$(fp)

\$fp has 10000000 (data frame)	
10000000	00000000
10000004	00000001
10000008	00000000? a after execution
1000000C	00000004 b
10000010	00000007 c

Instruction field format for add \$2,\$2,\$3  
 0000 0000 0100 0011 0001 0000 0010 0000 binary for 00431020 hex

vvvv vvss ssst tttt dddd dhhh hhvv vvvv 6,5,5,5,5,6 bit fields  
 0 | 2 | 3 | 2 | 0 | 32 decimal values of fields

Instruction field format for lw \$2,12(\$fp) \$fp is register 30  
 1000 1111 1100 0010 0000 0000 0000 1100 binary for 8FC2000C hex

vvvv vvxz xxxx dddd aaaa aaaa aaaa aaaa 6,5,5,16 bit fields  
 35 | 30 | 2 | 12 decimal values of fields

The person writing the assembler chose the format of an assembly language line. The person designing the ISA chose the format of the instruction. Why would you expect them to be in the same order?

ADD \$2,\$3,\$4 00641020<sub>16</sub>

0	3	4	2	0	32
31	2625	2120	1615	1110	6 5 0

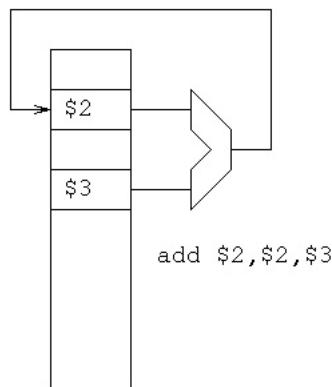
LW \$5,127(\$6) 8CC5007F<sub>16</sub>

35	6	5	127	0
31	2625	2120	1615	0

J 12345676 082F1853<sub>16</sub>

2	12345676	0
31	2625	0

A very simplified data flow of the add instruction. From the registers to the ALU and back to the registers.



The VHDL to use the ALU will be given to you as:

```

ALU: entity WORK.alu_32 port map(inA      => EX_A,
                                  inB      => EX_aluB,
                                  inst    => EX_IR,
                                  result  => EX_result);

```

We will call the upper input "A" and the lower input "B"

and the output "result".

The extra input, EX\_IR, not shown on the diagram above  
is the instruction the ALU is to perform, add, sub, etc.

The instructions we will use in this course are specifically:

[cs411\\_OPCODES.txt](#)

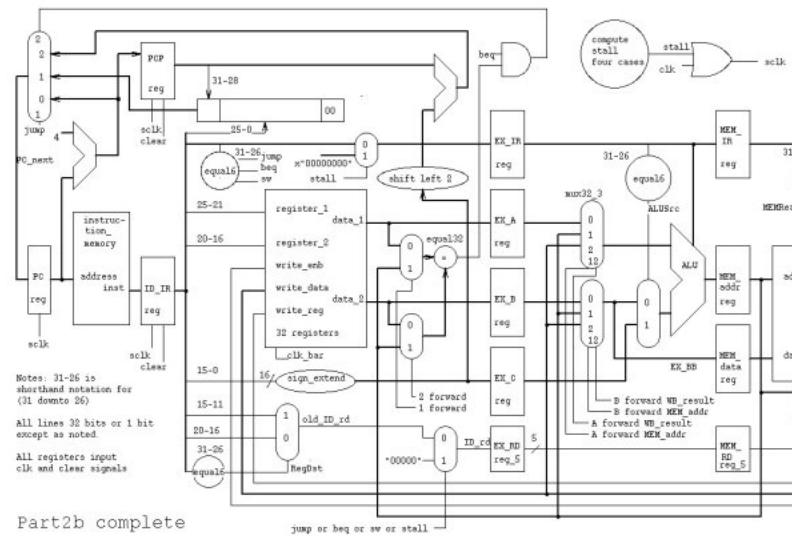
Each student needs to understand what the instructions are  
and the use of each field in each instruction.

(Note: a few have bit patterns different from the book and  
different from previous semesters in order to prevent copying.)

Our MIPS architecture computer uses five clocks to execute  
a load word instruction.

1 0 0 0 1 1 x x x x x r r r r r ---2's complement address----- lw r,adr(x)

1. Fetch the instruction from memory
2. Decode the instruction and read the value of register xxxx
3. Compute the memory address by adding the sign extended bottom  
16 bits of the instruction to the contents of register xxxx.
4. Fetch the data word from the memory address.
5. Write the data word from memory into the register rrrrr.



When we cover "pipelining" you will see why five clocks are  
used for every instruction, even though some instructions  
need less than five.

Computer languages come in many varieties.  
The information above applies to languages such  
as C, C++, Fortran, Ada and others.

Many languages abstract the concept of binary relocatable  
code, in what was originally called "crunch code".  
These languages use their own form of intermediate files.  
For example Pascal, Java, Python and others.

Other languages directly interpret the users source  
files, possibly with some preprocessing.  
For example SML, Haskell, Lisp, MatLab, Mathematica and

others.

With a completely new computer architecture, the first "language" would be an assembly language. From this, a primitive operating system would be built. Then, typically an existing C compiler would be modified for the new computer architecture. An alternative is to build a cross compiler from C and gas, to bootstrap existing code to the new architecture. From then on, "reuse" goes into full effect and millions of lines of existing software can be running on the new computer architecture.

#### For Homework 3

The computer `irix.gl.umbc.edu` is no longer available. This was a MIPS architecture using the same instructions as we are using. The MIPS architecture is studied because it is a much simpler and easier to understand architecture than the Intel X86, IA-32.

Thus, to see the instructions in RAM, we will use the `gdb` debugger on an Intel X86.

#### [HW3 information](#)

The information in `hex.out` will have lines similar to:

```
(gdb) disassemble
Dump of assembler code for function main:
RAM addr      offset    op code   address and register
0x08048384 <main+0>:    lea     0x4(%esp),%ecx
0x08048388 <main+4>:    and    $0xffffffff,%esp
0x0804838b <main+7>:    pushl   0xfffffff(%ecx)
```

End of assembler dump.

```
(gdb) x/60x main
Note: 16 bytes per line, 4 32-bit words
      but, these are X86 instructions, not MIPS !
0x08048384 <main>: 0x04244c8d 0xffff0e483 0x8955fc71 0x535657e5
0x08048394 <main+16>: 0x58ec8351 0x4589e089 0xe445c7cc 0x00000064
      ##          ##
<main+19>----|           <main+31>---|
0x8048397          0x80483A3
```

Because the MIPS architecture we are studying is a big endian machine, we will count bytes from left to right for homework 3.

In hexadecimal, 0x12345678 is stored big end first	12
	34
	56
	78

Little endian 0x12345678 is stored little end first	78
	56
	34

Each byte, 8 bits, is two hex digits 12

## Lecture 5, Instructions and Registers

Get paper handout, fill in values for registers and memory as we discuss the instructions in this lecture.  
The program starts with PC set to address zero.  
The instructions are defined on [cs411\\_opcodes.txt](#)

[part1.asm](#)  
[part1.abs](#)

`part1.abs`  
address instruction assembly language

```

00000000 8C010074    lw   $1,w1($0)
00000004 8C020078    lw   $2,w2($0)
00000008 8C03007C    lw   $3,w3($0)
0000000C 00000000    nop
00000010 00000000    nop
00000014 00232020    add  $4,$1,$3
00000018 00222822    sub  $5,$1,$2
0000001C 000133C2    sll  $6,$1,15
00000020 00023C03    srl  $7,$2,16
00000024 0003400B    cmpl $8,$3
00000028 0022480D    or   $9,$1,$2
0000002C 0023500F    and  $10,$1,$3
00000030 00435818    div   $11,$2,$3
00000034 0062601B    mul   $12,$3,$2
00000038 AC010080    sw   $1,w4($0)
0000003C 300D0074    addi $13,w1
00000040 00000000    nop
00000044 00000000    nop
00000048 8DAE0004    lw   $14,4($13)
0000004C 31AF0008    addi $15,8($13)
00000050 3C100010    lwim $16,16
00000054 00000000    nop
00000058 00000000    nop
0000005C ADE30008    sw   $3,8($15)
00000060 00000000    nop
00000064 00000000    nop
00000068 00000000    nop
0000006C 00000000    nop
00000070 00000000    nop
00000074 11111111    w1: word 0x11111111
00000078 22222222    w2: word 0x22222222
0000007C 33333333    w3: word 0x33333333
00000080 44444444    w4: word 0x44444444

```

After the CPU has executed the first instruction:

	General Registers	RAM memory	initial	final
\$0	00000000			
\$1	11111111	00000074	11111111	-----
\$2	-----	00000078	22222222	-----
\$3	-----	0000007C	33333333	-----
\$4	-----	00000080	44444444	-----
\$5	-----	00000084	xxxxxxxx	-----
\$6	-----			
\$7	-----			
\$8	-----			
\$9	-----			
\$10	-----			
\$11	-----			
\$12	-----			

This is part of your project: [part1.abs](#)  
and the result of running that small program part1.chk:

[part1.chk](#)

Note the large amount of information printed each clock time.

Note that it takes 5 clock cycles to finish an instruction.

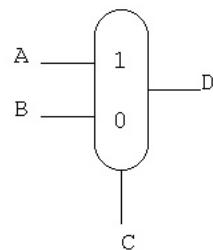
Basic MUX Truth Table and Schematic

### The Multiplexor      MUX

truth table

A	B	C	D
0	0	0	0
0	0	1	0
0	1	0	1
0	1	1	0
1	0	0	0
1	0	1	1
1	1	0	1
1	1	1	1

circuit symbol



concurrent VHDL statement

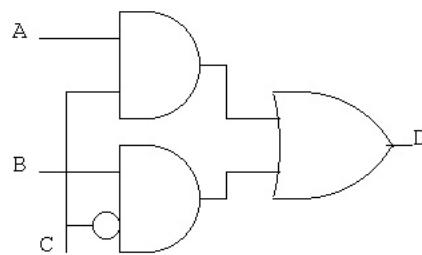
```
D <= A when C='1' else B;
```

```
D <= A after 2 ns when C='1'  
else B after 2 ns;
```

sequential VHDL statement

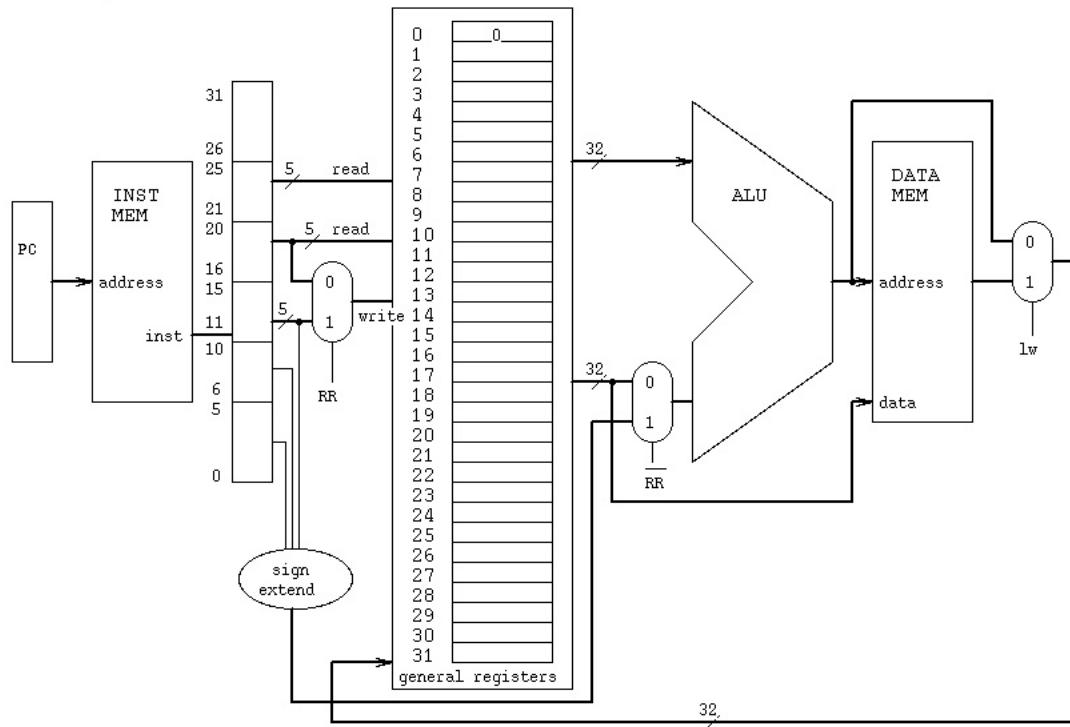
```
if C='1' then D <= A  
else D <= B end if;
```

a logic diagram



How MUX are used to route data

Multiplexors control flow of data



You can see much of the code for the above in the  
starter code for Proj1:

[part1\\_start.vhd](#)

There are basic design principles for computer architecture  
and many apply to broader applications.

Design Principle 1:

Simplicity is best achieved through regularity.

A few building blocks, used systematically, will have  
fewer errors, be available sooner and sell for less.  
A uniform instruction set allows better compilers.

Design Principle 2:

Smaller is faster:

Smaller feature size means signals can move faster.  
Shorter paths, less stages, allow completion sooner.

Design Principle 3:

Good design requires good compromises.

There are no perfect architectures. There are kluges.

Design Principle 4:

Make the common part fast.  
Amdahl's law, be sure you are maximizing speedup.

[Pentium 4 Hyper threading](#)

[Intel Core Duo](#)

[AMD quad core, one core shown](#)

[\\$329 for just 1/2 quad core processor](#)

[a 4 CPU, 8GB RAM configuration](#)

Now 12-core 16GB RAM, 3 hard drives, 2 DVD writers

Practice safe computing!

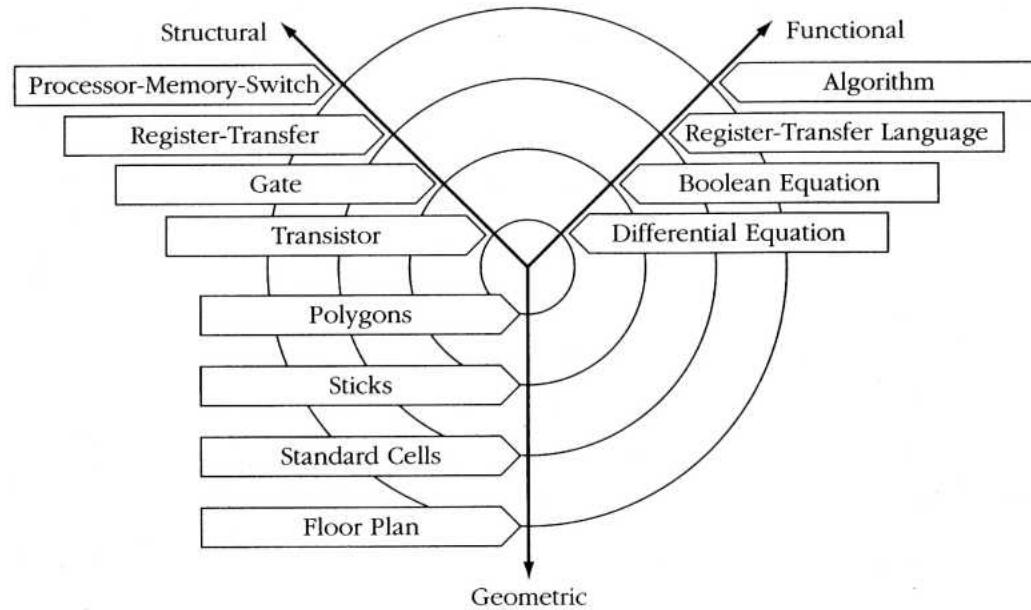
Beware Malware, Spyware and Adware:

Do everything you can to keep malware from infecting your systems, malware authors do all they can to keep their work from being detected and removed. By looking at the methods that malware uses to keep itself safe, you can better root it out and remove it before the damage is done. Downloading attachments is the primary way malware gets into your system.

[HW3 assigned](#)

## Lecture 6, VHDL introduction

**VHDL is used for structural and functional modeling of digital circuits.**



*Domains and levels of abstraction. The radial axes show the three different domains of modeling. The concentric rings show the levels of abstraction, with the more abstract levels on the outside and more detailed levels toward the center.*

The geometric modeling is handled by other Cadence programs.

First, simple VHDL statements for logic gates:

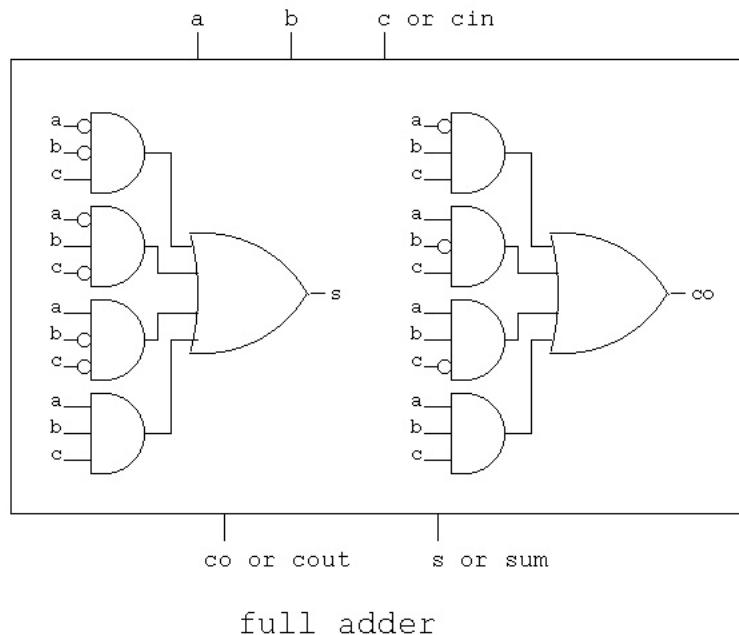
[logic gates and corresponding VHDL statements](#)

```
VHDL comments start with -- acting like C++ and Java //
VHDL like C++ and Java end statements with a semicolon ;
VHDL uses "library" and "use" where C++ uses #include Java uses import
VHDL uses ".all" where Java uses ".*"
VHDL uses names similar to Pascal, case insensitive, var is same as Var, VAR
```

VHDL has a two part basic structure for each circuit  
 that is more than one gate, the "entity" and the "architecture".  
 There needs to be a "library" and "use" for features that are used.

The word "port" is used to mean interface.  
 The term "std\_logic" is a type used for one bit.  
 The term "std\_logic\_vector" is a type used for more than one bit.  
 The time from an input changing to when the output may change  
 is optional. "after 1 ps" indicates 1 pico second.  
 "after 2 ns" indicates 2 nano seconds.

This circuit is coded as a full adder component in VHDL:



```

library IEEE;
use IEEE.std_logic_1164.all;

entity fadd is          -- full adder stage, interface
  port(a    : in  std_logic;
       b    : in  std_logic;
       cin  : in  std_logic;
       s    : out std_logic;
       cout : out std_logic);
end entity fadd;

architecture circuits of fadd is -- full adder stage, body
begin -- circuits of fadd
  s <= a xor b xor cin after 1 ps;
  cout <= (a and b) or (a and cin) or (b and cin) after 1 ps;
end architecture circuits; -- of fadd

```

Notice that `entity fadd is ... end entity fadd;` is a statement

Notice that `architecture circuits of fadd is ... end architecture circuits;`  
is a statement. The "of fadd" connects the architecture to the entity.

The arbitrary signal names `a`, `b`, `cin`, `s`, `cout` were required to be assigned a type, `std_logic` in this case, before being used. Typical for many programming languages.

Now, use a loop to combine 32 `fadd` into a 32 bit adder:

Note: to use `fadd`, a long statement must be used

```
a0: entity WORK.fadd port map(a(0), b(0), cin, sum(0), cout);
```

A unique label `a0` followed by a colon :

Then entity `WORK.fadd` naming the entity to be used in `WORK` library.

Then port map( with actual signals for `a`, `b`, `cin`, `s`, `cout` )

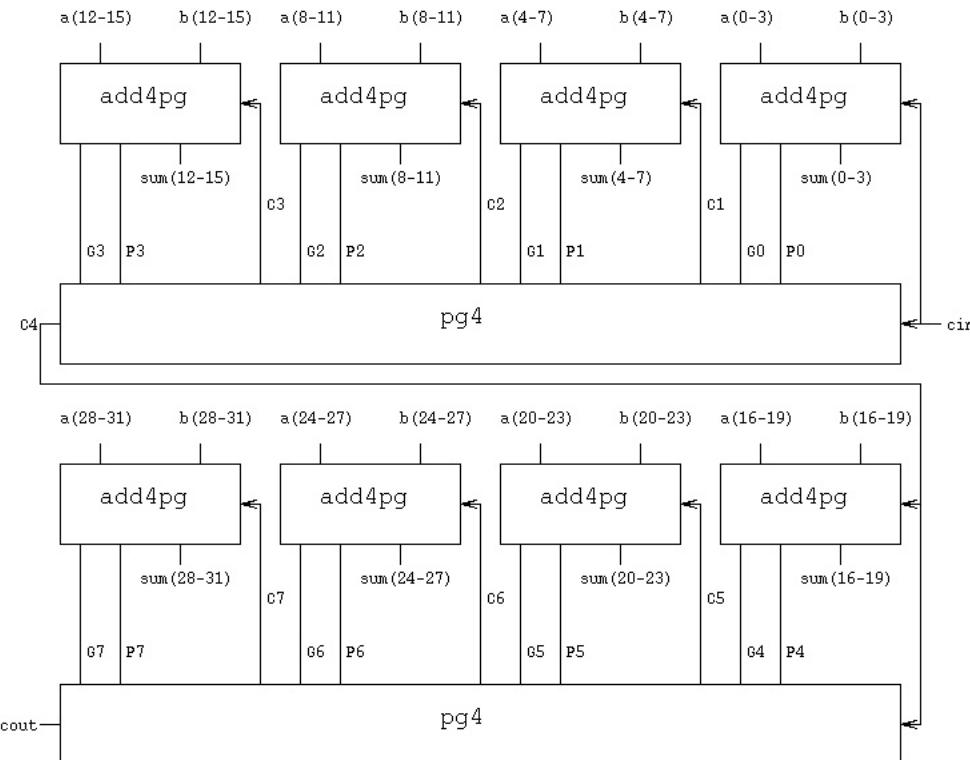
Note subscripts for bit numbers in parenthesis, not [ ] .

The first and last stage are slightly different from the 30 stages in the loop.

[add32.vhd1 using the fadd above](#)

Another variation of an adder, propagate generate.

add32 propagate - generate



#### [add32pg\\_start.vhdl for HW4](#)

A "main" entity to use the component add32 with test data.

Note: just structure of "entity" then big architecture

```
entity tadd32 is
end tadd32;
```

-- test bench for add32.vhd  
-- no requirement to use "main"

architecture circuits of tadd32 is ...

#### [tadd32.vhdl for main entity for HW4](#)

The output from the write statements in tadd32.vhd is:

#### [tadd32.chk output of tadd32.vhd](#)

The additional file tadd32.run was needed to tell the VHDL simulator how long to run:

#### [tadd32.run used to stop simulation](#)

After using `vhdl_cshrc` and `cds.lib`,  
the command line commands for using Cadence VHDL are:

```
ncvhdl -v93 add32.vhdl  
ncvhdl -v93 tadd32.vhdl  
ncelab -v93 tadd32:circuits  
ncsim -batch -logfile tadd32.out -input tadd32.run tadd32
```

#### [output of simulation](#)

The command line commands for using GHDL are:

```
ghdl -a --ieee=synopsys add32.vhdl  
ghdl -a --ieee=synopsys tadd32.vhdl  
ghdl -e --ieee=synopsys tadd32  
ghdl -r --ieee=synopsys tadd32 --stop-time=65ns > tadd32.gout
```

#### [output of simulation](#)

Use a Makefile for sets of commands. You will be running more than once to get homework and projects correct.

I provide a starter Makefile for VHDL and GHDL.

Browse and use as a reference for HW4, HW6, and Project.  
You must do the setup exactly as stated in [HW4](#)

#### [Sample designs and corresponding VHDL code](#)

#### [VHDL Language Compact Summary](#)

The setup for HW4, HW6 and Project will be covered in the next lecture.  
You will be using command lines in a terminal window on `linux.gl.umbc.edu`  
You are given a `cs411.tar` file that creates the needed directories for Cadence.  
`Makefile_ghdl` sets up Makefile for GHDL.  
You will be modifying a Makefile for HW4, HW6, and Project parts.  
The basic VHDL commands are shown in the starter Makefile's  
[Makefile\\_vhdl\\_for\\_Cadence](#)  
[Makefile\\_ghdl\\_for\\_GHDL](#)

The number system of interest in computer architecture re:  
 Sign Magnitude - binary magnitude with sign bit  
 Ones Complement - negative numbers have all bits inverted  
 Twos Complement - Ones Complement with one added to lsb

All number systems have the sign bit 0 for positive and 1 for negative. The msb is the sign bit and thus the word length is important.

Number systems, using 4-bit words

Hex	Binary	Sign	Ones	Twos
Digit	Bits	Magnitude	Complement	Complement
0	0000	0	0	0
1	0001	1	1	1
2	0010	2	2	2
3	0011	3	3	3
4	0100	4	4	4
5	0101	5	5	5
6	0110	6	6	6
7	0111	7	7	7
8	1000	-0	-7	-8 difference starts here
9	1001	-1	-6	-7
A	1010	-2	-5	-6
B	1011	-3	-4	-5
C	1100	-4	-3	-4
D	1101	-5	-2	-3
E	1110	-6	-1	-2
F	1111	-7	-0	-1

to negate: invert sign invert all bits and add one

math  $-(-N)=N$  OK OK  $-(-8)=-8$  YUK!

Addition	Sign	Ones	Twos
	Magnitude	Complement	Complement

$$\begin{array}{r}
 2 \\
 +3 \\
 \hline
 +5
 \end{array}
 \quad
 \begin{array}{rrr}
 0010 & 0010 & 0010 \\
 0011 & 0011 & 0011 \\
 \hline
 0101 & 0101 & 0101 \\
 \text{OK}
 \end{array}$$

$$\begin{array}{r}
 4 \\
 +5 \\
 \hline
 9
 \end{array}
 \quad
 \begin{array}{rrr}
 0100 & 0100 & 0100 \\
 0101 & 0101 & 0101 \\
 \hline
 1001 & 1001 & 1001 \\
 -1 & -6 & -7
 \end{array}$$

overflow gives wrong answer on fixed length, computer, numbers

Subtraction: negate second operand and add

$$\begin{array}{r}
 4 \\
 -5 \\
 \hline
 -1
 \end{array}
 \quad
 \begin{array}{rrr}
 0100 & 0100 & 0100 \\
 1101 & 1010 & 1011 \\
 \hline
 1001 & 1110 & 1111 \\
 -1 & -1 & -1
 \end{array}$$

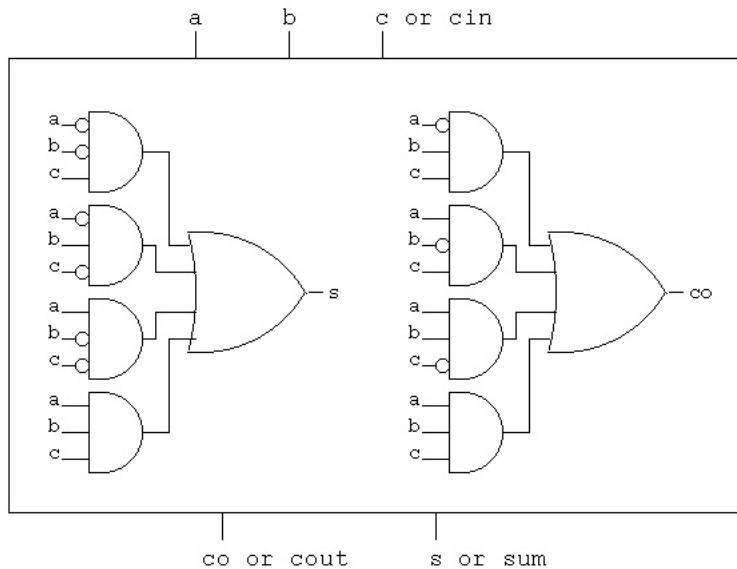
works, using correct definition of negate

Sign Magnitude bigger minus smaller, fix sign  
 Twos Complement, just add. Most computers today  
 Ones Complement, just add. e.g. Univac computers

It was discovered the "add one" was almost zero cost, thus most integer arithmetic is twos complement.

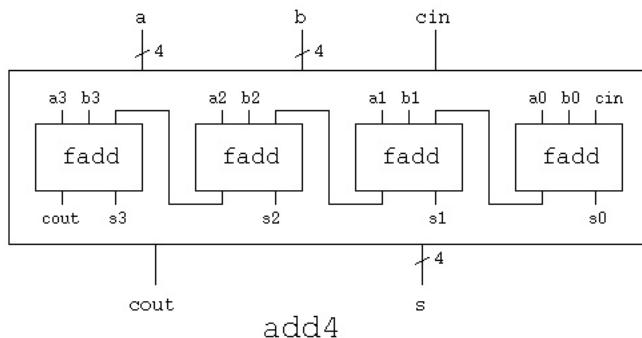
The hardware adder has a carry-in input that implements the "add one" by making this input a "1".

Basic one bit adder, called a full adder.

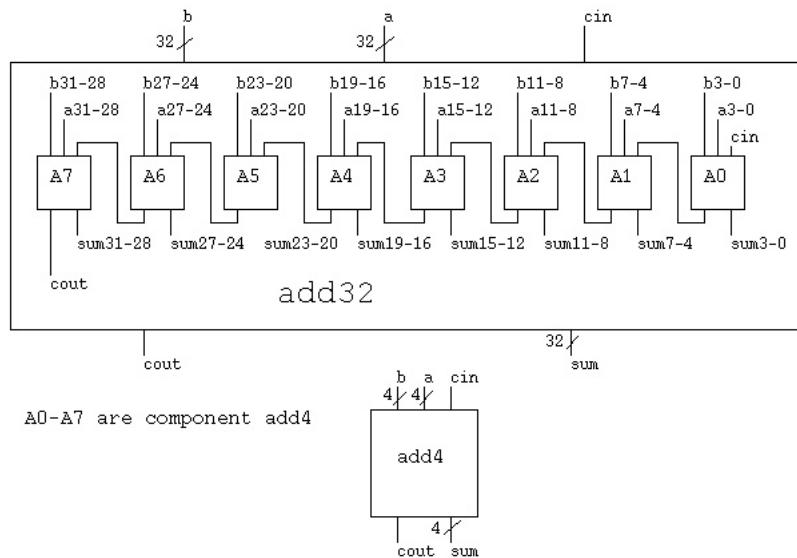


full adder

Combining four full adders to make a 4-bit adder.

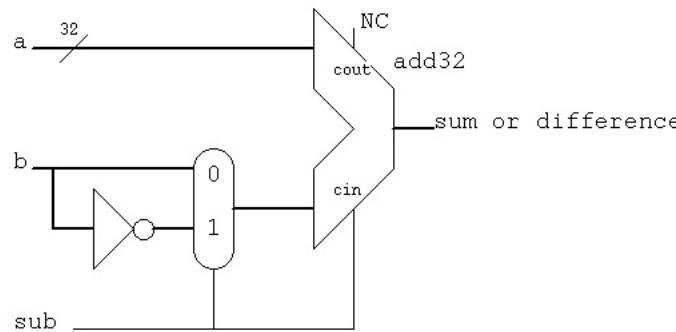


Combining eight 4-bit adders to make a 32-bit adder.



A quick look at VHDL that implements the above diagrams, with some optimization, is [an add32](#)

Using a multiplexor with 32-bit adder for subtraction.  
 "sub" is '1' for subtract, '0' for add.  
 (NC is no connection, use open in VHDL)

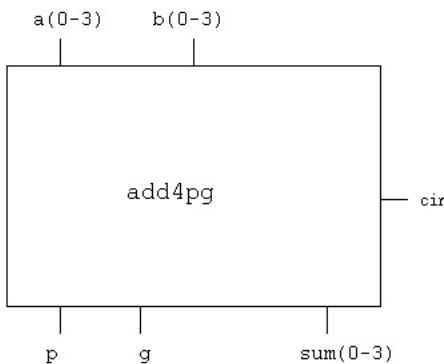


**add32** used for add or subtract

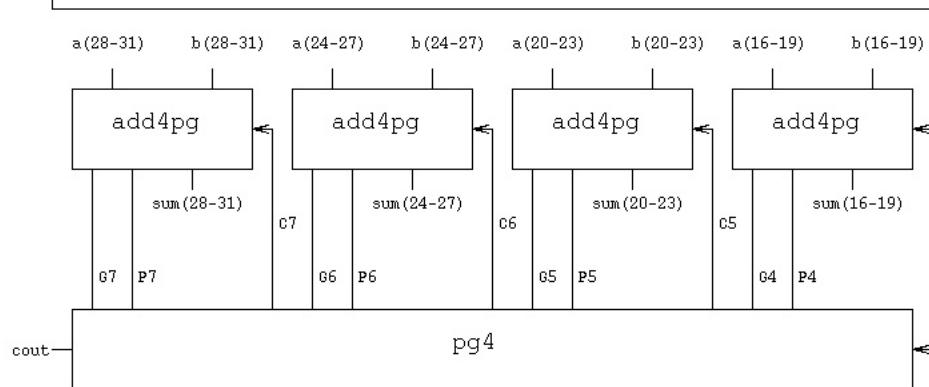
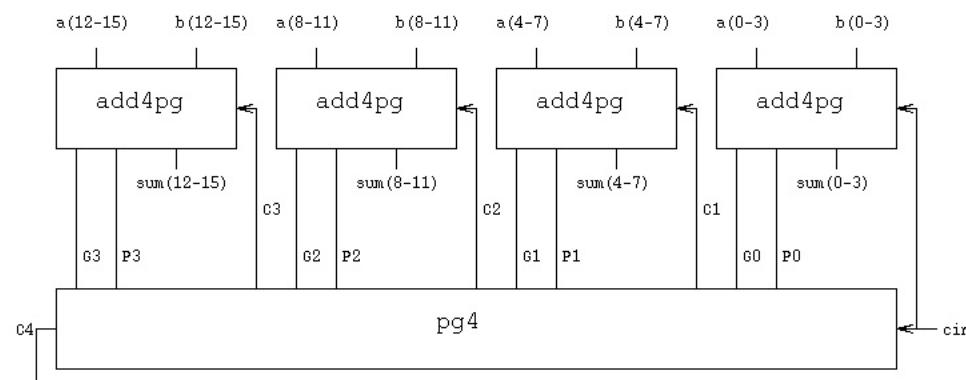
There are many types of adders. "Bit slice" will be covered in the next lecture on the ALU. First, related to Homework 4 is the "propagate generate" adder, then the "Square root N" adder for Computer Engineering majors.

The "Propagate Generate" PG adder has a propagation time

proportional to  $\log_2 N$  for  $N$  bits.



add32 propagate - generate



The "add4pg" unit has four full adders and extra circuits, defined by equations rather than logic gates:

```
-- add4pg.vhd      entity and architecture
--          for 4 bits of a propagate-generate, pg, adder
library IEEE;
```

```

use IEEE.std_logic_1164.all;
entity add4pg is
  port(a      : in  std_logic_vector(3 downto 0);
       b      : in  std_logic_vector(3 downto 0);
       cin   : in  std_logic;
       sum   : out std_logic_vector(3 downto 0);
       p     : out std_logic;
       g     : out std_logic );
end entity add4pg ;

architecture circuits of add4pg is
  signal c : std_logic_vector(2 downto 0);
begin -- circuits of add4pg
  sum(0) <= a(0) xor b(0) xor cin after 2 ps;
  c(0)   <= (a(0) and b(0)) or (a(0) and cin) or (b(0) and cin) after 2 ps;
  sum(1) <= a(1) xor b(1) xor c(0) after 2 ps;
  c(1)   <= (a(1) and b(1)) or
            (a(1) and a(0) and b(0)) or
            (a(1) and a(0) and cin) or
            (a(1) and b(0) and cin) or
            (b(1) and a(0) and b(0)) or
            (b(1) and a(0) and cin) or
            (b(1) and b(0) and cin) after 2 ps;
  sum(2) <= a(2) xor b(2) xor c(1) after 2 ps;
  c(2)   <= (a(2) and b(2)) or (a(2) and c(1)) or (b(2) and c(1)) after 2 ps;
  sum(3) <= a(3) xor b(3) xor c(2) after 2 ps;
  p     <= (a(0) or b(0)) and (a(1) or b(1)) and
            (a(2) or b(2)) and (a(3) or b(3)) after 2 ps;
  g     <= (a(3) and b(3)) or ((a(3) or b(3)) and
            ((a(2) and b(2)) or ((a(2) or b(2)) and
            ((a(1) and b(1)) or ((a(1) or b(1)) and
            ((a(0) and b(0))))))) after 2 ps;
end architecture circuits; -- of add4pg

```

The "PG4" box is defined by equations and thus no schematic:

```

-- pg4.vhdl    entity and architecture Carry-Lookahead unit
-- pg4 is driven by four add4pg entities
library IEEE;
use IEEE.std_logic_1164.all;
entity pg4 is
  port(p0  : in  std_logic;
       p1  : in  std_logic;
       p2  : in  std_logic;
       p3  : in  std_logic;
       g0  : in  std_logic;
       g1  : in  std_logic;
       g2  : in  std_logic;
       g3  : in  std_logic;
       cin : in  std_logic;
       c1  : out std_logic;
       c2  : out std_logic;
       c3  : out std_logic;
       c4  : out std_logic);
end entity pg4;

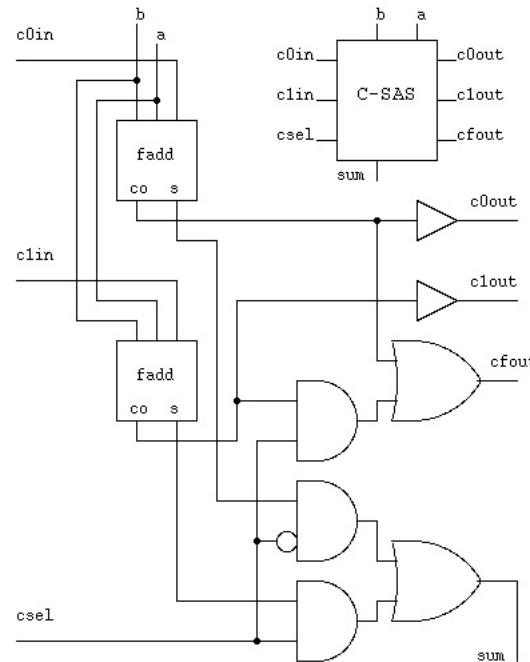
architecture circuits of pg4 is
begin -- circuits of pg4
  c1 <= g0 or (p0 and cin) after 2 ps;
  c2 <= g1 or (p1 and g0) or (p1 and p0 and cin) after 2 ps;
  c3 <= g2 or (p2 and g1) or (p2 and p1 and g0) or
            (p2 and p1 and p0 and cin) after 2 ps;
  c4 <= g3 or
            (p3 and g2) or
            (p3 and p2 and g1) or
            (p3 and p2 and p1 and g0) or
            (p3 and p2 and p1 and p0 and cin) after 2 ps;
end architecture circuits; -- of pg4

```

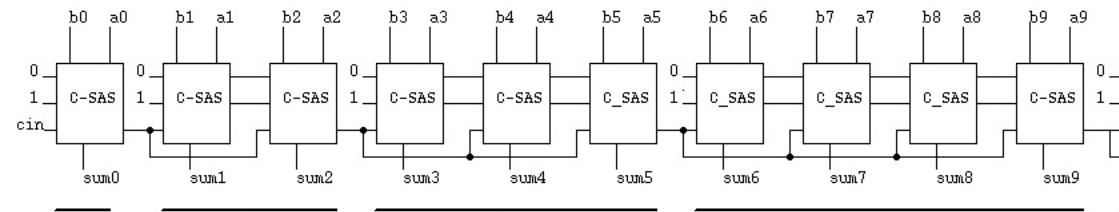
The "Carry Select" CS, adder gets increased speed from computing

[https://userpages.umbc.edu/~squire/cs411\\_lect.html](https://userpages.umbc.edu/~squire/cs411_lect.html)

the possible output with carry in to that stage being both '0' and '1'. The "Carry Select" adder has a propagation time proportional to  $\text{sqrt}(N)$  for N bits.



carry-select adder stage



ten stage carry-select adder

The above diagram has only 10 bits drawn.  
You need 32 bits. Thus you need additional group of 5,  
group of 6, group of 7, and a final group of 4.  
 $1+2+3+4+5+6+7+4=32$

If  $N = 64$ ,  $\log_2 N = 6$ ,  $\text{sqrt}(N) = 8$  speedup vs complexity (size)

Behavioral VHDL for our add32:

```
library IEEE;
use IEEE.std_logic_1164.all;
entity add32 is
  port(a : in std_logic_vector(31 downto 0);
```

```

b    : in std_logic_vector(31 downto 0);
cin : in std_logic;
sum : out std_logic_vector(31 downto 0);
cout : out std_logic;
end entity add32; -- same for all implementations

library IEEE;
use IEEE.std_logic_arith.all;
architecture behavior of add32 is
  signal temp : std_logic_vector(32 downto 0);
  signal vcin : std_logic_vector(32 downto 0) := X"00000000"&'0';
  signal va  : std_logic_vector(32 downto 0) := X"00000000"&'0';
  signal vb  : std_logic_vector(32 downto 0) := X"00000000"&'0';
  -- 33 bits (32 downto 0) needed to compute cout
begin -- circuits of add32
  vcin(0) <= cin;
  va(31 downto 0) <= a;
  vb(31 downto 0) <= b;
  temp <= unsigned(va) + unsigned(vb) + unsigned(vcin); -- 33 bit add
  cout <= temp(32) after 6 ps;
  sum  <= temp(31 downto 0) after 6 ps;
end architecture behavior; -- of add32

```

[Now go to Homework 4 and the setup commands.](#)

Expect errors. Nobody's perfect.

For many errors after typing 'make'

touch add32.vhd1

make |& more # hit space for next page, enter for next line

make >& add32.prt # results, including error go to a file

# use editor to read file, you can search

FIX THE FIRST ERROR !!!!

Yes, you can fix other errors also, but one error can cause  
a cascading effect and produce many errors.

Don't panic when there was only one error, you fixed that,  
then the next run you get 37 errors. The compiler has stages,  
it stops on a stage if there is an error. Fixing that error  
lets the compiler move to the next stage and check for other  
types of errors.

Don't give up. Don't make wild guesses. Do experiment with  
one change at a time. You may actually have to read some  
of the handouts :)

Cadence VHDL error message. (actually an extra semicolon)

```

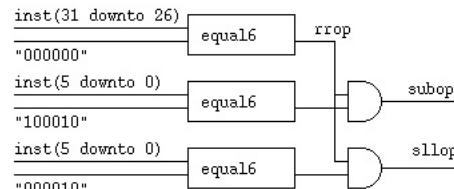
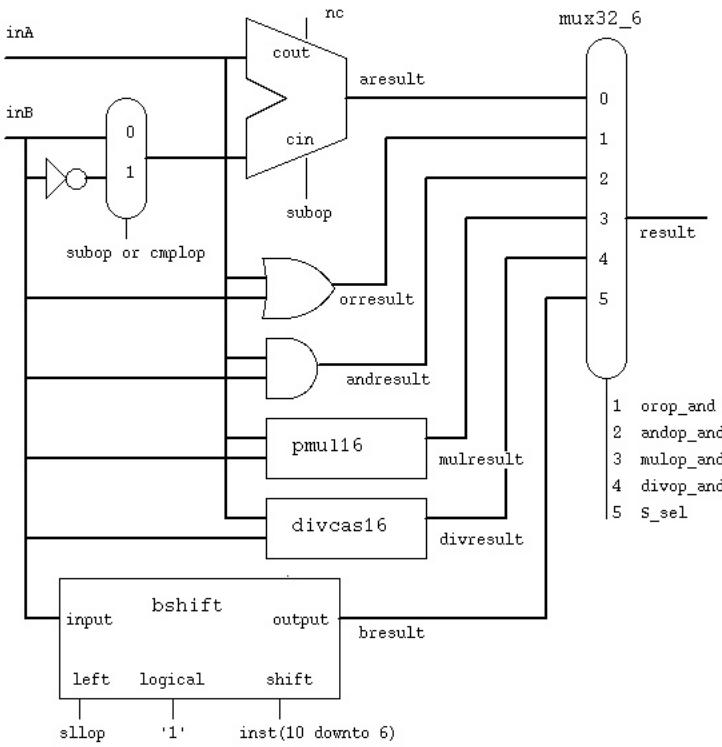
ncvhdl: 05.40-s011: (c) Copyright 1995-2005 Cadence Design Systems, Inc.
      OUTT : out std_logic;
      |
ncvhdl_p: *E,PORNKW (error.vhdl,10|28): identifier expected.
      OUTT : out std_logic;;

```

[Then to VHDL resource.](#)

## Lecture 8, ALU

The Arithmetic Logic Unit is the section of the CPU that actually performs add, subtract, multiply, divide, and, or, floating point and other operations. The choice of which operations are implemented is determined by the Instruction Set Architecture, ISA. Most modern computers separate the integer unit from the floating point unit. Many modern architectures have simple integer, complex integer, and an assortment of floating point units.



Schematic for partice\_div.vhdl

The ALU gets inputs from registers [reg\\_use.jpg](#)

Where did numbers such as 100010 for subop and 000010 for sllop come from? [cs411\\_OPCODES.TXT](#)

```
-- alu_start.vhdl
library IEEE;
use IEEE.std_logic_1164.all;

entity alu_32 is
port(inA : in std_logic_vector (31 downto 0);
```

```

inB    : in  std_logic_vector (31 downto 0);
inst   : in  std_logic_vector (31 downto 0);
result : out std_logic_vector (31 downto 0));
end entity alu_32;

architecture schematic of alu_32 is
  signal cin    : std_logic := '0';
  signal cout   : std_logic;
begin -- schematic
  --
  -- REPLACE THIS SECTION FOR PROJECT PART 1
  -- (add the signals you need above the "begin"
  -- add logic below the "begin")

  adder: entity WORK.add32 port map(a      => inA,
                                    b      => inB,      -- change
                                    cin   => cin,      -- change
                                    sum   => result,   -- change
                                    cout  => cout);

  -- examples of entity instantiations:

  -- bsh: entity WORK.bshift port map (left   => sllop,
  --                                     logical => '1',
  --                                     shift   => inst(10 downto 6),
  --                                     input   => inB,
  --                                     output  => bresult);

  -- r1: entity WORK.equal6 port map (inst  => inst(31 downto 26),
  --                                     test   => "000000",
  --                                     equal  => rrop);

  -- s1: entity WORK.equal6 port map (inst  => inst(5 downto 0),
  --                                     test   => "100010",      -- 34
  --                                     equal  => subop1);
  -- s1a: subop <= subop1 and rrop;

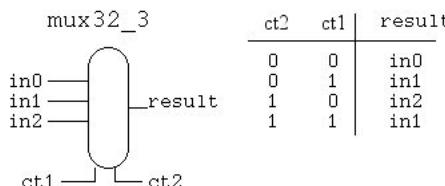
  --      S_sel <= sllop_or_srlop; -- for mux32_6

  -- much more

end architecture schematic; -- of alu_32

```

Many variations of subop, subop1, subop\_and, subopa  
Your starter partice\_start.vhdl uses subopa short for subop\_and.  
[partice\\_start.vhdl](#)



in1 has priority over in2

-- mux32\_3.vhdl

```

library IEEE;
use IEEE.std_logic_1164.all;
entity mux32_3 is
  port(in0    : in  std_logic_vector (31 downto 0);
       in1    : in  std_logic_vector (31 downto 0);
       in2    : in  std_logic_vector (31 downto 0);

```

```

ct1  : in  std_logic;      -- pass in1(has priority)
ct2  : in  std_logic;      -- pass in2
result : out std_logic_vector (31 downto 0);
end entity mux32_3;

architecture behavior of mux32_3 is
begin -- behavior -- no process needed with concurrent statements
  result <= in1 when ct1='1' else in2 when ct2='1' else in0 after 50 ps;
end architecture behavior; -- of mux32_3

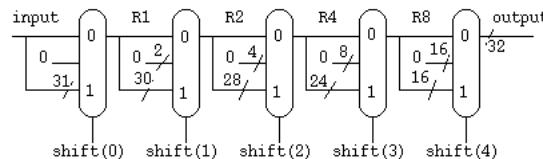
-- mux_32_6.vhd have only zero or one ctl = '1'

library IEEE;
use IEEE.std_logic_1164.all;

entity mux_32_6 is
  port(in0  : in  std_logic_vector (31 downto 0);
       in1  : in  std_logic_vector (31 downto 0);
       in2  : in  std_logic_vector (31 downto 0);
       in3  : in  std_logic_vector (31 downto 0);
       in4  : in  std_logic_vector (31 downto 0);
       in5  : in  std_logic_vector (31 downto 0);
       ctl1 : in  std_logic;
       ctl2 : in  std_logic;
       ctl3 : in  std_logic;
       ctl4 : in  std_logic;
       ctl5 : in  std_logic;
       result : out std_logic_vector (31 downto 0));
end entity mux_32_6;

architecture behavior of mux_32_6 is
begin -- behavior -- no process needed with concurrent statements
  result <= in1 when ctl1='1' else in2 when ctl2='1' else
    in3 when ctl3='1' else in4 when ctl4='1' else
    in5 when ctl5='1' else in0 after 10 ps;
end architecture behavior; -- of mux_32_6

```



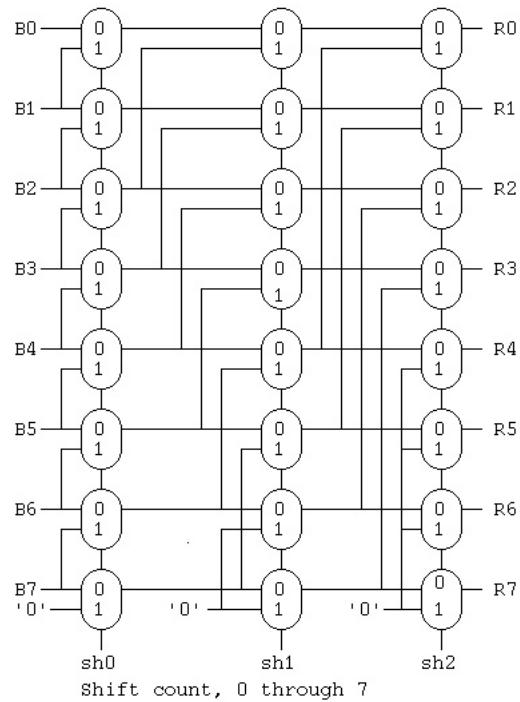
Partial schematic of shift right logical in circuits architecture

Note that bshift.vhd contains two different architectures for the same entity. A behavioral architecture using sequential programming and a circuits architecture using digital logic components.

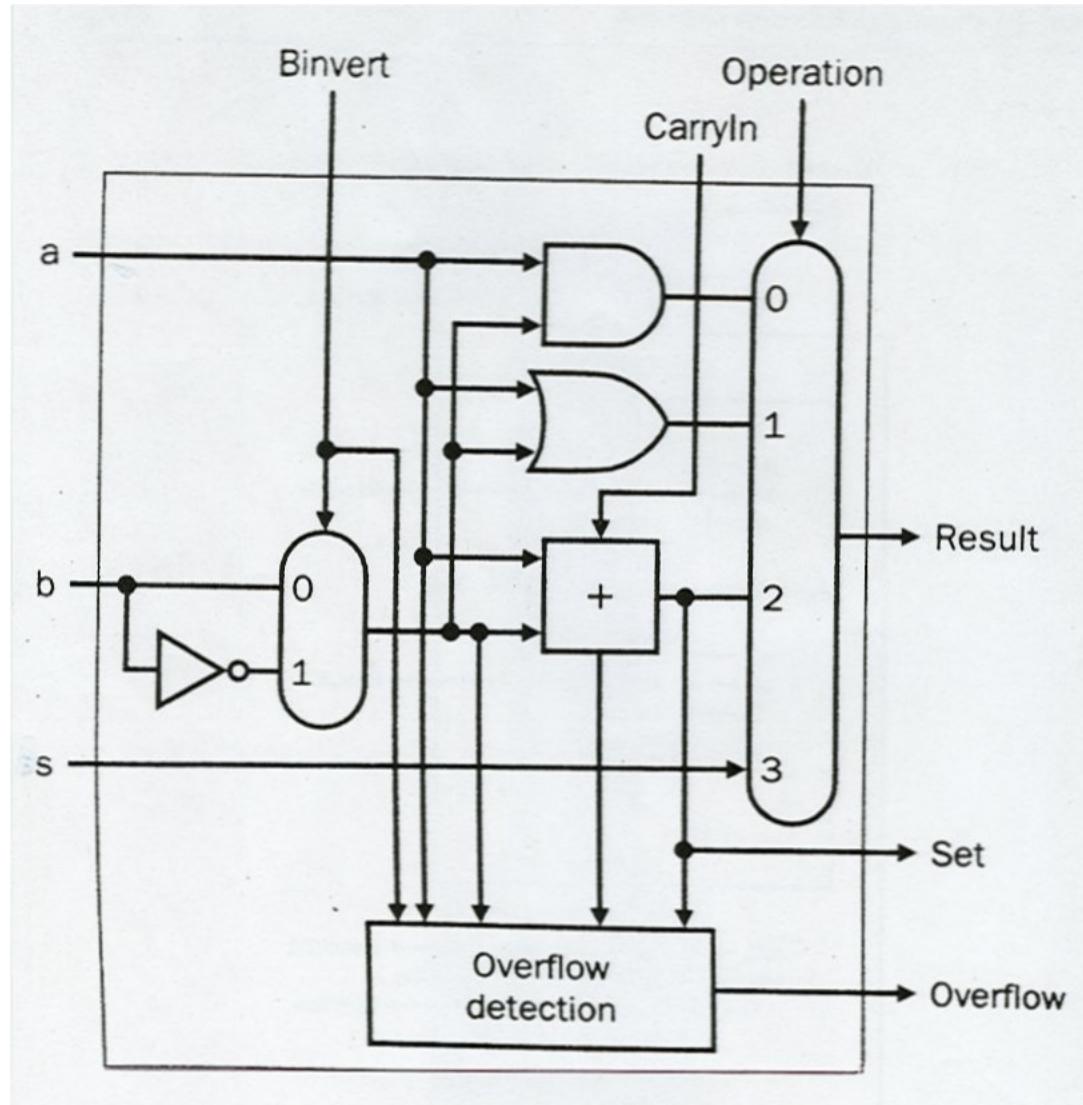
[bshift.vhd](#)

An 8-bit version of shift right logical, using single bit signals, three bit shift count, is:

8-bit barrel shift



There are many ways to build an ALU. Often the choice is based on mask making and requires a repeated pattern. The "bit slice" method uses the same structure for every bit. One example is:



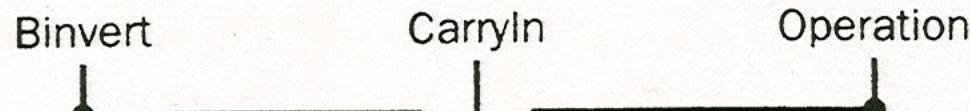
Note that 'Operation' is two bits, 0 for logical and, 1 for logical or, 2 for add or subtract, and 3 for an operation called set used for comparison.

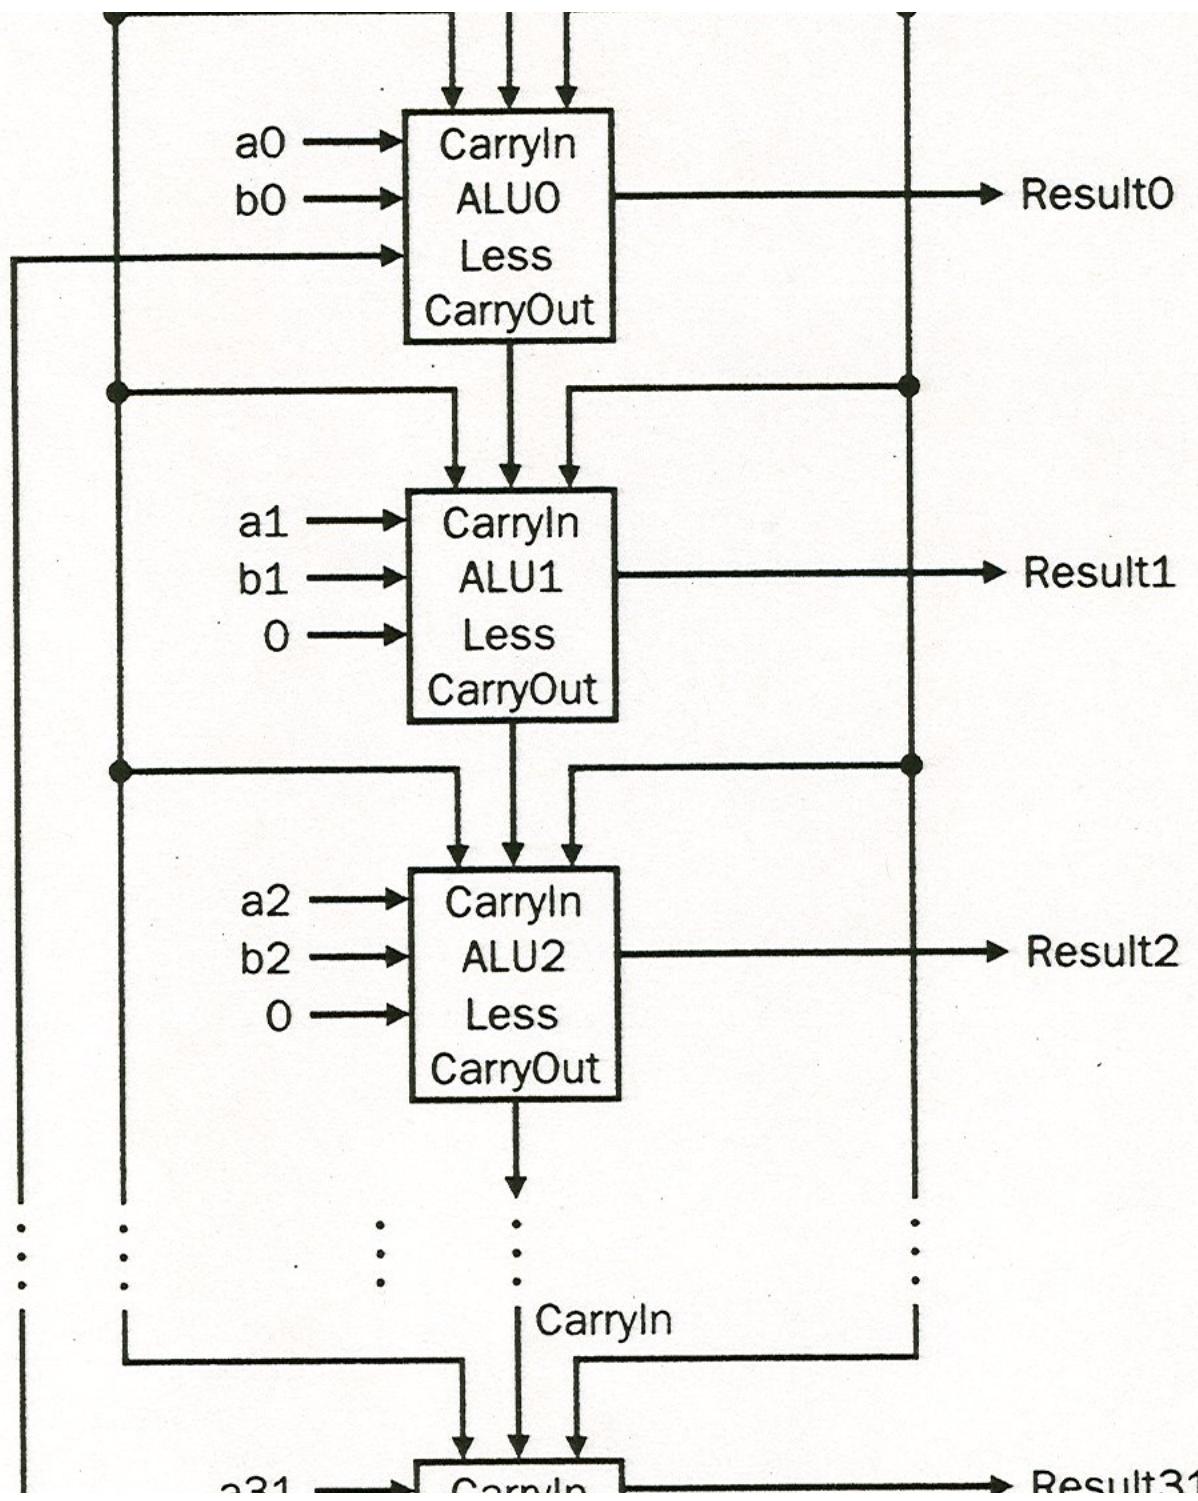
'Binvert' and 'CarryIn' would be set to '1' for subtract.

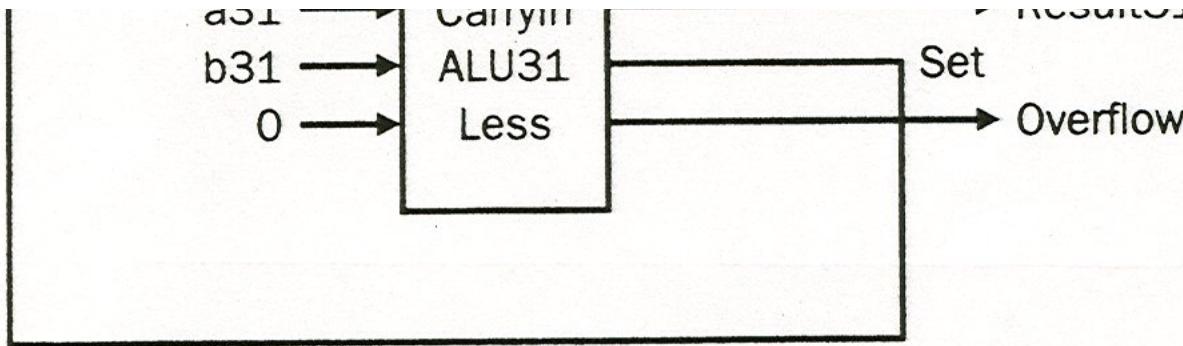
'Binvert' and 'a' set to '0' would be complement.

The overflow detection is in every stage yet only used in the last stage.

The bit slices are wired together to form a simple ALU:

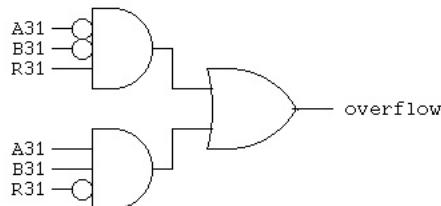






The 'set' operation would give non zero if 'a' < 'b' and zero otherwise. A possible condition status or register value for a "beq" instruction.

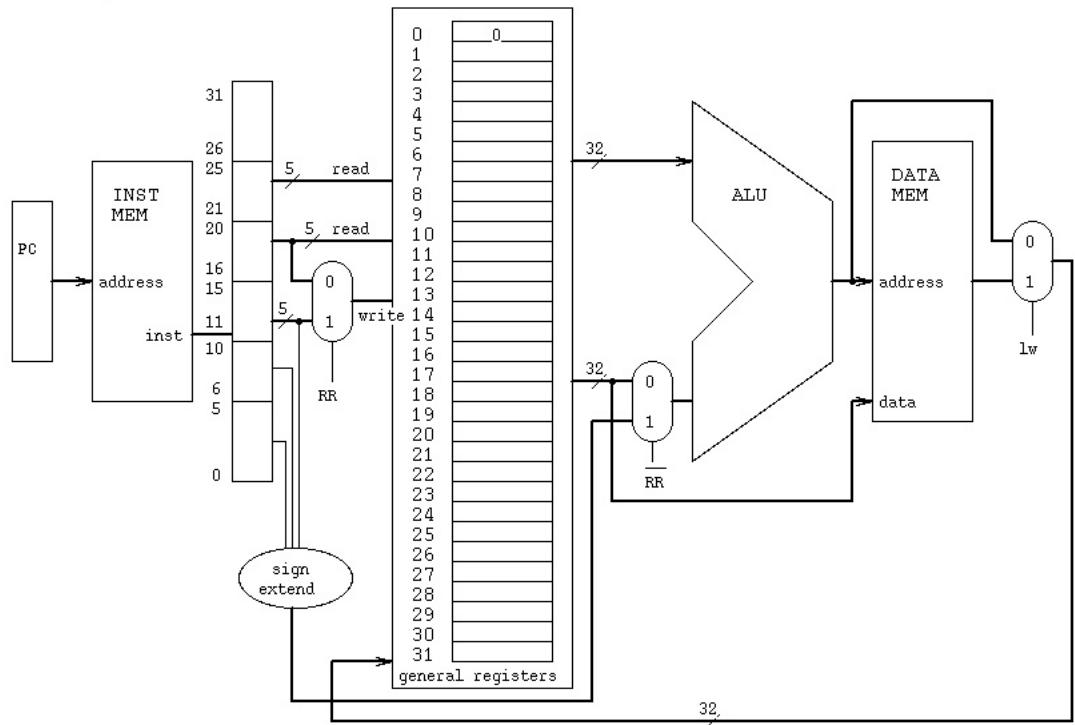
If overflow was to be detected, the circuit below uses the sign bit of the A and B inputs and the sign bit of the result to detect overflow on twos complement addition.



Two's complement addition overflow  
when signs of input same but  
sign of output is different.

The ALU fits into the machine architecture as shown below:

Multiplexors control flow of data



32-bit and 64-bit ALU architectures are available.

A 64-bit architecture, by definition, has 64-bit integer registers. Many computers have had 64-bit IEEE floating point for many years. The 64-bit machines have been around for a while as the Alpha and PowerPC yet have become popular for the desktop with the Intel and AMD 64-bit machines.

## Current 64-bit microprocessor architectures

64-bit microprocessor architectures (as of 2006) include:

- The DEC Alpha architecture (view [Digital Alpha timeline](#))
- Intel's IA-64 architecture (used in Intel's Itanium CPUs)
- AMD's [AMD64](#) architecture, a 64-bit version of the x86 architecture (used in AMD's Athlon 64, Opteron, Sempron, and Turion 64 CPUs).
  - Intel now uses the same instruction set, calling it EM64T.
- SPARC architecture (64-bit as of SPARC V9)
  - Sun's UltraSPARC architecture
  - Fujitsu's SPARC64 architecture
- IBM's POWER architecture (64-bit as of POWER3 and RS64 variants)
- IBM/Motorola's PowerPC architecture (64-bit PowerPC 620 and PowerPC 970 variants)
- IBM's z/Architecture, used by IBM zSeries and System z9 mainframes, a 64-bit version of the [ESA/390](#) architecture
- MIPS Technologies' MIPS IV, MIPS V, and MIPS64 architectures
- HP's PA-RISC family (64-bit as of PA-RISC 2.0)

Most 64-bit processor architectures can execute code for the 32-bit version of the architecture natively without any performance penalty.

This kind of support is commonly called *bisarch support* or more generally *multi-arch support*.

Software has been dragging well behind computer architecture.

The chaos started in 1979 with the following "choices."

The differences between the three models lies in the non-pointer data types. The table below details the data types for the above three data models and includes LP32 and ILP32 for comparison purposes.

Data Type	LP32	ILP32	ILP64	LLP64	LP64
<b>char</b>	8	8	8	8	8
<b>short</b>	16	16	16	16	16
<b>int32</b>			32		
<b>int</b>	16	32	64	32	32
<b>long</b>	32	32	64	32	64
<b>long long (int64)</b>				64	
<b>pointer</b>	32	32	64	64	64

When the width of one or more of the C data types is changed, applications may be affected in various ways. These effects fall into two main categories:

- Data objects, such as a structure, defined with one of the 64-bit data types will be different in size from those declared in an identical way on a 16 or 32-bit system.
- Common assumptions about the relationships between the fundamental data types may no longer be valid in a 64-bit data model. Applications which depend on those relationships often cease to work properly when compiled on a 64-bit platform. A typical assumption made by many application developers is that:

`sizeof(int) = sizeof(long) = sizeof(pointer)`

This relationship is not codified in any C programming standard, but it is valid for the ILP32 data model. However, it is not valid for two of the three 64-bit data models described above, nor is it valid for the LP32 data model.

The full whitepaper [www.unix.org/whitepapers/64bit.html](http://www.unix.org/whitepapers/64bit.html)

My desire is to have the compiler, linker and operating system be LLP64. All my code would work fine. I make no assumptions about word length. I use `sizeof(int)` `sizeof(size_t)` etc. when absolutely needed. On my 8GB computer I use a single array of over 4GB thus the subscripts must be 64-bit. The only option, I know of, for gcc is `-m64` and that just gives LP64. Yuk! I have to change my source code and use "long" everywhere in place of "int". If you get the idea that I am angry with the compiler vendors, you are correct!

Here are sample programs and output to test for 64-bit capability in gcc:

[Get sizeof on types and variables big.c](#)

[output from gcc -m64 big.c big.out](#)

[malloc more than 4GB big\\_malloc.c](#)

[output from big\\_malloc mac.out](#)

Newer Operating Systems and compilers (note 'sizeof' changed to long)

[Get sizeof on types and variables big12.c](#)

[output from gcc big12.c big12.out](#)

The early 64-bit computers were:

[DEC Alpha](#)

[DEC Alpha](#)

[IBM PowerPC](#)

Some history of 64-bit computers:

- 1961: IBM delivered the [IBM 7030 Stretch supercomputer](#). This used 64-bit data words and 32 or 64-bit instruction words.
- 1974: Control Data Corporation launched the [CDC Star-100 vector supercomputer](#), which used a 64-bit word architecture (previous CDC systems were based on a 60-bit architecture).
- 1976: Cray Research delivered the first [Cray-1 supercomputer](#). This was based on a 64-bit word architecture, which formed the basis for later Cray vector supercomputers.
- 1983: Elxsi launched the [Elxsi 6400 parallel minisupercomputer](#). The Elxsi architecture had 64-bit data registers but a 32-bit address space.
- 1991: [MIPS Technologies](#) produced the first 64-bit microprocessor, as the third revision of their [MIPS RISC](#) architecture, the [R4000](#). The CPU was used in [SGI](#) graphics workstations starting with the [IRIS Crimson](#). However, 64-bit support for the R4000 was not included in the [IRIX](#) operating system until IRIX 6.2, released in 1996.
- 1992: [Digital Equipment Corporation \(DEC\)](#) introduced the pure 64-bit [Alpha](#) architecture which was born from the [PRISM](#) project.
- 1993: DEC released the 64-bit [OSF/1 AXP](#) Unix-like operating system (later renamed [Tru64 UNIX](#)) and the [OpenVMS](#) operating system for Alpha systems.
- 1994: Intel announced plans for the 64-bit [IA-64](#) architecture (jointly developed with HP) as a successor to its 32-bit IA-32 processors. A 1998–1999 launch date was targeted. SGI released IRIX 6.0, with 64-bit support for [R8000](#) CPUs.
- 1995: Sun launched a 64-bit [SPARC](#) processor, the [UltraSPARC](#). Fujitsu-owned [HAL Computer Systems](#) launched workstations based on a 64-bit CPU, HAL's independently designed first generation SPARC64. IBM released 64-bit AS/400 systems, with the upgrade able to convert the operating system, database and applications. DEC released [OpenVMS Alpha 7.0](#), the first full 64-bit version of OpenVMS for Alpha.
- 1996: HP released an implementation of the 64-bit 2.0 version of their [PA-RISC](#) processor architecture, the [PA-8000](#). Nintendo introduces the [Nintendo 64](#) video game console, built around a low-cost variant of the MIPS R4000.
- 1997: IBM released their [RS64](#) full 64-bit PowerPC processors.
- 1998: IBM released their [POWER3](#) full 64-bit PowerPC/POWER processors. Sun released Solaris 7, with full 64-bit UltraSPARC support.
- 1999: Intel released the [instruction set](#) for the IA-64 architecture. First public disclosure of AMD's set of 64-bit extensions to IA-32, called [x86-64](#) (later renamed [AMD64](#)).
- 2000: IBM shipped its first 64-bit mainframe, the [zSeries z900](#), and its new [z/OS](#) operating system — culminating history's biggest 64-bit processor development investment and instantly wiping out 31-bit plug-compatible competitors Fujitsu/Amdahl and Hitachi. 64-bit Linux on [zSeries](#) followed almost immediately.
- 2001: Intel finally shipped its 64-bit processor line, now branded [Itanium](#), targeting high-end servers. It fails to meet expectations due to the repeated delays getting IA-64 to market, and becomes a flop. [Linux](#) was the first operating system to run on the processor at its release.
- 2002: Intel introduced the [Itanium 2](#) as a successor to the Itanium.
- 2003: AMD brought out its [AMD64](#)-architecture [Opteron](#) and [Athlon 64](#) processor lines. Apple also shipped 64-bit "G5" PowerPC 970 CPUs courtesy of IBM, along with an update to its Mac OS X operating system, that added partial support for 64-bit mode. Several Linux distributions released with support for AMD64. Microsoft announced that it would create a version of its Windows operating system for these AMD chips. Intel maintained that its Itanium chips would remain its only 64-bit processors.

[Java for 64-bit, source compatible](#)

[Then to VHDL resource, FPGA,  
get free GHDL](#)

## Lecture 9, Multiply

Standard decimal and binary multiplication could look like:

$$\begin{array}{r}
 234 \\
 \times 121 \\
 \hline
 234 \\
 468 \\
 234 \\
 \hline
 028314
 \end{array}
 \quad
 \begin{array}{r}
 01010 \\
 \times 00011 \\
 \hline
 01010 \\
 01010 \\
 00000 \\
 \hline
 00000
 \end{array}
 \quad
 \begin{array}{l}
 \text{multiplicand} \\
 \times \text{multiplier} \\
 \hline
 \text{product}
 \end{array}$$

```

|     0000011110 5-bits times 5-bits gives a 10-bit product,
|     in a computer leading zeros are kept.

3-digits times 3-digits gives a 6-digit product, yet in
decimal, we do not write the leading zeros.

```

We have covered how computer adders work and how they are built.  
Exactly two numbers are added to produce one sum, thus the binary multiply above needs to be rewritten as:

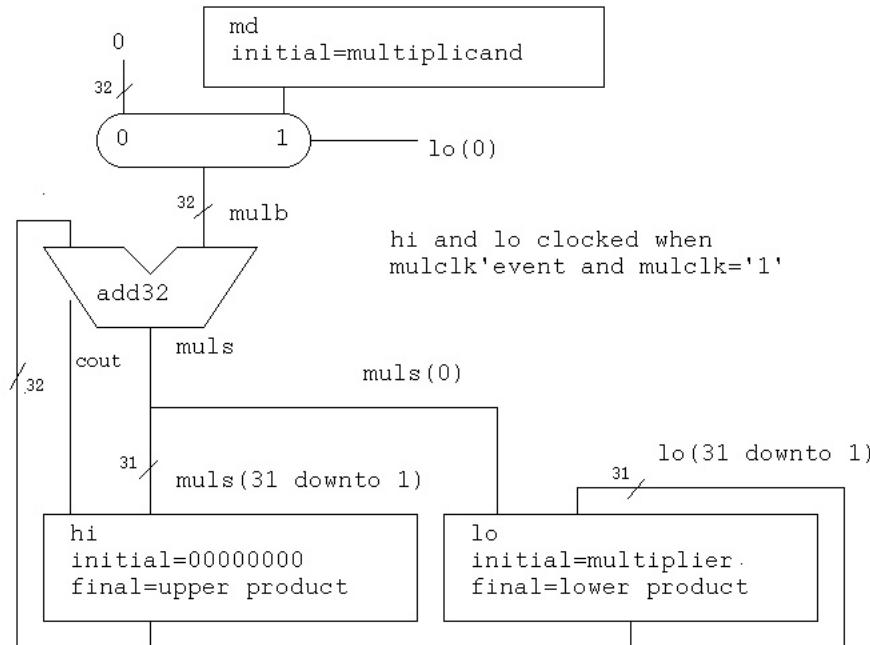
```

    01010
    x 00011
    -----
      001010 -- multiplier LSB anded with multiplicand
      + 01010 -- multiplier bit-1 anded with multiplicand
      -----
      0011110 -- partial sum, bottom bit passed down
      + 00000 -- multiplier bit-2 anded with multiplicand
      -----
      00011110 -- partial sum, bottom two bits passed down
      + 00000 -- multiplier bit-3 anded with multiplicand
      -----
      000011110 -- partial sum, bottom three bits passed down
      + 00000 -- multiplier bit-4 anded with multiplicand
      -----
      0000011110 -- final product, four bits passed down

```

Thus, by this simple method, with a 5-bit unsigned multiplier, there are four additions needed. A circuit that uses one adder and performs serial multiplication follows directly. This design chose to use a multiplexor rather than an 'and' operation to select the multiplicand or zero.

#### [How a register works](#)



The VHDL code that represents the above circuit is:

```
mula <= hi;
```

```

mulb <= md when (lo(0)='1') else x"00000000" after 50 ps;
adder:entity WORK.add32 port map(mula, mulb, '0', muls, cout);
hi <= cout & muls(31 downto 1) when mulclk'event and mulclk='1';
lo <= muls(0) & lo(31 downto 1) when mulclk'event and mulclk='1';

```

The signal "mulclk" runs for the number of clock cycles that their are bits in the multiplier, 32 for this example. For simplicity of design, zero is added in the first step. Note that "cout" is used when loading the "hi" register. The shifting is accomplished by wire routing.

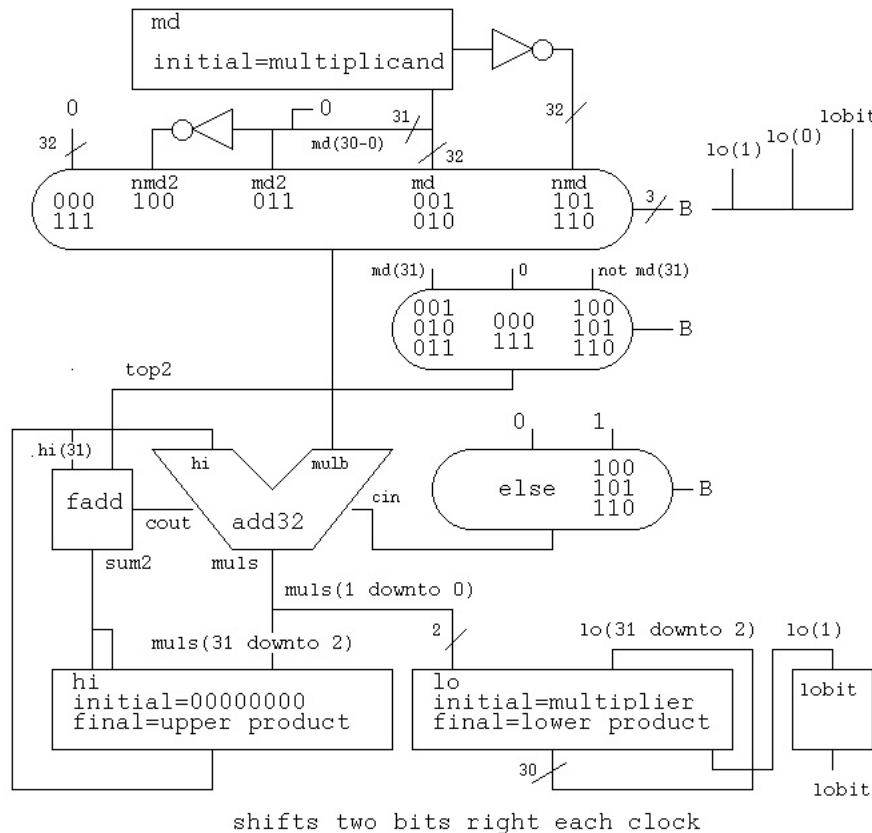
The VHDL test source code is [mul\\_ser.vhd](#)

The output from the test is [mul\\_ser.out](#)

P.S. The above was an introduction, never use that method or circuit.

A serial multiplier can be built using only half as many clock cycles. We use the technique developed by Mr. Booth. Two multiplier bits are used each clock cycle. Only one add operation is needed each cycle, yet the augend has several possible values as shown by the multiplexor in the schematic and the table in the VHDL source code.

### Booth serial multiplier



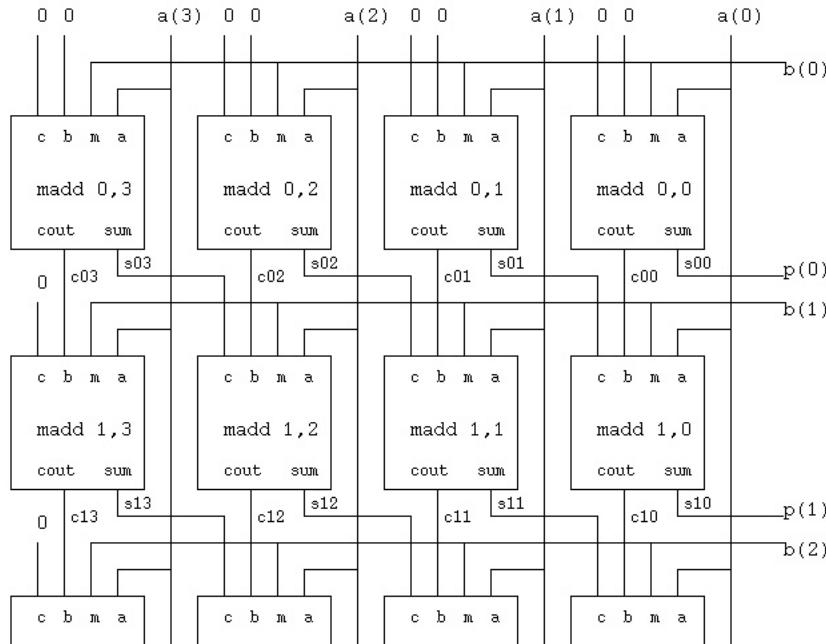
The VHDL test source code is [bmul\\_ser.vhd](#)

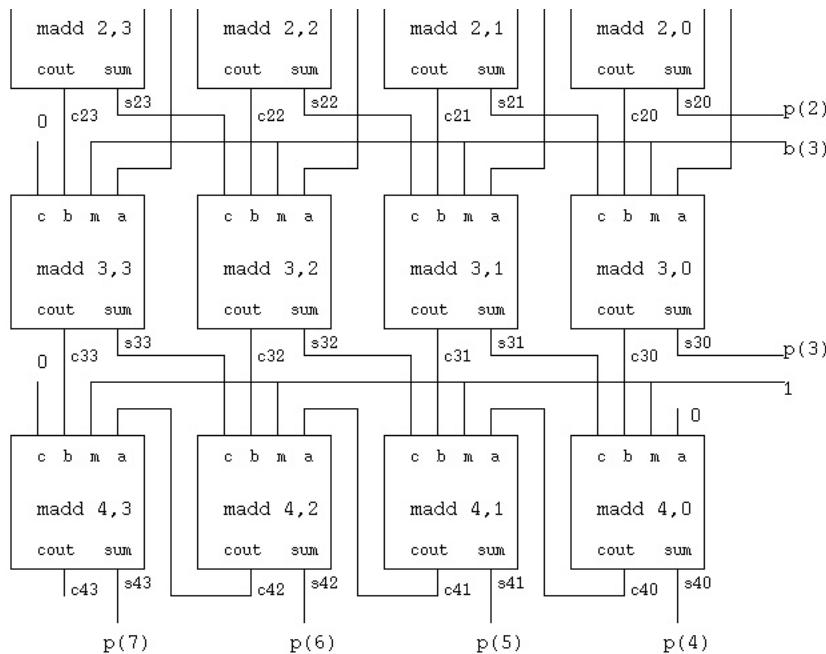
The output from the test is [bmul\\_ser.out](#)

Next, parallel multiplication with a carry-save design.  
Note there is no carry propagation except in the last stage.

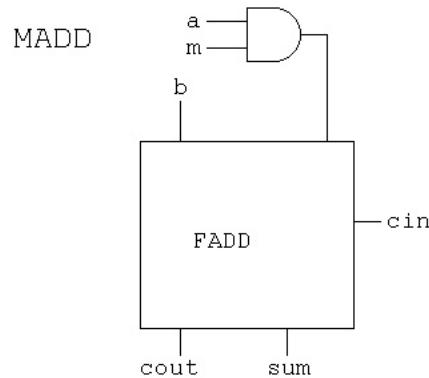
parallel multiplier (unsigned)

$$a(0-3) * b(0-3) = p(0-7)$$



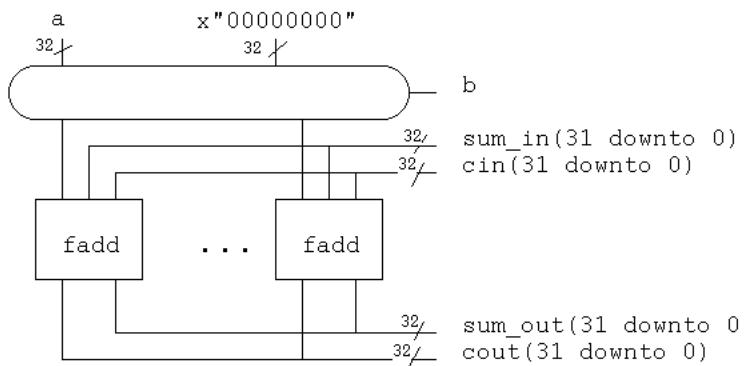
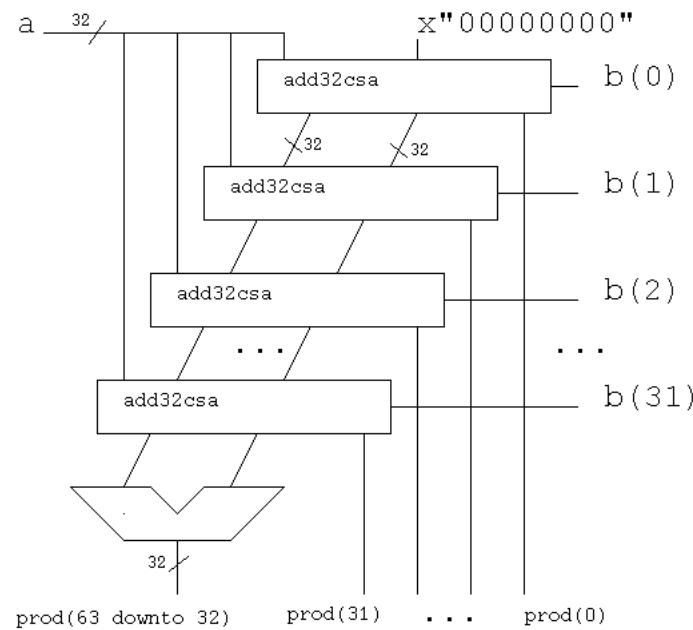


Notes: c00 is shorthand for  $c(0,0)$       1 is shorthand for '1'  
 s00 is shorthand for  $s(0,0)$       0 is shorthand for '0'  
 statements are needed for  $p(7) \leq s(4,3), \dots p(0) \leq s(0,0)$



Some fancy VHDL using double subscripting and "generate".  
[pmul4.vhd1](#)

A 32 bit design using an add32csa entity is:



Schematic of add32csa

The VHDL entity for the carry-save multiplier is [mul32c.vhd1](#)

The VHDL test source code is [mul32c\\_test.vhd1](#)

The output from the test is [mul32c\\_test.out](#)

We can now combine the Booth multiplication technique to reduce the number of stages in half, still using the parallel multiply. The VHDL was written without a diagram, thus no schematic exists, yet.

The VHDL entity for the carry-save multiplier is [bmul32.vhd1](#)

The VHDL test source code is [bmul32\\_test.vhd1](#)

The output from the test is [bmul32\\_test.out](#)

[Homework 5 is assigned](#)

## Lecture 10, Divide

Hopefully you understand decimal division:

```
49 quotient
divisor 47 / 2345 dividend
      188
      ---
      465
      423
      ---
      42 remainder
```

And check division by multiplication:

```
49 multiplicand is the quotient above
x 47 multiplier is the divisor above
-----
2303
+ 42 add the remainder above
-----
2345 final sum is the dividend above
```

A smaller case that is used below in binary:

```
12 quotient
divisor 7 / 85 dividend
      7
      --
      15
      14
      --
      1 remainder
```

Binary divide, conventional method and non restoring method

These examples are shown in a form that can be directly implemented in a computer architecture.

The divisor, quotient and remainder are each one word.  
The dividend is two words.

The equations  $\text{dividend} = \text{quotient} * \text{divisor} + \text{remainder}$   
and  $|\text{remainder}| < |\text{divisor}|$   
must be satisfied.

When a choice is possible, choose the sign of the remainder to be the same as the sign of the dividend.

Save the sign bits of the dividend and divisor, if necessary,  
negate the dividend and divisor to make them positive.  
Fix up the sign bits of the quotient and dividend after dividing.

Example: dividend = 85 , divisor = 7

Decimal divide 85 / 7 = quotient 12 , remainder 1

Restoring (conventional) binary divide, twos complement 4-bit numbers

```
1 1 0 0   quotient
0 1 1 1   / 0 1 0 1 0 1 0 1
      -0 1 1 1   may subtract by adding twos complement
                  - 0 1 1 1   is 1 0 0 1
5 - 7 = -2   1 1 1 0
negative, add 7   +0 1 1 1
```

```

restored      1 0 1 0
next bit     -0 1 1 1
              -----
10 - 7 = 3   0 1 1 1
quotient=1, next bit -0 1 1 1
              -----
7 - 7 = 0   0 0 0 0
quotient=1, next bit -0 1 1 1
              -----
0 - 7 = -7   1 0 0 1
negative, add 7 +0 1 1 1
quotient=0
restored, next bit 0 0 0 1
                    -0 1 1 1
              -----
1 - 7 = -6   1 0 1 0
negative, add 7 +0 1 1 1
quotient=0
restored, finished 0 0 0 1 final remainder
(8 cycles using adder)

```

Clock cycles can be saved by not performing the "restored" operation.

non-restoring binary divide, twos complement 4-bit numbers  
note: 7 = 0 1 1 1    -7 = 1 0 0 1

```

          1 1 0 0  quotient
0 1 1 1 / 0 1 0 1 0 1 0 1
pre shift      +1 0 0 1 adding twos complement of divisor
              -----
10 - 7 = 3   0 0 1 1 1
quotient=1    +1 0 0 1
next bit subtract 0 0 0 0 0
7 - 7 = 0   0 0 0 0 0
quotient=1    +1 0 0 1
next bit subtract 0 0 0 1 1
0 - 7 = -7   1 0 0 1 1
quotient=0    +0 1 1 1 adding divisor
next bit add 2 + 7 = 9 = -7   1 0 1 0
quotient=0    +0 1 1 1
correction add
final remainder 0 0 0 1 remainder
(5 cycles using adder)

```

Correcting signs:

dividend	divisor	quotient	remainder
+	+	+	+85 / +7 = +12 R +1
+	-	-	+85 / -7 = -12 R +1
-	+	-	-85 / +7 = -12 R -1
-	-	+	-85 / -7 = +12 R -1

Humans, not the computer, keeps track of the binary point.

Integers	Fractions	(fixed point)
qqqq.	.qqqq	q.qqq
ssss. / dddddddd.	.ssss / .ddddd	ss.ss / ddd.dddd
rrrr.	.0000rrrr	.0rrrr
<hr/>		
* qqqq. * ssss.	* .qqqq * .ssss	* q.qqq * ss.ss

$$\begin{array}{r}
 \text{ttttttt.} & \text{.ttttttt} & \text{ttt.ttttt} \\
 + \text{rrrr.} & + \text{.0000rrrr} & + \text{.0rrrr} \\
 \hline
 \text{ddddd.} & \text{.ddddd} & \text{ddd.dddd}
 \end{array}$$

for multiply, counting positions from the right, the binary point of the product is at the sum of the positions of the multiplicand and multiplier.

for divide, counting positions from the right, the binary point of the quotient is at the difference of the positions of the dividend and divisor. The binary point of the remainder is in the same position as the binary point of the dividend.

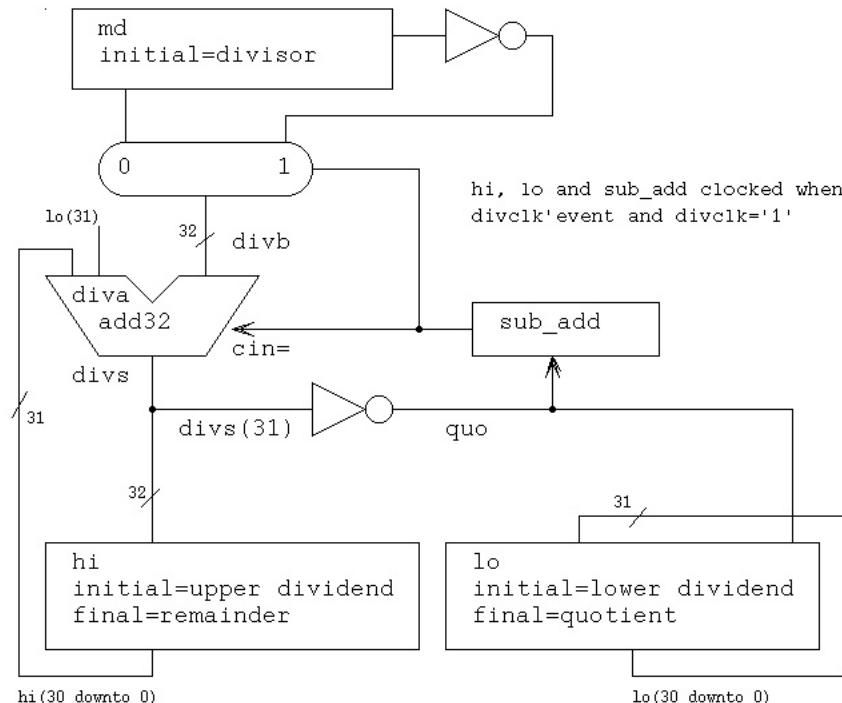
Overflow occurs when the top half of dividend is greater than or equal to the divisor, thus division by zero is always overflow.

No schematic or VHDL is provided for restoring division because it is never used in practice. The serial non restoring division is:

A possible design for a serial divide, does not include remainder correction:

```

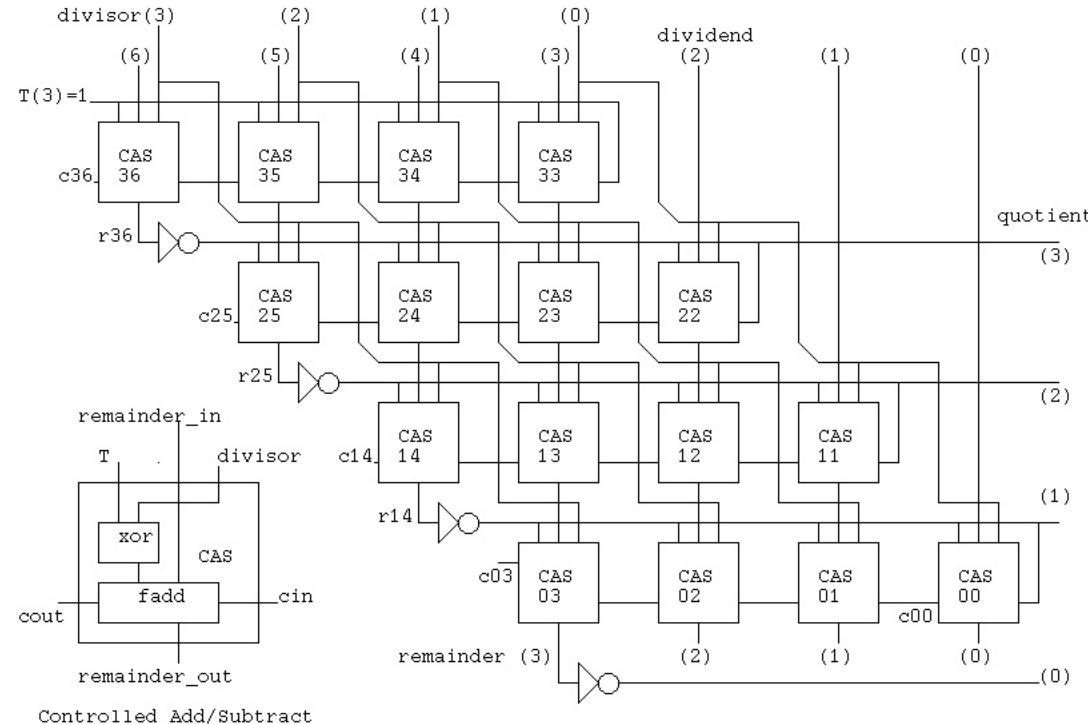
diva  <= hi(30 downto 0) & lo(31) after 50 ps; -- shift
divb  <= not md when sub_add='1' else md after 50 ps; -- subtract or add
adder:entity WORK.add32 port map(diva, divb, sub_add, divs, cout);
quo   <= not divs(31) after 50 ps; -- quotient bit
hi    <= divs when divclk'event and divclk='1';
lo    <= lo(30 downto 0) & quo when divclk'event and divclk='1';
sub_add <= quo when divclk'event and divclk='1';
  
```



The full VHDL code is [div\\_ser.vhd](#)  
with output [div\\_ser.out](#)

Note that the remainder is not corrected by this circuit.  
The FFFFFFFFA should have the divisor 00000007 added to it,  
making the remainder 00000001

Now that you understand how binary division works and understand  
how multiplication can be speeded up using parallel circuits,  
we show a parallel division circuit and its simulation.



[divcas4\\_test.vhd](#)

[divcas4\\_test.out](#)

Note that the output includes the time.  
Observe the first few lines of printout replacing 'U' undefined,  
meaning not computed, with zeros or ones. Unfortunately, if VHDL  
prints hexadecimal, any state except one is printed as zero.

For part1 project you are given divcas16.vhd  
This divides as 32 bit number by a 16 bit number and  
produces a 16 bit quotient and 16 bit remainder.

[divcas16.vhd](#)

It would be nice if I could have a 4-bit radix 2 or radix 4 SRT  
division schematic here. Parallel circuits that perform division  
may use (-2, -1, 0, 1, 2) values for intermediate signals.  
Two or more bits of the quotient may be computed at each stage,  
based on a table and a few bits of the divisor and partial  
remainder.

[SRT Divide, click on slide show .pdf](#)  
[SRT Divide .pdf local](#)

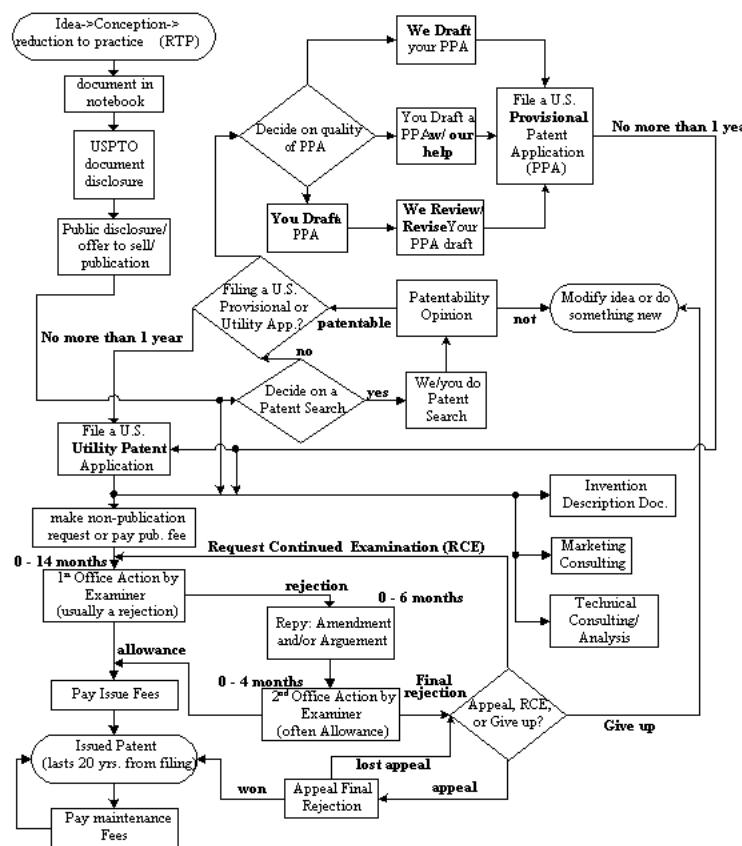
[freepatentsonline.com/5272660.html](http://freepatentsonline.com/5272660.html)

Software can be copyrighted. Just doing a physical embodiment makes you the owner of the copyright. Add Copyright year name to the document or computer file. If you want your copyright to stand up in a court of law, you need to file the copyright. Get the latest information, at one time there was a \$40.00 filing fee and the copyright was good for 28 years, renewable for 67 more years, for a total of 95 years.

There is a "fair use" clause that allows personal use of parts of a copyrighted document.

Software and hardware and processes may be patented. A utility patent is good for 20 years, a design patent is good for 14 years. The cost of completing the process of getting a patent is variable. 20 years ago the average cost was \$5,000.00 and today the average cost is about \$15,000.00. There are companies that can help you, do-it-yourself, with advertised cost starting from about \$1,500.00. (There may be additional maintenance fees at 3 1/2 years etc.) ((It may take a year or more to get a patent.))

One version of the process to get a patent is:



© 2002-2003 Bay Area Intellectual Property Group, LLC. All Rights Reserved.

There is no "fair use" clause on patents.

## Lecture 11, Floating Point

Almost all Numerical Computation arithmetic is performed using IEEE 754-1985 Standard for Binary Floating-Point Arithmetic. The two formats that we deal with in practice are the 32 bit and 64 bit formats. You need to know how to get the format you desire in the language you are programming. Complex numbers use two values.

C	Java	Fortran 95	older Fortran	Ada 95	MATLAB
32 bit	float	float	real	real	float N/A
64 bit	double	double	double precision	real*8	long_float 'default'

complex					
32 bit	'none'	'none'	complex	complex	complex N/A
64 bit	'none'	'none'	double complex	complex*16	long_complex 'default'

'none' means not provided by the language (may be available as a library)  
N/A means not available, you get the default.

IEEE Floating-Point numbers are stored as follows:

The single format 32 bit has  
1 bit for sign, 8 bits for exponent, 23 bits for fraction  
The double format 64 bit has  
1 bit for sign, 11 bits for exponent, 52 bits for fraction

There is actually a '1' in the 24th and 53rd bit to the left of the fraction that is not stored. The fraction including the non stored bit is called a significand.

The exponent is stored as a biased value, not a signed value.  
The 8-bit has 127 added, the 11-bit has 1023 added.  
A few values of the exponent are "stolen" for special values, +/- infinity, not a number, etc.

Floating point numbers are sign magnitude. Invert the sign bit to negate.

Some example numbers and their bit patterns:

decimal	stored hexadecimal	sign	exponent	fraction	significand in binary
					The "1" is not stored
					biased exponent
1.0	3F 80 00 00	0	0 1111111	00000000000000000000000000000000	1.0 * 2^(127-127)
0.5	3F 00 00 00	0	0 1111110	00000000000000000000000000000000	1.0 * 2^(126-127)
0.75	3F 40 00 00	0	0 1111110	10000000000000000000000000000000	1.1 * 2^(126-127)
0.9999995	3F 7F FF FF	0	0 1111110	11111111111111111111111111111111	1.1111* 2^(126-127)
0.1	3D CC CC CD	0	0 1111011	10011001100110011001101	1.1001* 2^(123-127)
					63 62..... 52 51 ..... 0
1.0	3F F0 00 00 00 00 00	0	0 1111111111	000 ... 000	1.0 * 2^(1023-1023)
0.5	3F E0 00 00 00 00 00	0	0 1111111110	000 ... 000	1.0 * 2^(1022-1023)
0.75	3F E8 00 00 00 00 00	0	0 1111111110	100 ... 000	1.1 * 2^(1022-1023)
0.999999999999995	3F EF FF FF FF FF FF	0	0 1111111110	111 ...	1.11111* 2^(1022-1023)

```

0.1
3F B9 99 99 99 99 9A   0 0111111011 10011..1010 1.10011* 2^(1019-1023)
                           sign   exponent      fraction
                           before storing subtract bias

```

Note that an integer in the range 0 to  $2^{23} - 1$  may be represented exactly. Any power of two in the range -126 to +127 times such an integer may also be represented exactly. Numbers such as 0.1, 0.3, 1.0/5.0, 1.0/9.0 are represented approximately. 0.75 is 3/4 which is exact. Some languages are careful to represent approximated numbers accurate to plus or minus the least significant bit. Other languages may be less accurate.

```

/* flt.c just to look at .o file with hdump */
void flt() /* look at IEEE floating point */
{
    float x1 = 1.0f;
    float x2 = 0.5f;
    float x3 = 0.75f;
    float x4 = 0.99999f;
    float x5 = 0.1f;

    double d1 = 1.0;
    double d2 = 0.5;
    double d3 = 0.75;
    double d4 = 0.99999999;           The "1" not stored
    double d5 = 0.1;                 in binary
}
31 30....23 22.....0 |
3F 80 00 00 00 00 00 00 00 00000000000000000000000000000000 1.0 * 2^(127-127)
3F 00 00 00 00 00 00 00 00 00000000000000000000000000000000 1.0 * 2^(126-127)
3F 40 00 00 00 00 00 00 00 00000000000000000000000000000000 1.1 * 2^(126-127)
3F 7F FF 58 00 0111110 1111111111111010101000 1.1111* 2^(126-127)
3D CC CC CD 0 0111011 10011001100110011001101 1.1001* 2^(123-127)

```

```

63 62.....52 51 .....0
3F F0 00 00 00 00 00 00 00 00 00000000000000000000000000000000 1.0 * 2^(1023-1023)
3F E0 00 00 00 00 00 00 00 00 00000000000000000000000000000000 1.0 * 2^(1022-1023)
3F E8 00 00 00 00 00 00 00 00 00000000000000000000000000000000 1.1 * 2^(1022-1023)
3F EF FF FF FA A1 9C 47 0 011111110 111 ... 1.11111* 2^(1022-1023)
3F B9 99 99 99 99 99 9A 0 0111111011 1001 ..1010 1.10011* 2^(1019-1023)
                           sign   exponent      fraction
                           before storing subtract bias
decimal          binary fraction / decimal exponent IEEE normalize
                           binary

```

Now, all the above is the memory, RAM, format. Upon a load operation of either float or double into one of the floating point registers, the format in the register extended to greater precision than double. All floating point arithmetic is performed at this greater precision. Upon a store operation, the greater precision is reduced to the memory format, possibly with rounding. From a programming viewpoint, always use double.

exponents must be the same for add and subtract!

A = 3.5 * 10^6	a = 11.1 * 2^6	1.11 * 2^7		
B = 2.5 * 10^5	b = 10.1 * 2^5	1.01 * 2^6		
A+B	3.50 * 10^6	a+b	11.10 * 2^6	1.110 * 2^7
	+ 0.25 * 10^6		+ 1.01 * 2^6	+ 0.101 * 2^7
	<hr/>		<hr/>	<hr/>
	3.75 * 10^6		100.11 * 2^6	10.011 * 2^7
			normalize 1.0011 * 2^8	
			IEEE	
A-B	3.50 * 10^6		normalize 0.10011 * 2^9	
	- 0.25 * 10^6		fraction	

```

-----
 3.25 * 10^6

A*B      3.50 * 10^6
          * 2.5 * 10^5
-----
 8.75 * 10^11

A/B    3.5 *10^6 / 2.5 *10^5 = 1.4 * 10^1

```

The mathematical basis for floating point is simple algebra

The common uses are in computer arithmetic and scientific notation

given: a number  $x_1$  expressed as  $10^{e_1} \cdot f_1$   
then 10 is the base,  $e_1$  is the exponent and  $f_1$  is the fraction  
example  $x_1 = 10^3 \cdot 1.234$  means  $x_1 = 123.4$  or  $1.234 \cdot 10^3$   
or in computer notation  $0.1234E3$

In computers the base is chosen to be 2, i.e. binary notation  
for  $x_1 = 2^{e_1} \cdot f_1$  where  $e_1=3$  and  $f_1 = .1011$   
then  $x_1 = 101.1$  base 2 or, converting to decimal  $x_1 = 5.5$  base 10

Computers store the sign bit, 1=negative, the exponent and the fraction in a floating point word that may be 32 or 64 bits.

The operations of add, subtract, multiply and divide are defined as:

Given  $x_1 = 2^{e_1} \cdot f_1$   
 $x_2 = 2^{e_2} \cdot f_2$  and  $e_2 \leq e_1$

$x_1 + x_2 = 2^{e_1} \cdot (f_1 + 2^{-(e_1-e_2)} \cdot f_2)$   $f_2$  is shifted then added to  $f_1$

$x_1 - x_2 = 2^{e_1} \cdot (f_1 - 2^{-(e_1-e_2)} \cdot f_2)$   $f_2$  is shifted then subtracted from  $f_1$

$x_1 \cdot x_2 = 2^{(e_1+e_2)} \cdot f_1 \cdot f_2$

$x_1 / x_2 = 2^{(e_1-e_2)} \cdot (f_1 / f_2)$

an additional operation is usually needed, normalization.  
if the resulting "fraction" has digits to the left of the binary point, then the fraction is shifted right and one is added to the exponent for each bit shifted until the result is a fraction.

We will use fraction normalization, not IEEE normalization:

if the resulting "fraction" has zeros immediately to the right of the binary point, then the fraction is shifted left and one is subtracted from the exponent for each bit shifted until there is a non zero digit to the right of the binary point.

Numeric examples using equations:

(exponents are decimal integers, fractions are decimal)  
(normalized numbers have  $1.0 > \text{fraction} \geq 0.5$ )  
(note fraction strictly less than 1.0, greater than or equal 0.5)

$x_1 = 2^4 \cdot 0.5$  or  $x_1 = 8.0$   
 $x_2 = 2^2 \cdot 0.5$  or  $x_2 = 2.0$

$x_1 + x_2 = 2^4 \cdot (.5 + 2^{-(4-2)} \cdot .5) = 2^4 \cdot (.5 + .125) = 2^4 \cdot .625$

$x_1 - x_2 = 2^4 \cdot (.5 - 2^{-(4-2)} \cdot .5) = 2^4 \cdot (.5 - .125) = 2^4 \cdot .375$   
not normalized, multiply fraction by 2, subtract 1 from exponent  
 $= 2^3 \cdot .75$

$x_1 \cdot x_2 = 2^{(4+2)} \cdot (.5 \cdot .5) = 2^6 \cdot .25$  not normalized  
 $= 2^5 \cdot .5$  normalized

$x_1 / x_2 = 2^{(4-2)} \cdot (.5 / .5) = 2^2 \cdot 1.0$  not normalized  
 $= 2^3 \cdot .5$  normalized

```

Numeric examples, people friendly:
(exponents are decimal integers, fractions are decimal)
(normalized numbers have 1.0 > fraction >= 0.5)

x1 = 0.5 * 2^4
x2 = 0.5 * 2^2

x1 + x2 = 0.500 * 2^4
+ 0.125 * 2^4 unnormalize to make exponents equal
-----
0.625 * 2^4 result is normalized, done.

x1 - x2 = 0.500 * 2^4
- 0.125 * 2^4 unnormalize to make exponents equal
-----
0.375 * 2^4 result is not normalized
0.750 * 2^3 double fraction, halve exponential

x1 * x2 = 0.5 * 0.5 * 2^2 * 2^4 = 0.25 * 2^6 not normalized
= 0.5 * 2^5 normalized

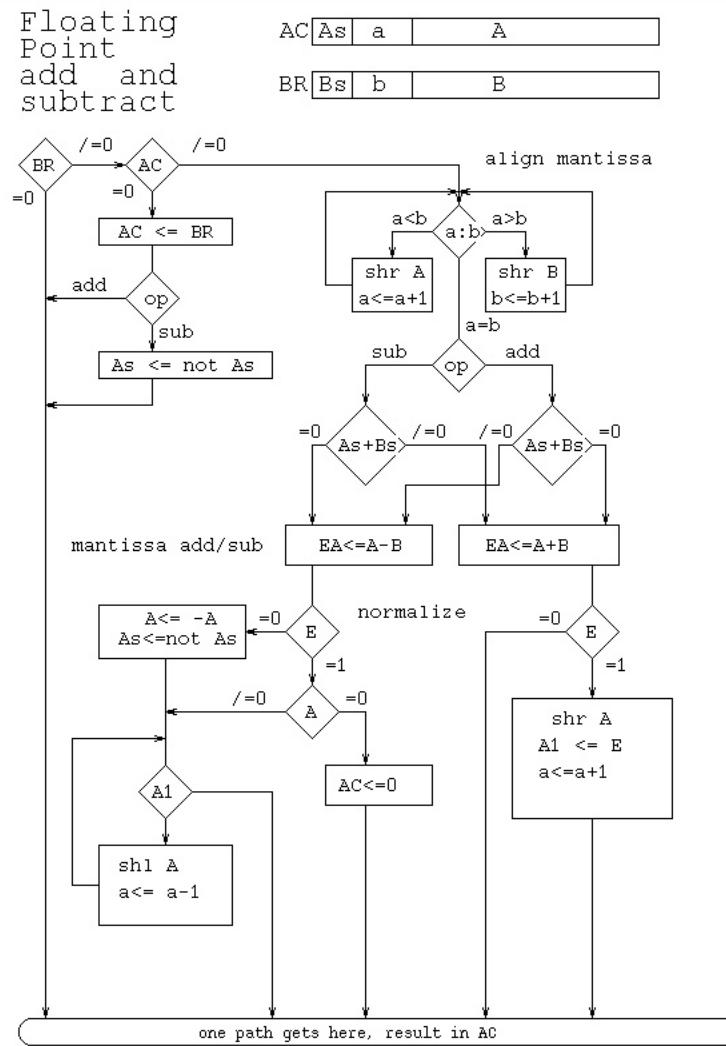
x1 / x2 = (.5/.5) * 2^4/2^2 = 1.0 * 2^2 not normalized
= 0.5 * 2^3 normalized
halve fraction, double exponential

```

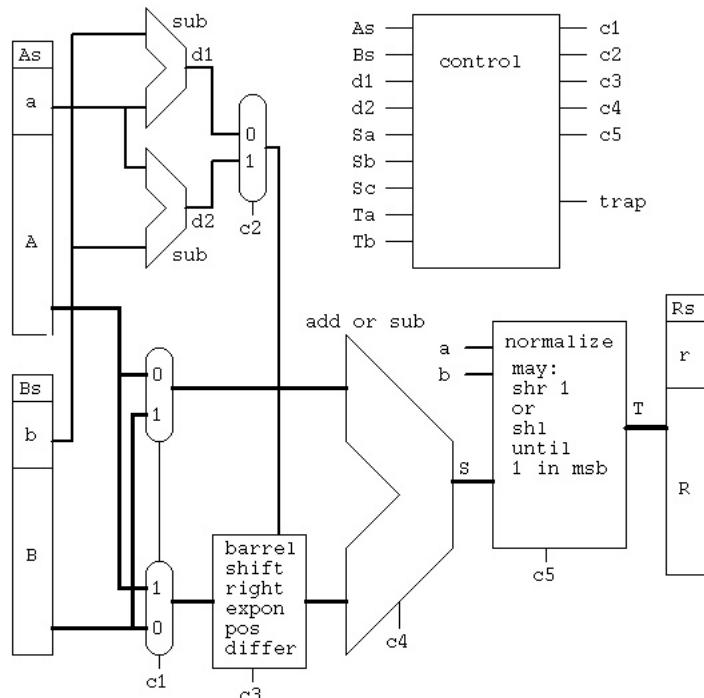
#### [IEEE 754 Floating Point Standard](#)

A few minor problems, e.g. the square root of all complex numbers are in the right half of the complex plane and thus the real part of the square root should never be negative. As a concession to early hardware, the standard define the  $\sqrt{-0}$  to be  $-0$  rather than  $+0$ . Several places the standard uses the word should. If a standard is specifying something, the word shall is typically used.

Basic decisions and operations for floating point add and subtract:



The decisions indicated above could be used to design the control component shown in the data path diagram below:



Floating point add, subtract components and data paths

A hint on normalization, using computer scientific notation:

```
1.0E-8 == 10.0E-9 == 0.01E-6 == 0.0000001 == 10ns == 0.01 microseconds
1.0E8 == 0.1E9 == 100.0E6 == 100,000,000 == 100MHz == 0.1 GHz
1.0/1.0GHz = 1ns clock period
```

Some graphics boards have large computing capacity and some are releasing the specs so programmers can use the computing capacity.

[nVidia example 2007](#)

[512-core by 2011, more today.](#)

Programming 512 cores or more with CUDA or OpenCL is quite a challenge.  
New languages are coming, not optimized yet.

Fortunately, CMSC 411 does not require VHDL for floating point, just the ability to manually do floating point add, subtract, multiply and divide. (Examples above and in class on board.)

## Lecture 12, VHDL - circuits and debugging

Debugging VHDL (or almost any computer input)

1) Expect errors. Nobody's perfect.

2) Automate to make it easy to re-run, e.g. Makefile

In Makefile Cadence VHDL, keep adding more tests

```
all: tadd32.out pmul16_test.out part1.out
# these are the final result files

# <tab> not spaces, precede ncvhdl, ncelab, ncsim

tadd32.out: add32.vhdl add32.vhdl tadd32.run # HW4
ncvhdl -v93 add32.vhdl
ncvhdl -v93 tadd32.vhdl
ncelab -v93 tadd32:circuits
ncsim -batch -logfile tadd32.out -input tadd32.run tadd32
# blank line ends commands

pmul16_test.out: pmul16.vhdl pmul16_test.vhdl pmul16_test.run # HW6
ncvhdl -v93 pmul16_test.vhdl # 8 for CSEE
ncelab -v93 pmul16_test:circuits
ncsim -batch -logfile pmul16_test.out \
-input pmul16_test.run pmul16_test

part1.out: part1.vhdl add32.vhdl bshift.vhdl divcas16.vhdl
pmul16.vhdl part1.chk part1.abs part1.run
ncvhdl -v93 add32.vhdl
ncvhdl -v93 bshift.vhdl
ncvhdl -v93 divcas16.vhdl # 16 for both
ncvhdl -v93 pmul16.vhdl # 8 for CSEE
ncvhdl -v93 part1.vhdl
ncelab -v93 part1:schematic
ncsim -batch -logfile part1.out -input part1.run part1
diff -iw part1.out part1.chk
```

In Makefile GHDL, keep adding more tests

```
all: tadd32.out pmul16_test.out part1.out
# these are the final result files

# <tab> not spaces, precede ghdl

tadd32.out: add32.vhdl tadd32.vhdl # HW4
ghdl -a --ieee=synopsys add32.vhdl
ghdl -a --ieee=synopsys tadd32.vhdl
ghdl -e --ieee=synopsys tadd32
ghdl -r --ieee=synopsys tadd32 --stop-time=160ns > tadd32.out

# blank line ends commands
```

then diff -iw tadd32.out tadd32.chkg

```
pmul16_test.out: pmul16.vhdl pmul16_test.vhdl # HW6
ghdl -a --ieee=synopsys pmul16.vhdl
ghdl -a --ieee=synopsys pmul16_test.vhdl
ghdl -e --ieee=synopsys pmul16_test
ghdl -r --ieee=synopsys pmul16_test --stop-time=8704ns > pmul16_test.out
```

then diff -iw pmul16\_test.out pmul16\_test.chkg

```
part1.out: part1.vhdl add32.vhdl bshift.vhdl divcas16.vhdl
pmul16.vhdl part1.chk part1.abs
ghdl -a --ieee=synopsys add32.vhdl
ghdl -a --ieee=synopsys bshift.vhdl
ghdl -a --ieee=synopsys divcas16.vhdl
ghdl -a --ieee=synopsys pmul16.vhdl
ghdl -a --ieee=synopsys part1.vhdl
ghdl -e --ieee=synopsys part1
ghdl -r --ieee=synopsys part1 --stop-time=280ns > pmul16_test.out
```

then diff -iw part1.out part1.chkg

```

3) Log onto linux.g1.umbc.edu to work on VHDL
cd vhdl
tcsh          # only for Cadence
source vhdl_cshrc # only for Cadence
make clean
touch part1.vhdl # not needed if you edit the file
make      # results come to screen gmake on Solaris, make on Linux
# "don't know how to make xxx.yyy" file missing or typo
# much output flying by on screen, or use:

make |& more # hit space for next page, enter for next line or:
make >& part1.prt # results, including error go to a file
# use editor to read file, you can search for "error"

make clean      # saves on quota and removes bad compilations

```

4) FIX THE FIRST ERROR !!!!  
 Yes, you can fix other errors also, but one error can cause  
 a cascading effect and produce many errors.

Don't panic when there was only one error, you fixed that,  
 then the next run you get 37 errors. The compiler has stages,  
 it stops on a stage if there is an error. Fixing that error  
 lets the compiler move to the next stage and check for other  
 types of errors. Go to step 3)

5) Don't give up. Don't make wild guesses. Do experiment with  
 one change at a time. You may actually have to read some  
 of the handouts :)

6) Your circuit compiles and simulates but the output is not  
 correct. Solution: find first difference, or add debug print.

Most circuits in this course have a print process. You can  
 easily add printout of more signals. Look for the existing  
 code that has 'write' and 'writeline' statements.  
 To print out some signal, xxx, after a 'writeline' statement add

```

write(my_line, string(" xxx=")); -- label printout
hwrite(my_line, xxx);           -- hex for long signals
write(my_line, string(" enb="));
write(my_line, enb);            -- bit for single values
writeline(output, my_line);    -- outputs line

```

7) You have a signal that seems to be wrong and you can not find  
 when it gets the wrong value. OK, create a new process to  
 print every change and when it occurs.

```

prtxxx: process (xxx)
  variable my_line : LINE; -- my_line needs to be defined
begin
  write(my_line, string("xxx="));
  write(my_line, xxx);           -- or hwrite for long signals
  write(my_line, string(" at="));
  write(my_line, now);           -- "now" is simulation time
  writeline(output, my_line);   -- outputs line
end process prtxxx;

```

When adding 'write' statements, you may need to add the  
 context clause in front of the enclosing design unit. e.g.

```

library STD;
use STD.textio.all; -- defines LINE, writeline, etc.
library IEEE;
use IEEE.std_logic_1164.all;
use IEEE.std_logic_textio.all; -- defines write on std_logic (_vector)

```

8) Read your code.

Every identifier must be declared before it is used.

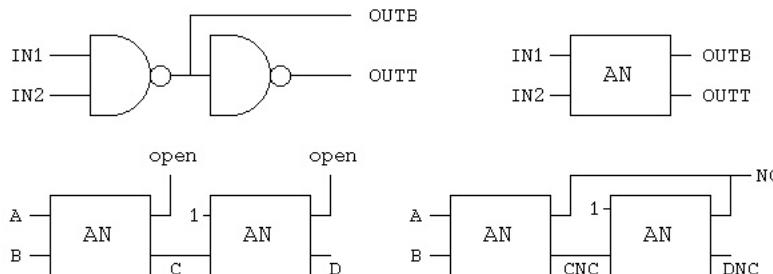
Every signal MUST be set exactly once, e.g.

```
xxx <= a;
xxx <= b; -- somewhere else, BAD !
-- all hardware runs all the time
-- the ordering of statements does not matter
```

```
a0: fadd port map(a(0), b(0), cin , sum(0), c(0));
a1: fadd port map(a(1), b(1), c(0), sum(1), c(0));
#####      BAD !
```

Signals must match in type and size. An error having "shape mismatch" means incompatible size. You can not put one bit into a 32 bit signal nor 32 bits into a one bit signal. "...type... error" Are you putting an integer into a std\_logic? You can not put an identifier of type std\_logic into std\_logic\_vector. a(31 downto 28) is of type std\_logic\_vector, a(31) is of type std\_logic.

Everywhere a specific signal name is used, these points are wired together. For VHDL simulation purposes, all points on a wire always have exactly the same value. Zero propagation delay through a wire. Be careful what you wire together. Use the VHDL reserved word 'open' for open circuits rather than NC for no connection.



Use 'open'  
Bad design,  
OUTB should have driver

Do not connect unused outputs  
it may change the logic

ncsim: 05.40-s011: (c) Copyright 1995-2005 Cadence Design Systems, Inc.

ncsim> run 7 ns

```
A= 1  B= 1  C= U  D= U  CNC= U  DNC= U  NC= U at time 0 ns
A= 1  B= 1  C= U  D= U  CNC= U  DNC= U  NC= U at time 1 ns
A= 1  B= 1  C= 1  D= U  CNC= U  DNC= U  NC= U at time 2 ns
A= 1  B= 1  C= 1  D= U  CNC= U  DNC= U  NC= U at time 3 ns
A= 1  B= 1  C= 1  D= 1  CNC= U  DNC= U  NC= U at time 4 ns
A= 1  B= 1  C= 1  D= 1  CNC= U  DNC= U  NC= U at time 5 ns
A= 1  B= 1  C= 1  D= 1  CNC= U  DNC= U  NC= U at time 6 ns
Ran until 7 NS + 0
```

ncsim> exit

!!! !!! never set due to connection

```
-- use_open.vhd1
library IEEE;
use IEEE.std_logic_1164.all;

entity AN is
port(IN1 : in std_logic;
      IN2 : in std_logic;
      OUTB : inout std_logic; -- because used internally, bad design
      OUTT : out std_logic);
end entity AN;
```

```

architecture circuits of AN is
begin -- circuits
    OUTB <= IN1 nand IN2 after 1 ns;
    OUTT <= not OUTB      after 1 ns;
end architecture circuits; -- of AN

library IEEE;
use IEEE.std_logic_1164.all;
use STD.textio.all;
use IEEE.std_logic_textio.all;

entity use_open is
end entity use_open;

architecture circuits of use_open is
    signal A : std_logic := '1';
    signal B : std_logic := '1';
    signal C, CNC : std_logic;
    signal D, DNC : std_logic;
    signal NC : std_logic := '1'; -- for no connection or tied off
begin
    my_print : process is
        variable my_line : line;
    begin
        write(my_line, string('A= "'));
        write(my_line, A);
        write(my_line, string(' B= "'));
        write(my_line, B);
        write(my_line, string(' C= "'));
        write(my_line, C);
        write(my_line, string(' D= "'));
        write(my_line, D);
        write(my_line, string(' CNC= "'));
        write(my_line, CNC);
        write(my_line, string(' DNC= "'));
        write(my_line, DNC);
        write(my_line, string(' NC= "'));
        write(my_line, NC);
        write(my_line, string(' at time "'));
        write(my_line, now);
        writeline(output, my_line);
        wait for 1 ns;
    end process my_print;

n01: entity WORK.AN port map(A, B, open, C);
n02: entity WORK.AN port map('1', C, open, D);
n03: entity WORK.AN port map(A, B, NC, CNC);
n04: entity WORK.AN port map('1', CNC, NC, DNC);

end architecture circuits; -- of use_open

```

#### Truth tables using type std\_logic

#### t\_table.vhdl

Now, some Cadence VHDL error messages.

```
-- error.vhdl  demonstrate VHDL compiler error messages
```

```

library IEEE;
use IEEE.std_logic_1164.all;

entity AN is
    port(IN1 : in std_logic;
         IN2 : in std_logic;
         OUTB : inout std_logic; -- because used internally
         OUTT : out std_logic););
end entity AN;

architecture circuits of AN is
    signal aaa : std_logic;
begin -- circuits

```

```

OUTB <= aa and IN1 and IN2 after 1 ns;
OUTT <= not OUTB      after 1 ns;
end architecture circuits; -- of AN

tcsh          # only Cadence
source vhdl_cshrc # only Cadence similar error messages from GHDL
ncsim -v93 error.vhdl >& error.out # note "|" and line numbers

ncvhdl: 05.40-s011: (c) Copyright 1995-2005 Cadence Design Systems, Inc.
        OUTT : out std_logic;);

ncvhdl_p: *E,PORNKW (error.vhdl,10|28): identifier expected.
        OUTT : out std_logic;);

ncvhdl_p: *E,MISCOL (error.vhdl,10|28): expecting a colon (':') 87[4.3.3] 93[4.3.2].
        OUTT : out std_logic;);

ncvhdl_p: *E,PORNKW (error.vhdl,10|31): identifier expected.
        OUTT : out std_logic;);

ncvhdl_p: *E,MISCOL (error.vhdl,10|31): expecting a colon (':') 87[4.3.3] 93[4.3.2].
end entity AN;

ncvhdl_p: *E,EXPRIS (error.vhdl,11|13): expecting the reserved word 'IS' [1.1].
        OUTB <= aa and IN1 and IN2 after 1 ns;

ncvhdl_p: *E,IDENTU (error.vhdl,16|11): identifier (AA) is not declared [10.3].

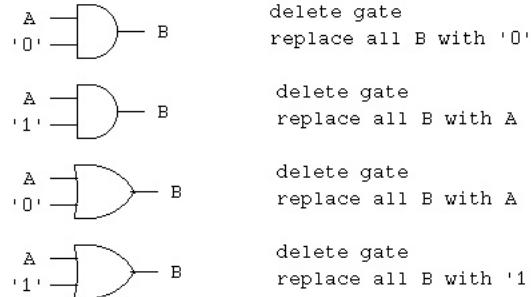
```

Now you are ready to tackle [Homework 6](#)

To simplify

Simplify circuit by eliminating gates

before                  after



Repeat until no gates eliminated

[sqrt examples of simplify](#)

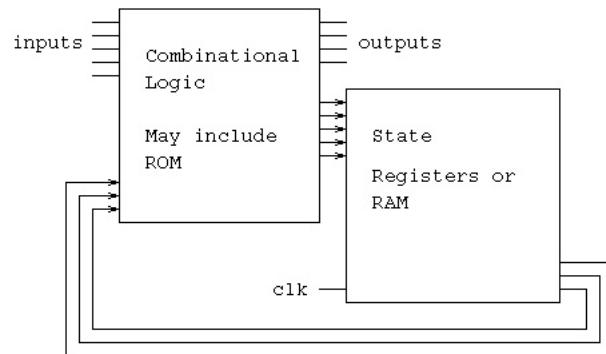
### Lecture 13, Microprogramming - review

Review is paper handout.  
 Following, microcontrollers, microprogramming and 64-bit.

A microcontroller may be a very small and inexpensive device.  
 The basic parts are Combinational Logic, logic gates, and some

type of storage, Sequential Logic.

## Simplified Microcontroller



Combinational Logic can not have loops.

Logic Gates is another name for Combinational logic

All types of storage must be clocked.

For students who have taken CMSC 451, this is the classic Deterministic Automata, a Finite State Machine.

A microcontroller may have Read Only Memory, ROM, that contains a microprogram to run the microcontroller. Micro assemblers and micro compilers may be used to generate the microprogram. The microprogram is manufactured in the microcontroller.

Micro instructions may be very long, 40 to 64 bits is common. Often there are bits to directly control multiplexors.

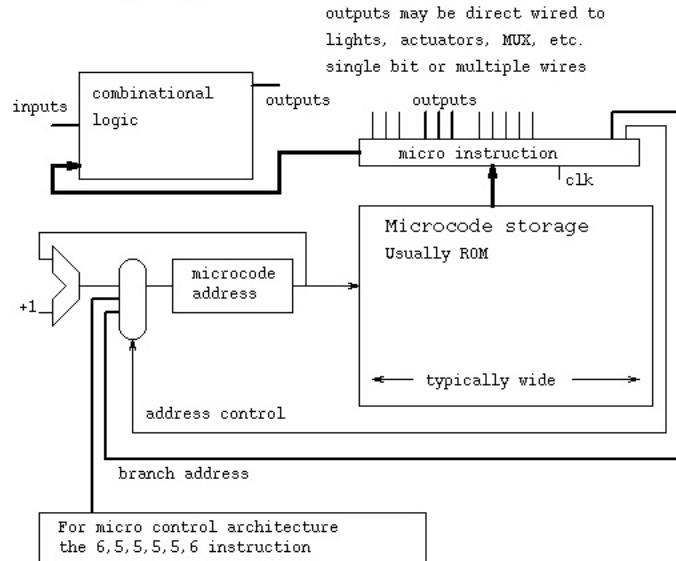
Often there are groups of bits to directly control other units such as the ALU.

There may be bits that go directly to outputs.

Every microinstruction may have a jump address.

The jump may be a conditional branch based on some state bits.

## Microprogrammed Microcontroller



The micro instruction typically has individual bits each output. There is also a branch address, optional, in each instruction.

From Wikipedia [wiki/Microcode](#)

Terminology: Combinational Logic is just gates. No storage.  
 Sequential Logic has storage, flipflop(s) or register(s).  
 A flipflop or register holds the output until changed.  
 A flipflop or register will have a clock and data  
 will only be input to change state on a clock edge.  
 There may be a clear or set input that does not  
 need a clock signal, typically used to initialize  
 a logic circuit to a known state.

This lecture also covers 64-bit machines (If not covered earlier)

A 64-bit architecture, by definition, has 64-bit integer registers.  
 Many computers have had 64-bit IEEE floating point for many years.  
 The 64-bit machines have been around for a while as the Alpha and  
 PowerPC yet have become popular for the desktop with the Intel and  
 AMD 64-bit machines.

## Current 64-bit microprocessor architectures

64-bit microprocessor architectures (as of 2006) include:

- The DEC Alpha architecture (view [Digital Alpha timeline](#))
- Intel's IA-64 architecture (used in Intel's Itanium CPUs)
- AMD's AMD64 architecture, a 64-bit version of the x86 architecture (used in AMD's Athlon 64, Opteron, Sempron, and Turion 64 CPUs).
  - Intel now uses the same instruction set, calling it EM64T.
- SPARC architecture (64-bit as of SPARC V9)
  - Sun's UltraSPARC architecture
  - Fujitsu's SPARC64 architecture
- IBM's POWER architecture (64-bit as of POWER3 and RS64 variants)
- IBM/Motorola's PowerPC architecture (64-bit PowerPC 620 and PowerPC 970 variants)
- IBM's z/Architecture, used by IBM zSeries and System z9 mainframes, a 64-bit version of the ESA/390 architecture
- MIPS Technologies' MIPS IV, MIPS V, and MIPS64 architectures
- HP's PA-RISC family (64-bit as of PA-RISC 2.0)

Most 64-bit processor architectures can execute code for the 32-bit version of the architecture natively without any performance penalty.

This kind of support is commonly called *bisarch support* or more generally *multi-arch support*.

Software has been dragging well behind computer architecture.

The chaos started in 1979 with the following "choices."

The differences between the three models lies in the non-pointer data types. The table below details the data types for the above three data models and includes LP32 and ILP32 for comparison purposes.

Data Type	LP32	ILP32	ILP64	LLP64	LP64
<b>char</b>	8	8	8	8	8
<b>short</b>	16	16	16	16	16
<b>int32</b>			32		
<b>int</b>	16	32	64	32	32
<b>long</b>	32	32	64	32	64
<b>long long (int64)</b>				64	
<b>pointer</b>	32	32	64	64	64

When the width of one or more of the C data types is changed, applications may be affected in various ways. These effects fall into two main categories:

- Data objects, such as a structure, defined with one of the 64-bit data types will be different in size from those declared in an identical way on a 16 or 32-bit system.
- Common assumptions about the relationships between the fundamental data types may no longer be valid in a 64-bit data model. Applications which depend on those relationships often cease to work properly when compiled on a 64-bit platform. A typical assumption made by many application developers is that:

`sizeof(int) = sizeof(long) = sizeof(pointer)`

This relationship is not codified in any C programming standard, but it is valid for the ILP32 data model. However, it is not valid for two of the three 64-bit data models described above, nor is it valid for the LP32 data model.

The full whitepaper [www.unix.org/whitepapers/64bit.html](http://www.unix.org/whitepapers/64bit.html)

My desire is to have the compiler, linker and operating system be ILP64. All my code would work fine. I make no assumptions about word length. I use `sizeof(int)` `sizeof(size_t)` etc. when absolutely needed. On my 8GB computer I use a single array of over 4GB thus the subscripts must be 64-bit. The only option, I know of, for gcc is `-m64` and that just gives LP64. Yuk! I have to change my source code and use "long" everywhere in place of "int". If you get the idea that I am angry with the compiler vendors, you are correct!

Here are sample programs and output to test for 64-bit capability in gcc:

[Get sizeof on types and variables big.c](#)

[output from gcc -m64 big.c big.out](#)[malloc more than 4GB big\\_malloc.c](#)[output from big\\_malloc mac.out](#)

Newer Operating Systems and compilers

[Get sizeof on types and variables big12.c](#)[output from gcc big12.c big12.out](#)

The early 64-bit computers were:

[DEC Alpha](#)[DEC Alpha](#)[IBM PowerPC note 5 clocks, similar to project](#)**review for midterm, handout****Lecture 14, mid-term exam**

open book, open note, download, edit, submit

Students with email user name starting a b c d e f g h i  
download and edit midterm32a.doc[download midterm32a.doc](#)Students with email user name starting j k l m n o p q  
download and edit midterm32b.doc[download midterm32b.doc](#)Students with email user name starting r s t u v w x y z  
download and edit midterm32c.doc[download midterm32c.doc](#)Follow instructions in exam, edit, then  
submit cs411 midterm midterm32...

You can do the exam on linux.gl.umbc.edu in your directory

```
cp /afs/umbc.edu/users/s/q/squire/pub/download/midterm32?.doc .
libreoffice midterm32?.doc
submit cs411 midterm midterm32?.doc
rm midterm32?.doc
```

Before Exam:

Review HW2, HW3, HW4 (VHDL) and HWS

Review WEB Lecture Notes 1 through 13.

There are 10 types of people:

Those who know binary.

Those who do not know binary.

Teach your children to count in the computer age:

zero

one

two

three

four

Computer bits are numbered from the bottom

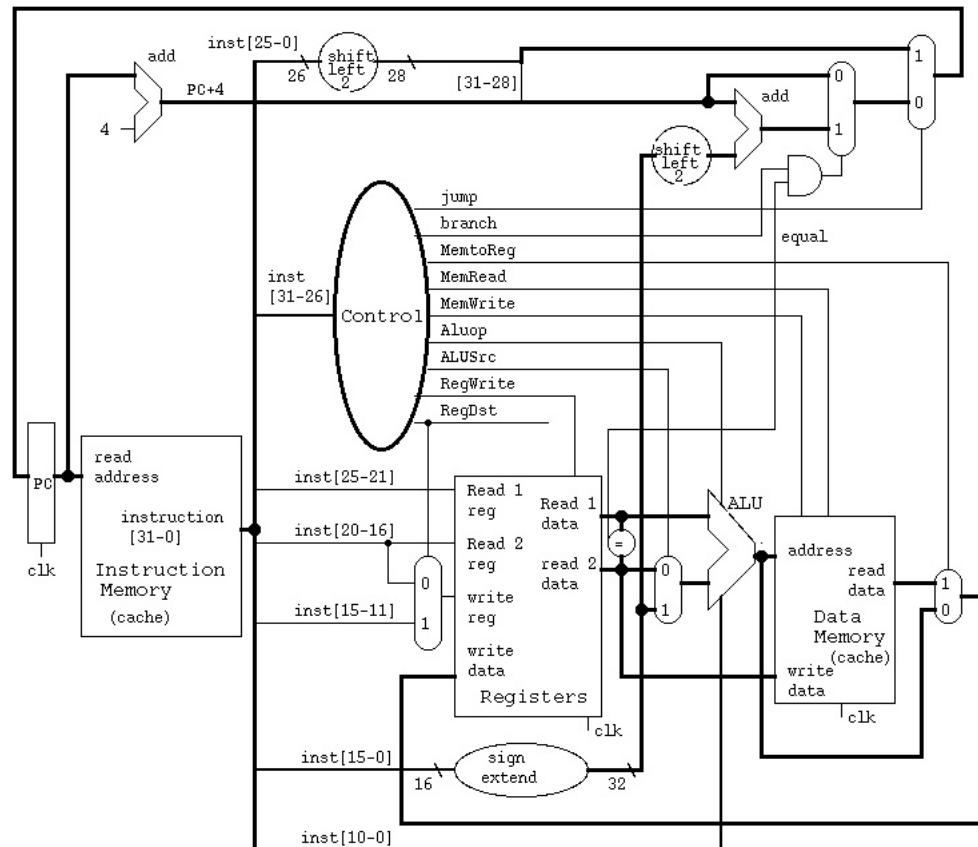
0	0	1	0	1	= 5
4	3	2	1	0	bit numbers (actually powers of 2)

## Lecture 15, Control Unit

We now start the second half of the semester, focusing on the five part project to simulate part of a real computer. Note that the hardware does not change. Only multiplexer control signals are needed to execute various instructions.

The first complete computer architecture is a single cycle design. On each clock cycle this computer executes one instruction. CPI=1 (The clock would be slow compared to pipeline computers in the next lecture.)

Signals are inputs to components on the left and outputs of components on the right. Wide lines are 32-bits. Narrow signals are one-bit unless otherwise indicated.



Every clock, we use the rising edge, the program counter register, PC, takes the 32 bit input from the left most signal on the diagram. The output of the PC is a memory address for an instruction.

The 32 bit instruction is "decoded" by routing various parts of the

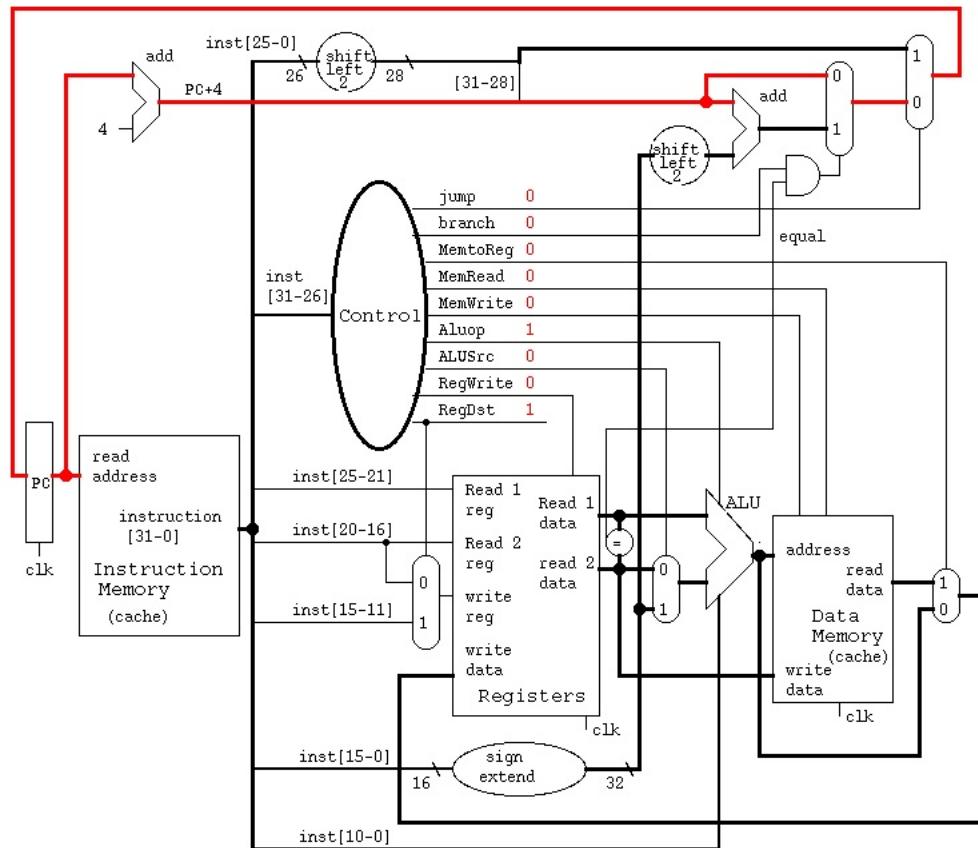
instruction to various places.  
 Bits 31 downto 26 of the instruction go to the control unit.  
 (The schematic of the control unit is shown below.)  
 Bits 10 downto 0 of the instruction go to the ALU, the shift count and  
 the ALU op code.  
 Bits 25 downto 21 are a register address that is read and the 32 bit  
 contents of that register are placed on read data 1.  
 Bits 20 downto 16 are a register address that is read and the 32 bit  
 contents of that register are placed on read data 2.  
 Bits 15 downto 11 are a register address that may be written with the  
 32 bit write data.  
 Bits 25 downto 0 go to the jump address computation.

The sequence of diagrams that follow will show the control signals  
 and the data paths for various instructions.

The bit patterns for our CMSC 411 machine are [cs411\\_OPCODES.TXT](#)  
[INSIDE THE ALU ENTITY](#).

The first instruction is the nop instruction.  
 This instruction shows the basic updating of the PC, while changing  
 no other registers or memory. All other instructions shown below,  
 except branch and jump , use this updating of the PC.

## nop



The PC plus 4 is the next sequential instruction address. The 32 bit instruction has four bytes. The bottom two bits of all instruction addresses are zero. The instructions are "aligned."

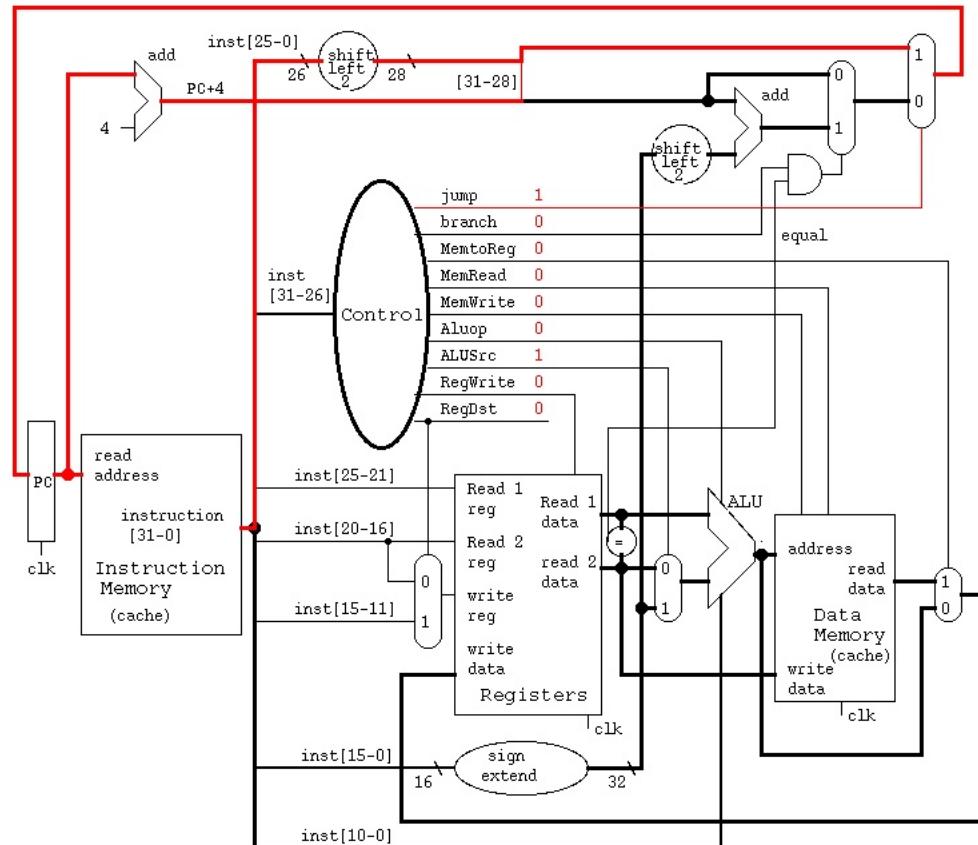
The critical control signals are:

```
jump    0
branch  0
MemWrite 0
RegWrite 0
```

The other control signals are shown for completeness.

The next instruction, jump, is just slightly more complex than nop.  
The bit pattern for jump in [cs411\\_OPCODES.TXT](#)

## jump



Note the wiring where instruction bits 25 downto 0 are shifter left two places. This provides a larger jump range and aligns the address on a quad byte boundary. The top four bits come from the incremented PC and the resulting 32 bit address is routed through the multiplexer back to the PC, ready for the next clock.

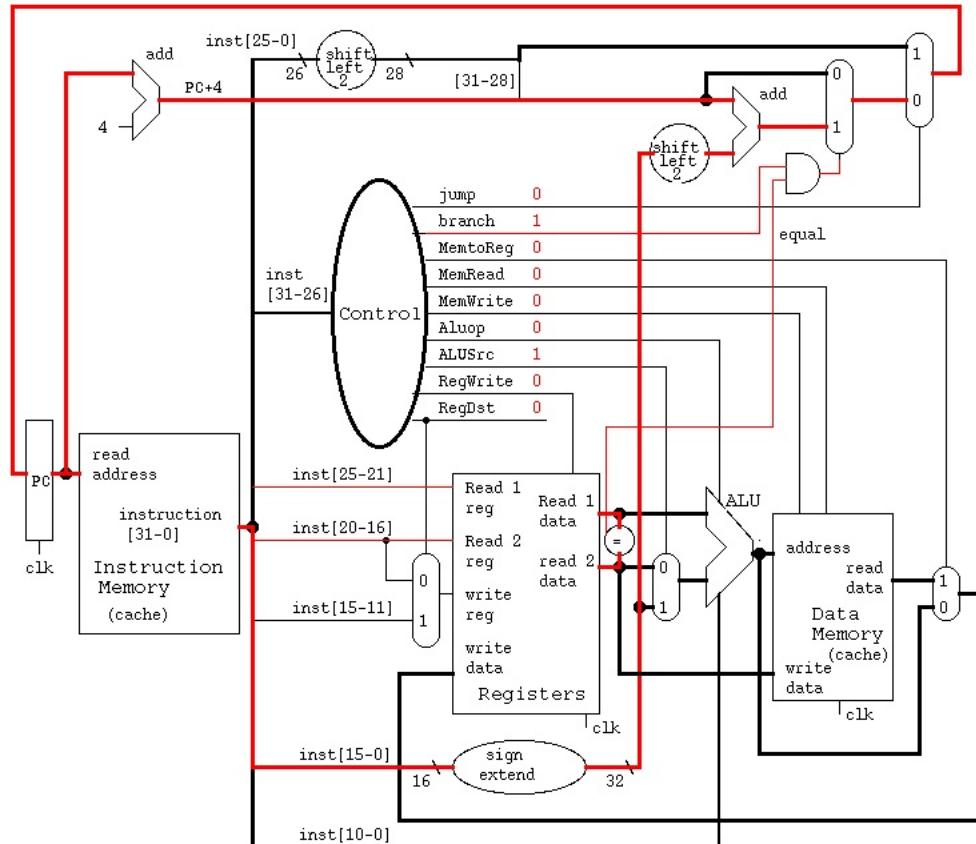
The critical control signals are:

```
jump 1
MemWrite 0
RegWrite 0
```

The other control signals are shown for completeness.

The next instruction, branch, uses the remainder of the upper schematic to compute a new instruction address relative to the incremented PC. Note that the assembler subtracts 4 from the branch address before generating the machine instruction.  
The bit pattern for beq in [cs411\\_OPCODES.TXT](#)

## branch



Note the equal comparator immediately next to the registers. This is the design we will use in the project because it provides better performance in the pipeline architecture. If the branch condition is not satisfied, the instruction becomes a nop. The branch condition for beq is that the contents of the registers are the same and a beq instruction is executing. Note the and gate driving the multiplexer.

The critical control signals are:  
jump 0

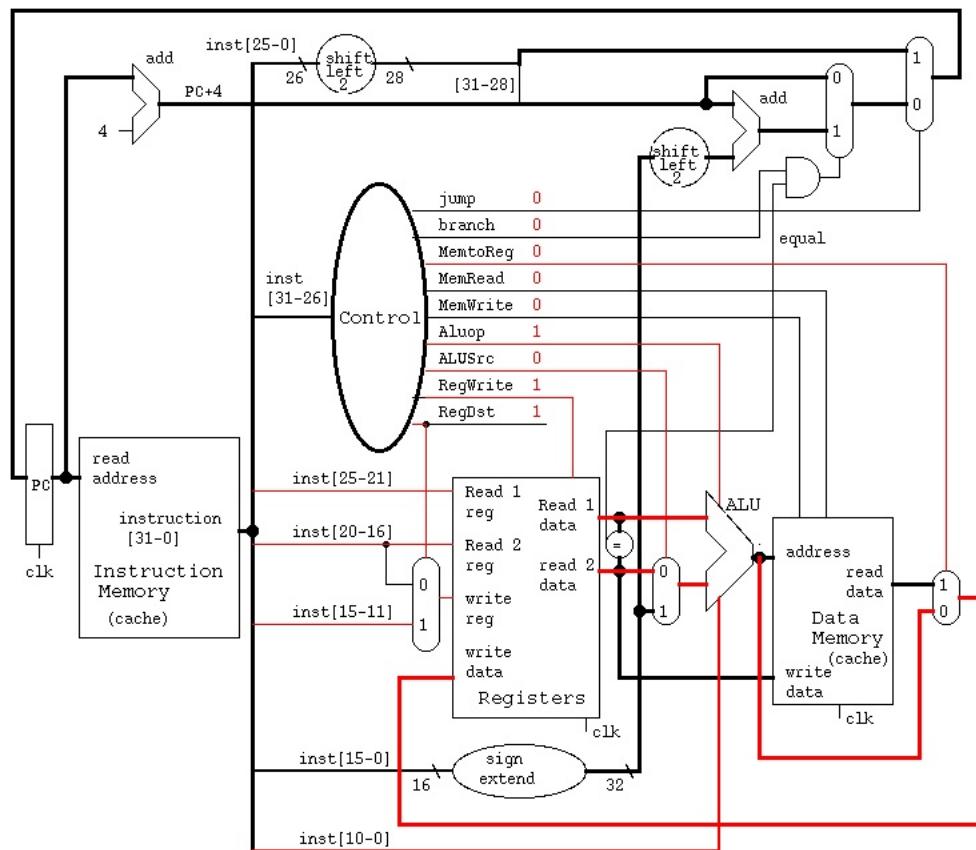
```

branch 1      and the equal comparison
MemWrite 0
RegWrite 0
The other control signals are shown for completeness.

```

The add instruction is shown with just the data paths and control paths for the instruction shown. The upper control to increment the PC is the same as shown for the nop instruction. The bit pattern for add in [cs411\\_opcodes.txt](#)

## add



The contents of two registers are combined in the ALU. The ALU op code in the instruction bits 5 downto 0 would have 100000 for add . Other instructions such as subtract, shift, and, etc follow the same data paths and control, executing the instruction coded in the instruction bits 5 downto 0. The output of the ALU is routed back to the registers and written on the falling edge of the clock, clk.

The critical control signals are:

```

jump 0
branch 0
MemtoReg 0

```

```

MemWrite 0
Aluop 1
ALUSrc 0
RegWrite 1
RegDst 1

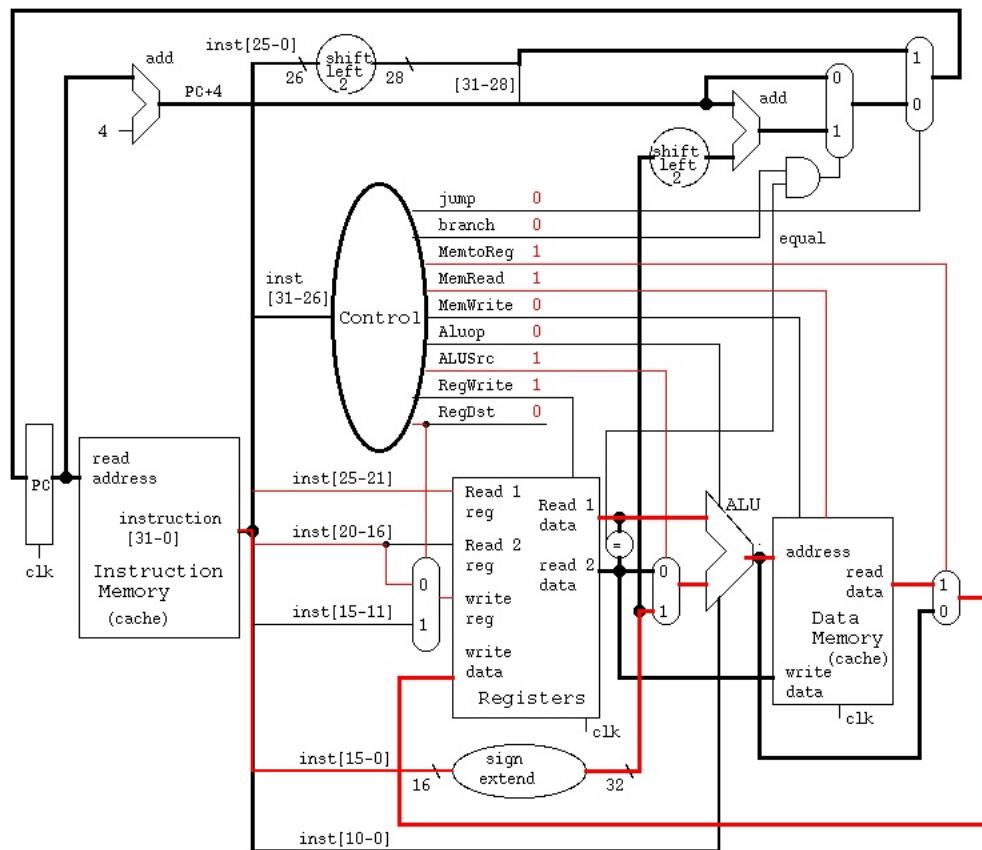
```

The other control signals are shown for completeness.

The load word, lw , instruction computes a memory address using the two's complement offset in the instruction bits 15 downto 0, sign extended to 32 bits and added to a register. The memory is read and the contents from memory is routed through the multiplexer into the destination register. The PC is incremented as shown in the nop instruction.

The bit pattern for lw in [cs411\\_opcodes.txt](#)

## load word, lw



The critical control signals are:

```

jump 0
branch 0
MemtoReg 1
MemRead 1
MemWrite 0

```

Aluop 0 the ALU performs an add when Aluop is zero

ALUSrc 1

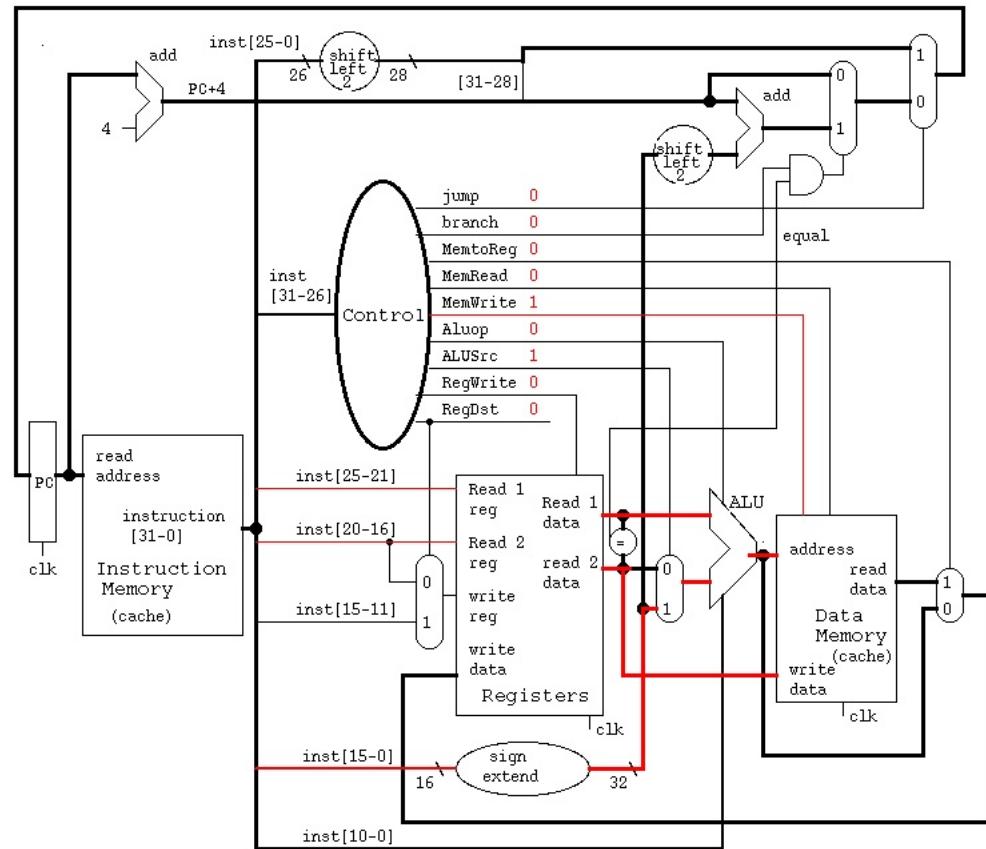
RegWrite 1

RegDst 0

The other control signals are shown for completeness.

The store word, sw , instruction computes a memory address using the two's complement offset in the instruction bits 15 downto 0, sign extended to 32 bits and added to a register. The read data 2 is stored in memory. The PC is incremented as shown in the nop instruction. The bit pattern for sw in [cs411\\_OPCODES.TXT](#)

## store word, sw



Note the data path around the ALU into the write data input to the memory

The critical control signals are:

jump 0

branch 0

MemRead 0

MemWrite 1

Aluop 0 the ALU performs an add when Aluop is zero

ALUSrc 1

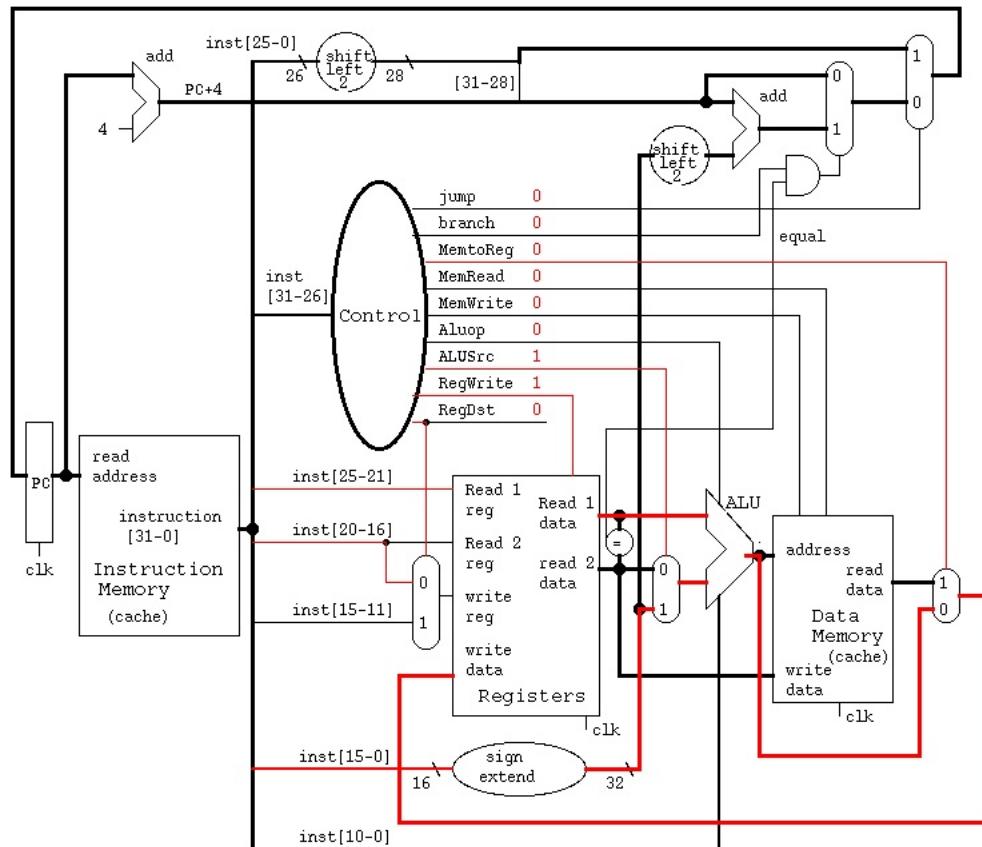
RegWrite 0

The other control signals are shown for completeness.

The add immediate, addi , instruction adds the two's complement bits 15 downto 0 of the instruction to a register and places the sum into the destination register. The PC is incremented as shown in the nop instruction.

The bit pattern for addi in [cs411\\_OPCODES.TXT](#)

## add immediate, addi



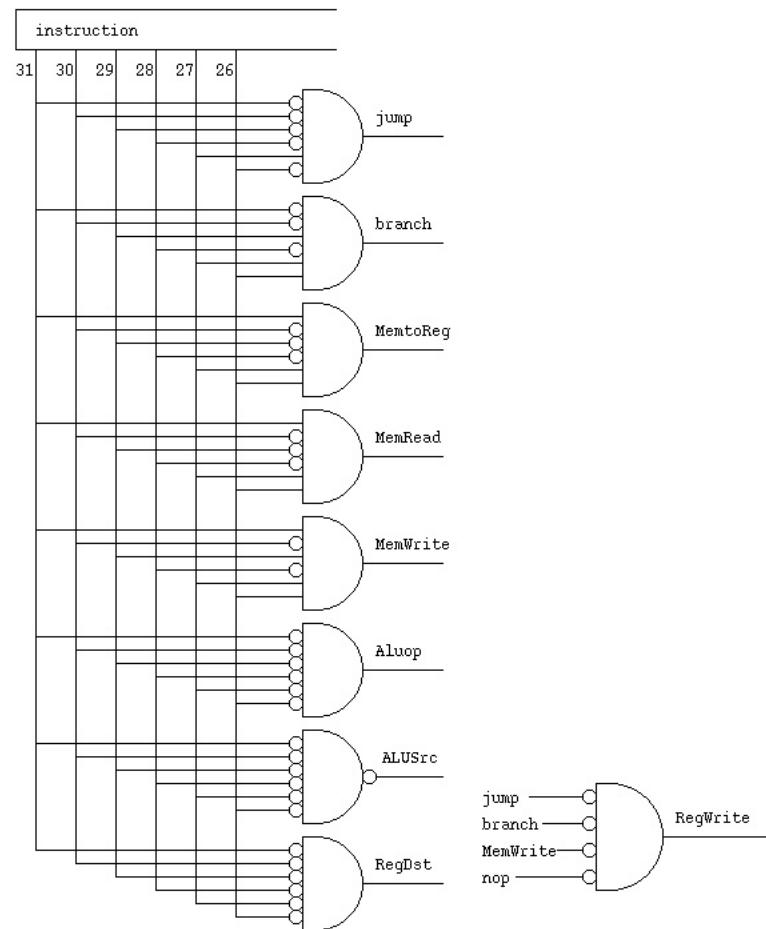
The critical control signals are:

```

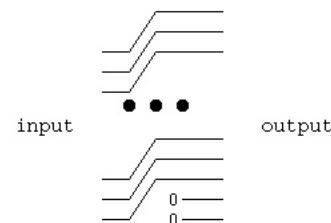
jump    0
branch  0
Memento 0
MemWrite 0
AluOp   0      the ALU performs an add when AluOp is zero
ALUSrc  1
RegWrite 1
RegDst  0
  
```

The other control signals are shown for completeness.

The control schematic for some specific instructions, possibly not this semester, for the one cycle architecture, is:

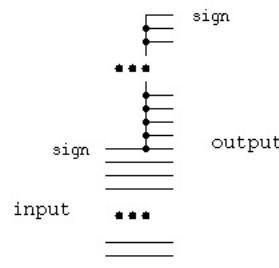


The shift left 2 circuit is just bent wires.  
The VHDL is   output <= input(29 downto 0) & "00";



The sign extend circuit is just wiring. The input is a 16 bit  
twos complement word and outputs a 32 bit twos complement word.

The VHDL is    output(15 downto 0) <= input;  
                    output(31 downto 16) <= (others => input(15));



cs411\_opcodes.txt different from Computer Organization and Design 1/8/2020

**rd** is register destination, the result, general register 1 through 31  
**rs** is the first register, A, source, general register 0 through 31  
**rt** is the second register, B, source, general register 0 through 31

--val--- generally a 16 bit number that gets sign extended  
--addr--- a 16 bit address, gets sign extended and added to (rx)  
"i" is generally immediate, operand value is in the instruction

Opcode    Operands    Machine code format  
                               6    5    5    5    5    6    number of bits in field

### Definitions:

**nop** no operation, no programmer visible registers or memory are changed, except  $PC \leq PC+4$

j adr bits 0 through 25 of the instruction are inserted into PC(27:2)

probably should zero bits PC(1:0) but should be zero already

**lw r,adr(x)** load word into register r from memory location (register x plus sign extended adr field)

**sw b,adr(x)** store word from register b into memory location (register x plus sign extended adr field)

**beq a,b,adr** branch on equal, if the contents of register a are equal to the contents of register b, add the, shifted by two, sign extended adr to the PC (The PC will have 4 added by then)

**lwim r,val(x)** add immediate, the contents of register x is added to the sign extended value and the result put into register r

**addi r,val(x)** add immediate, the contents of register x is added to the sign extended value and the result is added to register r

**add r,a,b** add register a to register b and put result into register r

**sub r,a,b** subtract register b from register a and put result into register r

**mul r,a,b** multiply register a by register b and put result into register r

**div r,a,b** divide register a by register b and put result into register r

**and r,a,b** and register a to register b and put result into register r

**or r,a,b** or register a to register b and put result into register r

**srl r,b,s** shift the contents of register b by s places right and put result in register r

**sll r,b,s** shift the contents of register b by s places left and put result in register r

**cmpl r,b** one's complement of register b goes into register r

Also: no instructions are to have side effects or additional "features"

last updated 1/8/2020 (slight difference in opcodes from previous semesters)

## Lecture 16, Pipelining 1

First, a few definitions:

**Pipelining** : Multiple instructions being executed, each in a different stage of their execution. A form of parallelism.

**Super Pipelining** : Advertising term, just longer pipelines.

**Super Scalar** : Having multiple ALU's. There may be a mix of some integer ALU's and some Floating Point ALU's.

**Multiple Issue** : Starting a few instructions every clock.  
The CPI can be a fraction, 4 issue gives a CPI of 1/4 .

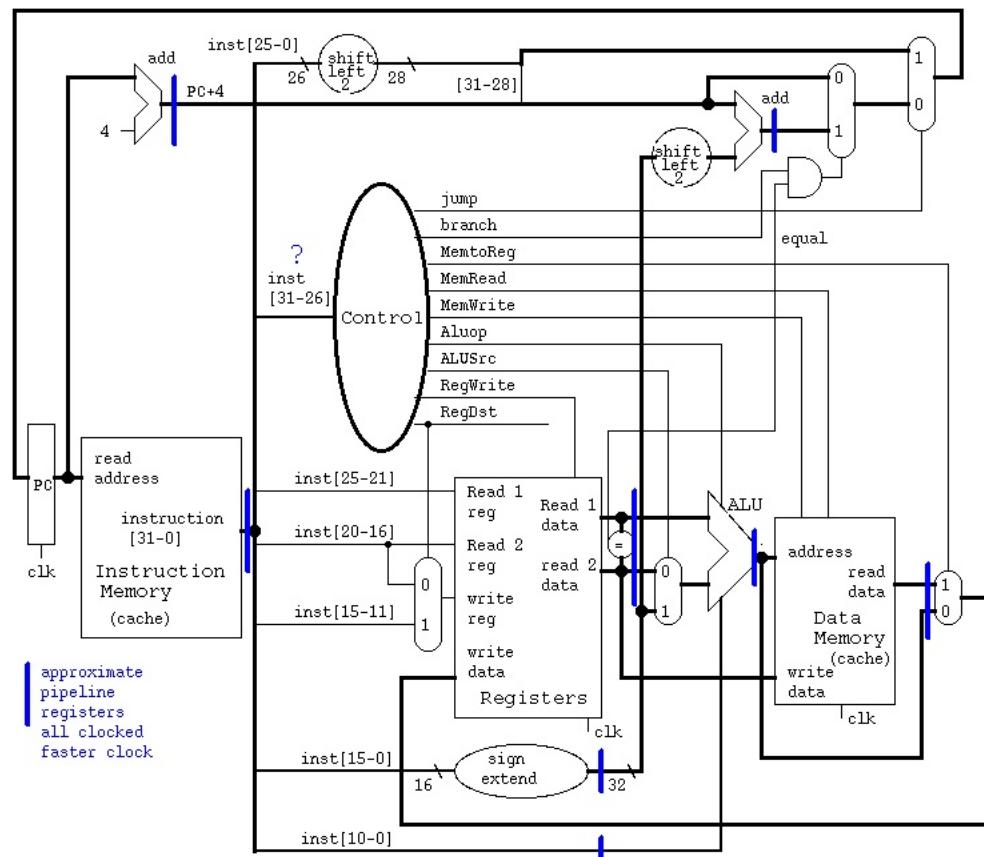
**Dynamic Pipeline** : This may include all of the above and also can reorder instructions, use data forwarding and hazard workarounds.

**Pipeline Stages** : For our study of the MIPS architecture,  
 IF Instruction Fetch stage  
 ID Instruction Decode stage  
 EX Execute stage  
 MEM Memory access stage  
 WB Write Back into register stage

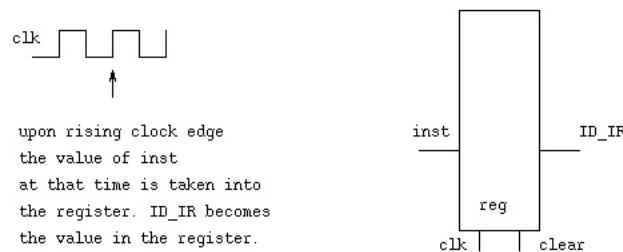
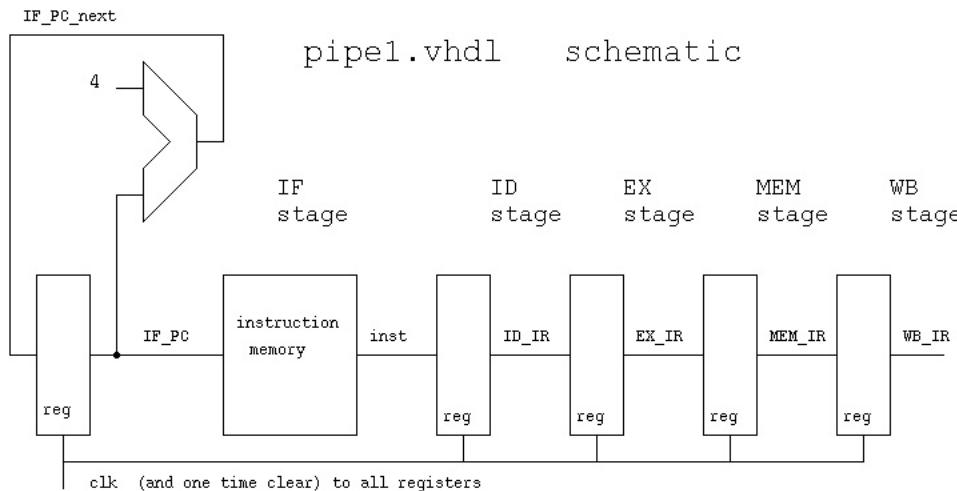
**Hyper anything** : Generally advertising terminology.

Consider the single cycle machine in the previous lecture.

The goal is to speed up the execution of programs, long sequences of instructions. Keeping the same manufacturing technology, we can look at speeding up the clock by inserting clocked registers at key points. Note the placement of blue registers that tries to minimize the gate delay time between any pair of registers. Thus, allowing a faster clock.



This is called approximate because some additional design must be performed, mostly on "control", that must now be distributed. The next step in the design, for our project, is to pass the instruction along the pipeline and keep the design of each stage of the pipeline simple, just driven by the instruction presently in that stage.

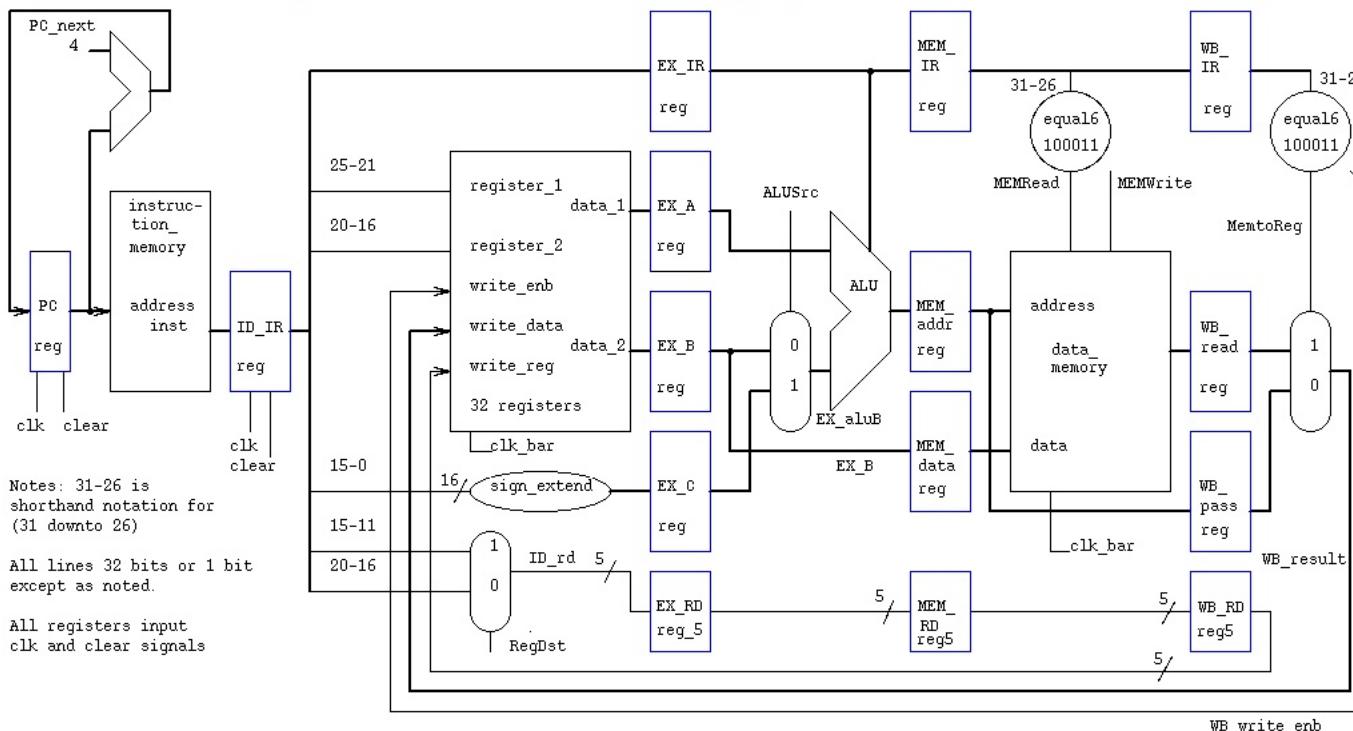


[pipe1.vhd1](#) implementation moves instruction  
note clock and reset generation  
look at register behavioral implementation  
instruction memory is preloaded

[pipe1.out](#) just numbers used for demonstration

## Pipelined Architecture with distributed control

## Pipelined design without branch and jump



[pipe2.vhd1](#) note additional entities  
 equal6 for easy decoding  
 data memory behavioral implementation

[pipe2.out](#) instructions move through stages

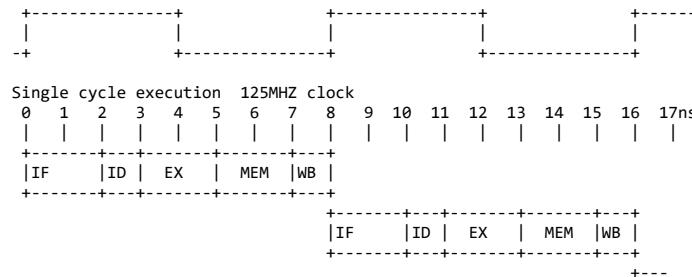
## Timing analysis

Consider four instructions being executed.

First on the single cycle architecture, needing 8ns per instruction.

The time for each part of the circuit is shown.

The clock would be:



| IF ... 24ns  
+---

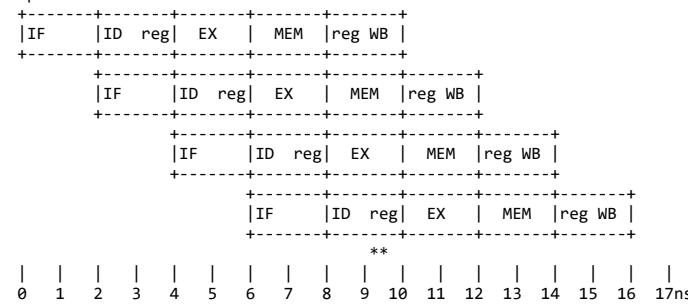
... 32n:

The four instructions finished in 32ns.  
An instruction started every 8ns.  
An instruction finished every 8ns.

Now, the pipelined architecture has the clock determined by the slowest part between clocked registers. Typically, the ALU. Thus use the same ALU time as above, the clock would be:

+---+ +---+ +---+ +---+ +---+ +---+ +---+ +---+ +---+

Pipelined Execution 500MHZ clock    \*\*



The four instructions finished in 16ns. (But, the speedup is not 2)

An instruction started every 2ns.

An instruction finished every 2ns. Thus, the speedup is  $8\text{ns}/2\text{ns} = 4$ .

Since an instruction finishes every 2ns for the pipelined architecture and every 8ns for the single cycle architecture, the speedup will be  $8ns/2ns = 4$ . The speedup would change with various numbers of instructions if the total time was used. Thus, the time between the start or end of adjacent instructions is used in computing speedup.

Note the \*\* above in the pipeline. The first of the four instructions may load a value in a register. This load takes place on the falling edge of the clock. The fourth instruction is the earliest instruction that could use the register loaded by the first instruction. The use of the register comes after the rising edge of the clock. Thus use of both halves of the clock cycle is important to this architecture and to many modern computer architectures.

Remember, every stage of the pipeline must be the same time duration. The system clock is used by all pipeline registers. The slowest stage determines this time duration and thus determines the maximum clock frequency.

The worse case delay that does not happen often because of optimizing compilers, is a load word, `lw`, instruction followed by an instruction that needs the value just loaded. The sequence of instructions, for this unoptimized architecture, would be:

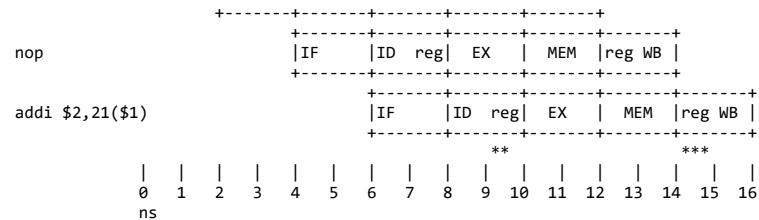
```
lw $1,val($0) load the 32 bit value at location val into register 1
nop
nop
addi $2,21($1) register 1 is available, add 21 and put result into reg 2
```

As can be seen in the pipelined timing below, `lw` would load register 1 by 9ns and register 1 would be used by `addi` by 10ns (\*\*). The actual add would be finished by 12 ns and register 2 updated sum by 15 ns (\*\*\*)

```

+-----+-----+-----+-----+-----+
lw $1, val($0) | IF      | ID     | reg   | EX    | MEM   | reg   WB |
+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+
nop          | IF      | ID     | reg   | EX    | MEM   | reg
+-----+-----+-----+-----+-----+

```



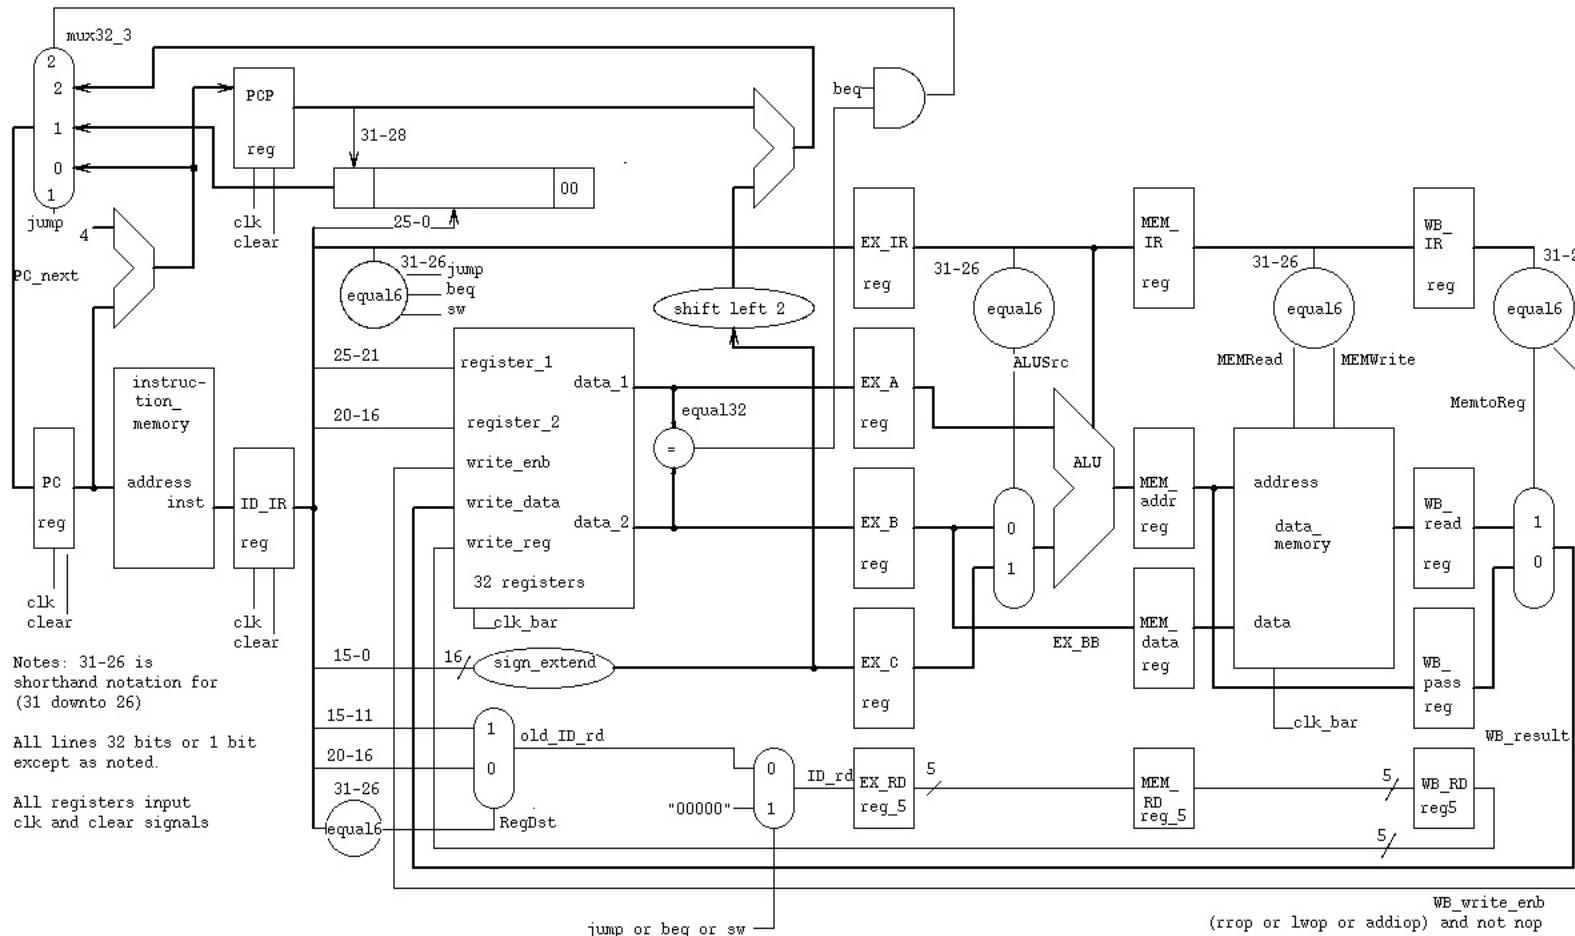
It is interesting to note some similarity to an IBM Power PC that came a few years after the MIPS R3000 architecture that is similar to the above design.

[IBM Power PC stages and clock usage](#)

[new IBM Power PC](#)  
Shipped 2012 at 5.5Ghz

## Lecture 17, Pipelining 2

The pipeline for this course with branch and jump optimized:



Note the three input mux replacing two mux in previous lecture.

Note the distributed control using the equal6 entity:

```
eq6j: entity WORK.equal6 port map(ID_IR(31 downto 26), "000010", jump);
      jumpaddr <= PCP(31 downto 28) & ID_IR(25 downto 0) & "00";
```

[cs411\\_OPCODES.txt](#) look at jump

In a later lecture, we will cover data forwarding to avoid nop's in arithmetic and automatic stall to avoid putting all nop's in source code.

For the basic machine above, we have the timing shown here.

The branch slot, programming to avoid delays (filling in nop's):

Note: beq and jump always execute the next physical instruction.  
This is called the "delayed branch slot", important for HW7.

```
if(a==b) x=3; /* simple C code */
else x=4;
y=5;
```

```
lw $1,a      # possible unoptimized assembly language
lw $2,b      # no ($0) shown on memory access
nop          # wait for b to get into register 2
```

```

nop      # wait for b to get into register 2
beq $1,$2,lab1
nop      # branch slot, always executed *****
addi $1,4   # else part
nop      # wait for 4 to get into register 1
nop      # wait for 4 to get into register 1
sw  $1,x    # x=4;
j   lab2
nop      # branch slot, always executed *****
lab1: addi $1,3   # true part
nop      # wait for 3 to get into register 1
nop      # wait for 3 to get into register 1
sw  $1,x    # x=3;
lab2: addi $1,5   # after if-else, always execute
nop      # wait for 5 to get into register 1
nop      # wait for 5 to get into register 1
sw  $1,y    # y=5;

```

Unoptimized, 20 instructions.

Now, a smart compiler would produce the optimized code:

```

lw  $1,a      # possible unoptimized assembly language
lw  $2,b      # no ($0) shown on memory access
addi $4,4     # for else part later
addi $3,3     # for true part later
beq $1,$2,lab1
addi $5,5     # branch slot, always executed, for after if-else
j   lab2
sw  $4,x    # x=4; in branch slot, always executed !! after jump
lab1: sw  $3,x    # x=3;
lab2: sw  $5,y    # y=5;

```

Optimized, 10 instructions.

The pipeline stage diagram for a==b true is:

	1	2	3	4	5	6	7	8	9	10	11	12	clock
lw \$1,a	IF	ID	EX	MM	WB								
lw \$2,b	IF	ID	EX	MM	WB								
addi \$4,4	IF	ID	EX	MM	WB								
addi \$3,3	IF	ID	EX	MM	WB								
beq \$1,\$2,L1	IF	ID	EX	MM	WB	assume equal, branch to L1							
addi \$5,5	IF	ID	EX	MM	WB	delayed branch slot							
j L2													
sw \$4,x													
L1:sw \$3,x	IF	ID	EX	MM	WB								
L2:sw \$5,y	IF	ID	EX	MM	WB								
	1	2	3	4	5	6	7	8	9	10	11	12	

The pipeline stage diagram for a==b false is:

	1	2	3	4	5	6	7	8	9	10	11	12	13	clock
lw \$1,a	IF	ID	EX	MM	WB									
lw \$2,b	IF	ID	EX	MM	WB									
addi \$4,4	IF	ID	EX	MM	WB									
addi \$3,3	IF	ID	EX	MM	WB									
beq \$1,\$2,L1	IF	ID	EX	MM	WB	assume not equal								
addi \$5,5	IF	ID	EX	MM	WB									
j L2						IF ID EX MM WB jumps to L2								
sw \$4,x						IF ID EX MM WB								
L1:sw \$3,x	IF	ID	EX	MM	WB									
L2:sw \$5,y	IF	ID	EX	MM	WB									
	1	2	3	4	5	6	7	8	9	10	11	12	13	

```

if(a==b) x=3; /* simple C code */
else
y=5;

```

Renaming when there are extra registers that the programmer can not assess (diagram in Alpha below) with multiple units there can be multiple issue (parallel execution of instructions)

The architecture sees the binary instructions from the following:

```

lw $1,a
lw $2,b
nop
sll $3,$1,8
sll $6,$2,8
add $9,$1,$2
sw $3,c
sw $6,d
sw $9,e
lw $1,aa
lw $2,bb
nop
sll $3,$1,8
sll $6,$2,8
add $9,$1,$2
sw $3,cc
sw $6,dd
sw $9,ee

```

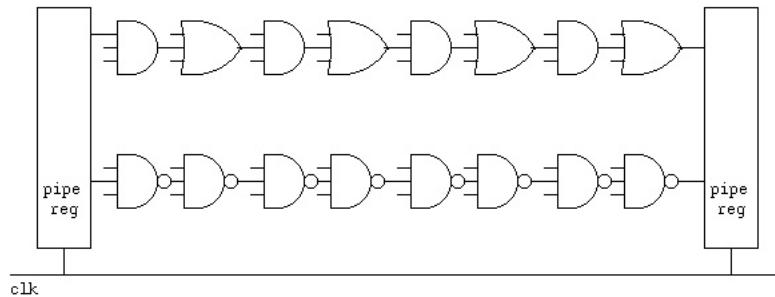
Two ALU's, each with their own pipelines, multiple issue, register renaming:  
The architecture executes two instruction streams in parallel.  
(Assume only 32 user programmable registers, 80 registers in hardware.)

lw \$1,a	lw \$41,aa
lw \$2,b	lw \$42,bb
nop	nop
sll \$3,\$1,8	sll \$43,\$41,8
sll \$6,\$2,8	sll \$46,\$42,8
add \$9,\$1,\$2	add \$49,\$41,\$42
sw \$3,c	sw \$43,cc
sw \$6,d	sw \$46,dd
sw \$9,e	sw \$49,ee

Out of order execution to avoid delays. As seen in the first example,  
changing the order of execution without changing the semantics of the  
program can achieve faster execution.

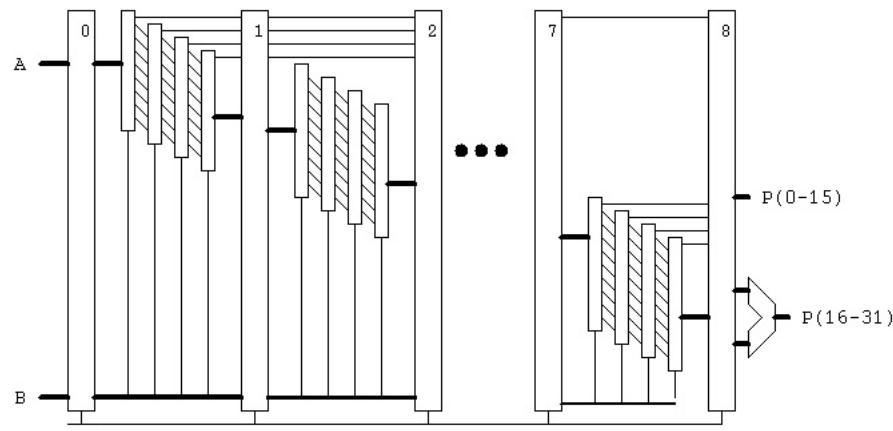
There can be multiple issue when there are multiple arithmetic and  
other units. This will require significant hardware to detect the  
amount of out of order instructions that can be issued each clock.

Now, hardware can also be pipelined, for example a parallel multiplier.  
Suppose we need to have at most 8 gate delays between pipeline  
registers.



Note that any and-or-not logic can be converted to use only nand gates  
or only nor gates. Thus, two level logic can have two gate delays.

We can build each multiplier stage with two gate delays. Thus we can  
have only four multiplier stages then a pipeline register. Using a  
carry save parallel 32-bit by 32-bit multiplier we need 32 stages, and  
thus eight pipeline stages plus one extra stage for the final adder.



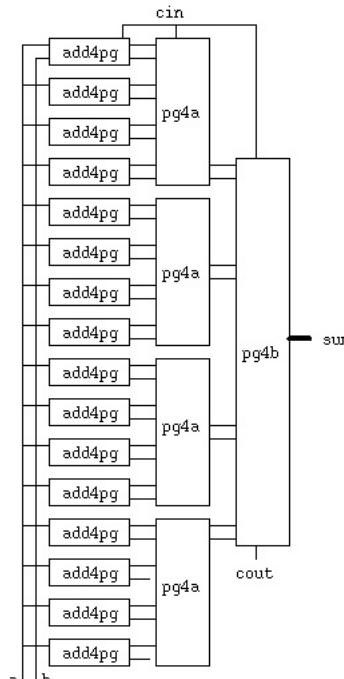
$P = A * B$     eight pipeline stages    32 carry save blocks

### Parallel Pipelined Multiplier

Note that a multiply can be started every clock. Thus a multiply can be finished every clock. The speedup including the last adder stage is 9 as shown in:

[pipemul\\_test.vhd1](#)  
[pipemul\\_test.out](#)  
[pipemul.vhd1](#)

A 64-bit PG adder may be built with eight or less gate delays.  
 The signals a, b and sum are 64 bits. See [add64.vhd1](#) for details.



64-bit PG adder

[add64.vhd1](#)

Any combinational logic can be performed in two levels with "and" gates feeding "or" gates, assuming complementation time can be ignored.  
Some designers may use diagrams but I wrote a Quine McClusky minimization program that computes the two level and-or-not VHDL statement for combinational logic.

[quine\\_mcclusky.c](#) logic minimization

[eqn4.dat](#) input data

[eqn4.out](#) both VHDL and Verilog output

[there are  \$2^{2^N}\$  possible functions of N bits](#)

Not as practical, I wrote a Myhill minimization of a finite state machine, a Deterministic Finite Automata, that inputs a state transition table and outputs the minimum state equivalent machine. "Not as practical" because the design of sequential logic should be understandable. The minimized machine's function is typically unrecognizable.

[myhill.cpp](#) state minimization  
[initial.dfa](#) input data  
[myhill.dfa](#) minimized output

A reasonably complete architecture description for the Alpha showing the pipeline is:

[basic Alpha](#)  
[more complete Alpha](#)

The "Cell" chip has unique architecture:

[Cell architecture](#)

Some technical data on Intel Core Duo (With some advertising.)

[Core Duo all on WEB](#)

From Intel, with lots of advertising:

power is proportional to capacitance \* voltage^2 \* frequency, page 7.

[tech overview](#)

[whitepaper](#)

[Intel quad core demonstrated](#)

[AMD quad core](#)

By 2010 AMD had a 12-core available and Intel had a 8-core available.

[and 24 core and 48 core AMD](#)

[IBM Power6 at 4.7GHz clock speed](#)

Intel I7 920 Nehalem 2.66GHz not quad	\$279.99
Intel I7 940 Nehalem 2.93GHz quad core	\$569.99
Intel I7 965 Nehalem 3.20GHz quad core	\$999.99
Prices vary with time, <a href="#">NewEgg.com search Intel I7</a>	

Motherboard [Asus products-motherboards-intel i7](#)

Intel socket 1366

[Supermicro.com motherboards, 12-core](#)

local, bad formatting, in case web page goes away. Good history.

[Core Duo 1](#)

[Core Duo 2](#)

[Core Duo 3](#)

[Core Duo 4](#)

[Core Duo 5](#)

[Core Duo 6](#)

[Core Duo 7](#)

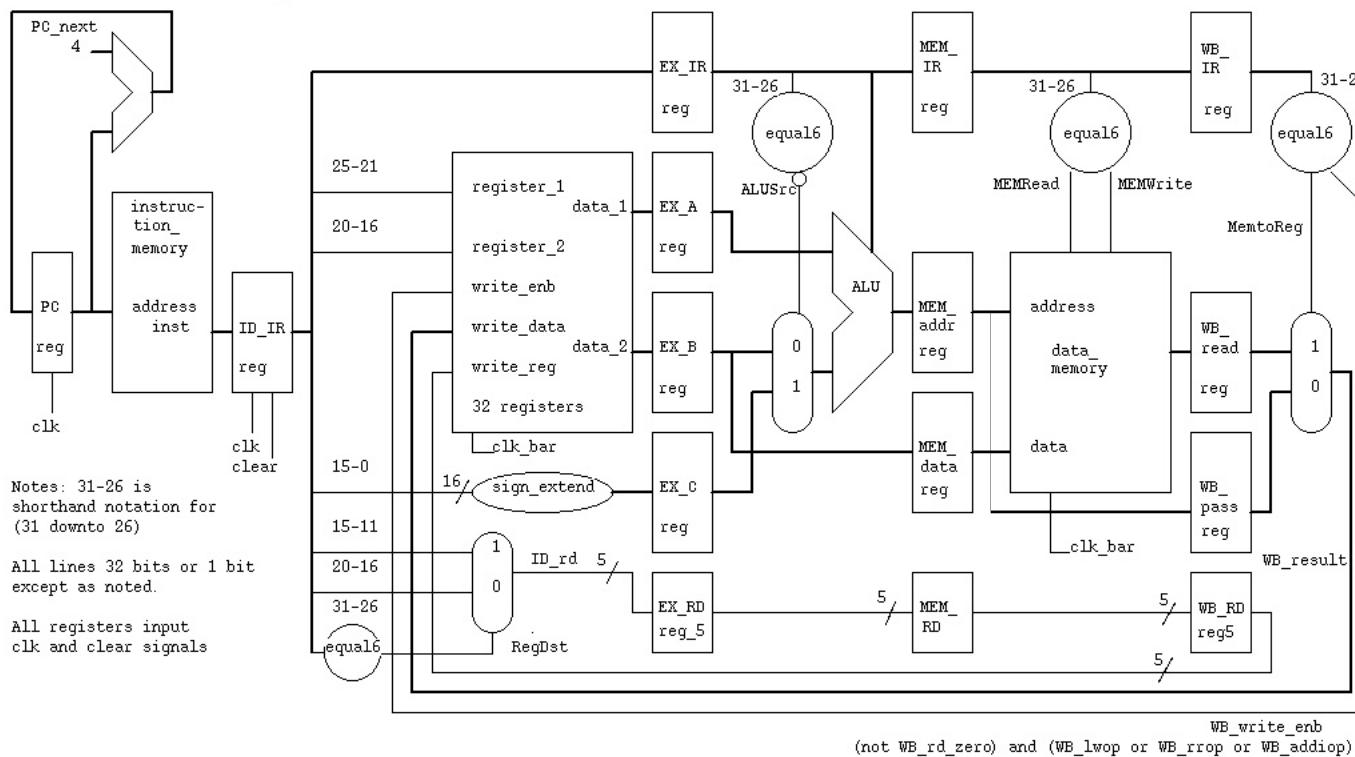
[Core Duo 8](#)

[HW7 is assigned](#)

## Lecture 18, Project Outline and VHDL

Project part1 starts with [\\_part1\\_start.vhd1](#).  
Search for "???" where you need to do some work.  
!!! remove ??? , ... , they are not legal VHDL.

## CMSC 411 Project Part1 schematic



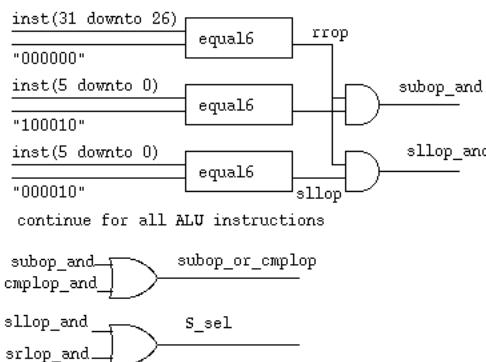
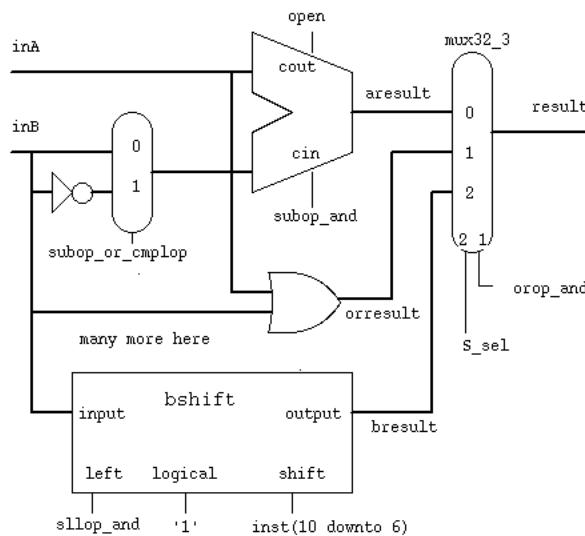
WB\_write\_enb  
(not WB\_rd\_zero) and (WB\_lwop or WB\_rrop or WB\_addiop)

WB\_write\_enb <= needs WB\_lwop or WB\_lwimop or ...

Above: RegDst WORK.equal6 ID\_IR(31 downto 26) , "00000"  
Similar for ALUSrc compare to "00000" get complement,  
ALUSrc <= not complement

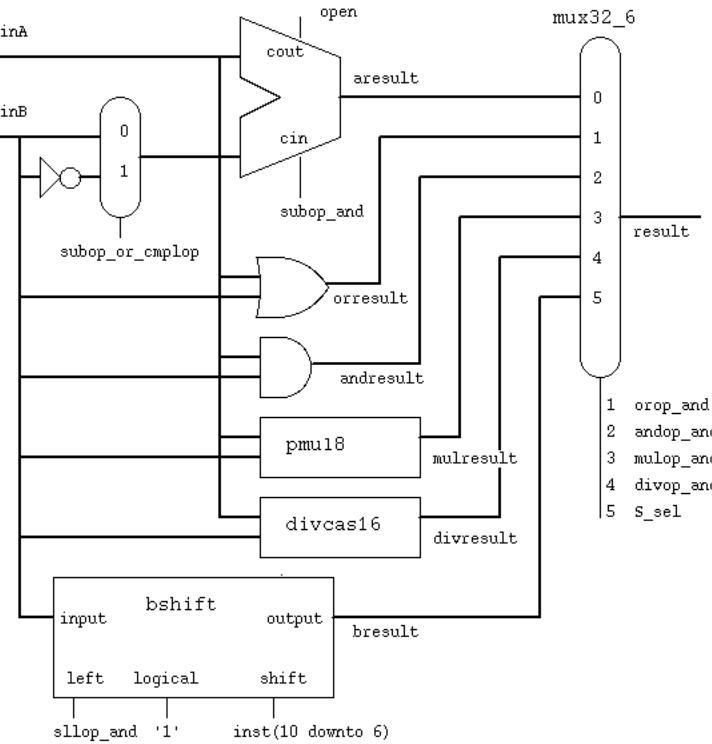
Below: need "not inB" signal, into WORK.mux\_32 and new output name that also goes into B side of ALU.

with ALU schematic for all, also see more on schematic below.



A possible ALU schematic

All include divide, divcas16 covered in Lecture 8 and provided.  
Use your *add32.vhdl* from HW4.  
Use your *pmul8.vhdl* from HW6.



Various versions have different signal names for same signal,  
`orop_and` may be just `orop`, result of anding `oropa` with `rrop`

`S_sel` may be shortened name for `sllop_or_srlop`  
`S_sel` <= `sllop_and` or `srlop_and`;

Remember from `cs411_OPCODES.txt`, `sll` instruction has bottom six bits "000010" and typical code would call that signal `sllop`. But, many instructions could have those bottom bits, thus to be sure the instruction is `sll` check top six bits, `Rrop`, equal to zero and call that signal `sllop_and`. Similar for all instructions. Some schematics use a short hand, just `sllop` meaning the instruction is an `sll`, yet VHDL code needs `sllop_and`.

Extracted code to indicate where you need to do some work "...":  
-- part1\_start.vhdl VHDL '93 version using entities from WORK library  
[part1\\_start.vhdl to modify](#)

```

library IEEE;
use IEEE.std_logic_1164.all;

entity alu_32 is -- given. Do not change this interface
port(inA      : in std_logic_vector (31 downto 0);
     inB      : in std_logic_vector (31 downto 0);
     inst     : in std_logic_vector (31 downto 0);
     result   : out std_logic_vector (31 downto 0));
end entity alu_32;

architecture schematic of alu_32 is
signal cin      : std_logic := '0';
signal cout    : std_logic;
signal RRop    : std_logic;
signal orop    : std_logic;

```

```

signal orop_and : std_logic;
signal andop   : std_logic;
signal andop_and : std_logic;
signal S_sel    : std_logic;
-- ??? insert other needed signals

signal mulop     : std_logic;
signal mulop_and : std_logic;
signal divop     : std_logic;
signal divop_and : std_logic;

signal aresult : std_logic_vector (31 downto 0);
signal bresult : std_logic_vector (31 downto 0);
signal orresult : std_logic_vector (31 downto 0);
signal andresult : std_logic_vector (31 downto 0);
signal mulresult : std_logic_vector (31 downto 0);
signal divresult : std_logic_vector (31 downto 0);
signal divrem : std_logic_vector (31 downto 0);

begin -- schematic
  --
  -- REPLACE THIS SECTION FOR PROJECT PART 1
  -- (add the signals you need above "begin"
  --
  ORR : entity WORK.equal6 port map(inst(31 downto 26), "000000", RRop);
  Oor: entity WORK.equal6 port map(inst(5 downto 0), "001101", orop);
  Omul: entity WORK.equal6 port map(inst(5 downto 0), "011011", mulop);
  Odiv: entity WORK.equal6 port map(inst(5 downto 0), "011000", divop);
-- ??? insert other xxop statements

  orop_and <=orop and RRop;
  mulop_and <=mulop and RRop;
  divop_and <=divop and RRop;
-- ??? insert other xxx_and statements

adder: entity WORK.add32 port map(a      => inA,
                                  b      => inB,
                                  cin   => cin,
                                  sum   => aresult,
                                  cout  => cout);

Mul:  entity WORK.pmul8 port map(inA(7 downto 0),
                                 inB(7 downto 0),
                                 mulresult(15 downto 0));

Div:  entity WORK.divcas16 port map(inA(31 downto 0),
                                    inB(15 downto 0),
                                    divresult(15 downto 0),
                                    divrem(15 downto 0));

Omux: entity WORK.mux32_6 port map(in0=>aresult,
                                     in1=>bresult,
                                     in2=>andresult,
                                     in3=>orresult,
                                     in4=>mulresult,
                                     in5=>divquo32,
                                     ctrl=>S_sel,
                                     ct2=>andop_and,
                                     ct3=>orop_and,
                                     ct4=>mulop_and,
                                     ct5=>divop_and,
                                     result=>result);
end architecture schematic; -- of alu_32

... big cut

-- put additional debug print here, if needed, delete before submit

end architecture schematic; -- of part1_start

```

```

Do a final search for ???
Oh! You need to compute WB_RRop.
You know RRop is register to register operations add, sub, ...
that has 6 zeros in instruction bits 31 downto 0.
WB write back stage instruction is WB_IR.
WBrrrop: entity WORK.equal6 port map( WB_IR(31 downto 26),"000000", WB_RRop);
similar statement for WB_addiop look up "-----"
Of course, you need to define the signals WB_RRop and WB_addiop and
put the or ... inside the )

```

The additional files needed are:

[part1.abs](#) the program to be executed

[part1.run](#) to stop execution, no halt instruction

[part1.chk](#) the expected output

[cs411\\_OPCODES.txt](#) opcode bit patterns

You will need to enter opcode bit patterns not in part1\_start.vhdl.

Modify your Makefile :

```
all: ... part1.out
```

```

part1.out: add32.vhdl bshift.vhdl pmul8.vhdl divcas16.vhdl \
           part1.vhdl part1.abs part1.run
    ncvhdl -v93 add32.vhdl
    ncvhdl -v93 bshift.vhdl
    ncvhdl -v93 pmul8.vhdl
    ncvhdl -v93 divcas16.vhdl
    ncvhdl -v93 part1.vhdl # renamed and modified part1ce_start.vhdl
    ncelab -v93 part1:schematic
    ncsim -batch -logfile part1.out -input part1.run part1

```

optional in Makefile, should be run on GL before submit cs411 part1 part1.vhdl

```
diff -iw part1.out part1.chk      should be no differences other
                                         than copyright or dates
```

[Now, work on the ALU](#)

The full project writeup:

[cs411\\_proj.shtml](#)

## Lecture 19, Pipelining Data Forwarding

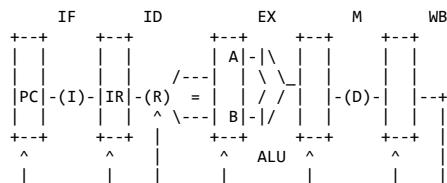
Data forwarding example CMSC 411 architecture

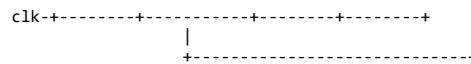
Consider the five stage pipeline architecture:

```

IF instruction fetch, PC is address into memory fetching instruction
ID instruction decode and register read out of two values
EX execute instruction or compute data memory address
M data memory access to store or fetch a data word
WB write back value into general register

```





Now consider the instruction sequence:

```

400  lw $1,100($0) load general register 1 from memory location 100
404  lw $2,104($0) load general register 2 from memory location 104
408  nop
40C  nop      wait for register $2 to get data
410  add $3,$1,$2 add contents of registers 1 and 2, sum into register 3
414  nop
418  nop      wait for register $3 to get data
41C  add $4,$3,$1 add contents of registers 3 and 1, sum into register 4
420  nop
424  nop      wait for register $4 to get data
428  beq $3,$4,-100 branch if contents of register 3 and 4 are equal to 314
42C  add $4,$4,$4 add ..., this is the "delayed branch slot" always exec.

```

The pipeline stage table with NO data forwarding is:

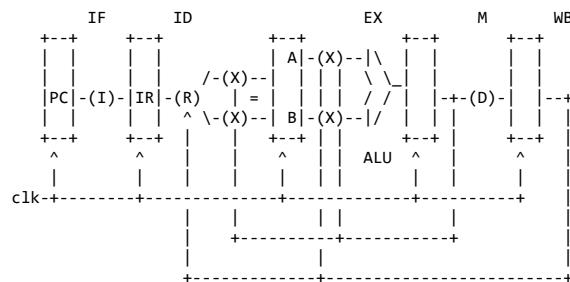
```

lw  IF ID EX M  WB
lw  IF ID EX M  WB
nop  IF ID EX M  WB
nop  IF ID EX M  WB
add  IF ID EX M  WB
nop  IF ID EX M  WB
nop  IF ID EX M  WB
nop  IF ID EX M  WB
add  IF ID EX M  WB
nop  IF ID EX M  WB
nop  IF ID EX M  WB
beq  IF ID EX M  WB
add  IF ID EX M  WB

```

time 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16

This can be significantly improved with the addition of four multiplexors and wiring.



The pipeline stage table with data forwarding is:

```

lw  IF ID EX M  WB
lw  IF ID EX M  WB
nop  IF ID EX M  WB      saved one nop
add  IF ID EX M  WB      $2 in WB and used in EX
add  IF ID EX M  WB      saved two nop's $3 used
nop  IF ID EX M  WB      saved one nop
beq  IF ID EX M  WB      $4 in MEM and used in ID
add  IF ID EX M  WB

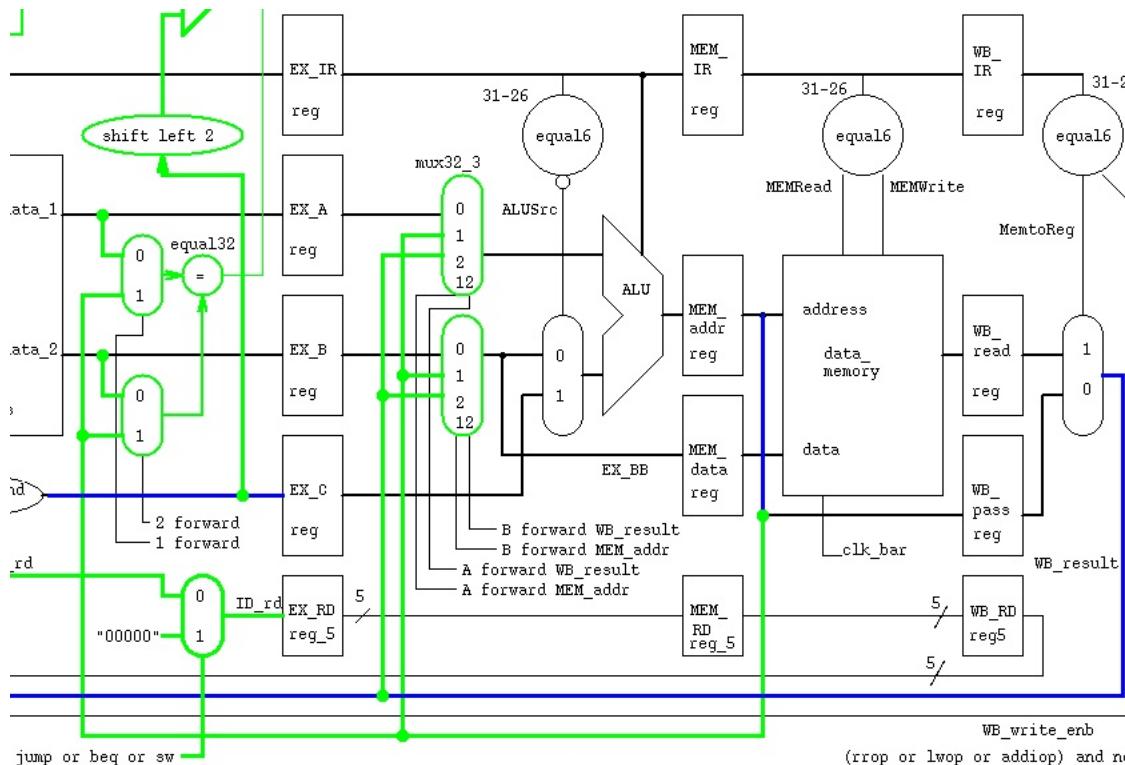
```

time 1 2 3 4 5 6 7 8 9 10 11 12

Note the required nop from using data immediately after a load.

Note the required nop for the beq in the ID stage using an ALU result.

The data forwarding paths are shown in green with the additional multiplexors. The control is explained below.

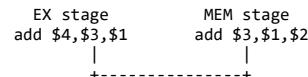


Green must be added to part2a.vhdl.

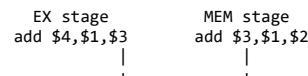
Blue already exists, used for discussion, do not change.

To understand the logic better, note that MEM\_RD contains the register destination of the output of the ALU and MEM\_addr contains the value of the output of the ALU for the instruction now in the MEM stage.

If the instruction in the EX stage has the MEM\_RD destination in bits 25 downto 21, then MEM\_addr must be routed to the A side of the ALU. (This is the A forward MEM\_addr control signal.)



If the instruction in the EX stage has the MEM\_RD destination in bits 20 downto 16, then MEM\_addr must be routed to the B side of the ALU. (This is the B forward MEM\_addr control signal.)



To understand the logic better, note that WB\_RD contains the register destination of the output of the ALU or Memory and WB\_result contains

the value of the output of the ALU or Memory for the instruction now in the WB stage.

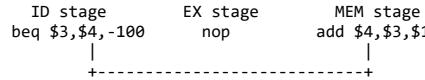
If the instruction in the EX stage has the WB\_RD destination in bits 25 downto 21, then WB\_result must be routed to the A side of the ALU. (This is the A forward WB\_result control signal.)

If the instruction in the EX stage has the WB\_RD destination in bits 20 downto 16, then WB\_result must be routed to the B side of the ALU. (This is the B forward WB\_result control signal.)

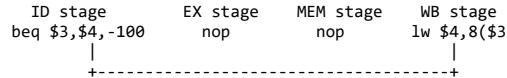
Note that a beq instruction in the ID stage that needs a value from the instruction in the WB stage does not need data forwarding.

A beq instruction in the ID stage has the MEM\_RD destination in bits 25 downto 21, then MEM\_addr must be routed to the top side of the equal comparator. (This is the 1 forward control signal.)

A beq instruction in the ID stage has the MEM\_RD destination in bits 20 downto 16, then MEM\_addr must be routed to the bottom side of the equal comparator. (This is the 2 forward control signal.)



A beq instruction in the ID stage has the WB\_RD destination in bits 20 downto 16, then WB\_result must be used by the bottom side of the equal comparator. (This happens by magic. Not really, two rules above apply.)



The data forwarding rules can be summarized based on the cs411 schematic, shown above.

ID stage beq data forwarding:

```

default with no data forwarding is ID_read_data_1
1 forward MEM_addr is ID_reg1=MEM_RD and MEM_rd/=0 and MEM_OP/=lw

default with no data forwarding is ID_read_data_2
2 forward MEM_addr is ID_reg2=MEM_RD and MEM_rd/=0 and MEM_OP/=lw
  
```

EX stage data forwarding:

```

default with no data forwarding is EX_A
A forward MEM_addr is EX_reg1=MEM_RD and MEM_RD/=0 and MEM_OP/=lw
A forward WB_result is EX_reg1=WB_RD and WB_RD/=0

default with no data forwarding is EX_B
B forward MEM_addr is EX_reg2=MEM_RD and MEM_RD/=0 and MEM_OP/=lw
B forward WB_result is EX_reg2=WB_RD and WB_RD/=0
  
```

Note: the entity mux32\_3 is designed to handle the above.

ID\_RD is 0 for ID\_OP= beq, j, sw (nop, all zeros, automatic zero in RD)  
thus EX\_RD, MEM\_RD, WB\_RD = 0 for these instructions  
Because register zero is always zero, we can use 0 for a destination for every instruction that does not produce a result in a register. Thus no data forwarding will occur for instructions that do not produce a value in a register.

```
note: ID_reg1 is ID_IR(25 downto 21)
      ID_reg2 is ID_IR(20 downto 16)
      EX_reg1 is EX_IR(25 downto 21)
      EX_reg2 is EX_IR(20 downto 16)
      MEM_OP is MEM_IR(31 downto 26)
      EX_OP is EX_IR(31 downto 26)
      ID_OP is ID_IR(31 downto 26)
```

These shorter names can be used with VHDL alias statements

```
alias ID_reg1 : word_5 is ID_IR(25 downto 21);
alias ID_reg2 : word_5 is ID_IR(20 downto 16);
alias EX_reg1 : word_5 is EX_IR(25 downto 21);
alias EX_reg2 : word_5 is EX_IR(20 downto 16);
alias MEM_OP : word_6 is MEM_IR(31 downto 26);
alias EX_OP : word_6 is EX_IR(31 downto 26);
alias ID_OP : word_6 is ID_IR(31 downto 26);
```

Why is the priority mux, mux32\_3 needed?

[mux32\\_3.vhd1 gives priority to ct1 over ct2](#)

Answer: Consider MEM\_RD with a destination value 3 and  
WB\_RD with a destination value 3.

What should add \$4,\$3,\$3 use? MEM\_addr or WB\_result ?

For this to happen, some program or some person would have  
written code such as:

```
sub $3,$12,$11
add $3,$1,$2
add $4,$3,$3 double the value of $3
```

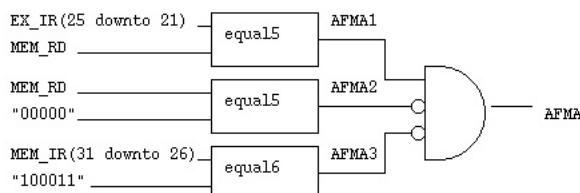
Well, rather obviously, the result of the sub is never used and  
thus the answer to our question is that MEM\_addr must be used. This  
is the closest prior instruction with the required result. The  
correct design is implemented using the priority mux32\_3 with the  
MEM\_addr in the in1 priority input.

The control signal A forward MEM\_addr may be implemented in VHDL as:

Example VHDL for part2a, data forwarding.

Given: A forward MEM\_addr is EX\_reg1=MEM\_RD and MEM\_RD/=0  
and MEM\_OP/=lw

The schematic for computing "A forward MEM\_addr", named AFMA is:



The VHDL needs statements such as the following:

```
signal AFMA, AFMA1, AFMA2, AFMA3: std_logic;
eq17: entity WORK.equal5 port map(EX_IR(25 downto 21), MEM_RD, AFMA1);
eq18: entity WORK.equal5 port map(MEM_RD, "00000", AFMA2);
eq19: entity WORK.equal6 port map(MEM_OP(31 downto 26), "100011", AFMA3);
AFMA <= AFMA1 and not AFMA2 and not AFMA3;
```

btw: 100011 in any\_IR(31 downto 26) is the lw opcode in this example,  
be sure to check this semester's [cs411\\_opcodes.txt](#)

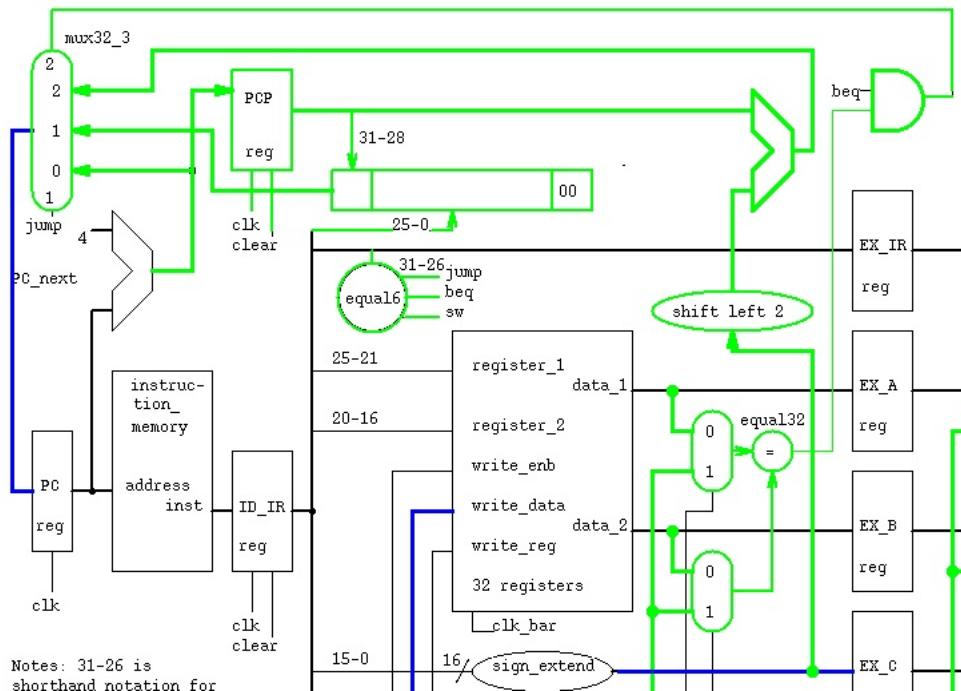
Here is where you may want to add a debug process. Replace AFMA  
with any signal name of interest:

```
prtAFMA: process (AFMA)
  variable my_line : LINE; -- my_line needs to be defined
begin
  write(my_line, string("AFMA="));
  write(my_line, AFMA); -- or hwrite for long signals
  write(my_line, string(" at="));
  write(my_line, now); -- "now" is simulation time
  writeline(output, my_line); -- outputs line
end process prtAFMA;
```

[part2a.chk has the RD signals and values](#)

[cs411\\_opcodes.txt for op code values](#)

Now, to finish part2a.vhdl, the jump and branch instructions must be implemented. This is shown in green on the upper part of the schematic.



The signal out of the jump address box would be coded in VHDL as:

```
jump_addr <= PCP(31 downto 28) & ID_IR(25 downto 0) & "00";
```

The adder symbol is just another instance of your Homework 4, add32.

The "shift left 2" is a simple VHDL statement:

```
shifted2 <= ID_sign_ext(29 downto 0) & "00";
```

The project writeup: [part2a](#)

For more debugging, uncomment print process and diff against:

[part2a\\_print.chk](#)

## Lecture 20, Hazard and Stalls

Our design goal is to eliminate the need for `nop` instructions. The design method is to detect the need for a `nop` and stall the IF and ID stages of the pipeline, inserting a `nop` into the execution stage instruction register, `EX_IR`.

The initial instruction sequence was:

```
400 lw $1,100($0) load general register 1 from memory location 100
404 lw $2,104($0) load general register 2 from memory location 104
408 nop
40C nop      wait for register $2 to get data
410 add $3,$1,$2 add contents of registers 1 and 2, sum into register 3
414 nop
418 nop      wait for register $3 to get data
41C add $4,$3,$1 add contents of registers 3 and 1, sum into register 4
420 nop
424 nop      wait for register $4 to get data
428 beq $3,$4,-100 branch if contents of register 3 and 4 are equal to 314
42C add $4,$4,$4 add ..., this is the "delayed branch slot" always exec.
```

The pipeline stage table with data forwarding and automatic hazard elimination reduces to:

400 lw \$1,100(\$0)	IF	ID	EX	M	WB							
404 lw \$2,104(\$0)		IF	ID	EX	M	WB						
408 add \$3,\$1,\$2		IF	ID	EX	M	WB	--					
40C add \$4,\$3,\$1		IF	IF	ID	EX	M	WB					
410 beq \$3,\$4,-100			IF	ID	ID	EX	M	WB				
414 add \$4,\$4,\$4				IF	IF	ID	EX	M	WB			

time 1 2 3 4 5 6 7 8 9 10 11 12

(actually clock count)

On any clock there can be only one instruction in each pipeline stage.  
Empty stages do not need to be shown, they have an inserted `nop`.  
(useful for Homework 8)

Note that the -- indicates that IF stage and ID stage have stalled.  
The -- also indicates a `nop` instruction has automatically been inserted into the EX stage.

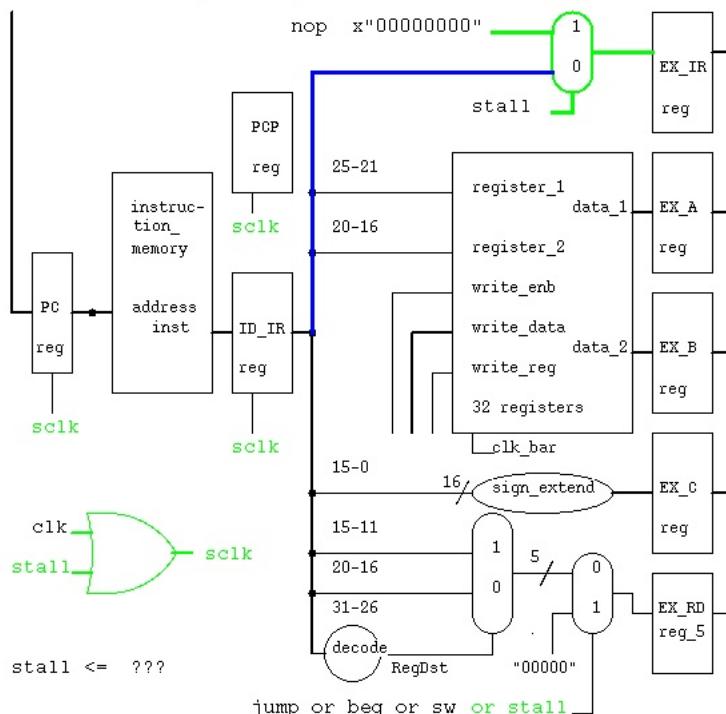
A new instruction can not move into the ID stage when an instruction is stalled there. A new instruction can not move into the IF stage when an instruction is stalled there. No column may have more than one instruction in each stage. Any unlisted stage has a `nop`.

The compiler may now generate compressed code for the computer architecture, saving on memory bandwidth because `nop` instructions are not needed in the executable memory image. (Except a rare `nop` instruction after a branch or jump instruction.)

The primary task will be the implementation of a "stall" signal for the project `part2b.vhdl`. The "stall" signal will then be used to prevent clocking of the instruction fetch, IF stage and instruction decode, ID stage by using a new clock signal "sclk". The explanation for generating "sclk" is presented below.  
Note that when the `nop` instruction is muxed into `EX_IR` then the `EX_RD` must be set to zero along with the existing `beq`, `sw` and `jump`.

The changes in `part2b.vhdl` are in the IF and ID stages.  
Green must be added. The signal "stall" is computed from the information presented below.

## CS411 Project part2b additions



A "hazard" is a condition in the pipeline when a stage of the pipeline would not perform the correct processing with the available data. To be a hazard, the action of data forwarding, covered in the previous lecture, must be taken into account.

Some cases where hazards would occur are:

```

lw $1,100($0)
add $2,$1,$1

      EX stage      MEM stage
add $2,$1,$1    lw $1,100($0)  hazard!
                  value for $1 not available

```

Thus hold add \$2,\$1,\$1 in ID stage, insert nop in EX, this is a stall.

ID stage	EX stage	MEM stage	
add \$2,\$1,\$1	nop	lw \$1,100(\$0)	no hazard

ID stage	EX stage	MEM stage	WB stage	
add \$2,\$1,\$1	nop	lw \$1,100(\$0)	lw \$1,100(\$0)	no hazard

| | +-----+ data forwarding

```

add $4,$3,$1
beq $3,$4,-100

      ID stage      EX stage      MEM stage
beq $3,$4,-100    add $4,$3,$1    hazard!
                  value for $4 not available

```

```

beq $3,$4,-100      nop      add $4,$3,$1      no hazard
|                   |                   |
+-----+           data forwarding

```

```

lw $5,40($1)
beq $5,$4,L2

```

ID stage beq \$5,\$4,L2	EX stage lw \$5,40(\$1)	hazard! value for \$5 not available
----------------------------	----------------------------	--

ID stage beq \$5,\$4,L2	EX stage nop	MEM stage lw \$5,40(\$1) hazard! value for \$5 not available
----------------------------	-----------------	--

ID stage beq \$5,\$4,L2	EX stage nop	MEM stage nop	WB stage lw \$5,40(\$1) no hazard
			+-----+ normal lw

Cases for stall hazards (taking into account data forwarding)  
based on cs411 schematic. This is NOT VHDL, just definitions.

Note: ( OP stands for opcode, bits (31 downto 26)  
 lw stands for load word opcode "100011"  
 addi stands for add immediate opcode "001100" etc.  
 rr\_op stands for OP = "000000" )

```

lw $a, ...
op $b, $a, $a where op is rr_op, beq, sw

```

stall\_lw is EX\_OP=lw and EX\_RD/=0 and  
 (ID\_reg1=EX\_RD or ID\_reg2=EX\_RD)  
 and ID\_OP/=lw and ID\_OP /=addi and ID\_OP /=j

(note: the above handles the special cases where  
 sw needs both registers. sll, srl, cmpl have a zero in unused register.  
 no stall can occur based on EX\_RD, MEM\_RD or WB\_RD = 0)

```

lw $a, ...
lw $b,addr($a) or addi $b,addr($a)

```

stall\_lwlw is EX\_OP=lw and EX\_RD/=0 and  
 (ID\_OP=lw or ID\_OP=addi) and  
 ID\_reg1=EX\_RD

```

lw $a ...
beq $a,$a, ...

```

stall\_mem is ID\_OP=beq and MEM\_RD/=0 and MEM\_OP=lw and  
 (ID\_reg1=MEM\_RD or ID\_reg2=MEM\_RD)

```

op $a, ... where op is rr_op and addi
beq $a,$a, ...

```

stall\_beq is ID\_OP=beq and EX\_RD/=0 and  
 (ID\_reg1=EX\_RD or ID\_reg2=EX\_RD)

ID\_RD is 0 for ID\_OP= beq, j, sw, stall (nop automatic zero)  
 thus EX\_RD, MEM\_RD, WB\_RD = 0 for these instructions

rr\_op is "000000" for add, sub, cmpl, sll, srl, and, mul, ...

stall is stall\_lw or stall\_lwlw or stall\_mem or stall\_beq

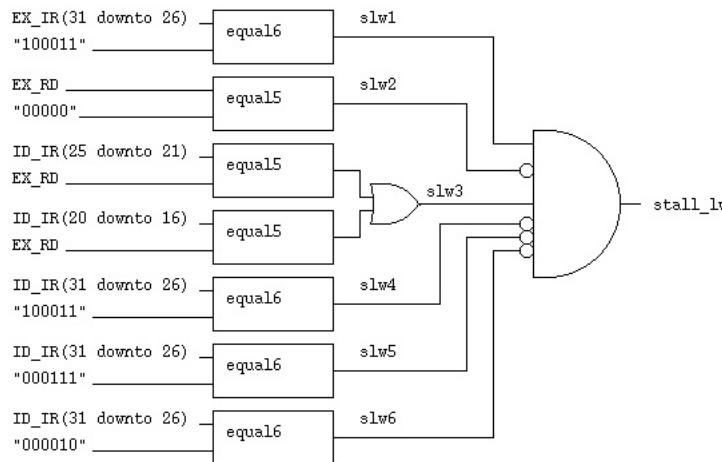
Be sure to use this semesters cs411\_opcodes.txt, it changes every semester.  
[cs411\\_opcodes.txt for op codes](#)

An partial implementation of stall\_lw is:

Example VHDL schematic for part2b, stall\_lw

Given: stall\_lw is EX\_OP=lw and EX\_RD/=0 and  
 (ID\_reg1=EX\_RD or ID\_reg2=EX\_RD)  
 and ID\_OP/=lw and ID\_OP/=addi and ID\_OP/=j

A schematic computing the signal, stall\_lw, could be:



VHDL signals must be defined for, slw1, etc

signal slw1: std\_logic; -- etc

The final "and" gate is coded:

```

stall_lw <= slw1 and not slw2 and slw3 and not slw4 and
          not slw5 and not slw6;
  
```

to get slw5 use "001100" for addiop per cs411\_opcodes.txt

To check on the "stall" signal, you may need to add:

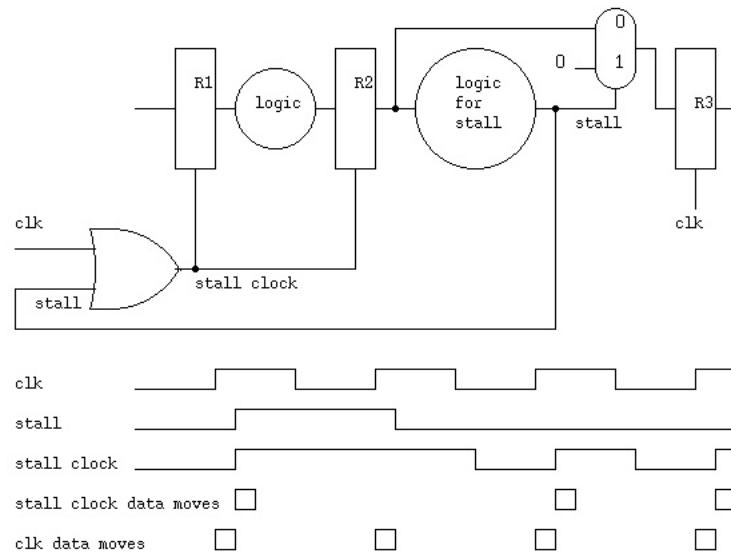
```

prtstall: process (stall)
  variable my_line : LINE; -- my_line needs to be defined
begin
  write(my_line, string'("stall="));
  write(my_line, stall); -- or hwrite for long signals
  write(my_line, string'(" at="));
  write(my_line, now); -- "now" is simulation time
  writeline(output, my_line); -- outputs line
end process prtstall;
  
```

stall clock, sclk, is:

for raising edge registers clk or stall (our circuit)

Stalling a pipeline architecture  
with registers accepting data on rising clock.



For checking your results:  
[part2b.chk](#) look for inserted nop's

[part2b.jpg](#) complete schematic as jpeg image  
[part2b.ps](#) complete schematic as postscript image

Project writeup [part2b](#)

Why is eliminating nop from the load image important?  
Answer: memory bandwidth. RAM memory has always been slower than the CPU. Often by a factor of 10. Thus, the path from RAM memory into the CPU has been made wide. a 64 bit wide memory bus is considered small today. 128 bit and 256 bit memory input to the CPU is common.

Many articles have been written that say "adding more RAM to your computer will give more performance improvement than adding a faster CPU." This is often true because of the complex interaction of the operating system, application software, computer architecture and peripheral equipment. Adding RAM to most computers is easy and can be added by non experts. The important step in adding more RAM is to get the correct Dual Inline Memory Modules, DIMM's. There are speed considerations, voltage considerations, number of pins and possible pairing considerations. The problem is that there are many choices. The following table indicates some of the choices yet does not include RAM size.

Type	Memory Bus	Symbol	Module Bandwidth	DIMM Pins	Nominal Voltage	Memory clock
DDR4	1700Mhz	PC4-2133	25.6GB/sec	288	1.2 volt	
DDR3	1600Mhz	PC3-12800	12.8GT/sec 38.4GB/sec	240	1.6 volt	200Mhz
DDR3	1333Mhz	PC3-10600	10.7GT/sec	240	1.6 volt	166Mhz triple
DDR3	1066Mhz	PC3-8500	8.5GT/sec	240	1.6 volt	133Mhz channel

DDR3	800Mhz	PC3-6400	6.4GT/sec	240	1.6 volt	100MHz (10ns)
------	--------	----------	-----------	-----	----------	---------------

DDR2	1066MHz	PC2-8500	17.0GB/sec	240	2.2 volt	two channel
DDR2	1000MHz	PC2-8000	16.0GB/sec	240	2.2 volt	
DDR2	900MHz	PC2-7200	14.4GB/sec	240	2.2 volt	
DDR2	800MHz	PC2-6400	12.8GB/sec	240	2.2 volt	
DDR2	667MHz	PC2-5300	10.6GB/sec	240	2.2 volt	
DDR2	533MHz	PC2-4200	8.5GB/sec	240	2.2 volt	
DDR2	400MHz	PC2-3200	6.4GB/sec	240	2.2 volt	

DDR	556MHz	PC-4500	9.0GB/sec	184	2.6 volt	
DDR	533MHz	PC-4200	8.4GB/sec	184	2.6 volt	
DDR	500MHz	PC-4000	8.0GB/sec	184	2.6 volt	
DDR	466MHz	PC-3700	7.4GB/sec	184	2.6 volt	
DDR	433MHz	PC-3500	7.0GB/sec	184	2.6 volt	
DDR	400MHz	PC-3200	6.4GB/sec	184	2.6 volt	
DDR	366MHz	PC-3000	5.8GB/sec	184	2.6 volt	
DDR	333MHz	PC-2700	5.3GB/sec	184	2.6 volt	
DDR	266MHz	PC-2100	4.2GB/sec	184	2.6 volt	
DDR	200MHz	PC-1600	3.2GB/sec	184	2.6 volt	

Pre DDR had 168 pin 3.3 volt DIMM's.  
Older machines had 72 pin RAM

Then, there is the size of the DIMM in bytes.  
(may need 2 DDR2 or 3 DDR3 in parallel, minimum 6GB DDR3)

128MB						
256MB						
512MB						
1024MB	1GB					
2048MB	2GB					
4096MB	4GB					

Then, there is a choice of NON-ECC or ECC, Error Correcting Code  
that may be desired in commercial systems.

Then, possibly a choice of buffered or unbuffered.

Then, a choice of response CL3, CL4, CL5 clock waits.  
(in detail may read 7-7-7-20 notation)

Then, shop by price or manufacturers history of reliability.

Some systems require DIMM's of the same size and speed be installed  
in pairs. Read your computers manual or check for information on  
WEB sites. I have used the following sites to get information and  
purchase more RAM.

[www.crucial.com](http://www.crucial.com)

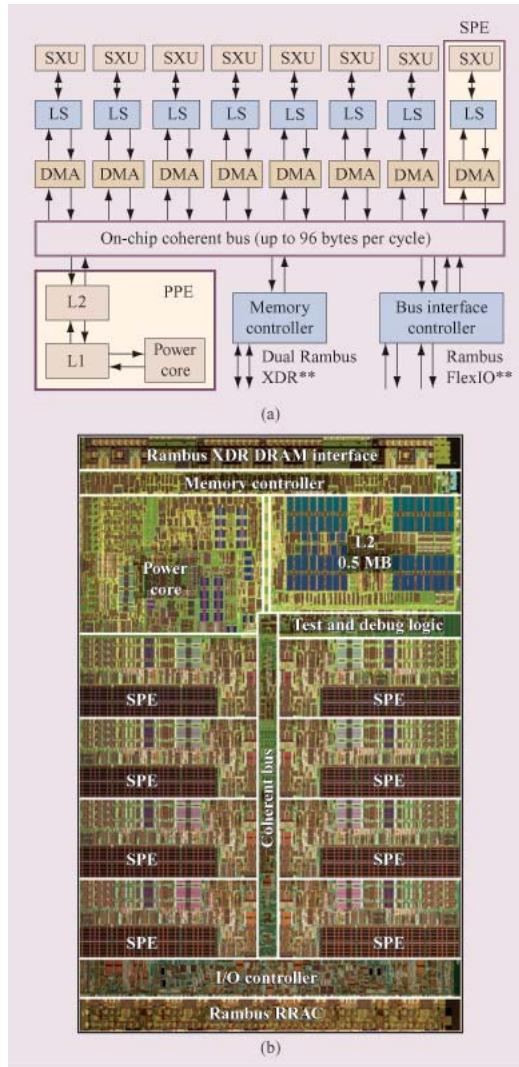
You may search by your computers make and model, or by  
DDR2 and see specification to find what is available.

[www.kingston.com](http://www.kingston.com)

[www.kingston.com/KHX8500](http://www.kingston.com/KHX8500)

[www.valueram.com/datasheets/KHX8500D2\\_1G.pdf](http://www.valueram.com/datasheets/KHX8500D2_1G.pdf)

Now, how can an architecture best make use of the combination of  
pipelines and memory. IBM Cell Processor uses an architecture of  
a general purpose CPU on chip with eight additional pipeline  
processors.

**Figure 1**

(a) Cell processor block diagram and (b) die photo. The first generation Cell processor contains a power processor element (PPE) with a Power core, first- and second-level caches (L1 and L2), eight synergistic processor elements (SPEs) each containing a direct memory access (DMA) unit, a local store memory (LS) and execution units (SXUs), and memory and bus interface controllers, all interconnected by a coherent on-chip bus. (Cell die photo courtesy of Thomas Way, IBM Burlington.)

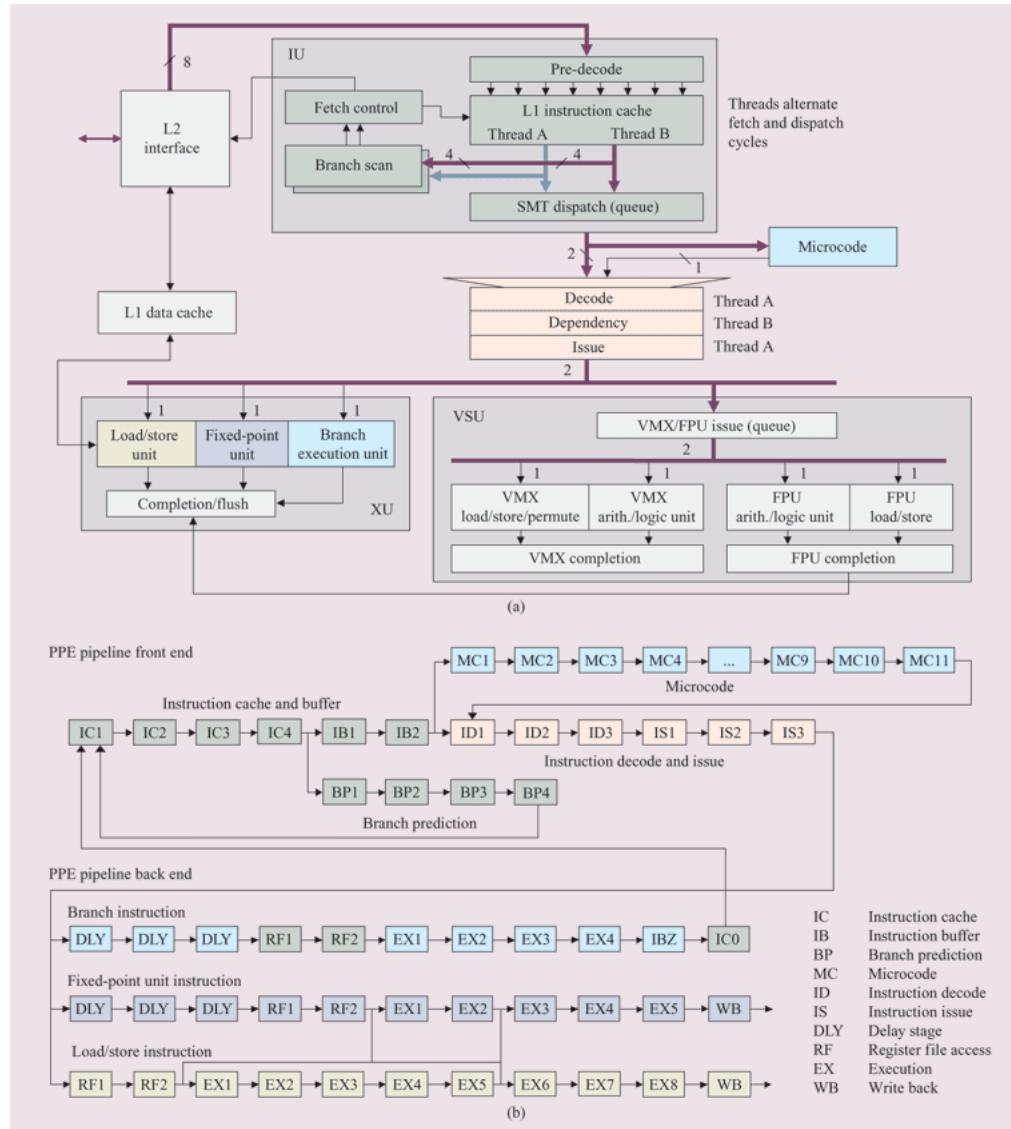


Figure 2

Power processor element (a) major units and (b) pipeline diagram. Instruction fetch and decode fetches and decodes four instructions in parallel from the first-level instruction cache for two simultaneously executing threads in alternating cycles. When both threads are active, two instructions from one of the threads are issued in program order in alternate cycles. The core contains one instance of each of the major execution units (branch, fixed-point, load/store, floating-point (FPU), and vector-media (VMX)). Processing latencies are indicated in part (b) [color-coded to correspond to part (a)]. Simple fixed-point instructions execute in two cycles. Because execution of fixed-point instructions is delayed, load to use penalty is limited to one cycle. Branch miss penalty is 23 cycles and is comparable to the penalty in designs with a much lower operating frequency.

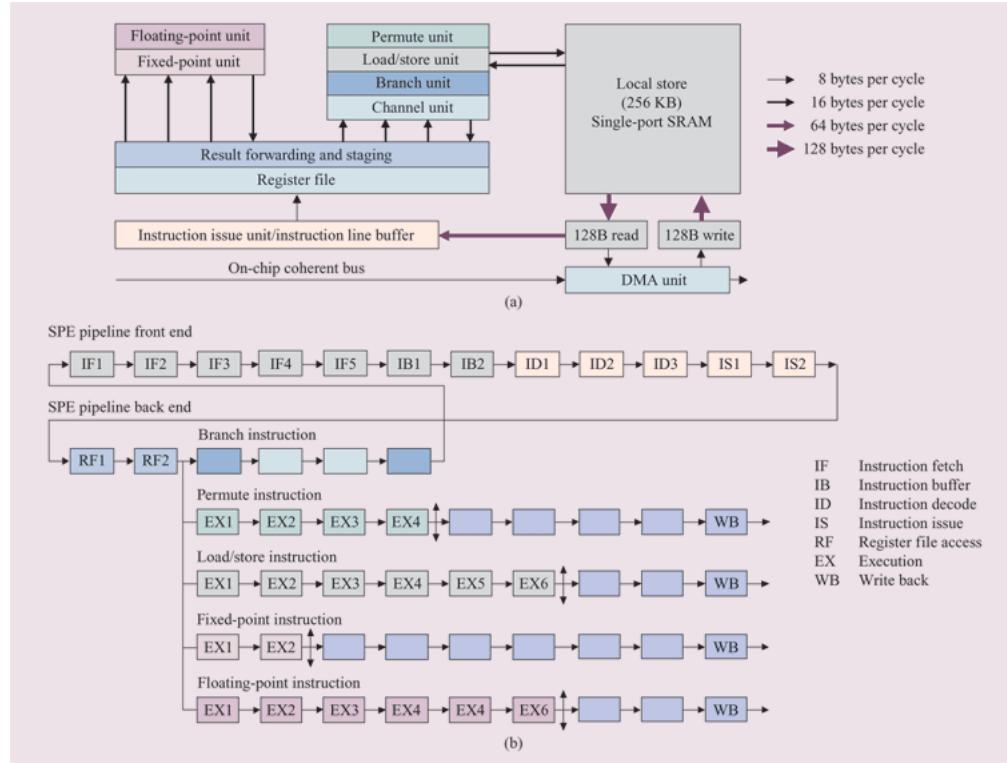
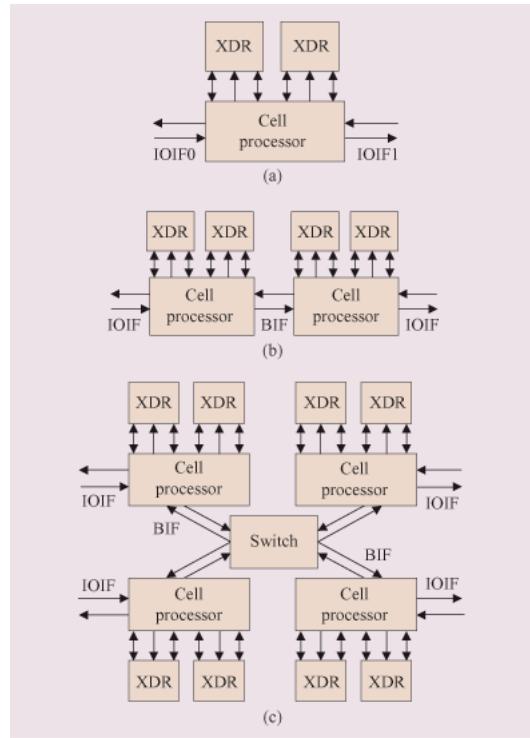
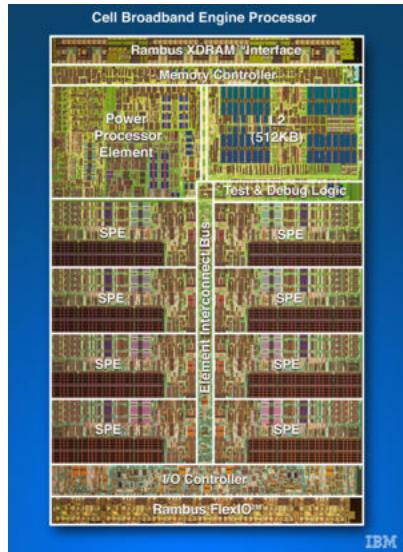


Figure 3

Synergistic processor element (a) organization and (b) pipeline diagram. Central to the synergistic processor is the 256-KB local store SRAM. The local store supports both 128-byte access from direct memory access (DMA) read and write, as well as instruction fetch, and a 16-byte interface for load and store operations. The instruction issue unit buffers and pre-fetches instructions and issues up to two instructions per cycle. A 6-read, 2-write port register file provides both execution pipes with 128-bit operands and stores the results. Instruction execution latency is two cycles for simple fixed-point instructions and six cycles for both load and single-precision floating-point instructions. Instructions are staged in an operand-forwarding network for up to six additional cycles; all execution units write their results in the register file in the same stage. The penalty for mispredicted branches is 18 cycles.

**Figure 4**

Cell system configuration options: (a) Base configuration for small systems. (b) “Glueless” two-way symmetric multiprocessor. (c) Four-way symmetric multiprocessor. (IOIF: Input–output interface; BIF: broadband interface.)



[Cell-tutorial.pdf](#)

[HW8 is assigned](#)

[part2b is assigned](#)

For more debugging, uncomment print process and diff against:

[part2b\\_print.chk](#)

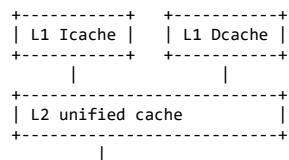
## Lecture 21, Cache

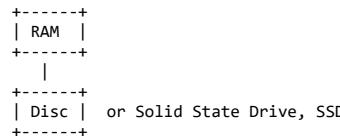
The "cache" is very high speed memory on the CPU chip. Typical CPU's can get words out of the cache every clock. In order to be as fast as the logic on the CPU, the cache can not be as large as the main memory. Typical cache sizes are hundreds of kilobytes to a few megabytes.

There is typically a level 1 instruction cache, a level 1 data cache. These would be in the blocks on our project schematic labeled instruction memory and data memory.

Then, there is typically a level 2 unified cache that is larger and may be slower than the level 1 caches. Unified means it is used for both instructions and data.

Some computers have a level 3 cache that is larger and slower than the level 2 cache. Multi core computers have at least a L1 instruction cache and a L1 data cache for every core. Some have a L3 unified cache that is available to all cores. Thus data can go from one core to another without going through RAM.





The goal of the architecture is to use the cache for instructions and data in order to execute instructions as fast as possible. Typical RAM requires 5 to 10 clocks to get an instruction or data word. A typical CPU does prefetching and branch prediction to bring instructions into the cache in order to minimize stalls waiting for instructions. You will simulate a cache and the associated stalls in part 3 of your project.

Intel IA-64 cache structure, page 3

[IA-64 Itanium](#)

An approximate hierarchy is:

	size	response
CPU		0.5 ns 2 GHz clock
L1 cache	.032MB	0.5 ns one for instructions, another for data
L2 cache	4MB	1.0 ns
RAM	4000MB	4.0 ns
disk	500000MB	4.0 ms = 4,000,000 ns

A program is loaded from disk, into RAM, then as needed into L2 cache, then as needed into L1 cache, then as needed into the CPU pipelines.

- 1) The CPU initiates the request by sending the L1 cache an address. If the L1 cache has the value at that address, the value is quickly sent to the CPU.
- 2) If the L1 cache does not have the value, the address is passed to the L2 cache. If the L2 cache has the value, the value is quickly passed to the L1 cache. The L1 cache passes the value to the CPU.
- 3) If the L2 cache does not have the value at the address, the address is passed to a memory controller that must access RAM in order to get the value. The value passes from RAM, through the memory controller to the L2 cache then to the L1 cache then to the CPU.

This may seem tedious yet each level is optimized to provide good performance for the total system. One reason the system is fast is because of wide data paths. The RAM data path may be 128-bits or 256-bits wide. This wide data path may continue through the L2 cache and L1 cache. The cache is organized in blocks (lines or entries may be used in place of the word blocks) that provide for many bytes of data to be accessed in parallel. When reading from a cache, it is like combinational logic, it is not clocked. When writing into a cache it must write on a clock edge.

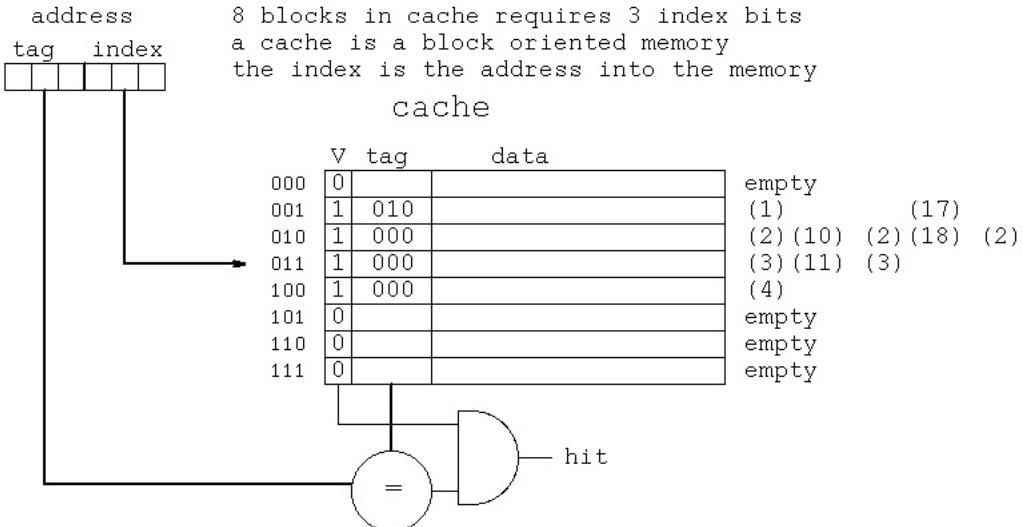
A cache receives an address, a computer address, a binary number. The parts of the cache are all powers of two. The basic unit of an address is a byte. For our study, four bytes, one word, will always be fetched from the cache. When working the homework problems be sure to read the problem carefully to determine if the addresses given are byte addresses or word addresses. It will be easiest and less error prone if all addresses are converted to binary for working the homework.

The basic elements of a cache are:

- A valid bit: This is a 1 if values are in the cache block
- A tag field: This is the upper part of the address for the values in the cache block.
- Cache block: The values that may be instructions or data

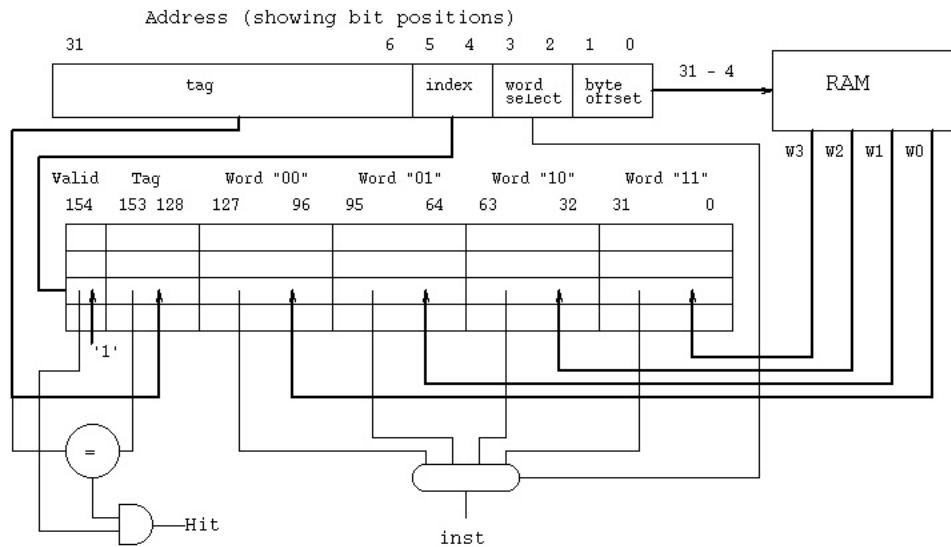
In order to understand a simple cache, follow the sequence of word addresses presented to the following cache.

## Direct Mapped Cache with 8 words, 1 word per block



## Sequence of addresses and cache actions

decimal	binary	hit/miss	action
tag	index		
1	000 001	miss	set valid, load data
2	000 010	miss	set valid, load data
3	000 011	miss	set valid, load data
4	000 100	miss	set valid, load data
10	001 010	miss	wrong tag, load data
11	001 011	miss	wrong tag, load data
1	000 001	hit	no action
2	000 010	miss	wrong tag, load data
3	000 011	miss	wrong tag, load data
17	010 001	miss	wrong tag, load data
18	010 010	miss	wrong tag, load data
2	000 010	miss	wrong tag, load data
3	000 011	hit	no action
4	000 100	hit	no action



Notes: Dark paths for loading cache block upon a miss

Index is the address of the block in the cache

Lines are drawn for block 2 selected by Index="10"

### Instruction Cache for project part 3a

#### Sequence of addresses and cache actions

decimal	binary	hit/miss	action
tag	index	word	
1	00	00 01	miss set valid, load data (0)(1)(2)(3)
2	00	00 10	hit no action
3	00	00 11	hit no action
4	00	01 00	miss set valid, load data (4)(5)(6)(7)
10	00	10 10	miss set valid, load data (8)(9)(10)(11)
11	00	10 11	hit no action
1	00	00 01	hit no action
2	00	00 10	hit no action
3	00	00 11	hit no action
17	01	00 01	miss wrong tag, load data (16)(17)(18)(19)
18	01	00 10	hit no action
2	00	00 10	miss wrong tag, load data (0)(1)(2)(3)
3	00	00 11	hit no action
4	00	01 00	hit no action

There are many cache organizations. The ones you should know are:

A direct mapped cache: the important feature is one tag comparator.

An associative cache: the important feature is more than one tag comparator. "Two way associative" means two tag comparators. "Four way associative" means four tag comparators.

A fully associative cache: Every tag slot has its own comparator.  
This is expensive, typically used for TLB's.

For each organization the words per block may be some power of 2.

For each organization the number of blocks may be some power of 2.

The size of the address that the cache must accept is determined by the CPU. Note that the address is partitioned starting with the low order bits. Given a byte address, the bottom two bits do not go to the cache. The next bits determine the word. If there are 4 words per block, 2-bits are needed, if there are 8 words per block, 3-bits are needed, if there are 16 words per block 4-bits are needed.  $2^{n-1}$  or number of bits is log base 2 of number of words. The next bits are called the index and basically address a block. For  $2^n$  blocks,  $n$  bits are needed. The top bits, whatever is not in the byte, word or index are the tag bits.

Given a 32-bit byte address, 8 words per block, 4096 blocks you would have:

byte    2-bits	
word    3-bits	
index 12-bits	
tag    15-bits	
total 32-bits      tag   index   word   byte   address +---+---+---+---+ 15    12    3    2	

To compute the number of bits in this cache:

4096 x 8 words at 32 bits per word = 1,048,576	
4096 x 15 bits tags	= 61,440
4096 x 1 bits valid bits	= 4,096
----- total bits = 1,114,112 (may not be a power of 2)	

Each cache block or line or entry, for this example has:

valid tag    8 words data or instructions	
+-- +---+ +-----+  1    15     8*32=256 bits	total 272 bits
+-- +---+ +-----+	

then 12 bit index means  $2^{12}=4096$  blocks.  $4096 * 272 = 1,114,112$  bits.

Cache misses may be categorized by the reason for the miss:

Compulsory miss: The first time a word is used and the block that contains that word has never been used.

Capacity miss: A miss that would have been a hit if the cache was big enough.

Conflict miss: A miss that would have been a hit in a fully associative cache.

The "miss penalty" is the time or number of clocks that are required to get the data value.

Data caches have two possible architectures in addition to all other variations. Consider the case where the CPU is writing data to RAM, our store word instruction. The data actually is written into the L1 data cache by the CPU. There are now two possibilities:

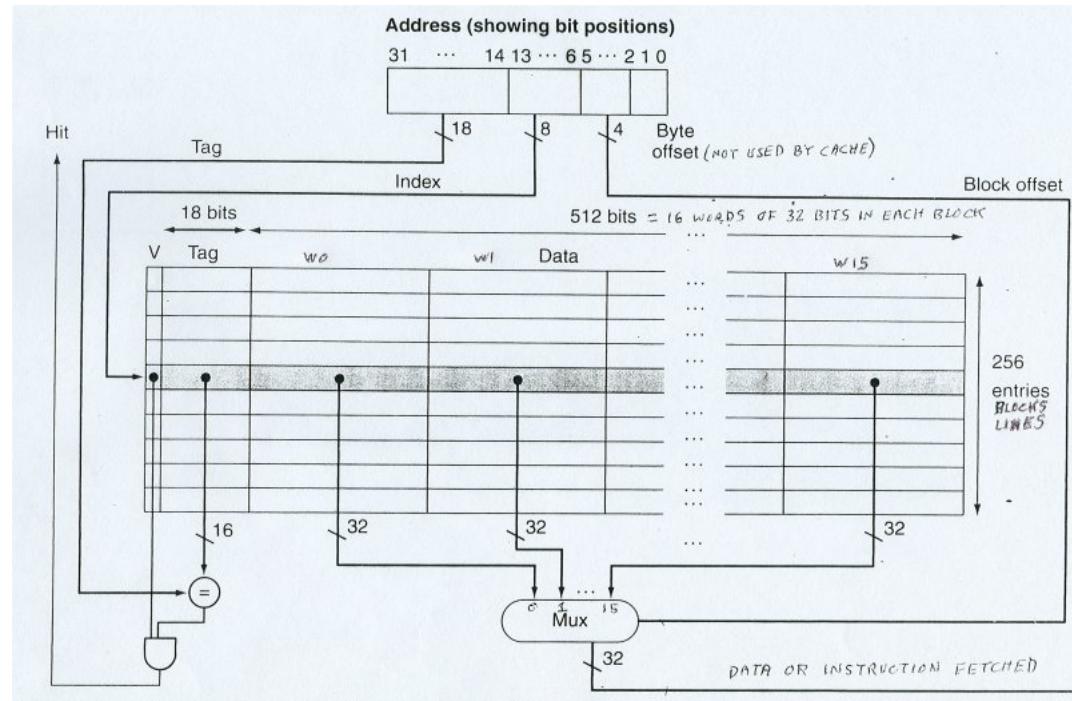
Write back cache: the word is written to the cache. No memory access is made until the block where the word is written is needed, at which time the entire block is written to RAM. It is possible the word could be written, and read, many times before any memory access.

Write through cache: the word is written to the cache and the single word is sent to the RAM memory. This causes to RAM memory to be accessed on every store word but there is no block write when the block is needed for other data. Most of the memory bandwidth is wasted on a wide 128 or 256 bit memory bus.

Tradeoff: Some motherboards have a jumper that you can change to

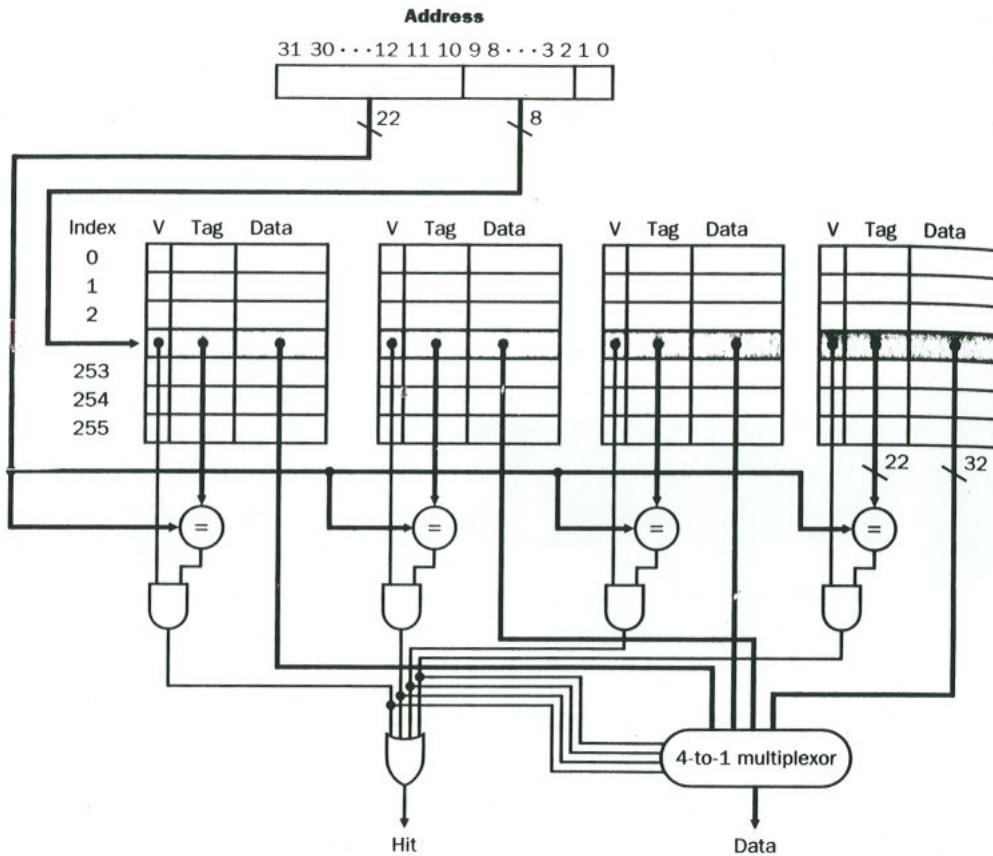
have a write back or write through cache. My choice is a write back cache because I find it gives my job mix better performance.

16 words per block. Note partition of address bits.



**FIGURE 7.9** The 16 KB caches in the Intrinsity FastMATH each contain 256 blocks with 16 words per block. The tag field

A four way associative cache. Note four comparators.  
Each of the four caches could be any of the above architectures  
and sizes.



#### [Homework 9 on cache](#)

The motherboard is essential to support the CPU, RAM and other devices.

#### [Battle of the MotherBoards](#)

#### [An Asus motherboard example](#)

#### [Asus motherboards](#)

#### [2007 Mother Boards, note RAM and hard drive capability](#)

#### [Graphics Cards for mother boards without enough power](#)

Latest high speed IBM Power6, 448 cores at 4.7Ghz

#### [Water cooled](#)

## Lecture 22, Cache Performance

Cache "miss rate" is used as a measure of cache performance.

Given 10 accesses to a cache, 9 hits and 1 miss,  
the miss rate =  $1/10 = 10\%$

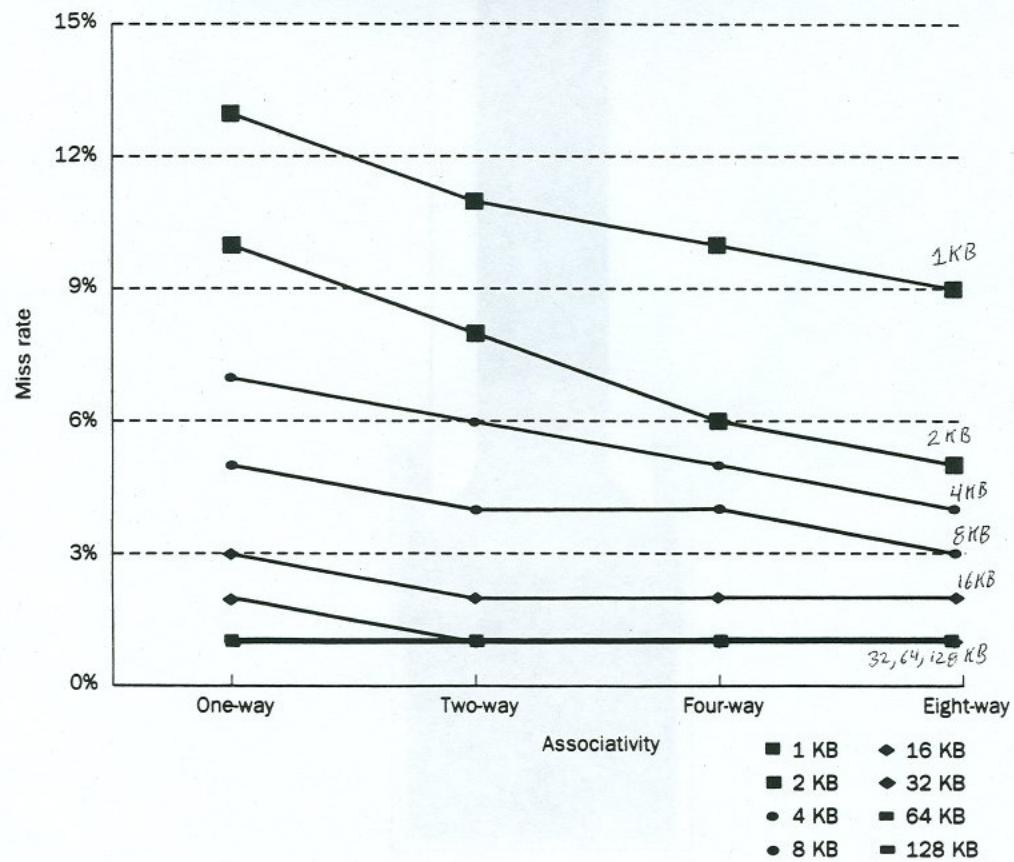
Because there must always be compulsory misses, the miss rate can never be zero. On some plots below, the miss rate is 1% meaning a 99% hit rate.

The importance of the plots is not the numbers, rather the trends. Note that this was based on SPEC92, over 20 years ago. Programs were much smaller back then, yet the trend for performance is the same today. Caches are scaled up today, 1MB and 2MB caches are common and 8MB caches are available.

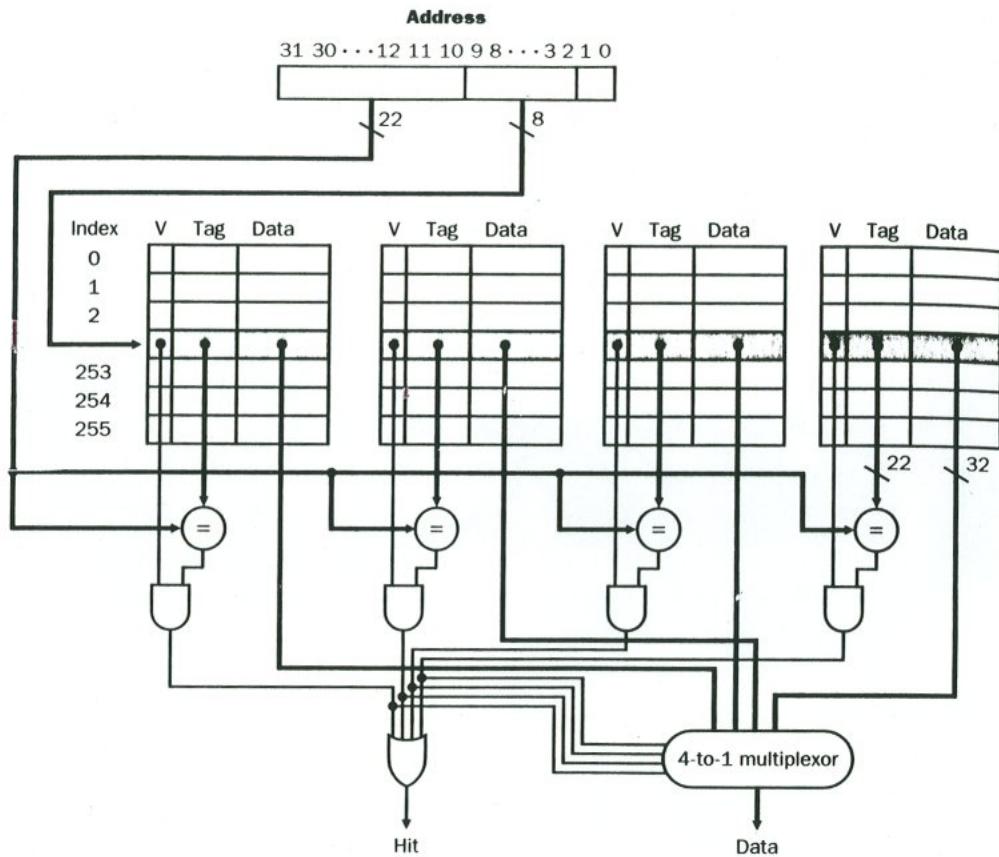
Cache performance based on two factors:

- 1) Cache size (bigger is better)
- 2) Cache associativity (more is better)

The advantage of increasing the degree of associativity is that it usually decreases the miss rate. The improvement in miss rate comes from reducing misses that compete for the same location. We will examine both of these in more detail shortly. First, let's look at how much improvement is gained. Figure 7.29 shows the data for a workload consisting of the SPEC92 benchmarks with caches of 1 KB to 128 KB, varying from direct mapped to eight-way set associative. The largest gains are obtained in going from direct mapped to two-way.

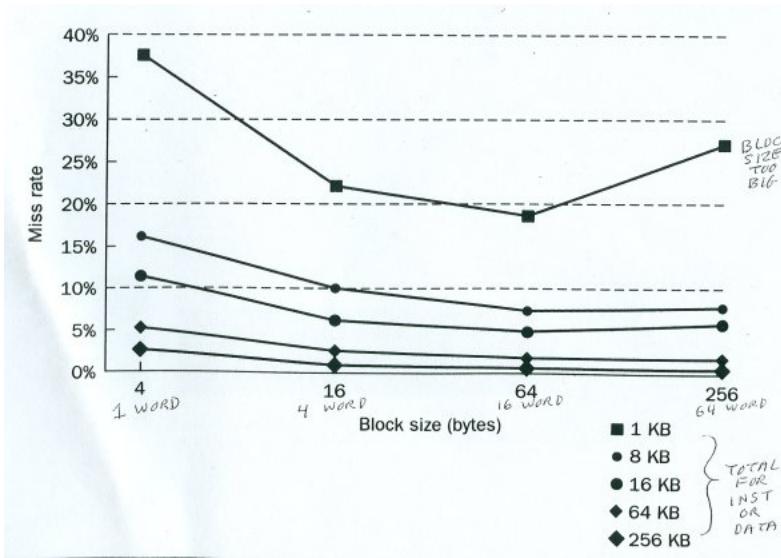


A 4 way associative cache. Count tag equal comparators.



Cache performance based on two factors:

- 1) cache size (bigger is better)
- 2) block size (more is usually better, but not for small caches!)

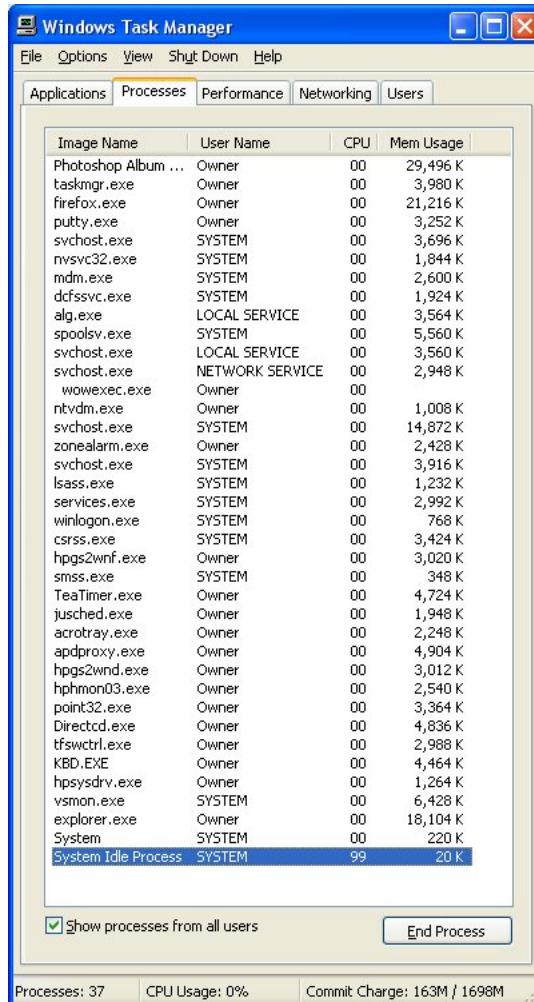


Caches hold a small part of memory in the CPU for fast access.  
The following two sets of memory usage are from my computers and  
show the size of some programs on Windows and Linux.

Memory usage on Windows XP:

37 processes		
Windows Explorer	18,104 KB	18 MB too big for cache
Firefox	21,216 KB	
Photoshop	29,496 KB	
etc.		
total	163,000 KB	163MB of 512MB used.

You would want good performance by keeping most of a program  
in cache. Thus, the need for caches in the megabytes.



## Memory usage on RedHat Linux:

```
83 processes, 3 running
X           38,119 KB  way too big for cache
Firefox      20,083 KB
Gimp         5,402 KB  with extras running
etc
```

```
running   top   reports:
                  306 MB memory used
                  195 MB memory free
                  14 MB memory buff
```

```
From: ps -Al          ## memory size in KB
F S   UID   PID  PPID  C PRI  NI ADR  SZ WCHAN  TTY      TIME CMD
4 S   0     1    0  1 75  0 - 345 schedu ?      00:00:04 init
1 S   0     2    1  0 75  0 - 0 contex ?      00:00:00 keventd
1 S   0     3    1  0 75  0 - 0 schedu ?      00:00:00 kapmd
1 S   0     4    1  0 94 19 - 0 ksofti ?      00:00:00 ksoftirqd_C
1 S   0     9    1  0 85  0 - 0 bdflush ?      00:00:00 bdflush
1 S   0     5    1  0 75  0 - 0 schedu ?      00:00:00 kswapd
1 S   0     6    1  0 75  0 - 0 schedu ?      00:00:00 kscand/DMA
```

```

1 S 0 7 1 0 75 0 - 0 schedu ? 00:00:00 kscand/Norm
1 S 0 8 1 0 75 0 - 0 schedu ? 00:00:00 kscand/High
1 S 0 10 1 0 75 0 - 0 schedu ? 00:00:00 kupdated
1 S 0 11 1 0 85 0 - 0 md_thr ? 00:00:00 mdrecoveryd
1 S 0 15 1 0 75 0 - 0 end ? 00:00:00 kjournald
1 S 0 73 1 0 85 0 - 0 end ? 00:00:00 khubd
1 S 0 1012 1 0 75 0 - 0 end ? 00:00:00 kjournald
1 S 0 1137 1 0 85 0 - 0 end ? 00:00:00 kjournald
1 S 0 3676 1 0 84 0 - 524 schedu ? 00:00:00 dhclient
5 S 0 3727 1 0 75 0 - 369 schedu ? 00:00:00 syslogd
5 S 0 3731 1 0 75 0 - 344 do_sys ? 00:00:00 klogd
5 S 32 3749 1 0 75 0 - 388 schedu ? 00:00:00 portmap
5 S 29 3768 1 0 75 0 - 391 schedu ? 00:00:00 rpc.statd
1 S 0 3812 1 0 75 0 - 0 end ? 00:00:00 rpciod
1 S 0 3813 1 0 85 0 - 0 schedu ? 00:00:00 lockd
5 S 0 3825 1 0 84 0 - 343 schedu ? 00:00:00 apmd
5 S 0 3841 1 0 85 0 - 5014 schedu ? 00:00:00 ypbnd
1 S 0 3945 1 0 75 0 - 372 pipe_w ? 00:00:00 automount
1 S 0 3947 1 0 75 0 - 372 pipe_w ? 00:00:00 automount
1 S 0 3949 1 0 75 0 - 372 pipe_w ? 00:00:00 automount
5 S 0 3968 1 0 85 0 - 879 schedu ? 00:00:00 sshd
5 S 38 3989 1 0 75 0 - 601 schedu ? 00:00:00 ntpd
1 S 0 4013 1 0 75 0 - 0 schedu ? 00:00:00 afs_rxliste
1 S 0 4015 1 0 75 0 - 0 end ? 00:00:00 afs_callbac
1 S 0 4017 1 0 75 0 - 0 schedu ? 00:00:00 afs_rxevent
1 S 0 4019 1 0 75 0 - 0 schedu ? 00:00:00 afs_fsd
1 S 0 4021 1 0 75 0 - 0 schedu ? 00:00:00 afs_checkse
1 S 0 4023 1 0 75 0 - 0 end ? 00:00:00 afs_backgro
1 S 0 4025 1 0 75 0 - 0 end ? 00:00:00 afs_backgro
1 S 0 4027 1 0 75 0 - 0 end ? 00:00:00 afs_backgro
1 S 0 4029 1 0 75 0 - 0 end ? 00:00:00 afs_cachetr
5 S 0 4037 1 0 75 0 - 354 schedu ? 00:00:00 gpm
1 S 0 4046 1 0 75 0 - 358 schedu ? 00:00:00 crond
5 S 43 4078 1 0 76 0 - 1226 schedu ? 00:00:00 xfs
1 S 2 4087 1 0 85 0 - 355 schedu ? 00:00:00 atd
4 S 0 4306 1 0 82 0 - 340 schedu tty1 00:00:00 mingetty
4 S 0 4307 1 0 82 0 - 340 schedu tty2 00:00:00 mingetty
4 S 0 4308 1 0 82 0 - 340 schedu tty3 00:00:00 mingetty
4 S 0 4309 1 0 82 0 - 340 schedu tty4 00:00:00 mingetty
4 S 0 4310 1 0 82 0 - 340 schedu tty5 00:00:00 mingetty
4 S 0 4311 1 0 82 0 - 340 schedu tty6 00:00:00 mingetty
4 S 0 4312 1 0 75 0 - 616 schedu ? 00:00:00 kdm
4 S 0 4325 4312 1 75 0 - 38119 schedu ? 00:00:02 X
5 S 0 4326 4312 0 77 0 - 877 wait4 ? 00:00:00 kdm
4 S 12339 4352 4326 0 85 0 - 1143 rt_sig ? 00:00:00 csh
0 S 12339 4393 4352 0 79 0 - 1034 wait4 ? 00:00:00 startkde
1 S 12339 4394 4393 0 75 0 - 785 schedu ? 00:00:00 ssh-agent
1 S 12339 4436 1 0 75 0 - 5012 schedu ? 00:00:00 kdeinit
1 S 12339 4439 1 0 75 0 - 5440 schedu ? 00:00:00 kdeinit
1 S 12339 4442 1 0 75 0 - 5742 schedu ? 00:00:00 kdeinit
1 S 12339 4444 1 0 75 0 - 9615 schedu ? 00:00:00 kdeinit
0 S 12339 4454 4436 0 75 0 - 2149 schedu ? 00:00:00 artsd
1 S 12339 4474 1 0 75 0 - 10689 schedu ? 00:00:00 kdeinit
0 S 12339 4481 4393 0 75 0 - 341 schedu ? 00:00:00 kwrapper
1 S 12339 4483 1 0 75 0 - 9466 schedu ? 00:00:00 kdeinit
1 S 12339 4484 4436 0 75 0 - 9772 schedu ? 00:00:00 kdeinit
1 S 12339 4486 1 0 75 0 - 9908 schedu ? 00:00:00 kdeinit
1 S 12339 4488 1 0 75 0 - 10299 schedu ? 00:00:00 kdeinit
1 S 12339 4489 4436 0 75 0 - 5085 schedu ? 00:00:00 kdeinit
1 S 12339 4493 1 0 75 0 - 9698 schedu ? 00:00:00 kdeinit
0 S 12339 4494 4436 0 75 0 - 2942 schedu ? 00:00:00 pam-panel-i
4 S 0 4495 4494 0 75 0 - 389 schedu ? 00:00:00 pam_timesta
1 S 12339 4496 4436 0 75 0 - 9994 schedu ? 00:00:00 kdeinit
1 S 12339 4497 4436 0 75 0 - 10010 schedu ? 00:00:00 kdeinit
1 S 12339 4500 1 0 75 0 - 9503 schedu ? 00:00:00 kalarmd
0 S 12339 4501 4496 0 75 0 - 1165 rt_sig pts/2 00:00:00 csh
0 S 12339 4502 4497 0 75 0 - 1159 rt_sig pts/1 00:00:00 csh
0 S 12339 4546 4501 0 85 0 - 1039 wait4 pts/2 00:00:00 firefox
0 S 12339 4563 4546 0 85 0 - 1048 wait4 pts/2 00:00:00 run-mozilla
0 S 12339 4568 4563 1 75 0 - 20083 schedu pts/2 00:00:01 firefox-bin
0 S 12339 4573 1 0 75 0 - 1682 schedu pts/2 00:00:00 gconfd-2
0 S 12339 4583 4502 0 75 0 - 5402 schedu pts/1 00:00:00 gimp
0 S 12339 4776 4583 0 85 0 - 2140 schedu pts/1 00:00:00 script-fu
1 S 12339 4779 4436 1 75 0 - 9971 schedu ? 00:00:00 kdeinit

```

```
0 S 12339 4780 4779 0 75 0 - 1155 rt_sig pts/3 00:00:00 csh
0 R 12339 4803 4780 0 80 0 - 856 - pts/3 00:00:00 ps
```

A benchmark that was designed to note discontinuity in time as the data size increased exceeding the L1 cache, L2 cache. It would take hours if the program exceeded RAM and went to virtual memory on disk!

The basic code, a simple matrix times matrix multiply:

```
/* matmul.c 100*100 matrix multiply */
#include <stdio.h>
#define N 100
int main()
{
    double a[N][N]; /* input matrix */
    double b[N][N]; /* input matrix */
    double c[N][N]; /* result matrix */
    int i,j,k;

    /* initialize */
    for(i=0; i<N; i++){
        for(j=0; j<N; j++){
            a[i][j] = (double)(i+j);
            b[i][j] = (double)(i-j);
        }
    }
    printf("starting multiply \n");

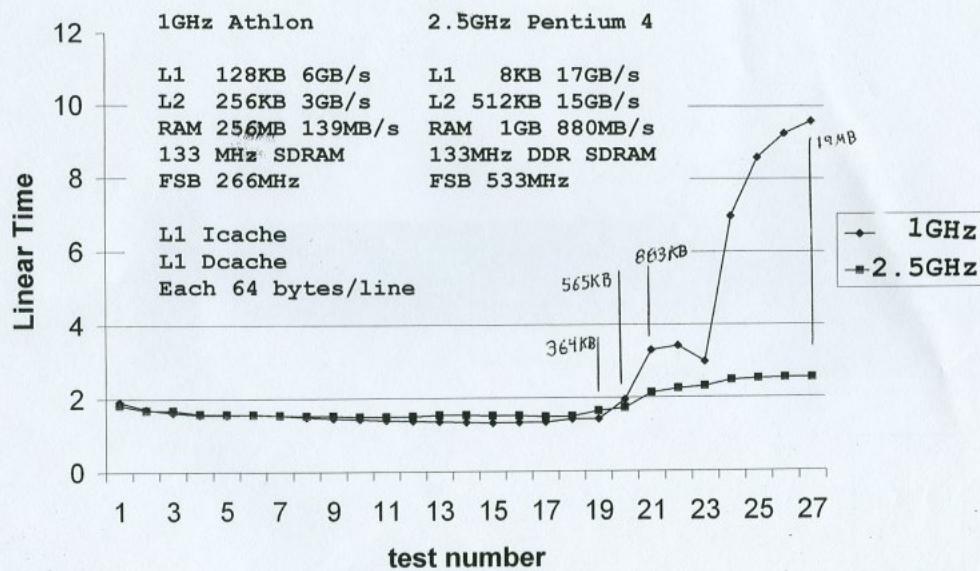
    for(i=0; i<N; i++){
        for(j=0; j<N; j++){
            c[i][j] = 0.0;
            for(k=0; k<N; k++){
                /* how many instructions are in this loop? */
                c[i][j] = c[i][j] + a[i][k]*b[k][j]; /* most time spent here! */
                /* this statement is executed one million times */
            }
        }
    }
    printf("a result %g \n", c[7][8]); /* prevent dead code elimination */
    return 0;
}
```

The actual code:

[time\\_matmul.c](#)  
and results:  
[time\\_matmul\\_1ghz.out](#)  
[time\\_matmul\\_p4\\_25.out](#)  
[time\\_matmul\\_2100.out](#)

Test results on two computers using same executable:

## matmul



A fact you should know about memory usage:

If your program gets more memory while running, e.g. using malloc, then tries to release that memory when not needed, e.g. free, the memory still belongs to your process. The memory is not given back to the operating system for use by another program. Thus, some programs keep growing in size as they run. Hopefully, internally, reusing any memory they previously freed.

On Linux you can use `cat /proc/cpuinfo` to see brief cache size  
[CS machine cpuinfo](#)  
[source code time\\_mp8.c](#)  
[measured time mp8.out](#)

We have seen the Intel P4 architecture, and here is a view of the AMD Athlon architecture circa 2001.

9 pipelines, possibly 9 instruction issued per clock, 3 is typical.

### Key Architectural Features of the AMD Athlon™ XP Processor Include:

QuantiSpeed™ Architecture for enhanced performance

- Nine-issue superpipelined, superscalar x86 processor microarchitecture designed for high performance
- Multiple parallel x86 instruction decoders
- Three out-of-order, superscalar, fully pipelined floating point execution units, which execute x87 (floating point), MMX™

and 3DNow!™ instructions

- Three out-of-order, superscalar, pipelined integer units
- Three out-of-order, superscalar, pipelined address calculation units
- 72-entry instruction control unit
- Advanced hardware data prefetch
- Exclusive and speculative Translation Look-aside Buffers
- Advanced dynamic branch prediction

3DNow!™ Professional technology for leading-edge 3D operation

- 21 original 3DNow!™ instructions—the first technology enabling superscalar SIMD
- 19 additional instructions to enable improved integer math calculations for speech or video encoding and improved data movement for Internet plug-ins and other streaming applications
- 5 DSP instructions to improve soft modem, soft ADSL, Dolby Digital surround sound, and MP3 applications
- 52 SSE instructions with SIMD integer and floating point additions offer excellent compatibility with Intel's SSE technology
- Compatible with Windows® XP, Windows 2000, Windows ME, and Windows 98 operating systems

333MHz and 266Mhz AMD Athlon™ XP processor system bus enables excellent system bandwidth for data movement-intensive applications

- Source synchronous clocking (clock forwarding) technology
- Peak data rate of 2.7GB/s
- Support for 64-bit bi-directional data

The AMD Athlon™ XP processor with performance-enhancing cache memory features 64K instruction and 64K data cache for a total of 128K L1 cache. 256K of integrated, on-chip L2 cache for a total of 384K full-speed, on-chip cache.

Socket A infrastructure designs are based on high-performance platforms and are supported by a full line of optimized infrastructure solutions (chipsets, motherboards, BIOS).

- Available in Pin Grid Array (PGA) for mounting in a socketed infrastructure
- Electrical interface compatible with 333MHz AMD Athlon XP system buses, based on Alpha EV6™ bus protocol

Die size: approximately 37.6 million transistors on 84mm<sup>2</sup>. Manufactured using AMD's state-of-the-art 0.13-micron copper process technology at AMD's Fab 30 wafer fabrication facility in Dresden, Germany.

You can find out your computer's cache sizes and speeds:

[www.memtest86.com](http://www.memtest86.com)

Get the .bin file to make a bootable floppy  
Get the .iso file to make a bootable CD

As part of the output, you do not have to run the memory test,  
you will see cache sizes and bandwidth values. (Shown on plot above.)

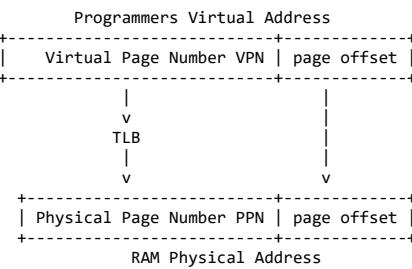
[part3a is assigned](#)

## Lecture 23, Virtual Memory 1

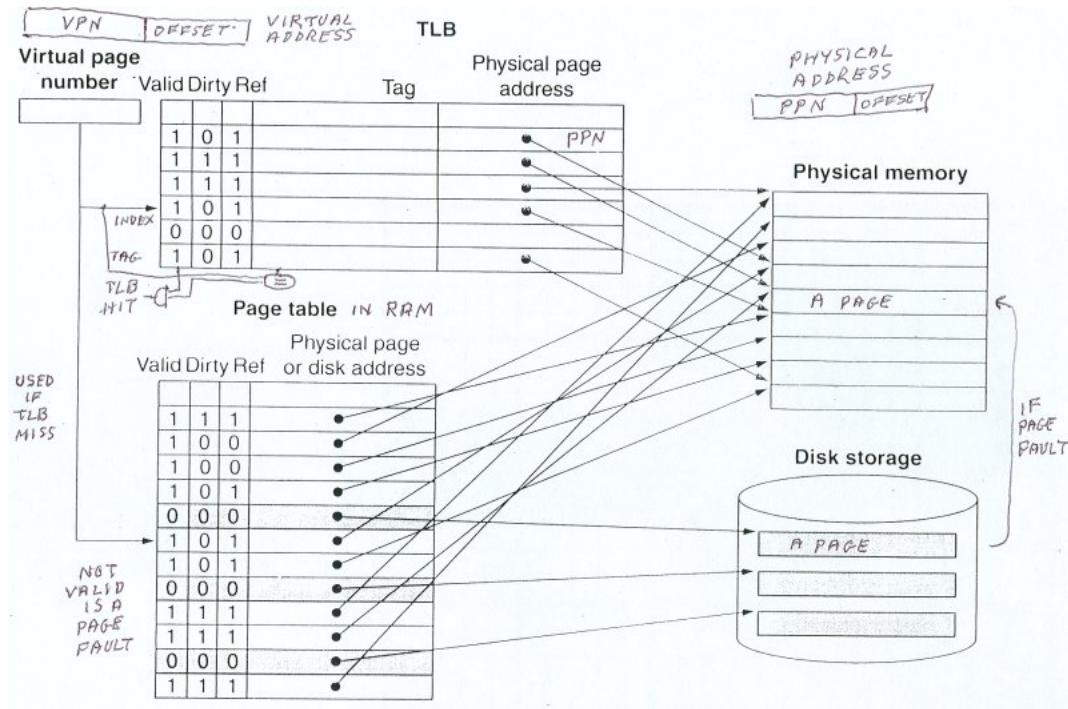
Most modern computers use the programmers addresses as virtual addresses. The virtual addresses must be converted to physical addresses in order to access data and instructions in RAM.

The RAM is divided into many pages. A page is some number of bytes that is a power of 2. A page could be as small as  $2^{12}=4096$  bytes up to  $2^{16}=65536$  bytes or larger. The page offset is the address within a specific page. The offset is 12-bits for a 4096 byte page and 16-bits for a 65536 byte page.

The virtual address and physical address do not necessarily have to be the same number of bits. The operation of virtual memory is to convert a virtual address to a physical address:



TLB is the acronym for Translation Lookaside Buffer. The TLB is the hardware on the CPU that converts the virtual page number to a physical page number. The Operating System is the resource manager and thus assigns each process the physical page numbers that the process may use. The virtual page numbers come from the programmers source code through compiler, assembler and loader onto disk. The addresses you saw in HW3 were virtual addressees. Not the address your program actually ran in RAM.



Two programs, p1 and p2, with code segments p1c and p2c, and data segments p1d and p2d. The operating system runs a simple program as a process. Now, each segment is divided into pages. p1c0, p1c1, p1c2 are the first three pages of program 1 code segment. These are virtual pages. These pages may be loaded into any physical pages in RAM. Each segment is consecutive as stored on disk as an executable program.

disk pages, each line is a page

```
...
p1c0  executable program 1
p1c1
p1c2
p1d0
p1d1
...
p2c0  executable program 2
p2c1
p2d0
p2d1
p2d2
```

There are also other types of segments.  
You may recall from Homework 3, the address of "main" was 0x08048390 28 bits of virtual address.

The page size may be chosen by the operating system author or in some computer architectures the page size is determined by the hardware, as shown below.

As time goes on, the operating system allocates and frees physical pages.  
Physical memory could look like this at some time:  
(Each line is a page, e.g. 8192 bytes)

```
os0  operating system pages
os1
```

```
...
osn
p2d3 somewhat randomly scattered pages
p1c2
empty
p2c0
p2c1
p1d5
p1c4
etc
```

Pages for a program may not be contiguous.

Pages for a segment of a program may not be contiguous.

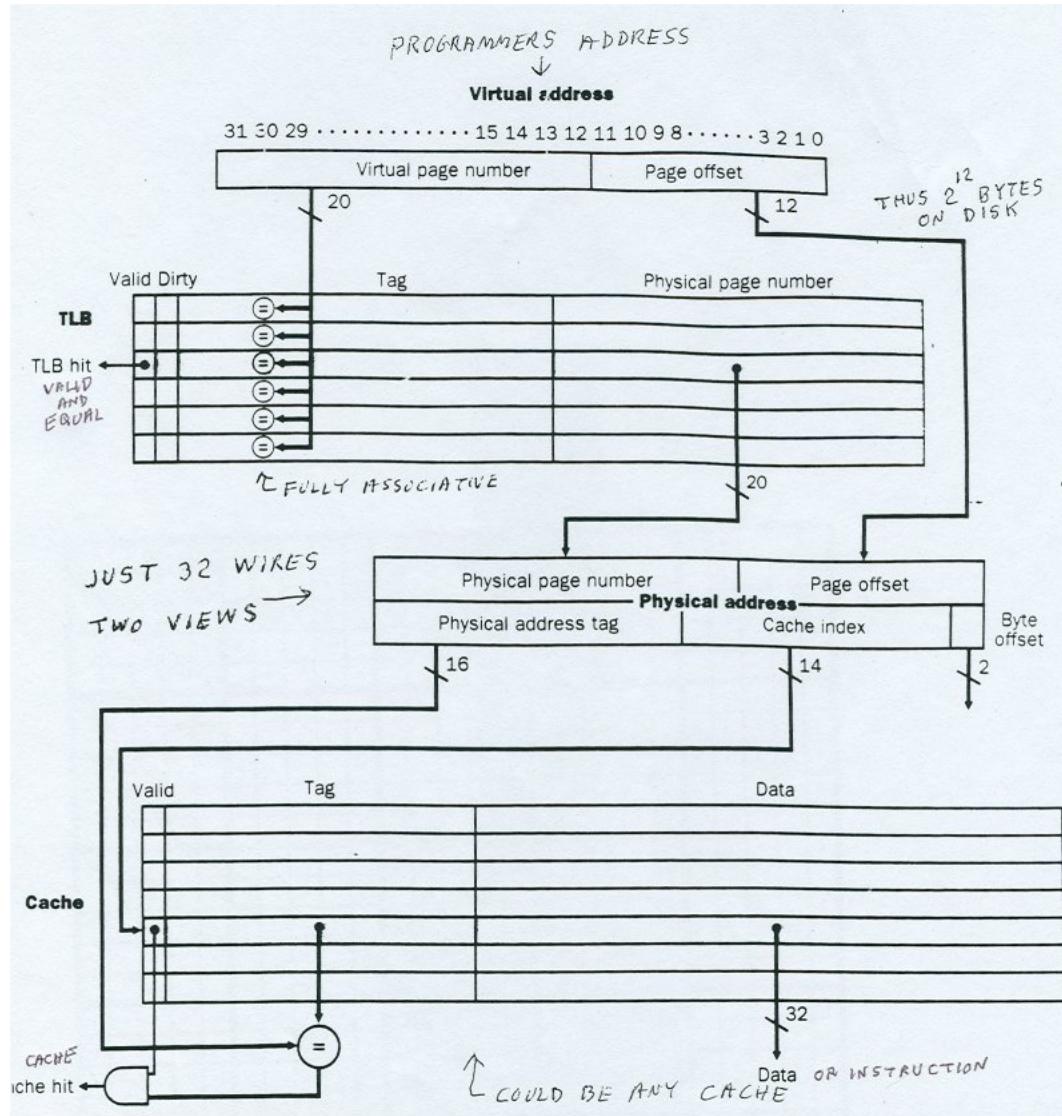
Basically, any virtual page can be in any physical page.

Code and data segments may not all be in physical memory.

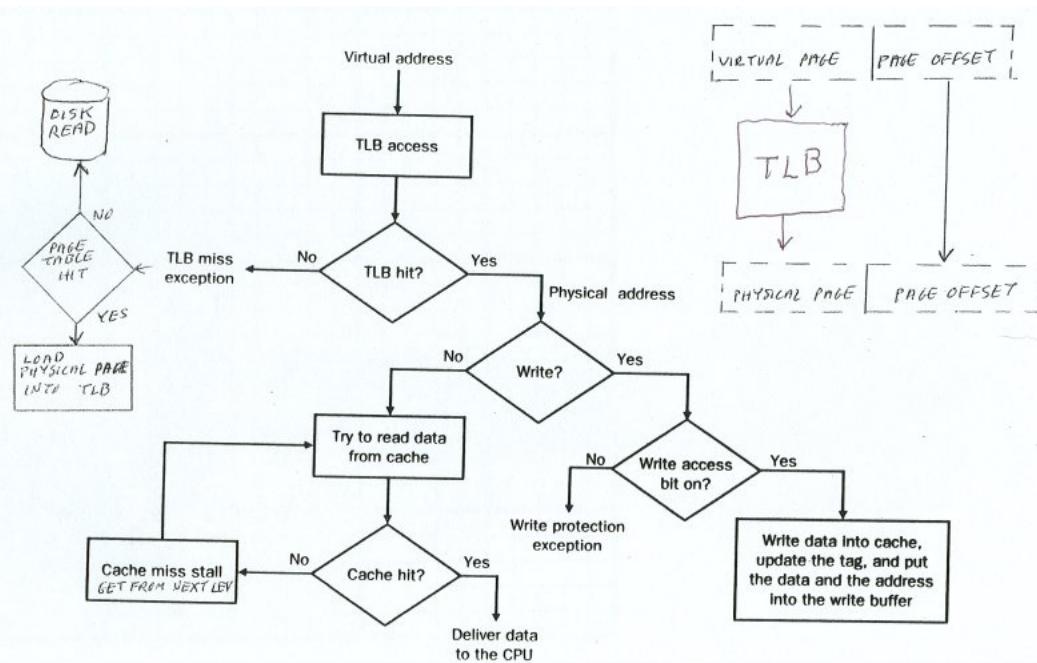
A TLB attached to a cache. Any cache could be used,

a simple one word per block cache is shown.

Note that the TLB is fully associative.



A flow diagram showing the logical steps to get from an executable programs virtual address to a physical address that can access RAM.



Note that a TLB is a cache yet it typically has some extra complexity. In addition to the valid bit there may be a "read only" bit that can easily prevent a store operation into a page. Another bit may be an "execute only" bit for instruction pages that prevents both load and store operations.

A required bit is a "dirty" bit. Consider a page that is referenced: The page must be loaded from disk or may be a created page of zeros in RAM. Then eventually that page in RAM is needed for some process. If any store operation changed that page in RAM, the page must be written out to disk. The page is "dirty" meaning changed. If the page in RAM is not dirty, the new page information just over writes the physical page in RAM with some other page.

A significant performance requirement for the operating system is to efficiently handle paging. If there are no physical pages on the OS free page list, a Least Recently Used, LRU, strategy is typically used to choose a page to over write.

The specific architecture of the TLB must be known in order to compute the number of bits of storage needed.

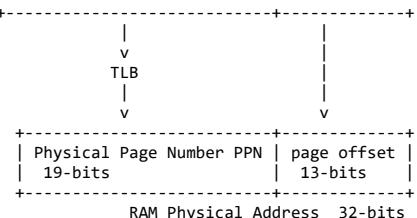
Given: a 36-bit virtual address,  
a 32-bit physical address,  
a 8192 byte page:

Compute:

$\log_2 8192 = 13$ -bit page offset. ( $2^{13}=8192$ )  
Thus the VPN is 36-13 = 23-bits  
the PPN is 32-13 = 19-bits

or, drawn

Programmers Virtual Address 36-bits	
+-----+-----+	
Virtual Page Number VPN   page offset	
23-bits   13-bits	



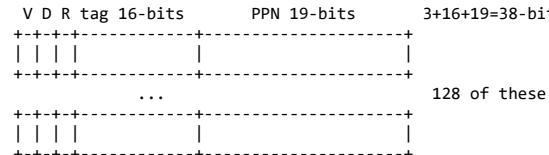
Given 128 blocks in the TLB,  
3 bits for valid, dirty and ref

Compute:

$$\log_2 128 = 7 \text{ bits in TLB index}$$

$$\text{VPN} = 23 - 7 \text{ bit index gives 16 bits in TLB tag}$$

or drawn



thus  $38 * 128 = 4864$  bits in TLB.

Now, given a simple page table is used, indexed by VPN,  
the page table has  $2^{\text{VPN}} = 2^{23} = 8,388,608$  entries.

Given a page table with three control bits V,D and R  
and a Physical Page Number then the page table needs

$$1 + 1 + 1 + 19 = 22 \text{ bits.}$$

$$\text{Total bits } 22 * 8,388,608 = 184,549,376 \text{ bits.}$$

Using power of 10,  $184 * 10^6$ , 184 million bits, Each process requires a page table. Fortunately, the OS uses intelligence and only builds a page table big enough for the size of the program or possibly for only the pages that are actually used. The page table itself is in a page and may, of course, be paged out. :)

A reminder on bits in address vs size of storage:

bits	size	approximate
10	kilobyte	$2^{10}$
20	megabyte	$2^{20}$
30	gigabyte	$2^{30}$
40	terabyte	$2^{40}$
50	petabyte	$2^{50}$
60	exabyte	$2^{60}$

Actually, modern computers use a hierarchy of page tables.

# Alpha VM Mapping

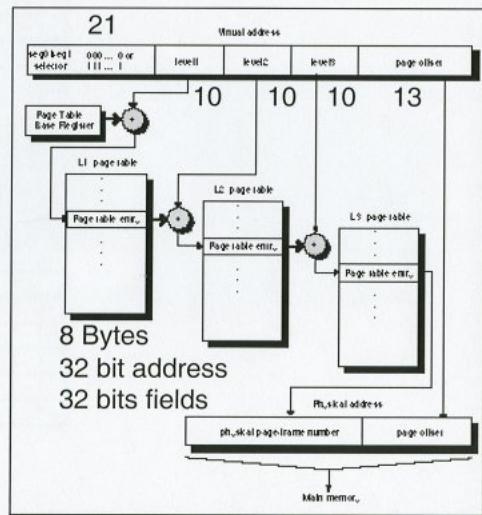
- “64-bit” address divided into 3 segments

- seg0 (bit 63=0) user code/heap
- seg1 (bit 63 = 1, 62 = 1) user stack
- kseg (bit 63 = 1, 62 = 0)  
kernel segment for OS

- Three level page table, each one page

- Alpha only 43 unique bits of VA
- (future min page size up to 64KB => 55 bits of VA)

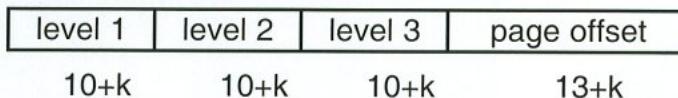
- PTE bits; valid, kernel & user read & write enable (No reference, use, or dirty bit)



# Expanding Alpha Address Space

## Increase Page Size

- Increasing page size 2X increases virtual address space 16X
  - 1 bit page offset, 1 bit for each level index

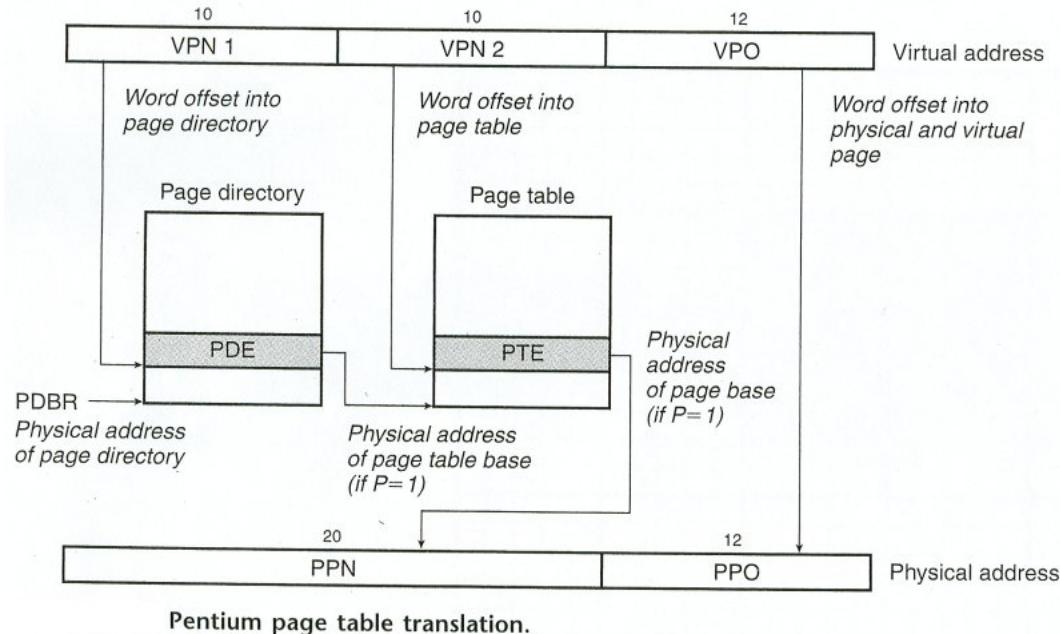


## Physical Memory Limits

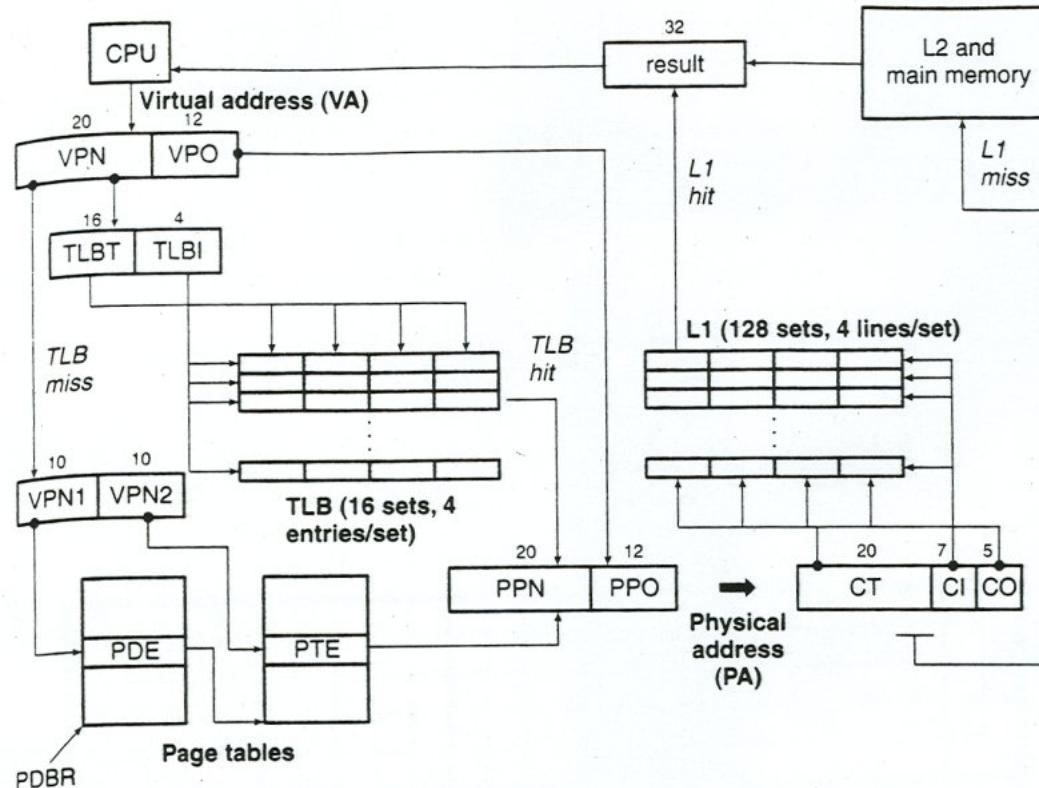
- Cannot be larger than kseg
  - VA bits  $-2 \geq PA$  bits
- Cannot be larger than 32 + page offset bits
  - Since PTE only has 32 bits for PPN

## Configurations

Page Size	8K	16K	32K	64K
VA Size	43	47	51	55
PA Size	41	45	47	48



Pentium page table translation.



Summary of Pentium address translation.

See [Homework 10](#)

## Lecture 24, Virtual Memory 2

This lecture covers the software interface to the computer architecture. Note that Unix was around many years before MS DOS and MS Windows, thus similar capability.

Just a little history from the current man page for `gcc`.  
 Note: The term "text" and "text segment" are instructions, executable code.

From `man gcc` then `/segment`

`-f writable-strings`  
 Store string constants in the writable data segment and don't unique them. This is for compatibility with old programs which assume they can write into string constants.

Writing into string constants is a very bad idea; "'constants'" should be constant.

This option is deprecated.

**-fconserve-space**

Put uninitialized or runtime-initialized global variables into the common segment, as C does. This saves space in the executable at the cost of not diagnosing duplicate definitions. If you compile with this flag and your program mysteriously crashes after "main()" has completed, you may have an object that is being destroyed twice because two definitions were merged.

This option is no longer useful on most targets, now that support has been added for putting variables into BSS without making them common.

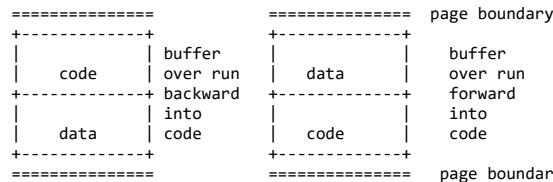
**-msep-data**

Generate code that allows the data segment to be located in a different area of memory from the text segment. This allows for execute in place in an environment without virtual memory management. This option implies -fPIC.

**-mno-sep-data**

Generate code that assumes that the data segment follows the text segment. This is the default.

in same page, better      more likely bad



Best if code and data not in same page

The page can then be "read-only" or "execute-only"

**-mid-shared-library**

Generate code that supports shared libraries via the library ID method. This allows for execute in place and shared libraries in an environment without virtual memory management. This option implies -fPIC.

We will see -fPIC is used directly, below.

Now, consider an operating system that allocated physical pages, via the TLB:

- 1) that contained only code - set to execute only or read only
- 2) that contained constant data - set to read only
- 3) that contained variables, including stack and heap - writable

Any virus or Trojan that tried to overwrite code would be trapped. No possible "buffer overrun" or other malicious action could occur.

But, today's operating systems may put both code and variables into the same physical page. This is most common with .so and .dll files. Thus, the hacker can cause data to be written over your programs instructions. What is written are the harmful instructions to erase your hard drive or do other damage. This is a legacy OS code problem that dates back to small core memory systems. There does not seem to be a willingness to fix this, currently, dangerous situation.

e.g. How could displaying a .jpg image allow a virus?

Oh! Because some idiot believed the size in the header and kept reading data that overwrote instructions.

Double Yuk! 1) Not checking size 2) code and data in same segment

Thus, they helped create cybercrime and thus cyberdefense.

As a part of MS Windows is DOS, now often called a command window or command prompt. Just typing "help" lists most available commands.

Different names for similar file types and commands are:

Unix, Linux, MacOSX	MS Windows	description
.o <no extension>	.obj	relocatable object file
.so	.exe	executable file
.a	.dll	shared object, dynamic link load
	.lib	library of relocatable object files
.c	.c	"C" source file
gcc -c xxx.c	cl /C xxx.c	just make relocatable object file
ar -crv libxxxx.a	cl /LD xxxx.lib	build a library file of many relocatable object files
-lxxxx	xxxx.lib	use library file

An example of building a self contained executable from a .a library and an executable that needs a shared object .so available:

A self contained executable can be distributed as a single file for a specific operating system.

An executable file that links to .so or .dll files will be much smaller and only one copy of the .so or .dll file needs to be in RAM, even when many executable programs need them.

The .so or .dll files must be distributed with the executable file.

First, the main programs and the four little C library functions that print their name in execution:

```
/* ax.c for libax.a test */
#include <stdio.h>
int main()
{
    printf("In ax main \n");
    abc();
    xyz();
    return 0;
}

/* abc.c for libax.a test */
#include <stdio.h>
void abc()
{ printf("In abc \n"); }

/* xyz.c for libax.a test */
#include <stdio.h>
void xyz()
{ printf("In xyz \n"); }

/* ab.c for libab.so test */
#include <stdio.h>
int main()
{
    printf("In ab main \n");
    aaa();
    bbb();
    return 0;
}

/* aaa.c for libab.so test */
#include <stdio.h>
void aaa()
{ printf("In aaa \n"); }

/* bbb.c for libab.so test */
#include <stdio.h>
void bbb()
{ printf("In bbb \n"); }

Then, the Makefile_so
# Makefile_so demo ar and ld and shared library .so
```

```

all: ax ab

ax : ax.c abc.c xyz.c
      gcc -c abc.c          # compile for library
      gcc -c xyz.c
      ar crv libax.a abc.o xyz.o # build library
      ranlib libax.a
      rm -f *.o
      gcc -o ax ax.c -L. -lax    # use library libax.a
      ./ax                         # execute

ab : ab.c aaa.c bbb.c
      gcc -c -fpic -shared aaa.c # compile for library
      gcc -c -fpic -shared bbb.c
      ld -o libab.so -shared aaa.o bbb.o -lm -lc
      rm -f *.o
      gcc -o ab ab.c -L. -lab   # use links to library
      ./ab # need LD_LIBRARY_PATH to include this directory
            # many users have "." meaning "here" "this directory" in path

abg : ab.c aaa.c bbb.c # uses /usr/local/lib needs root priv
      gcc -c -fpic -shared aaa.c
      gcc -c -fpic -shared bbb.c
      ld -o libab.so -shared aaa.o bbb.o -lm -lc
      rm -f *.o
      cp libab.so /usr/local/lib   # install for all users
      rm -f libab.so
      ldconfig
      gcc -o abg ab.c -lab       # any user can get libab.so
      ./abg # any user has access to libab.so

clean:
      rm -f ax
      rm -f ab
      rm *.a
      rm *.so

```

To see what is inside, gcc -S -g3 ax.c

[ax.s](#)

Here are some examples of addressing as seen in assembly code and .o or .obj files. Then in executable a.out or .exe files as seen through the debugger. The "relocatable" addresses are converted to "virtual" addresses then during execution converted to "physical" or RAM addresses. Coming soon to a WEB page near you.

To get memory map, yuk, output, add -Ml,-M to gcc -o ... command

[ax.map](#)

Remember, those huge addresses are virtual addresses.  
Your program may run with much smaller physical memory.

## Information that might help with Project part3

Some are ready to implement part3 of the project.

[Part3 description](#).

You may use a complete behavioral solution, just code the hit/miss process you did by hand in Homework 9, 2a. This may be based on the code below.

Put the caches inside the instruction memory, part3a, and and data memory, part3b, components (entity and architecture).

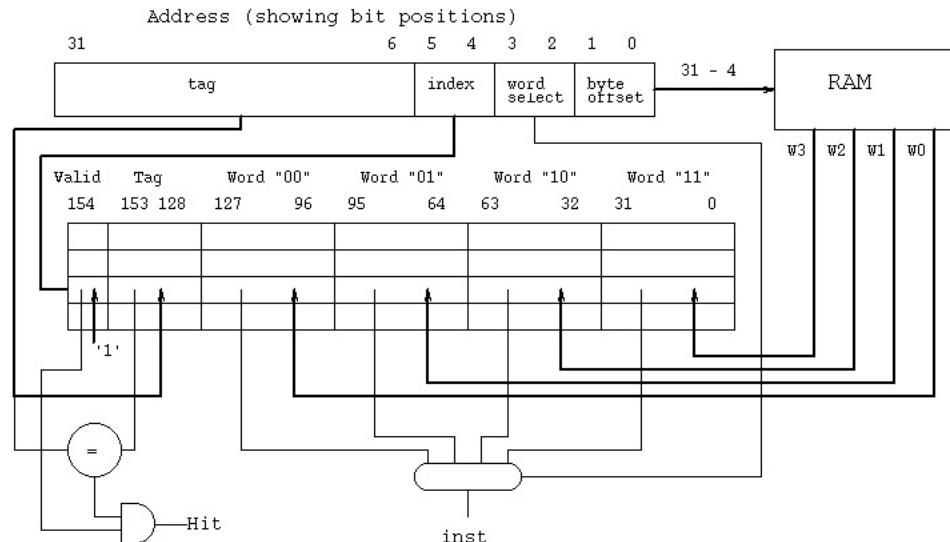
(you will need to pass a few extra signals in and out)

Use the existing shared memory data as the main memory.  
 Make a miss on the instruction cache cause a three cycle stall.  
 Make a miss on the data cache cause a three cycle stall.  
 Previous stalls from part2b must still work.

Both instruction cache and data cache hold 16 words  
 organized as four blocks of four words. Remember vhdl  
 memory is addressed by word address, the MIPS/SGI memory  
 is addressed by byte address and a cache is addressed by  
 block number.

The cache schematic for the instruction cache was handed out  
 in class and shown in: [icache.jpg](#)

The cache may be implemented using behavioral VHDL, basically  
 writing sequential code in VHDL or by connecting hardware.



### Instruction Cache for project part 3a

Possible behavioral, not required, VHDL to set up the start of a cache:  
 (no partial credit for just putting this in your cache.)

```
-- add in or out signals to entity instruction_memory as needed
-- for example, 'clk' 'clear' 'miss'

architecture behavior of instruction_memory is
  subtype block_type is std_logic_vector(154 downto 0);
  type cache_type is array (0 to 3) of block_type;
  signal cache : cache_type := (others=>(others=>'0'));
  -- now we have a cache memory initialized to zero
begin -- behavior
  inst_mem:
  process ... -- whatever, does not have to be just 'addr'
    variable quad_word_address : natural; -- for memory fetch
    variable cblock : block_type; -- the shaded block in the cache
    variable index : natural; -- index into cache to get a block
    variable word : natural; -- select a word
  begin
    if (addr <= 15) then
      cblock := cache(index);
      word := cblock(word_index);
      if (word <= 15) then
        quad_word := word;
      else
        quad_word := 0;
      end if;
    else
      quad_word := 0;
    end if;
  end process;
end;
```

```

variable my_line : line;      -- for debug printout
variable W0 : std_logic_vector(31 downto 0);
...
begin
  ...
  index := to_integer(addr(5 downto 4));
  word  := to_integer(addr(3 downto 2));
  cblock := cache(index);    -- has valid (154), tag (153 downto 128)
  -- W0 (127 downto 96), W1(95 downto 64)
  -- W2(63 downto 32), W3 (31 downto 0)
  -- cblock is the shaded block in handout
  ...
  quad_word_address := to_integer(addr(13 downto 4));
  W0 := memory(quad_word_address*4+0);
  W1 := memory(quad_word_address*4+1); -- ...
  -- fill in cblock with new words, then
  cache(index) <= cblock after 30 ns; -- 3 clock delay
  miss <= '1', '0' after 30 ns;       -- miss is '1' for 30 ns
  ...
  -- the part3a.chk file has 'inst' set to zero while 'miss' is 1
  -- not required but cleans up the "diff"

debug: process -- used to show cache
  variable my_line : LINE;      -- not part of working circuit
begin
  wait for 9.5 ns;           -- just before rising clock
  for I in 0 to 3 loop
    write(my_line, string("line="));
    write(my_line, I);
    write(my_line, string(" V="));
    write(my_line, cache_ram(I)(154));
    write(my_line, string(" tag="));
    hwrite(my_line, cache_ram(I)(151 downto 128)); -- ignore top bit
    write(my_line, string(" w0="));
    hwrite(my_line, cache_ram(I)(127 downto 96));
    write(my_line, string(" w1="));
    hwrite(my_line, cache_ram(I)(95 downto 64));
    write(my_line, string(" w2="));
    hwrite(my_line, cache_ram(I)(63 downto 32));
    write(my_line, string(" w3="));
    hwrite(my_line, cache_ram(I)(31 downto 0));
    writeline(output, my_line);
  end loop;
  writeline(output, my_line); -- blank line
  wait for 0.5 ns;           -- rest of clock
end process debug;

end architecture behavior; -- of cache_memory

For debugging your cache, you might find it convenient to add
this 'debug' print process inside the instruction_memory architecture:
Then diff -iw part3a.out part3a_print.chk

debug: process -- used to print contents of I cache
  variable my_line : LINE;      -- not part of working circuit
begin
  wait for 9.5 ns;           -- just before rising clock
  for I in 0 to 3 loop
    write(my_line, string("line="));
    write(my_line, I);
    write(my_line, string(" V="));
    write(my_line, cache(I)(154));
    write(my_line, string(" tag="));
    hwrite(my_line, cache(I)(151 downto 128)); -- ignore top bits
    write(my_line, string(" w0="));
    hwrite(my_line, cache(I)(127 downto 96));
    write(my_line, string(" w1="));
    hwrite(my_line, cache(I)(95 downto 64));
    write(my_line, string(" w2="));
    hwrite(my_line, cache(I)(63 downto 32));
    write(my_line, string(" w3="));
    hwrite(my_line, cache(I)(31 downto 0));
    writeline(output, my_line);
  end loop;

```

```
end loop;
wait for 0.5 ns;      -- rest of clock
end process debug;
```

see [part3a.print.chk](#) with debug

You may print out signals such as 'miss' using `prtmiss` from.  
[debug.txt](#)

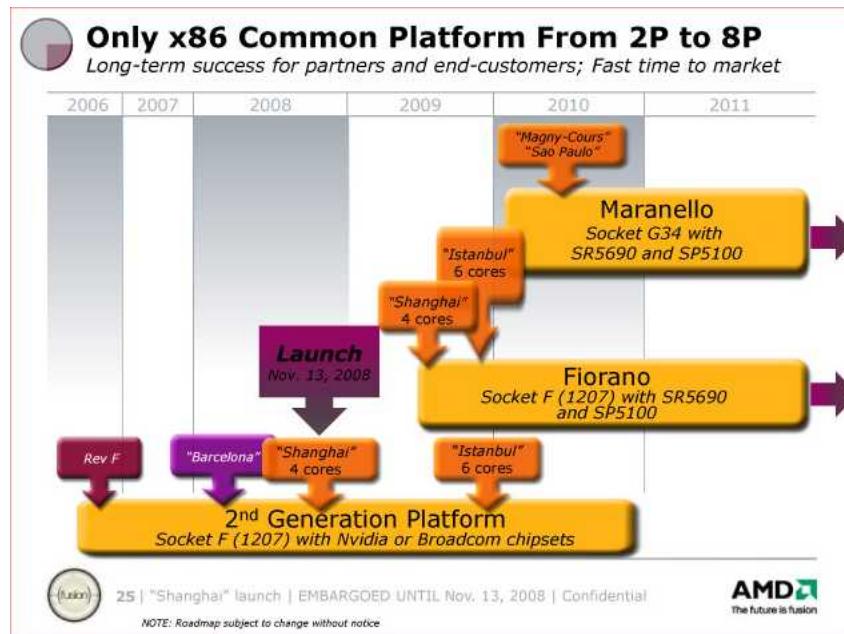
Change `MEMread : std_logic := '1';` to  
`MEMread : std_logic := '0';` for part3b.

You submit on GL using: `submit cs411 part3 part3a.vhd`

Do a write through cache for the data memory.  
 (It must work to the point that results in main memory are correct at the end of the run and the timing is correct, partial credit for partial functionality)  
 You submit this as `part3b.vhd`

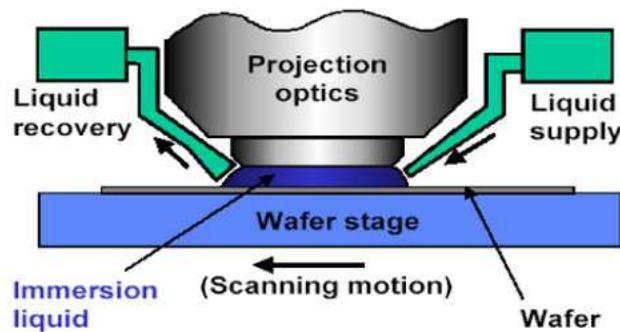
Cache hierarchy on a multiple core CPU.

AMD quad core to six core to shared memory.  
 17.6 GBs front side bus, DDR-800 RAM



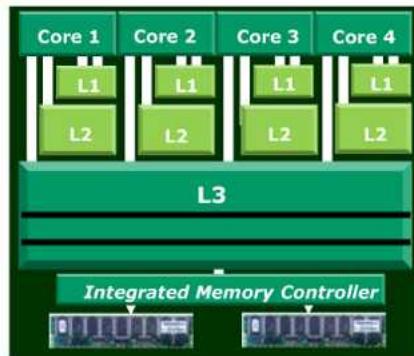
The AMD Shanghai is the first AMD chip built on the company's 45-nanometer manufacturing process, and chips in this series will have higher clock speeds, larger Level 3 cache sizes and better power efficiency compared with the previous generation of 65-nm Opteron processors. The first of these 45-nm Opteron processors is slated for two- and four-way systems. New 45-nm processors for single-socket servers and desktops will follow in 2009.

# 45nm with Immersion Lithography



Greater Frequency<sup>1</sup>

## L3 Cache Index Disable<sup>3</sup>

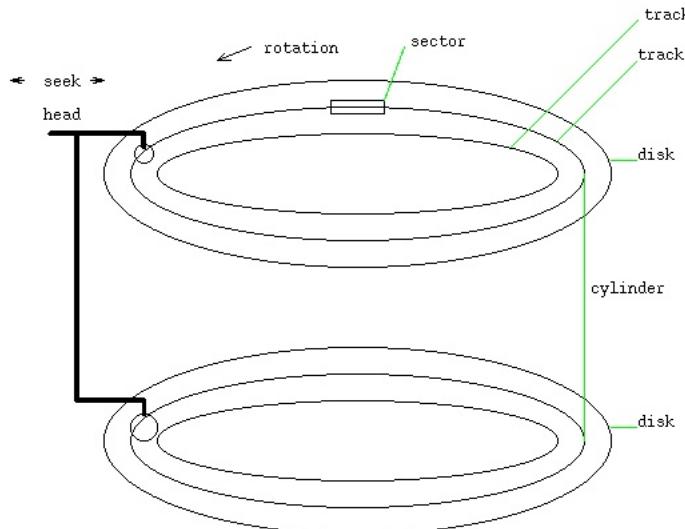


Superior Data Integrity

[part3b](#)

## Lecture 25, I/O types and performance

Take a look inside the hard drive being passed around.



Mine is bigger than yours.

#### How fast can you read a block of data?

There are four time components that must be known to answer this question.

- 1) The time for the read head to get to the required track.  
This is seek time.
- 2) The time for the disk to rotate to start reading the first byte.  
This is the rotational delay time.
- 3) The time to transfer the data from the disk to your RAM.  
This is the transfer time.

#### Gigabytes galore

Another hard drive milestone recedes into the rearview with the \$399 [Hitachi Deskstar 7K1000](#). The first one-terabyte (1TB) desktop-PC hard drive, this 3.5-inch drive spins at 7,200rpm and features a 3Gbps Serial ATA interface, as well as 32MB of buffer memory. At 40 cents per gigabyte, the cost per unit of storage is competitive with those of much-lower-capacity drives. The only question: How long will it take you to fill it up? [Hitachi Global Storage Technologies](#), [www.hitachigst.com](http://www.hitachigst.com)

- 4) Overhead that can be from software, application, OS or drivers.  
 This is overhead time.

### Seek time

The head may be on any track, thus there is seek time before any data can be read. The manufacturers published average seek time is standardized at the time to go from track 0 to the middle track, measured in milliseconds. In the 1990's the size of disk had become large enough such that the measured average seek time was 1/4 the published average seek time. We use 1/4 the published average seek time for our homework and exams. For your computer, having a hard drive with capacity over 120GB, I suggest using 1/8 the published average seek time for your estimates. The reason is that the files you are working with tend to cluster, thus you rarely will have a seek traveling 1/4 the tracks on the disk. For my example below, the published average seek time was 5.4 ms and thus  $5.4/4 = 1.4$  ms is used.

### Rotational delay time

The disk is spinning at a known Revolutions Per Minute, RPM. We deal in seconds, thus divide the RPM by 60 to get Revolutions Per Second, RPS.

How long, on average, does it take for the read head to reach data? This is the rotational delay time and only depends on the RPS. On average the time will be the time for 1/2 of a revolution, thus  $1/2 * 1/\text{RPS}$ . Typically expressed in milliseconds, ms. Some values are:

RPM	RPS	$1/2 * 1/\text{RPS}$	seconds	milliseconds
3600	60	0.00833	8.33	
5400	90	0.00555	5.55	
7200	120	0.00416	4.16	
10,025	167	0.00299	3.00	
15,000	250	0.00200	2.00	

### Transfer time

The time to transfer data depends on the bandwidth, typically given in Megabytes per second. The disk drive has internal RAM and usually can deliver a continuous stream of bytes at near the maximum transfer rate. The transfer may be slowed by your computers system bus or your RAM or other contention for the system bus to RAM path. The example below uses an 80MB/s transfer rate. Thus 80MB can be transferred in one second.

### Overhead time

The overhead time is estimated. 0.6ms

### Example

How long does it take to read a file from disk? (example calculation)

```
time = average seek time +
      average rotational delay +
      transfer time +
```

```

overhead

published average seek = 5.4 ms
"average" seek = 5.4/4      = 1.4ms

10,025 RPM or 167 RPS
1/2 * 1/167 = .00299 sec   = 3.0ms

Overhead assumed           = 0.6ms

Size independent delay, sum= 5.0ms

At 80 MB/sec transfer rate:

10KB   100KB   1MB   10MB

0.125  1.25   12.5  125. transfer time in ms
5.0     5.0    5.0   5.0
____   ____   ____   ____

5.125  6.25   17.5  130.0 ms

```

This is a one block "first read"  
The next read could be buffered

Notice that on small files, the latency (times 1) 2) and 3) dominate.  
On large files the transfer time dominates. Today, files in the  
tens of megabytes are common. Many years ago most files were around  
10 kilobytes. Today 1 to 10 megabyte is typical.

A benchmark I ran on reading 1KB, 10KB, 100KB, and 1MB of data  
from a 10MB file.

```

/* time_io.c check how much is cached in ram          */
/*          assumed pre-existing data file time_io.dat */
/*          created by running time_io_init            */
#include <stdio.h>
#include <time.h>
int main()
{
    FILE * handle;
    int i;
    int j;
    double cpu;
    char buf[1000000]; /* 1MB */
    int check;
    int n = 10000; /* number of reads on 10MB file for buf1*/
    int k = 1000; /* number of bytes read per read */

    printf("time_io.c 10MB file, read 1KB, 10KB, 100KB, 1MB \n");
    handle = fopen("time_io.dat","rb");
    printf("On rebooted machine, first read \n");
    cpu = (double)clock()/(double)CLOCKS_PER_SEC;
    for(i=0; i<n; i++)
    {
        check = fread(buf, k, 1, handle);
        if(check != buf[1]) printf("check failed \n");
    }
    cpu = (double)clock()/(double)CLOCKS_PER_SEC - cpu;
    fclose(handle);
    printf("first read time %g seconds \n", cpu);

    for(n=10000; n>=10; n=n/10)
    {
        printf("more reads, cached? consistent? \n");

        for(j=2; j<10; j++)
        {
            handle = fopen("time_io.dat","rb");
            cpu = (double)clock()/(double)CLOCKS_PER_SEC;
            for(i=0; i<n; i++)
            {
                check = fread(buf, k, 1, handle);
                if(check != buf[1]) printf("check failed \n");

```

```

    }
    cpu = (double)clock()/(double)CLOCKS_PER_SEC - cpu;
    fclose(handle);
    printf("%d read time %g seconds for %dKB block \n", j, cpu, k/1000);
}
k = k*10;
}
return 0;
} /* end time_io.c */

```

One computers output:  
time\_io.c 10MB file, read 1KB, 10KB, 100KB, 1MB  
On rebooted machine, first read  
first read time 0.12 seconds  
more reads, cached? consistent?  
2 read time 0.06 seconds for 1KB block  
3 read time 0.06 seconds for 1KB block  
4 read time 0.06 seconds for 1KB block  
5 read time 0.06 seconds for 1KB block  
6 read time 0.06 seconds for 1KB block  
7 read time 0.06 seconds for 1KB block  
8 read time 0.06 seconds for 1KB block  
9 read time 0.05 seconds for 1KB block  
more reads, cached? consistent?  
2 read time 0.05 seconds for 10KB block  
3 read time 0.05 seconds for 10KB block  
4 read time 0.04 seconds for 10KB block  
5 read time 0.05 seconds for 10KB block  
6 read time 0.05 seconds for 10KB block  
7 read time 0.05 seconds for 10KB block  
8 read time 0.05 seconds for 10KB block  
9 read time 0.05 seconds for 10KB block  
more reads, cached? consistent?  
2 read time 0.08 seconds for 100KB block  
3 read time 0.07 seconds for 100KB block  
4 read time 0.09 seconds for 100KB block  
5 read time 0.07 seconds for 100KB block  
6 read time 0.07 seconds for 100KB block  
7 read time 0.06 seconds for 100KB block  
8 read time 0.08 seconds for 100KB block  
9 read time 0.08 seconds for 100KB block  
more reads, cached? consistent?  
2 read time 0.09 seconds for 1000KB block  
3 read time 0.09 seconds for 1000KB block  
4 read time 0.09 seconds for 1000KB block  
5 read time 0.09 seconds for 1000KB block  
6 read time 0.11 seconds for 1000KB block  
7 read time 0.10 seconds for 1000KB block  
8 read time 0.10 seconds for 1000KB block  
9 read time 0.10 seconds for 1000KB block

Why did I reboot to run a file read test?  
On a computer that is not shut down, a file could  
remain in RAM and even partially in cache for days  
to weeks, if you were not using the computer.

By now you should know that I do a lot of benchmarking.  
I ran the above program on two computers each with two  
operating systems with three disk types.

Block	2.5GHz	2.5GHz	1GHz	1GHz
Size	P4 ATA 100	P4 ATA 100	ATA 66	SCSI 160
Windows XP	Linux	Windows 98	Linux	
1KB	0.0000015	0.000001	0.000016	0.000004
10KB	0.000015	0.000010	0.000060	0.000035
100KB	0.000150	0.000100	0.000500	0.000300
1MB	0.003100	0.002000	0.005000	0.004000

Fine print: CPU time in seconds, most frequent value of eight  
measurements after first read. Using fopen, fread, binary  
block read. Each measurement read 10MB. e.g 10 blocks read  
at 1MB, 100 blocks read at 100KB, 10,000 blocks read at 1KB.  
Other than the first number that is 1.5 microseconds, the

numbers can be read as integer microseconds.

As expected the SCSI disk was faster than the ATA disk. Note that the faster system clock can allow the actual transfer rate to be near the maximum while a slower clock speed can limit the transfer rate. The operating system, drivers and libraries have some impact on total time. This is lumped into "overhead."

Where do you find the disk specifications? Both the manufacturer and some retailers publish the disk specifications, and some prices.

e.g.

[evolution specs](#)

[2007 hard drives, note cache, RPM, transfer rate](#)

## Then SATA replaced ATA

Serial ATA changed the wiring and protocol. ATA had wide flat cables. Driven by PC manufacturers Dell, Gateway, HP, etc, they needed thinner cables. Thus higher speed transfer over fewer wires. Typical SATA bus maximum transfer rate is 3GB/s, 3 gigabytes per second.

Similar latency, similar seek, faster transfer rate.

A single drive with 500GB of storage became available at reasonable cost.

A terabyte of disk storage became practical for a desktop PC. Now multiple terabyte 6Gb/s disks are available.



Seagate Barracuda 3 TB HDD SATA 6 Gb/s NCQ 64MB Cache 3.5-Inch Internal  
by [Seagate](#)  
 (1,190 customer reviews)

List Price: \$269.99  
Price: **\$132.36 & FREE Shipping.** [Details](#)  
You Save: \$137.63 (51%)

**In Stock.**  
Ships from and sold by [Amazon.com](#) in certified [Frustration-Free Packaging](#). Gift-wrap available.

Want it tomorrow, May 2? Order within **3 hrs 32 mins** and choose **One-Day Shipping** at checkout. [Details](#)

Size: 3TB

Still too slow!

## Now, SSD, Solid State Disks

Replace the rotating disk drive with NAND Flash digital logic storage.

[Technology explanation](#)

[Performance comparisons](#)

[One technical specification](#)

Transcend 128GB \$229.99 TS128GSSD25S-M

[enclosure was needed for desktops, initially](#)

Check for latest size, speed, cost  
[computer-SSD-search SSD](#)

Reworking the example above for time to read a file:

#### Transfer time

The time to transfer data depends on the bandwidth, typically given in Megabytes per second. The example below uses an 80MB/s transfer rate. Thus 80MB can be transferred in one second.

#### Overhead time

The overhead time is estimated. 0.6ms

**No seek time, no rotational delay time, for SSD**

#### Example

How long does it take to read  
 a file from disk? (example calculation)

```
time = transfer time +
      overhead
```

At 80 MB/sec transfer rate:

10KB	100KB	1MB	10MB
0.125	1.25	12.5	125. trans
0.6	0.6	0.6	0.6
—	—	—	—
0.725	1.85	13.1	125.6 ms

This is a one block "first read"  
 The next read could be buffered

Notice that on very small files, the overhead time dominates.  
 On large files the transfer time dominates. Today, files in the tens of megabytes are common. Many years ago most files were around 10 kilobytes.

The SSD has a speedup of 7.07 for a 10KB file.  
 The SSD has a speedup of 1.03 for a 10MB file.

Your mileage may vary.

A typical desktop is executing 4,000,000 instructions per ms, millisecond.

[Homework 11](#)

## Lecture 26, DVR, DVD-RW, CDR, CD-RW

This lecture covers device characteristics and formats  
 of CD's and DVD's

It also covers aspects that bring together technology, business, teaming and public buying patterns.

There are many "ports" that allow CD and DVD connection to a common PC.

Parallel Port, IEEE 1284, about 2.5MB/sec

USB2, Universal Serial Bus, 60MB/sec

USB3, Universal Serial Bus, 600MB/sec some available in 2011

PCI, Peripheral Component Interconnect (bus) 528MB/sec

Firewire, IEEE 1394, 50MB/sec

Firewire, IEEE 1394b, 400MB/sec

Firewire, IEEE 1394c, 800MB/sec

SCSI, Small Computer System Interconnect, 320MB/sec

SCSI, up to 640MB/sec

ATA, Advanced Technology Attachment (commands) 160MB/sec

SATA 150MB/sec to 300MB/sec

SATA 3 to 750MB/sec = 6Gbit/sec

Unfortunately, the fastest DVD's are much slower.

CD and DVD drives can be found for many of these ports.

The "media" is the physical disk and typical names are:

CD a pre recorded disk

CDR a blank disk that can be recorded once

CDRW a blank disk that can be recorded many times

DVD a pre recorded disk

DVD-R a blank disk, dash media, that can be recorded once

DVD+R a blank disk, plus media, that can be recorded once

DVD-RW a blank disk, dash media, that can be recorded many times

DVD+RW a blank disk, plus media, that can be recorded many times

DVD-RAM a blank disk, RAM media, that can be recorded many times

Blu Ray DVD pre recorded or recordable

HD DVD pre recorded or recordable

There are many formats that can be used for CD's

Most of the varieties are audio formats.

There is a VCD, Video CD format.

The digital format is UDF, ISO 9660 compatible

DVD's chose to have only the UDF format

The information on a DVD or CD using UDF is directories

and files similar to any computer file system.

Movies use a set of files in MPEG format within the UDF file system.

In Windows, Windows Explorer or prompt command dir

or in Linux or any Unix, the command ls can be

used to look at the directory structure of the UDF file system.

Here is one such listing. Note required directory name video\_ts  
and required file name video\_ts for a DVD to play a movie.

```
Volume in drive E is ITALIAN_JOB
Volume Serial Number is 4E8F-DF0F
```

```
Directory of E:\
```

```
08/12/2003 03:13 AM
```

```
VIDEO_TS
  0 File(s)      0 bytes
```

```
Directory of E:\VIDEO_TS
```

08/12/2003 03:13 AM

08/12/2003 03:13 AM

```

08/12/2003 03:13 AM          20,480 VIDEO_TS.BUP
08/12/2003 03:13 AM          20,480 VIDEO_TS.IFO
08/12/2003 03:13 AM          909,312 VIDEO_TS.VOB
08/12/2003 03:13 AM          18,432 VTS_01_0.BUP
08/12/2003 03:13 AM          18,432 VTS_01_0.IFO
08/12/2003 03:13 AM          268,288 VTS_01_0.VOB
08/12/2003 03:13 AM          10,248 VTS_01_1.VOB
08/12/2003 03:13 AM          22,528 VTS_02_0.BUP
08/12/2003 03:13 AM          22,528 VTS_02_0.IFO
08/12/2003 03:13 AM          16,521,216 VTS_02_0.VOB
08/12/2003 03:13 AM          387,725,312 VTS_02_1.VOB
08/12/2003 03:13 AM          28,672 VTS_03_0.BUP
08/12/2003 03:13 AM          28,672 VTS_03_0.IFO
08/12/2003 03:13 AM          760,942,592 VTS_03_1.VOB
08/12/2003 03:13 AM          79,872 VTS_04_0.BUP
08/12/2003 03:13 AM          79,872 VTS_04_0.IFO
08/12/2003 03:13 AM          103,512,064 VTS_04_0.VOB
08/12/2003 03:13 AM          1,073,709,056 VTS_04_1.VOB
08/12/2003 03:13 AM          1,073,709,056 VTS_04_2.VOB
08/12/2003 03:13 AM          1,073,709,056 VTS_04_3.VOB
08/12/2003 03:13 AM          1,073,709,056 VTS_04_4.VOB
08/12/2003 03:13 AM          1,073,709,056 VTS_04_5.VOB
08/12/2003 03:13 AM          18,653,184 VTS_04_6.VOB
08/12/2003 03:13 AM          38,912 VTS_05_0.BUP
08/12/2003 03:13 AM          38,912 VTS_05_0.IFO
08/12/2003 03:13 AM          1,073,709,056 VTS_05_1.VOB
08/12/2003 03:13 AM          343,238,656 VTS_05_2.VOB
08/12/2003 03:13 AM          14,336 VTS_06_0.BUP
08/12/2003 03:13 AM          14,336 VTS_06_0.IFO
08/12/2003 03:13 AM          136,196,096 VTS_06_1.VOB
30 File(s) 8,210,677,760 bytes

```

Total Files Listed:

```

30 File(s) 8,210,677,760 bytes
3 Dir(s)      0 bytes free

```

The speeds of CD's and DVD have a large range. Generally they became faster as time passed and more were sold.

CD	DVD
1X = 150KB/sec	1X = 1.38MB/sec
2X = 300KB/sec	2X = 2.76MB/sec
10X = 1.5MB/sec	4X = 5.52MB/sec
20X = 3.0MB/sec	8X = 11 MB/sec
40X = 6.0MB/sec	16X = 22 MB/sec

Much slower than hard drives.

Most drives can read at a higher speed than they can write.

#### Capacity:

The disk capacity for CD's is from 74 to 80 minutes of music or 650MB to 700MB of digital storage in UDF file system.

DVD's have a wider range of storage from 2 to 4 hours of movies or  
 4.7GB single sided single layer  
 7.9GB single sided double layer  
 9.4GB double sided single layer  
 15.9GB double sided double layer

Blu Ray and HD DVD are aiming for 20 to 40 hours of conventional movies or 4 to 8 hours of HDTV, high definition TV, 1080i or hundreds of Gigabytes. The market was not stable for a time, and some technology, business, teaming and buying patterns are covered to show where we are now.

Technical information on CD's and DVD's

[DVD and CD Writing Technology](#)

[cont..](#)

[Reviews](#)

[Burn DVD using Linux](#)

[DVD-RW](#)

[Protecting, who?](#)

[Sony and friends vs. Toshiba and friends](#)

[Blue Ray vs. HD DVD](#)

## Blu-ray Burner Boasts Solid Software Bundle

### OPTICAL DRIVE

#### TEST THE PARADE OF Center

rewritable Blu-ray Disc drives continues with the introduction of Plextor's PX-B900A. I looked at a shipping version of this drive, and it is impressive. At \$999, it costs more than competing models, but you get a well-rounded software bundle in the box.

The Plextor offering writes to most flavors of CD, DVD,

#### PX-B900A

Plextor

PCW Rating  Very Good

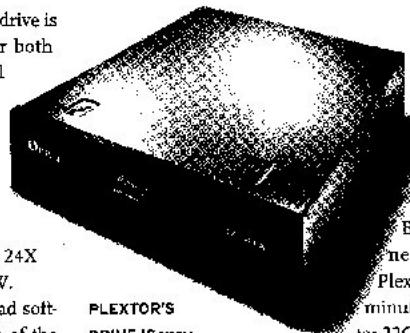
Internal Blu-ray burner offers an excellent software bundle, but its high price limits its appeal.

List: \$999

[find.pcworld.com/55623](http://find.pcworld.com/55623)

and Blu-ray Disc. The drive is rated at 2X speed for both single-layer 25GB and double-layer 50GB write-once BD-R and rewritable BD-RE. It also can handle 8X DVD±R, 8X DVD+RW, 6X DVD-RW, 5X DVD-RAM, 24X CD-R, and 16X CD-RW.

The InterVideo/Ulead software collection is one of the best bundles I've seen accompanying a Blu-ray Disc burner. You get WinDVD BD for watching Blu-ray movie discs, WinDVD for watching standard DVDs, Ulead Video-Recorder 10SE for video editing and disc authoring, DVD MovieFactory 5SE for author-



PLEXTOR'S  
DRIVE IS very  
simple to install.

The drive's write-once performance was in keeping with what we have seen from the other two Blu-ray Disc burners we've reviewed thus far, I-O Data's BRD-UM2/U and Pioneer's BDR-101A: The Plextor PX-B900A took 44 minutes, 35 seconds to master 22GB of data to BD-R, and 99 minutes to format a BD-RE and packet-write 22GB to disc—about the same as the other drives we've tested.

If you're eager for a Blu-ray Disc burner, the Plextor PX-B900A drive is a solid—albeit expensive—choice.

—Melissa J. Perenson

A prototype TDK 200GB blue laser disc would be able to hold a full 18 hours of high-definition video, the company said.

## Visible Light Spectrum

	Wavelength nano meters	Frequency Tera Hertz = $10^{12}$ Hz
Red		620-750
Orange		590-620
Yellow		570-590
Green		495-570
Blue		450-495
Violet		380-450

multiply nano meters by 10 to get Angstroms

1/5/2007

First Combo High-Def DVD Player Announced

Oh man, is CES going to be good! Lots of disruptive products out there, and I'm particularly excited about a new one from LG. The company promises to show off the first combo/hybrid drive for Blu-ray and HD DVD, possibly putting an end to the whole war for good. That's good news for consumers, who have mostly ignored the new discs. Our story wraps up what we know about LG's CES announcements, and also provides analysis as to what it means, and why this is so cool. Check it out, and stay tuned to our CES coverage all next week at [www.pcmag.com/ces](http://www.pcmag.com/ces) for all the breakthroughs.

[Get software to make your own DVD's](#)

[More HV vs Blu ray, gamers view](#)

[Blu ray vs HD DVD origin](#)

**And a final straw. Walmart said it would only carry Blu ray. HD\_DVD faded into oblivion.**

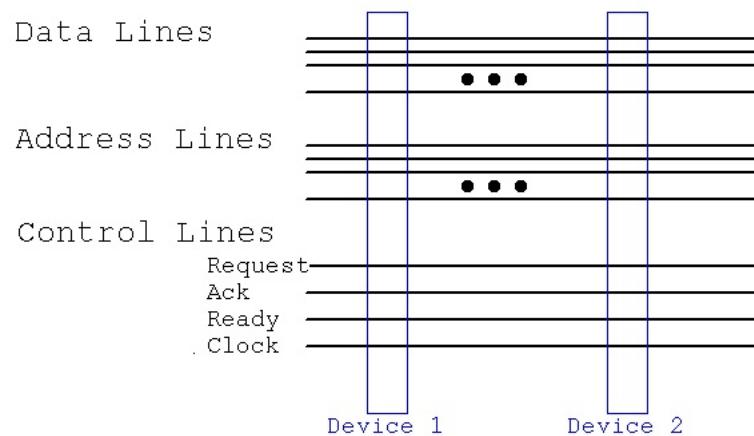
## Lecture 27, Busses, I/O-processor connection

A "bus" is just a number of wires in parallel used to transfer information from one device to another device. The wires may be built into a printed wiring board, PWB, or may be in a flexible cable.

The most important specification for a bus is its protocol. The protocol defines the method for accessing the bus, read requests, write requests, address and data sequencing, etc.

There may be many devices on a bus. In order for all the devices to work together, all must follow the protocol.

A possible bus may have the following sets of lines.



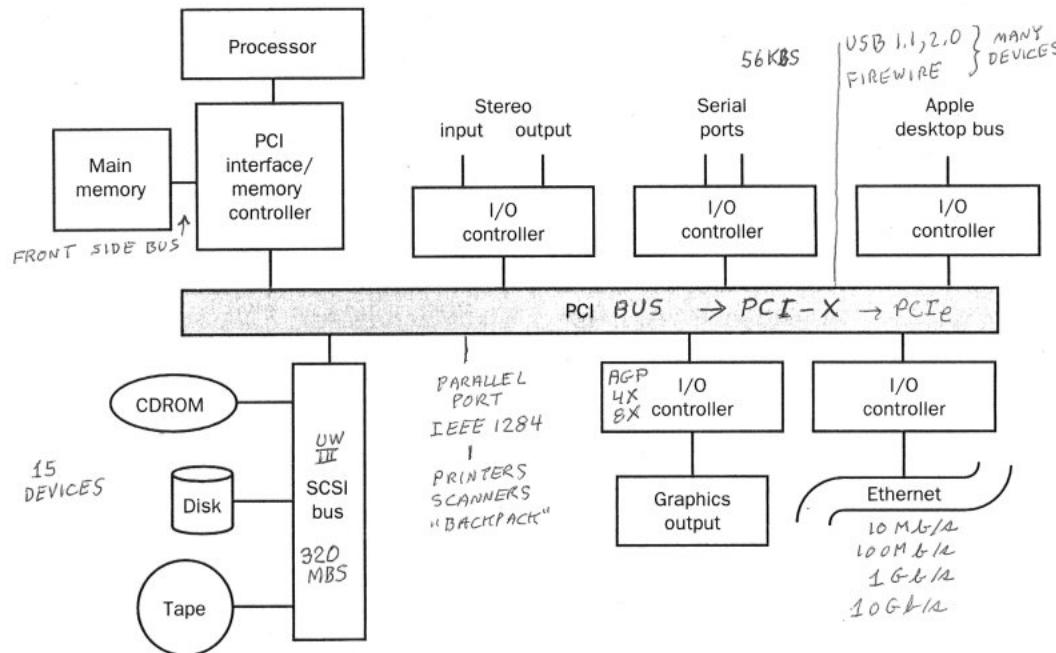
The Control lines are used to implement the protocol.

There may be a bus master, hardware, that arbitrates when two devices want to get on the bus at the same time.

When a bus has a clock, the bus is called synchronous.  
All signals change on rising edge, falling edge or both.

An asynchronous bus is driven at the speed of the device currently driving the bus.

Diagram showing how busses might be connected in a computer:



The bandwidth, speed, of a bus may be measured in bits per second, bps Mbps is  $10^6$  bps, not  $2^{20}$  bps bytes per second, Bps communication is typically powers of 10 megahertz, MHz words per second transactions per second

A transaction is a complete protocol sequence.

An example with time progressing down:

```

Device 1           Device 2
wait for bus available
put address on bus
set request to 1
wait for Ack = 1, acknowledge
wake up because request = 1
save address from bus
set Ack to 1
wait for request = 0

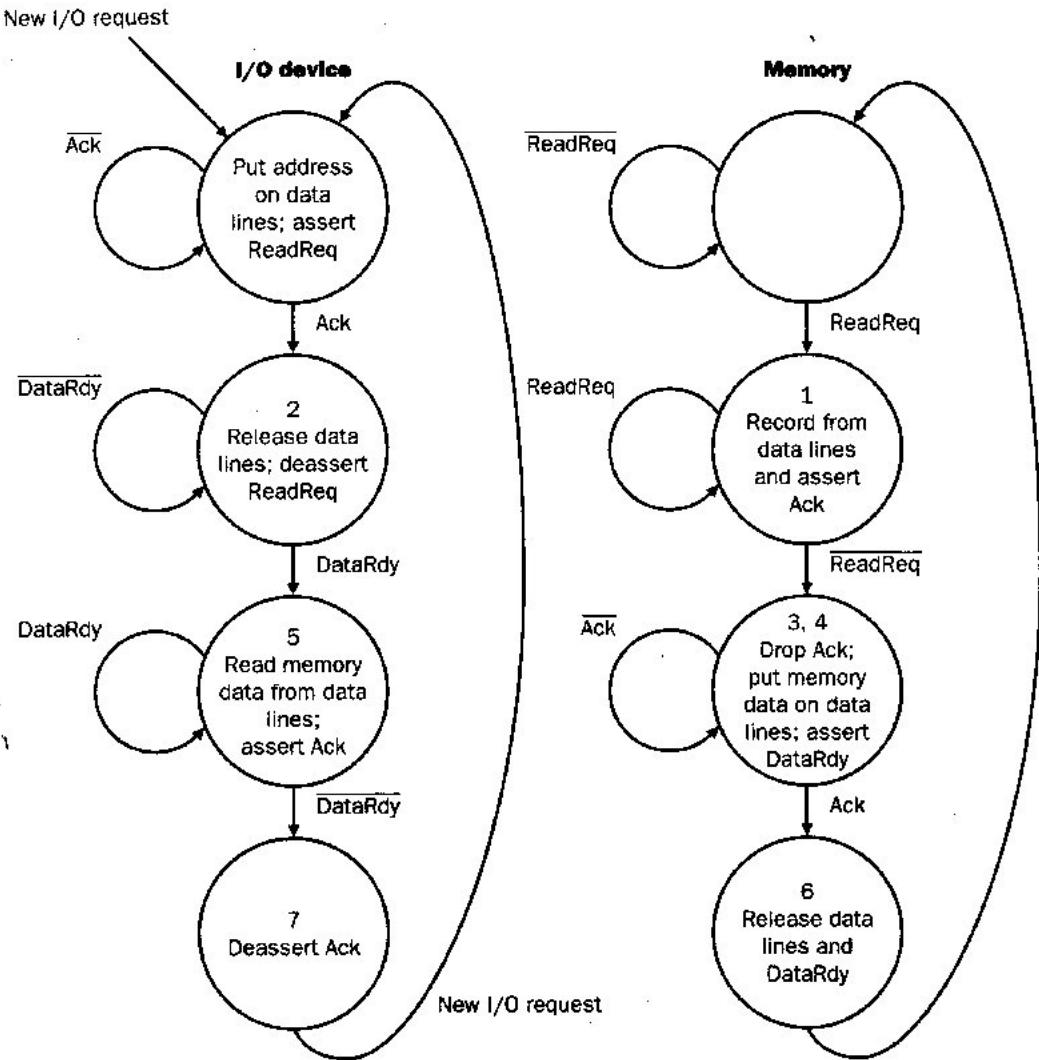
wake up because Ack = 1
release address lines
set request to 0
wait for ready = 1
wake up because request = 0
set Ack to 0
put data on the bus
set ready to 1
wait for Ack = 1

wake up because ready = 1
save data from bus
set Ack to 1
wait for ready = 0
wake up because Ack = 1
release data lines
set ready to 0
finished this transaction

wake up because ready = 0
set Ack to 0
finished this transaction
bus is available

```

Often, the bus protocol is implemented as a Deterministic Finite Automata, DFA. The state diagram for the above protocol could be shown as:



#### Examples of Busses circa 2012 including older (changes with time)

Bus name	Max Mbytes per sec	Max MHz per sec	width	comment	
front side	17,024 34,048 19,200 85,248 136,448 204,800 225,280 256,000 320,000	2,128 4,256 2,400 10,656 17,056 25,600 26,160 32,000 40,000	133 133 150 333 533 800 880 1,000 1,250	128 256 128 256 256 256 256 256 256	many possible

(PPC Mac G5)

	307,200	38,400	1,600	192	(I7 extreme, 1 channel)
AGP	2,112	264	66	32	
AGP8X	17,056	2,132	533	32	
PCI	1,056	132	33	32	
PCI	2,112	264	33	64	
PCI	2,112	264	66	32	
PCI	4,224	528	66	64	
PCI	4,224	528	133	32	
PCI	8,448	1,056	133	64	
PCIX	17,056	2,132	533	32	extended, compatible
PCIe	64,000	8,000	2000	32	express, one way, full duplex 1,2,4,8,12,16 or 32 lanes
ATA 100	800	100	25	32	
ATA 133	1064	133	33	32	
ATA 160	1280	160	40	32	
SATA 150	1200	150	600	2	one way, full duplex
SATA std	1500	187	1500	1	one way, full duplex limited by motherboard
SATA II 300	2400	300	1200	2	
SATA II std	3000	375	3000	1	no forcing to build standard
SATA 3.0	6000	750	6000	1	
SCSI 1	40	5	5	8	
SCSI 2	160	20	10	16	
SCSI 3	1280	160	80	16	
SCSI UW3	2560	320	160	16	
SCSI 320	5120	640	320	16	has cable terminators
Firewire1394	400	50	400	1	
Firewire1394b	800	100	800	1	many video cameras
Firewire S16	1600	200	1600	1	
Firewire S32	3200	400	3200	1	
Firewire S80	6400	800	6400	1	
USB 1.1	12	1.5	12	1	slow
USB 2	480	60	480	1	new cable
USB 3	3200	400	1600	2	new cable, dual differential
	5000	625	2500	2	new connectors, optional speed
	6400	800	3200	2	micro, mini, connectors etc.
Fiberchannel	1000	125	1000	1	1062.5
Fiberchannel	2000	250	2000	1	>mile
Fibre 16GFC	3200	14000			full duplex 10Km
Fibre 20GFC	5100	21000			full duplex
Ethernet 10	10	1.25	10	1	
Ethernet 100	100	12.5	100	1	
Ethernet 1Gig	1000	125	1000	1	
Ethernet 10G	10000	1,250	10000	1	
ISA	400	50	25	16	really old
IEEE 1284 ECP	2.5	0.31	0.31	8	half duplex printer port
V.90 56	0.056	0.005	0.056	1	modem, one way, full duplex
OC-48	2,500				optical cross country
OC-192 STM64	10,000				Optical Carrier
OC-768 STM256	40,000	5,000	light		
	Mbps	Mbps	MHz		

The speed of light limits the amount of information that can be sent over a given distance. Many busses have length restrictions.

Light can travel about

300,000,000 meters per second

300,000 meters per millisecond

300 meters per microsecond

0.3 meters per nanosecond (about 1 foot)

Unchanged in last few decades. (slower inside integrated circuit)

[Pentium 4 busses and PCI-X vs PCIe](#)

Note one example of AGP being replaced by PCI-e and the mention of many "busses" in the advertisement:

**Dual-Screen PCI-E GeForce® 7600 GT Graphics Card from NVIDIA®**



The NVIDIA GeForce 7600 GT PCI Express graphics card provides extreme high-definition gaming and home theater-quality video. Its blazing-fast RAMDACs support dual QXGA displays using resolutions up to 2048x1536 at 85Hz. The card's fourth-generation CineFX engine creates the most advanced visual effects for PC games and other visual applications.

Pixel Shaders.....	12
Vertex Shaders.....	5
Texture Units.....	12
ROPs.....	8
Memory Bus.....	128-bit
Core Speed (MHz).....	560
Mem. Speed (MHz).....	700 (1400 Dual)

**Includes DVI-I to VGA adapter, HDTV adapter and software on CD-ROM**

**Was \$169.  
ON SALE!  
\$144.**

PCI Express Video Card  
BFG Tech GeForce 7600 GT, NVIDIA  
202 0479

[SCSI and printer port, wave forms](#)

For HW12, read the directions carefully. Every bus is different.

[Example of HW12 solution method](#)

[Now you can do HW 12](#)

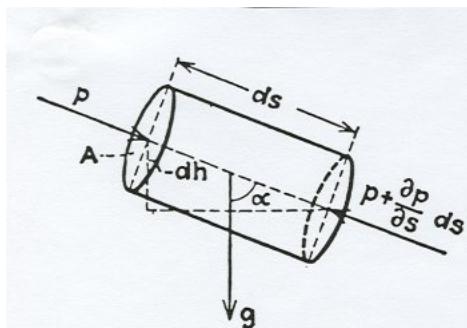
## Lecture 28, Multiprocessors

Classic problems that require multiprocessors:

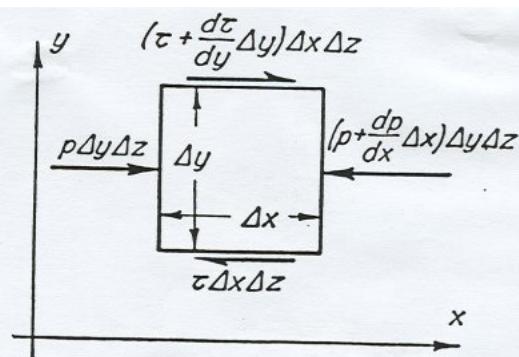
The general differential equations of fluid motion including the effect of viscosity are known as “the equations of Navier-Stokes.” A derivation of them can be found in Chap. XV, “Fundamentals.”<sup>1</sup> For the  $x$ -direction the equation is

$$\frac{\partial u}{\partial t} + u \frac{\partial u}{\partial x} + v \frac{\partial u}{\partial y} + w \frac{\partial u}{\partial z} = X - \frac{1}{\rho} \frac{\partial p}{\partial x} + \frac{\mu}{\rho} \left( \frac{\partial^2 u}{\partial x^2} + \frac{\partial^2 u}{\partial y^2} + \frac{\partial^2 u}{\partial z^2} \right).$$

Two corresponding equations hold for the  $y$ - and  $z$ -directions. It is seen that for no viscosity ( $\mu = 0$ ), the equation of Navier-Stokes (5) reduces to Euler’s equation (3).



—Forces on an element of ideal fluid.



Definition sketch for fluid pressure and shear.

Massive four dimensional arrays and many teraflops of computational power are needed to numerically solve non-trivial fluid problems.

Each cell of fluid has three physical dimensions and a time dimension.

#### Maxwell's Equations

Gauss' Law

$$\vec{\nabla} \cdot \vec{E} = \frac{\rho}{\epsilon_0}$$

Faraday's Law

$$\vec{\nabla} \times \vec{E} = -\frac{\partial \vec{B}}{\partial t}$$

$$\vec{\nabla} \cdot \vec{B} = 0$$

Ampere's Law

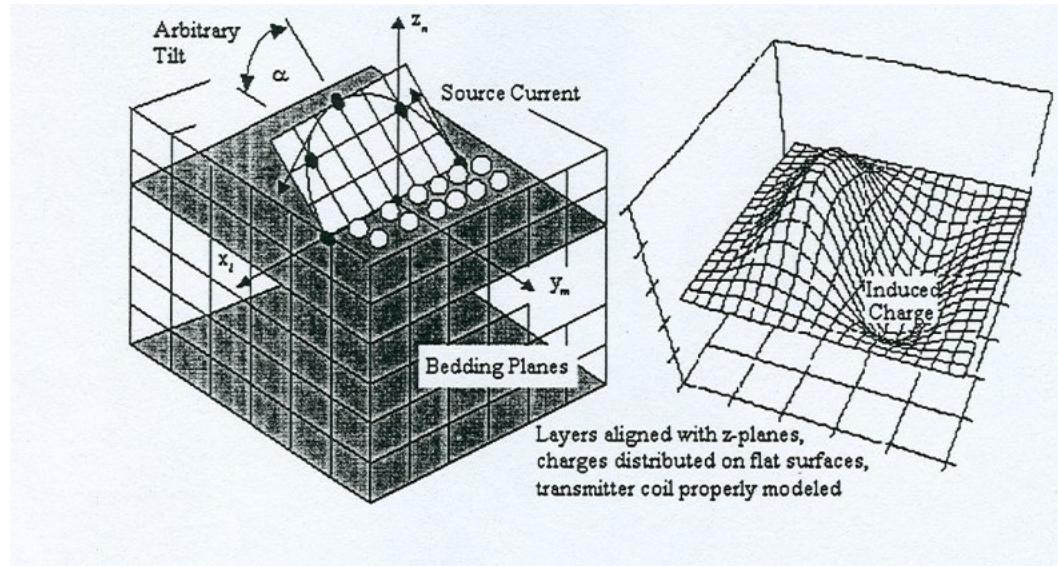
$$\vec{\nabla} \times \vec{B} = \mu_0 \vec{J} + \mu_0 \epsilon_0 \frac{\partial \vec{E}}{\partial t}$$

$$\oint \vec{E} \cdot d\vec{r} = -\frac{d}{dt} \iiint_S \vec{B} \cdot d\vec{A}$$

$$\oint \vec{B} \cdot d\vec{r} = \mu_0 (I + I_d)$$

$$\iint_S \vec{E} \cdot d\vec{A} = Q/\epsilon_0$$

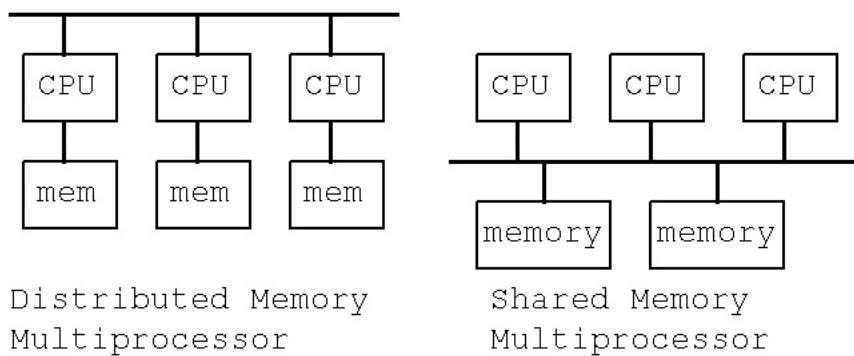
$$\iint_S \vec{B} \cdot d\vec{A} = 0$$



The numerical solution of Maxwell's Equations for electro-magnetic fields may use a large four dimensional array with dimensions X, Y, Z, T. Three spatial dimensions and time. Relaxation algorithms map well to a four dimensional array of parallel processors.

#### A 4D 12,288 node supercomputer

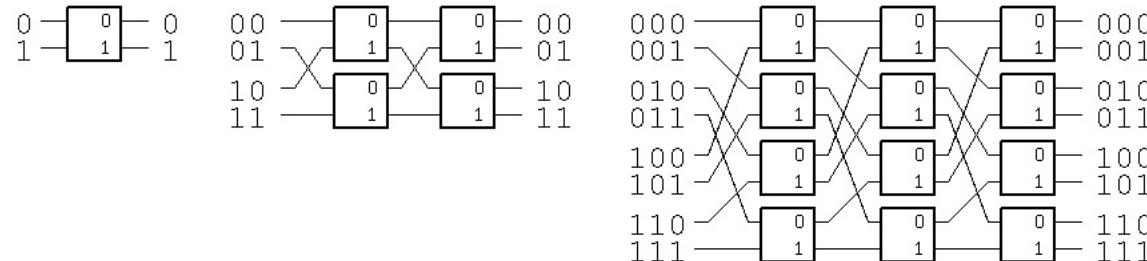
A multiprocessor may have distributed memory, shared memory or a combination of both.



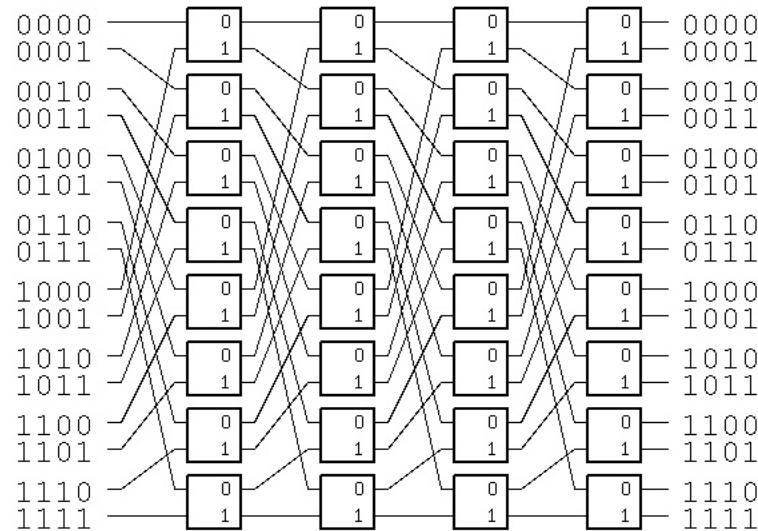
For the distributed memory and the shared memory multiprocessors, one possible connection, shown as a line above, is to use an omega network. The basic building block of an omega network is a switch with two inputs and two outputs. When a message arrives at this switch, the first bit is stripped off and the switch is set to: straight through if the bit is '0' on the top input or '1' on the bottom input else cross connected. Note that only one message can pass, the other being blocked, if two messages arrive and the exclusive or of the first bits is not '1'.



Then omega networks for connecting two devices, four devices or eight devices are built from this switch are shown below. The messages are sent with the most significant bit of the destination first.

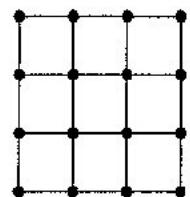


For 16 devices connected to the same or different 16 devices, the omega network is built from the primitive switch as:

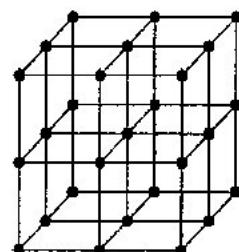


Note that connecting  $N$  devices requires  $N \log_2(N)$  switches. Given a set of random connections of  $N$  devices to  $N$  devices with an omega network, this is mathematically a permutation, then statistically  $1/2 N$  connections may be made simultaneously.

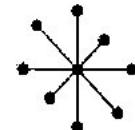
Then, we can call a CPU-memory pair a node, reduce the drawing of a node to a dot, and show a few connection topologies for multiprocessors



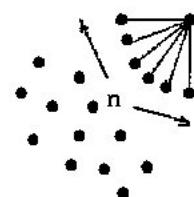
2D N=16  
X, Y  
Ports = 4  
Max path 6



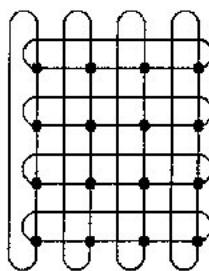
3D N=27  
X, Y, Z  
Ports = 6  
Max path 6



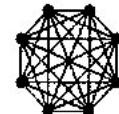
4D N = 9  
X, Y, Z, T  
Ports = 8  
Max path 1



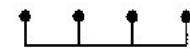
N-Cube N=2  
Ports = n  
Max path n



2D Torus N=16  
X, Y  
Ports = 4  
Max path 4



Star N=8  
Ports = N-1  
Max path 1



Bus N = 4  
Ports = 1  
Max path 1



Ring N = 4  
Ports = 2  
Max path 2

"Ports" is the number of I/O ports the node must have.

"Max path" is the maximum number of hops a message must take in order to get from one node to the farthest node. A message may be as small as a Boolean signal or as large as a big matrix.

The actual interconnect technology for those lines between the nodes has great variety. The lowest cost is Gigabit Ethernet while the best performance is with Myrinet and Infiniband.

## Interconnect

Interconnect	Count	Share %	Rmax Sum (GF)	Rpeak Sum (GF)	Processor Sum
Myrinet	54	10.80 %	282548	417651	68828
Gigabit Ethernet	256	51.20 %	794581	1511328	249994
Infiniband	35	7.00 %	219001	313513	52492
Quadrics	14	2.80 %	131249	173218	39108
Multi-stage crossbar	4	0.80 %	49852	55480	6136
Crossbar	7	1.40 %	54195	79566	7978
Myrinet/Infiniband	2	0.40 %	10858	16301	2804
NUMalink	9	1.80 %	55420	61028	9728
SP Switch	3	0.60 %	16714	25152	16768
Federation	33	6.60 %	246738	346643	42612
Colony	6	1.20 %	17646	34611	6656
HyperPlex	33	6.60 %	86738	142950	31718
Proprietary	26	5.20 %	629527	806058	291296
Fireplane	1	0.20 %	2054	3053	672
Hyper crossbar	1	0.20 %	10350	14336	2560
Numalink/Infiniband	2	0.40 %	58179	67872	11312
NUMalink/GigEthernet	1	0.20 %	2887	7680	1280
XT3 Internal Interconnect	9	1.80 %	109965	133803	28988
RapidArray	3	0.60 %	8388	10374	2357
InfiniPath	1	0.20 %	2073	2534	576
<b>Totals</b>	<b>500</b>	<b>100%</b>	<b>2788963.27</b>	<b>4223151.62</b>	<b>873863</b>

Now, the change 6 years later November 2012

Interconnect Top 500 Count Share (%)

Gigabit Ethernet	159	31.8
Infiniband QDR	106	21.2
Infiniband	59	11.8
Custom Interconnect	46	9.2
Infiniband FDR	45	9.0
10G Ethernet	30	6.0
Cray Gemini interconnect	15	3.0
Proprietary	11	2.2
Infiniband DDR	9	1.8
Aries interconnect	4	0.8
Infiniband DDR 4x	4	0.8
XT4 Internal Interconnect	4	0.8
Tofu interconnect	3	0.6
Myrinet 10G	3	0.6
Infiniband QDR Sun M9	1	0.2 new 100Gb/sec Ethernet
Mellanox 100G		

One measure of a multiprocessors communication capability is "bisection bandwidth". Optimally choose to split the processors into two equal groups and measure the maximum bandwidth that may be obtained between the groups.

Many modern multiprocessors are "clusters." Each node has a CPU, RAM, hard drive and communication hardware. The CPU may be dual or quad core and each CPU is considered a processor that may be assigned tasks. There is no display, keyboard, sound or graphics. The physical form factor is often a "blade" about 2 inches thick, 8 inches high and 12 inches deep with slide in connectors on the back. A blade may have multiple CPU chips each with multiple cores. 40 or more blades may be on one rack. Upon power up, each blade loads its operating system and applications from its local disk.

There is still a deficiency in some multiprocessor and multi core operating systems. The OS will move a running program from one CPU to another rather than leave a long running program and its

cache contents on one processor. Communication between multiprocesses may actually go out of a communication port and back into a communication processor when the processors are physically connected to the same RAM, rather than use memory to memory communication.

Another classification of multiprocessors is:  
 SISD Single Instruction Single Data (e.g. old computer)  
 SIMD Single Instruction Multiple Data (e.g. MASSPAR, CELL, GPU)  
 MIMD Multiple Instruction Multiple Data (e.g. cluster)

GPU stands for graphics processing unit, e.g. your graphics card that may have as many as 500 cores. Some of these cards have full IEEE double precision floating point in every core. There may be groups of cores that are SIMD and thus a group may be MIMD.

There are three main problems with massively parallel multiprocessors: software, hardware and software.

The operating systems are marginally useful for multiprogramming where a single program is to be run on a single data set using all the nodes and all the memory. Today, the OS is almost no help and the programmer must plan and program each node and every data transfer between nodes.

The programming languages are of little help. Java threads and Ada tasks are not guaranteed to run on individual processors. Posix threads are difficult to use and control. MPI and PVM libraries allow the programmer to specifically allocate tasks to nodes and control communication at the expense of significant programming effort.

Then there are programming classifications:

SPSD Single program Single Data (Conventional program)  
 SPMD Single Program Multiple Data (One program with "if" on all processors)  
 MPMD Multiple Program Multiple Data (Each processor has a unique program)

MPI Message Passing Interface is one of the SPMD toolkits that make programming distributed memory multiprocessors practical, yet still not easy.

There is a single program that runs on all processors with the allowance for if-then-else code dependent on processor number. The processor number may also be used for index and other calculations.

[My CMSC 455 lecture on MPI](#)

For shared memory parallel programming, threads are used, with one thread typically assigned to each cpu.

Only a small percent of application are in the class of "embarrassingly parallel". Most applications require significant design effort to obtain significant "speedup".

Yes, Amdahl's law applies to multiprocessors.  
 Given a multiprocessor with N nodes, the maximum speedup to be expected compared to a single processor of the same type as the node, is N. That would imply that 100% of the program could be made parallel.

Given 32 processors and 50% of the program can be made fully parallel, 25% of the program can use half the processors and the rest of the program must run sequentially, what is the speedup over one sequential processor?

Time sequentially is 100%	100%
50%    25%    25%	speedup = ----- = 3.55
Time multiprocessor is --- + --- + --- = 28.125%	28.125%
32    16    1	

far from the theoretical maximum of 32!

Note: "fully parallel" means the speedup factor is the number of processors.  
 "half the processors" in this case is 32/2 = 16.  
 the remaining 25% is sequential, thus factor = 1

Given 32 processors and 99% of the program can be fully parallel,

Time sequentially is 100%	100%
99%      1%	speedup = ----- = 24.4
Time multiprocessing is    --- + --- = 4.1%	4.1%
32            1	

about 3/4 the theoretical maximum of 32!

These easy calculations are only considering processing time.  
 In many programs there is significant communication time to  
 get the data to the required node and get the results to  
 the required node. A few programs may require more communication  
 time than computation time.

Consider a  $1024 = 2^{10}$  node multiprocessor.  
 Add  $1,048,576 = 2^{20}$  numbers as fast as possible on this multiprocessor.

Assume no communication cost (very unreasonable)

step	action
1	add $2^{10}$ numbers to $2^{10}$ numbers getting $2^{10}$ partial sums
2	add $2^{10}$ numbers to $2^{10}$ numbers getting $2^{10}$ partial sums

$\dots$

$2^9=512$  add  $2^{10}$  numbers to  $2^{10}$  numbers getting  $2^{10}$  partial sums

(so far fully parallel, now have only  $2^{19}$  numbers to add)

$2^{9+1}$  add  $2^{10}$  numbers to  $2^{10}$  numbers getting  $2^{10}$  partial sums

$2^{9+2}$  add  $2^{10}$  numbers to  $2^{10}$  numbers getting  $2^{10}$  partial sums

$\dots$

$2^{9+2^8}$  add  $2^{10}$  numbers to  $2^{10}$  numbers getting  $2^{10}$  partial sums

(so far fully parallel, now have only  $2^{18}$  numbers to add)

see the progression:

$2^9 + 2^8 + 2^7 + \dots 2^2 + 2^1 + 2^0 = 1023$  time steps  
 and we now have  $2^{10}$  partial sums, thus only  $2^9$  or 512  
 processors can be used on the next step

1024 add  $2^9$  numbers to  $2^9$  numbers getting  $2^9$  partial sums  
 (using 1/2 the processors)

1025 add  $2^8$  numbers to  $2^8$  numbers getting  $2^8$  partial sums  
 (using 1/4 the processors)

$\dots$

1033 add  $2^{0=1}$  number to  $2^{0=1}$  number to get the final sum  
 (using 1 processor)

sequential time 1,048,575  
 Thus our speedup is ----- = ----- = 1015  
 parallel time        1033

The percent utilization is  $1015/1024 * 100\% = 99.12\%$

Remember: Every program has a last, single, instruction to execute.  
 Jack Dongarra, an expert in the field of multiprocessor programming  
 says "It just gets worse as you add more processors."

Top 500 multiprocessors:

These have been and are evaluated by the Linpack Benchmark.  
 Heavy duty numerical computation. This Benchmark is close to  
 "embarrassingly parallel" and thus there is the start of a move  
 to the Graph 500 Benchmark that more fully measures the  
 interconnection capacity of the highly parallel machine.

[Graph500](#)

Some history of the top500:

[www.top500.org/lists/2006/06](http://www.top500.org/lists/2006/06)  
[www.top500.org/list/2007/11/100](http://www.top500.org/list/2007/11/100)  
[www.top500.org/lists/2008/11](http://www.top500.org/lists/2008/11)  
[www.top500.org/list/2015/06](http://www.top500.org/list/2015/06)

Over 1 million cores, over 12 megawatts of power.

[exascale](#)

[Gemini interconnect trying to solve the biggest problem](#)

[Latest VA Tech Machine](#)

Test your dual core, quad core, 8, 12 to be sure your operating system is assigning threads to different cores.

[time\\_mp2.c](#)

[time\\_mp4.c](#)

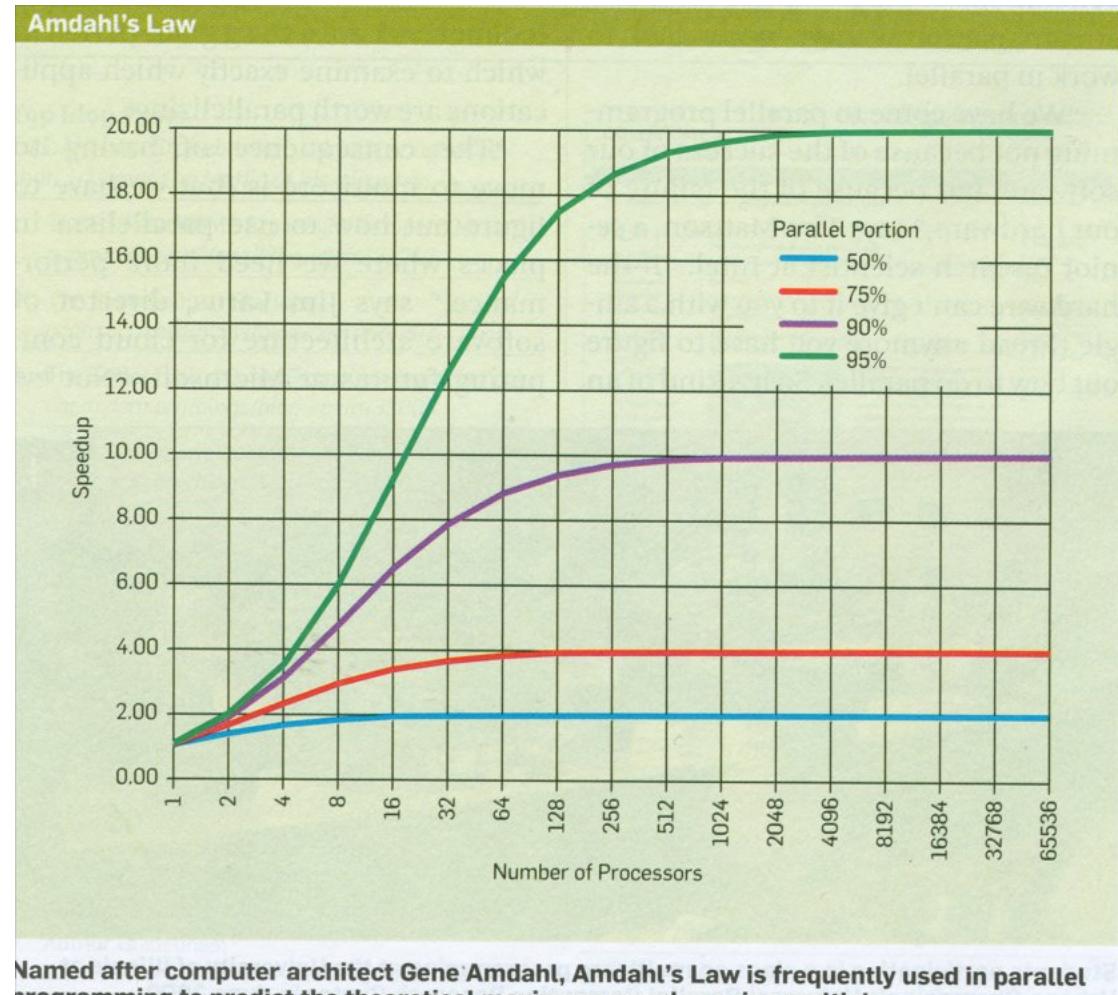
[time\\_mp8.c](#)

[time\\_mp12.c](#)

[time\\_mp12.c.out](#)

Here is a graph of Amdahl speedup for increasing number of processors, for 50%, 75%, 90% and 95% parallel execution.

As the curves flatten out, more processors or cores are useless.



[Tabular data](#)

[Project part3a hints](#)[diff1.png](#)[diff2.png](#)**Lecture 29, Review**

Covered on web: Previous Final Exam and Answers

Read over course WEB pages. (some have been updated)

Work all homeworks. (some similar problems on exam)

Do project at least through part2b. (some questions on exam)

**Lecture 30, Final Exam**

open book, open note, download, edit, submit  
 Edit by placing an x after a) b) c) that is your answer.  
 Edit with Microsoft Word on Windows, libreoffice on linux.g1

Finish homework and projects.

Students with email user name starting a b c d e f g h i  
 download and edit final32a.doc  
[download final32a.doc](#)

Students with email user name starting j k l m n o p q  
 download and edit final32b.doc  
[download final32b.doc](#)

Students with email user name starting r s t u v w x y z  
 download and edit final32c.doc  
[download final32c.doc](#)

Follow instructions in exam, edit, then  
 submit cs411 final final32?.doc

You can do the exam on linux.g1.umbc.edu in your directory

```
cp /afs/umbc.edu/users/s/q/squire/pub/download/final32?.doc .
libreoffice final32?.doc
submit cs411 final final32?.doc
rm final32?.doc
due by May 18, 2020
```

Before Exam:

Review HW2, HW3, HW4 (VHDL) and HW5  
 Review WEB Lecture Notes 14 through 29.

There are 10 types of people:  
 Those who know binary.  
 Those who do not know binary.

Teach your children to count in the computer age:  
 zero  
 one  
 two  
 three  
 four

Computer bits are numbered from the bottom

0	0	1	0	1	= 5
4	3	2	1	0	bit numbers (actually powers of 2)

last updated May 4, 2020

Last updated 10/28/09